

Omni-Compass Python Engine, Controllers and Harnesses: Source Code

Copyright © 2026 The Omni-Compass LLC. All rights reserved.

The Python engine of Omni-Compass (omnicompass/), its Kubernetes and GPU controllers and muscles (omni_controller/), the hardware and power-budget harnesses (hardware/), the tools (tools/), the benchmarks (benchmarks/) and the verification entry point (verify.py).

Version: git commit cd3f8426f3b68fdb128e2b4a9bb6021eaf3a895b

Date: 2026-09-28

Files: 54

Lines: 9,842

Use of this code is governed by the LICENSE file distributed with it.

Contents

omnicompass/adapter.py	310 lines	4
omnicompass/closure.py	251 lines	8
omnicompass/compass.py	164 lines	12
omnicompass/configurations.py	66 lines	15
omnicompass/conveyance.py	102 lines	16
omnicompass/core.py	241 lines	18
omnicompass/mathdrive.py	82 lines	21
omnicompass/muscles.py	155 lines	23
omnicompass/nervous.py	228 lines	25
omnicompass/nervous_system.py	159 lines	28
omnicompass/pools.py	34 lines	31
omnicompass/shield.py	107 lines	32
omnicompass/speed.py	92 lines	34
omnicompass/stack_sim.py	292 lines	36
omnicompass/storage.py	58 lines	40
omni_controller/controller.py	566 lines	41
omni_controller/gpu_governor.py	506 lines	49
omni_controller/muscles.py	681 lines	56
hardware/cpufreq.py	174 lines	66
hardware/node_exchange.py	157 lines	69
hardware/plant.py	188 lines	72
hardware/schedutil.py	98 lines	75
hardware/site_exchange.py	94 lines	77
tools/abc_report.py	446 lines	79
tools/abc_table.py	41 lines	86
tools/buyer_report.py	271 lines	87
tools/closed_loop_estimate.py	65 lines	91
tools/code_fingerprint.py	23 lines	92
tools/cpufreq_receipt.py	32 lines	93
tools/demo_nervous_system.py	75 lines	94
tools/full_report.py	791 lines	96
tools/gpu_baseline.py	71 lines	109
tools/gpu_physics_sim.py	177 lines	110
tools/gpu_pipeline_sim.py	156 lines	113
tools/gpu_reps.py	541 lines	116
tools/gpu_workload.py	145 lines	124
tools/live_reps.py	94 lines	126
tools/mechanism.py	113 lines	128
tools/mechanism_identity.py	165 lines	130
tools/protocol_bench.py	199 lines	133
tools/release_manifest.py	124 lines	136
tools/seal.py	106 lines	138
tools/source_book.py	135 lines	140

tools/strip_reference.py	34 lines	142
tools/three_way.py	93 lines	143
tools/tower_off_on.py	19 lines	145
tools/tracking_bounds.py	104 lines	146
tools/wall_meter.py	66 lines	148
benchmarks/core_evidence.py	63 lines	149
benchmarks/fleet_overhead.py	23 lines	150
benchmarks/multiplicity.py	65 lines	151
benchmarks/savings.py	61 lines	152
benchmarks/stack_benchmark.py	424 lines	153
verify.py	315 lines	159

omnicompass/adapter.py

310 lines SHA-256 57ca59da4024ccea2769083998fdeb7b7c9c05b5c0dbca638375350be7dd93ac

```
1  """Omni-Compass stack governor.
2
3  sense      : stack telemetry -> normalized observation vector
4  assimilate : observation blended into the six-state engine state
5  evolve     : one macro step of equations (1)-(7) with u = 0
6  allocate   : engine state -> directive (capacity, change gating, power, rollback, routing, alerting, paging)
7  authority  : OBSERVE | AUTOPILOT; kill() revokes authority and returns control to the native managers
8  """
9  from __future__ import annotations
10
11 import math
12 from collections import deque
13 from dataclasses import dataclass, asdict, field
14 from typing import Any, Dict, List, Optional
15
16 from .core import State, Params, macro_step, v_eff, control_command, U_AUTHORITY
17
18 OBSERVE = "observe"
19 AUTOPILOT = "autopilot"
20
21 # Engine parameters used for stack governance (fixed, not sampled).
22 STACK_PARAMS = Params(alpha=4.2, beta_int=0.0, beta_ext=0.0, k=1.7, sigma_l=0.38, delta=0.705, gamma_l=1.35,
23                       gamma_c=1.0, lambda_0=-0.40, lambda_1=1.55, lambda_2=1.15, c=1.0, E_max=2.0,
24                       omega_B=1.0, Q_B=2.5, alpha_s=0.12, beta_s=0.10, mu=3.782, alpha_E=3.487, alpha_U=4.2,
25                       lambda_I=0.719, lambda_U=0.0)
26 INITIAL_STATE = State(E=0.12, U=0.82, I_U=0.0, S=0.0, B=0.0, B_dot=0.0)
27 ASSIMILATION = 0.339
28
29
30 def clamp(v: float, lo: float, hi: float) -> float:
31     return lo if v < lo else hi if v > hi else v
32
33
34 def observe_vector(obs: Dict[str, float], conflicts: int) -> Dict[str, float]:
35     """Normalize telemetry. Inputs are the stack's delayed observations."""
36     return {
37         "q": clamp(float(obs.get("queue_ratio", 0.0)), 0.0, 2.0),
38         "load": clamp(float(obs.get("load_ratio", 0.0)), 0.0, 2.0),
39         "power": clamp(float(obs.get("power_stress", 0.0)), 0.0, 1.5),
40         "thermal": clamp(float(obs.get("thermal", 0.0)), 0.0, 1.5),
41         "network": clamp(float(obs.get("network_stress", 0.0)), 0.0, 1.5),
42         "drift": clamp(float(obs.get("drift_ratio", 0.0)), 0.0, 1.5),
43         "stale": clamp(float(obs.get("stale", 0.0)), 0.0, 1.0),
44         "security": clamp(float(obs.get("security_block", 0.0)), 0.0, 1.0),
45         "conflict": clamp(conflicts / 4.0, 0.0, 1.0),
46     }
47
48
49 def assimilate(x: State, o: Dict[str, float], p: Params) -> State:
50     """Blend the observation into the engine state and set the forcing terms."""
51     q, load, pw, th, nw, dr, st, sec, cf = (o[k] for k in
52         ("q", "load", "power", "thermal", "network", "drift", "stale", "security", "conflict"))
53     e_obs = clamp(0.25 * q + 0.18 * max(0.0, load - 0.85) + 0.16 * pw + 0.13 * th + 0.10 * nw + 0.08 * dr
54                 + 0.06 * cf + 0.04 * st, 0.0, 1.5)
55     u_obs = clamp(1.0 - (0.27 * q + 0.18 * pw + 0.16 * th + 0.12 * nw + 0.12 * dr + 0.10 * cf + 0.05 * st), 0.0,
56                 1.0)
57     s_obs = clamp(0.36 * th + 0.28 * pw + 0.18 * nw + 0.10 * sec + 0.08 * st - 0.20 * q, -1.0, 1.0)
58     b_obs = clamp(0.52 * pw + 0.26 * th + 0.22 * dr, -1.0, 1.5)
59     i_obs = clamp(q + dr + cf, 0.0, 2.0)
60     a = ASSIMILATION
61     p.beta_int = clamp(0.58 * q + 0.42 * max(0.0, load - 0.75), 0.0, 1.0)
62     p.beta_ext = clamp(0.33 * pw + 0.27 * th + 0.20 * nw + 0.12 * st + 0.08 * sec, 0.0, 1.0)
63     return State(E=(1 - a) * x.E + a * e_obs, U=(1 - a) * x.U + a * u_obs, I_U=(1 - a) * x.I_U + a * i_obs,
64                 S=(1 - a) * x.S + a * s_obs, B=(1 - a) * x.B + a * b_obs, B_dot=(1 - a) * x.B_dot + a * (b_obs -
65                 x.B))
66
67 @dataclass(frozen=True)
68 class AllocationLaw:
69     """Engine state -> governing directive.
70
71     release capacity removed only while push <= push_release (the engine's controller reports convergence)
72     add      without backlog, capacity added only while push >= push_add (the engine's controller reports need)
73     push     equation (2) control law evaluated on the evolved state with target sigma = +1:
74             u = clip(-f_U(x) + KP (1 - U), -U_AUTHORITY, U_AUTHORITY); push = u / U_AUTHORITY
75             (f_U contains the double-well drift and -(dE/dt)/E_max from equation (1))
76     headroom rho* = clamp(rho0 - kI*I_U - kE*E, rho_min, rho0)
77     sizing   load_s = load * cap_now if size_at_full_cap else load
78     capacity n_req = ceil(n * load / rho*) + ceil(kq * q * n); delta = clamp(n_req - n, -down_max, +up_max)
79     guard    no capacity removal while q >= guard_queue
80     band     capacity removed only when n_req <= n - down_band
81     dwell    capacity removed only after surplus persists for down_dwell intervals and no earlier than
82             down_after_add intervals after the last addition
```

```

82     envelope    no capacity addition while observed power >= power_ceiling (power_ceiling_backlog while q >=
guard_queue)
83     security    no capacity addition while a security block is observed
84     change      Terraform plans only while U >= U_gate, only in the direction of the live capacity decision
85     power       q < guard_queue and U >= U_gate: cap = clamp(cap_now * load * (1 + margin), cap_min, 1)  (trim
to demand)
86               q >= guard_queue at the power envelope: cap = max(cap_min, cap_now - spread_step)      (spread)
87               q >= guard_queue below the envelope: cap held (capacity added as nodes; 1 at max_nodes)
88               otherwise cap = 1; always cap <= 1 - cap_gain * max(0, B - B_cap)
89     rollback    a proposed rollback is authorised if S >= S_rollback, U is falling, or security blocks change
90     replicas    direct actuation: desired = ceil(replicas * load / rho*) (the HPA formula with the engine-set
target),
91               increases unless capacity is being reduced, decreases only while reducing;
92               otherwise autoscaler increases approved unless reducing, decreases only while reducing
93     consistency a step that expands capacity never tightens the power cap
94     paging      none: the governor holds every action a paged operator would take
95     """
96     rho0: float = 0.914
97     rho_min: float = 0.627
98     kI: float = 0.252
99     kE: float = 0.297
100    kq: float = 0.795
101    up_max: int = 2
102    down_max: int = 1
103    guard_queue: float = 0.05
104    power_ceiling: float = 0.95
105    power_ceiling_backlog: float = 1.00
106    spread_step: float = 0.05
107    backlog_release: bool = False
108    size_at_full_cap: bool = False
109    push_release: float = 0.005
110    push_add: float = -9.0
111    U_gate: float = 0.481
112    margin: float = 0.006
113    cap_min: float = 0.706
114    B_cap: float = 0.40
115    cap_gain: float = 0.60
116    S_rollback: float = 0.55
117    route_network: float = 0.20
118    down_band: int = 4
119    down_dwell: int = 4
120    down_after_add: int = 4
121    min_nodes: int = 8
122    max_nodes: int = 40
123
124
125 @dataclass
126 class Governor:
127     law: AllocationLaw = field(default_factory=AllocationLaw)
128     mode: str = AUTOPILOT
129     killed: bool = False
130     x: State = field(default_factory=lambda: State(**asdict(INITIAL_STATE)))
131     p: Params = field(default_factory=lambda: Params(**asdict(STACK_PARAMS)))
132     t: float = 0.0
133     current_cap: float = 1.0
134     nodes: int = 20
135     surplus_streak: int = 0
136     since_add: int = 99
137     last_push: float = 0.0
138     log: deque = field(default_factory=lambda: deque(maxlen=4096))
139     evolve: bool = True
140
141     def kill(self) -> None:
142         """Revoke authority; the engine continues to compute and log directives."""
143         self.killed = True
144
145     def set_mode(self, mode: str) -> None:
146         if mode not in (OBSERVE, AUTOPILOT):
147             raise ValueError(mode)
148         self.mode = mode
149
150     @property
151     def has_authority(self) -> bool:
152         return self.mode == AUTOPILOT and not self.killed
153
154     def step(self, obs: Dict[str, float], conflicts: int) -> Dict[str, Any]:
155         """Sense, assimilate, evolve, allocate."""
156         o = observe_vector(obs, conflicts)
157         U_before = self.x.U
158         xa = assimilate(self.x, o, self.p)
159         if self.evolve:
160             xn, _ = macro_step(xa, self.p, self.t, target=None)
161         else:
162             xn = xa
163         self.x = xn
164         self.t += 0.1

```

```

165     L = self.law
166     n = int(self.nodes)
167     u_push, _ = control_command(xn, self.p, self.t, +1)
168     push = u_push / U_AUTHORITY
169     self.last_push = push
170     rho = clamp(L.rho0 - L.kI * xn.I_U - L.kE * xn.E, L.rho_min, L.rho0)
171     converged = push <= L.push_release
172     load_sized = o["load"] * self.current_cap if L.size_at_full_cap else o["load"]
173     n_req = math.ceil(n * load_sized / rho) + math.ceil(L.kq * o["q"] * n)
174     delta = max(-L.down_max, min(L.up_max, n_req - n))
175     if delta < 0 and not converged:
176         delta = 0
177     if delta > 0 and o["q"] < L.guard_queue and push < L.push_add:
178         delta = 0
179     if delta < 0 and (o["q"] >= L.guard_queue or n_req > n - L.down_band):
180         delta = 0
181     self.surplus_streak = self.surplus_streak + 1 if delta < 0 else 0
182     if delta < 0 and (self.surplus_streak < L.down_dwell or self.since_add < L.down_after_add):
183         delta = 0
184     ceiling = L.power_ceiling_backlog if o["q"] >= L.guard_queue else L.power_ceiling
185     if delta > 0 and (o["power"] >= ceiling or n >= L.max_nodes):
186         delta = 0
187     if o["security"] > 0.5:
188         delta = min(delta, 0)
189     self.since_add = 0 if delta > 0 else self.since_add + 1
190     if delta < 0:
191         self.surplus_streak = 0
192     at_envelope = o["power"] >= ceiling
193     if o["q"] < L.guard_queue and xn.U >= L.U_gate:
194         cap = clamp(self.current_cap * o["load"] * (1.0 + L.margin), L.cap_min, 1.0)
195     elif o["q"] >= L.guard_queue and L.spread_step > 0.0:
196         if at_envelope:
197             cap = clamp(self.current_cap - L.spread_step, L.cap_min, 1.0)
198         else:
199             cap = 1.0 if L.backlog_release else (self.current_cap if n < L.max_nodes else 1.0)
200     elif o["q"] >= L.guard_queue and L.backlog_release:
201         cap = 1.0
202     else:
203         cap = 1.0
204     cap = min(cap, clamp(1.0 - L.cap_gain * max(0.0, xn.B - L.B_cap), L.cap_min, 1.0))
205     if delta > 0:
206         cap = max(cap, self.current_cap)
207     command = clamp(0.5 * (v_eff(xn, self.p) / max(self.p.c, 1e-9) + 1.0), 0.0, 1.0)
208     d = {
209         "node_delta": int(delta),
210         "change_permitted": xn.U >= L.U_gate and o["security"] <= 0.5,
211         "power_cap": cap,
212         "rollback_authorized": xn.S >= L.S_rollback or xn.U < U_before or o["security"] > 0.5,
213         "route_shift": o["network"] > L.route_network,
214         "state": {"E": xn.E, "U": xn.U, "I_U": xn.I_U, "S": xn.S, "B": xn.B},
215         "demand": rho,
216         "load": o["load"],
217         "command": command,
218     }
219     self.log.append({"mode": self.mode, "killed": self.killed, **{k: v for k, v in d.items() if k !=
"state"}},
220                    **d["state"]})
221     return d
222
223
224 def directive_to_actions(d: Dict[str, Any], stack_state: Dict[str, Any], proposals: List[Dict[str, Any]],
225                          cfg, direct_actuation: bool = False) -> List[Dict[str, Any]]:
226     """Directive -> executable stack actions. direct_actuation: capacity is owned by the governor; no Terraform
plans."""
227     tag = "Omni-Compass"
228     out: List[Dict[str, Any]] = []
229     cur = int(stack_state["actual_nodes"])
230     target = int(max(cfg.minimum_nodes, min(cfg.maximum_nodes, cur + int(d["node_delta"]))))
231     if target != cur:
232         out.append({"manager": tag, "action": "nodes", "target": target, "direction": 1 if target > cur else -1})
233     if not direct_actuation and d["change_permitted"] and any(p["action"] == "terraform_plan" for p in
proposals):
234         tf_dir = 1 if target > stack_state["terraform_desired_nodes"] else -1
235         live_dir = 1 if target > cur else -1 if target < cur else 0
236         if abs(target - stack_state["terraform_desired_nodes"]) >= 2 and live_dir in (0, tf_dir):
237             out.append({"manager": tag, "action": "terraform_plan", "target": target,
238                        "direction": 1 if target > stack_state["terraform_desired_nodes"] else -1})
239     reps = [int(p["target"]) for p in proposals if p["action"] == "replicas"]
240     if direct_actuation:
241         cur_r = int(stack_state["replicas"])
242         rt = max(20, min(500, math.ceil(cur_r * float(d["load"]) / max(float(d["demand"]), 1e-9))))
243         if (rt > cur_r and d["node_delta"] >= 0) or (rt < cur_r and d["node_delta"] < 0):
244             out.append({"manager": tag, "action": "replicas", "target": rt, "direction": 1 if rt > cur_r else -
1})
245     reps = []
246     if reps:

```

```

247         cur_r = int(stack_state["replicas"])
248         rt = None
249         if d["node_delta"] >= 0 and max(reps) > cur_r:
250             rt = max(reps)
251         elif d["node_delta"] < 0 and min(reps) < cur_r:
252             rt = min(reps)
253         if rt is not None:
254             out.append({"manager": tag, "action": "replicas", "target": rt,
255                        "direction": 1 if rt > cur_r else -1})
256         cap = float(d["power_cap"])
257         if any(a["direction"] > 0 and a["action"] in ("nodes", "terraform_plan", "replicas") for a in out):
258             cap = max(cap, float(stack_state["power_cap"]))
259         if abs(cap - stack_state["power_cap"]) > 1e-9:
260             out.append({"manager": tag, "action": "power_cap", "target": cap,
261                        "direction": -1 if cap < stack_state["power_cap"] else 1})
262         if d["route_shift"]:
263             out.append({"manager": tag, "action": "route_shift", "target": 1.0, "direction": 1})
264         if any(p["action"] == "rollback" for p in proposals) and d["rollback_authorized"]:
265             out.append({"manager": tag, "action": "rollback", "target": 1.0, "direction": -1})
266         alerts = [p for p in proposals if p["action"] == "alert"]
267         if alerts:
268             out.append({"manager": tag, "action": "alert", "target": max(int(p["target"]) for p in alerts),
269                        "direction": 1})
270         return out
271
272 THROUGHPUT_MODE = dict(backlog_release=True, power_ceiling=9.0, power_ceiling_backlog=9.0, spread_step=0.0,
273                        B_cap=9.0,
274                        size_at_full_cap=True, kI=0.05, kE=0.40, cap_min=0.65, down_band=1, down_dwell=10,
275                        down_after_add=3, kq=0.30, U_gate=0.45, push_release=0.01)
276
277 FLEET_MODE = dict(THROUGHPUT_MODE, rho0=0.829, rho_min=0.769, kI=0.072, kE=0.487, kq=0.635, down_band=3,
278                  down_dwell=1, down_after_add=5, push_release=0.053, cap_min=0.695, margin=0.131, U_gate=0.85, up_max=8)
279 FLEET_BALANCED_MODE = dict(FLEET_MODE, push_release=0.0098, push_add=0.0098)
280 FLEET_WEAR_MODE = dict(FLEET_MODE, push_release=0.0035, push_add=0.0035)
281 FLEET_DECISION_SECONDS = 60
282
283 def mode_law(mode: str, base: "AllocationLaw" = None) -> "AllocationLaw":
284     """Operating modes.
285
286     power_protect the site power limit is enforced (envelope, spread-and-throttle, bath cap; shield I4 active)
287     fleet         throughput mode with constants selected on the 15-second fleet harness (fleet/), decisions
288                  every 60 s;
289     fleet_balanced energy-first: the equation (2) release gate is lenient and the add gate is off
290     fleet_wear     fleet with the equation (2) gate active in both directions at 0.0098 (release and add)
291     throughput     fleet with the equation (2) gate active in both directions at 0.0035 (release and add)
292     backlog        the power envelope is not enforced (same power rules as the native stack); caps released on
293                  backlog;
294                  capacity sized at full power (load * cap_now), so power trimming does not trigger machine
295                  starts;
296                  constants selected by constrained search on development seeds (minimum energy subject to
297                  backlog,
298                  health, recovery, reversal and invariant constraints)
299
300     """
301     from dataclasses import replace
302     base = base or AllocationLaw()
303     if mode == "power_protect":
304         return base
305     if mode == "throughput":
306         return replace(base, **THROUGHPUT_MODE)
307     if mode == "fleet":
308         return replace(base, **FLEET_MODE)
309     if mode == "fleet_balanced":
310         return replace(base, **FLEET_BALANCED_MODE)
311     if mode == "fleet_wear":
312         return replace(base, **FLEET_WEAR_MODE)
313     raise ValueError(mode)

```

omnicompass/closure.py

251 lines SHA-256 cc91029928c57b97871ab5e4189a8eb998897753bb9c1cebbef8d70095413b93

```
1  """Closure governor: machines governed by the Omni-Compass closure law, preemptive coherence and the dual-bath
   tracker.
2
3  Sources (the owner's manuscript): Section 5 (5c preemptive coherence: forward projection  $\tau_{\text{fwd}} = \Phi_{\text{dt}}(x)$  and
4  correction  $-\lambda \text{grad } \psi$ ), Chapter 20 (master closure law  $F = G_0 + G_c$ ,  $G_c = -\rho \beta \text{Dh}$ , zero in the
   interior,
5  inward near the boundary with margin  $\delta$ ), Chapters 29-30 (the wheel: expansion, turning point  $dE/dt = 0$ ,
   compression;
6  release only after the turn, ignition before the threshold), Chapter 31 (dual-bath exchange  $dU/dt = -\alpha E + \beta S$ ,
    $dE/dt = \alpha U - \gamma B$ ).
7
8
9  Cluster realisation (one pool)
10 state      r(t) requested cores (what the pods ask for); n machines; c = cores x ALLOC per machine
11 dual bath  level L and rate v of r, the exchange pair of Chapter 31 in discrete form:
12            innovation  $e = r - (L + v \text{ dt})$ 
13             $L \leftarrow L + v \text{ dt} + a e$  (coherence bath absorbs the observation)
14             $v \leftarrow (1 - g) v + b e / \text{dt}$  (deviation bath exchanges with it; g = redistribution damping)
15 projection  $r_{\text{fwd}}(\tau) = L + v \tau$  for  $\tau$  in  $[0, H]$  (5c:  $\Phi$  over the look-ahead horizon)
16 admissible  $h(\tau) = r_{\text{fwd}}(\tau) / (n c) - \rho_{\text{max}} \leq 0$  (Chapter 20: the admissible domain)
17 closure  $G_c$  if  $\max\{\tau \leq H_{\text{add}}\} h(\tau) > -\delta$  (projected boundary approach, within one boot of lead):
18            add  $k = \text{ceil}((\max r_{\text{fwd}} / (\rho_{\text{max}} - \delta) - n c) / c)$  machines (beta sized to
   restore delta)
19 interior  $G_0$  energy descent: remove one machine only if it stays admissible with margin for the whole release
   horizon,
20             $\max\{\tau \leq H_{\text{rel}}\} r_{\text{fwd}}(\tau) / ((n - 1) c) \leq \rho_{\text{max}} - \delta$ ,
21            and only after the turn ( $v \leq 0$ : past the peak, the metric flip), and only while the six-state
   engine's
22            equation (2) push  $\leq$  push_release
23 interior  nothing else happens: inside the domain the platform physics runs unmodified ( $G_c = 0$ )
24 """
25 from __future__ import annotations
26
27 import math
28 from dataclasses import dataclass
29
30
31 @dataclass(frozen=True)
32 class ClosureLaw:
33     rho_max: float = 0.95 # admissibility boundary: requested cores / allocatable cores
34     delta: float = 0.05 # inwardness margin
35     H_add: int = 8 # look-ahead for the closure correction (ticks);  $\geq$  the boot delay
36     H_rel: int = 40 # look-ahead a release must survive (ticks)
37     a: float = 0.5 # dual-bath coupling: level absorption
38     b: float = 0.1 # dual-bath coupling: rate exchange
39     g: float = 0.05 # redistribution damping of the rate bath
40     push_release: float = 0.2 # engine gate on release
41     turn: bool = True # release only after the turning point ( $v \leq 0$ )
42     delta_rel: float = -1.0 # release band: a release must fit at  $\rho_{\text{max}} - \delta_{\text{rel}}$  (hysteresis;  $< 0$  means
   = delta)
43     z: float = 0.0 # deviation-bath band: a release must survive z standard deviations of the
   tracker's
44     noise; # innovation over the release horizon (Chapter 31: the deviation bath carries the
45     # the band is wide where demand is noisy, tight where it is calm)
46     tone: bool = False # muscle tone: release = park (alive, low power, instant wake), not power-off;
   parked
47     # machines beyond the reserve needed within tone_H are powered off
48     site: bool = False # whole body (multi-cluster sites): the law runs on the site total, one warm
   reserve for
49     # the site, and traffic shift lets a cluster use another's already-powered
   machines
50     tone_H: int = 960 # reserve horizon (ticks): the largest demand seen over this window sets the warm
   reserve
51     coord: bool = False # nervous-system coordination: pods decided first, machines on those pods, and
   the two
52     # organs never move in opposite directions in one decision
53     wake_first: bool = False # whole body: a site that needs a machine wakes a parked one in any cluster
   before it cold-
54     # boots one (traffic shift carries the load; a wake costs no start, a boot costs
   one)
55     dwell: int = 0 # release only after this many consecutive calm decisions (resource-aware
   envelope,
56     # Proposition 2: no release outside the calm set; 0 = no dwell requirement)
57     turn_rise: float = 0.0 # the turning-point veto binds only for a rise larger than this share of one
   machine over
58     # the release horizon ( $v * H_{\text{rel}} > \text{turn\_rise} * c$ ); 0 = any rise vetoes (the
   benchmarked
59     # law). A rise smaller than the release band cannot carry demand through the band,
   so it
60     # is no reason to hold a machine (one small rise must not hold every machine for
   a run)
61
```

```

62
63 class ClosureNodes:
64     def __init__(self, law: ClosureLaw = ClosureLaw()):
65         self.L = None; self.v = 0.0; self.law = law; self.calm = 0; self.s2 = 0.0; self.hist = []
66
67     def observe(self, r: float, dt: float = 1.0) -> None:
68         a, b, g = self.law.a, self.law.b, self.law.g
69         if self.law.tone:
70             self.hist.append(r)
71             if len(self.hist) > self.law.tone_H:
72                 self.hist.pop(0)
73         if self.L is None:
74             self.L = r; return
75         pred = self.L + self.v * dt
76         e = r - pred
77         self.s2 = 0.95 * self.s2 + 0.05 * e * e
78         self.L = pred + a * e
79         self.v = (1.0 - g) * self.v + b * e / dt
80
81     def reserve(self, c: float) -> int:
82         """Machines the law expects to need within the tone horizon (largest recent demand at the working
boundary)."""
83         if not self.hist:
84             return 0
85         return int(math.ceil(max(self.hist) / ((self.law.rho_max - self.law.delta) * c) - 1e-9))
86
87     def fwd_max(self, H: int) -> float:
88         return max(self.L + self.v * tau for tau in range(0, H + 1)) if self.L is not None else 0.0
89
90     def decide(self, n: int, c: float, push: float, n_min: int, n_max: int) -> int:
91         """Target machine count for this tick."""
92         L = self.law
93         if self.L is None:
94             return n
95         peak_add = max(self.fwd_max(L.H_add), 0.0)
96         if n <= 0 or peak_add / (n * c) - L.rho_max > -L.delta:
97             self.calm = 0
98             need = math.ceil(peak_add / ((L.rho_max - L.delta) * c) - 1e-9)
99             return int(min(n_max, max(n_min, need, n)))
100         peak_rel = max(self.fwd_max(L.H_rel), 0.0) + L.z * math.sqrt(self.s2 * max(1, L.H_rel)) * L.a
101         band = L.delta if L.delta_rel < 0 else L.delta_rel
102         calm = n - 1 >= n_min and peak_rel / ((n - 1) * c) <= L.rho_max - band \
103             and (not L.turn or self.v * L.H_rel <= L.turn_rise * c) and push <= L.push_release
104         self.calm = self.calm + 1 if calm else 0
105         if calm and self.calm > L.dwell:
106             self.calm = 0
107         return n - 1
108     return n
109
110
111 # -----
112 # Self-calibrating closure law: every margin is derived, none is tuned per workload.
113 #
114 # tracker      dual bath (level L, rate v); the gains are chosen by the data: a bank of trackers from slow-
trend to fast-
115 #              trend (a, b) runs side by side and the one with the smallest running one-step prediction error
drives the
116 #              forecast (model selection by prediction error; the fastest member is the critically damped a =
1/2,
117 #              b = a^2 / (2 - a))
118 # noise        sigma^2 = running variance of the tracker's one-step innovation (the deviation bath of Chapter
31)
119 # add horizon  H_add = boot delay + 1 tick: a machine ordered now is serving when the projection arrives
120 # release hor. H_rel = boot delay / park fraction: the energy break-even between keeping a machine parked and
booting
121 #              it cold again (a parked machine draws park_frac x idle per tick, a cold start draws boot x idle)
122 # boundary     requested cores <= allocatable cores (rho_max = 1, the physical limit); no fixed margin
123 # margin       z(t) sigma sqrt(H): the forecast error over the horizon at confidence z;
124 #              z(t) = z95 x max(1, S / S*), S the six-state engine's stress (equation 6) and S* its
equilibrium,
125 #              delta - alpha_s S* - (3/4) beta_s S*^2 = 0: the engine widens the margin when it is stressed
126 # Gc (add)     if projected peak over H_add + margin > n c: add ceil((peak + margin) / c) - n machines
127 # G0 (release) one machine, only if projected peak over H_rel + margin over max(H_rel, T*) <= (n - 1) c (the
machine is
128 #              not expected back within the learned break-even T* below), after the turn (v <= 0), and while
129 #              the engine's push <= push_release; the released machine is parked
130 # power-off    learned break-even (ski rental with measured return times): every release at level k opens a
record
131 #              that closes when demand needs k machines again; X = the time it took. A parked machine is
powered off
132 #              when it has been parked T* ticks, T* = argmin_T sum_i cost_i(T), cost_i = park_frac X_i if X_i
<= T,
133 #              else park_frac T + boot (energy in idle-machine ticks; still-open records are right-censored at
their
134 #              age). Before any record exists T* = boot / park_frac, the deterministic ski-rental rule (never

```

```

    more than
135 #         twice the energy of a controller that knows the future). Nothing is tuned: T* is computed from
    the
136 #         workload's own return times, per workload, while it runs
137 # -----
138 @dataclass(frozen=True)
139 class AutoClosureLaw:
140     boot: int = 6                # boot delay of a cold machine (ticks): a property of the plant
141     park_frac: float = 0.25      # parked power / idle power: a property of the hardware
142     z95: float = 1.6449         # one-sided 95% normal quantile
143     a: float = 0.5
144     push_release: float = 0.2
145     site: bool = False
146     # read by the plant (same interface as ClosureLaw)
147     tone: bool = True
148     rho_max: float = 1.0
149     delta: float = 0.0
150
151     @property
152     def b(self) -> float:
153         return self.a * self.a / (2.0 - self.a)
154
155     @property
156     def H_add(self) -> int:
157         return self.boot + 1
158
159     @property
160     def H_rel(self) -> int:
161         return int(math.ceil(self.boot / self.park_frac))
162
163     @property
164     def tone_H(self) -> int:
165         return self.H_rel
166
167
168 def stress_equilibrium(delta: float, alpha_s: float, beta_s: float) -> float:
169     """Positive root of  $\delta - \alpha_s S - (3/4) \beta_s S^2 = 0$  (the engine's stress equilibrium, equation 6).
    """
170     qa = 0.75 * beta_s
171     return (-alpha_s + math.sqrt(alpha_s * alpha_s + 4.0 * qa * delta)) / (2.0 * qa) if qa > 0 else delta /
    alpha_s
172
173
174 class AutoClosureNodes:
175     def __init__(self, law: AutoClosureLaw = AutoClosureLaw(), s_eq: float = None):
176         self.law = law; self.L = None; self.v = 0.0; self.s2 = 0.0; self.n_obs = 0
177         self.S = 0.0; self.s_eq = s_eq if s_eq else stress_equilibrium(0.5, 0.12, 0.10)
178         self.hist = []; self.ages = []; self.n_last = 0
179         self.t = 0; self.open = []; self.closed = [] # release records: (level, t0) open; X closed
180
181     BANK = ((0.2, 0.002), (0.2, 0.02), (0.5, 0.005), (0.5, 0.05), (0.5, 1.0 / 6.0))
182
183     def observe(self, r: float, dt: float = 1.0) -> None:
184         if self.L is None:
185             self.L = r; self.bank = [[r, 0.0, 0.0] for _ in self.BANK]; return
186         self.n_obs += 1; self.t += 1
187         self.ages = [x + 1 for x in self.ages]
188         w = max(1.0 / self.n_obs, 0.02) # running mean first, then a 50-tick exponential window
189         for (a, b), m in zip(self.BANK, self.bank):
190             pred = m[0] + m[1] * dt
191             e = r - pred
192             m[2] = (1.0 - w) * m[2] + w * e * e
193             m[0] = pred + a * e
194             m[1] = m[1] + b * e / dt
195         best = min(self.bank, key=lambda m: m[2])
196         self.L, self.v, self.s2 = best
197
198     def z(self) -> float:
199         return self.law.z95 * max(1.0, self.S / self.s_eq) if self.s_eq > 0 else self.law.z95
200
201     def peak(self, H: int, H_noise: int = None) -> float:
202         """Projected peak over H ticks (level + trend) plus the forecast-error margin accumulated over H_noise
    ticks."""
203         if self.L is None:
204             return 0.0
205         Hn = H if H_noise is None else H_noise
206         return max(self.L + self.v * tau for tau in range(0, H + 1)) + self.z() * math.sqrt(self.s2 * Hn)
207
208     GRID = (0, 6, 12, 24, 48, 96, 192, 384, 768, 1440)
209
210     def threshold(self) -> int:
211         L = self.law
212         if not self.closed and not self.open:
213             return L.H_rel
214         cens = [self.t - t0 for _, t0 in self.open]

```

```

215     def cost(T):
216         c = sum(L.park_frac * x if x <= T else L.park_frac * T + L.boot for x in self.closed)
217         return c + sum(L.park_frac * T + L.boot if a > T else L.park_frac * a for a in cens)
218     return min(self.GRID, key=lambda T: (cost(T), T))
219
220     def reserve(self, c: float) -> int:
221         """Machines to keep powered (running + parked younger than the learned break-even T*); older parked go
off."""
222         T = self.threshold()
223         self.ages = [x for x in self.ages if x < T]
224         return self.n_last + len(self.ages)
225
226     def decide(self, n: int, c: float, push: float, n_min: int, n_max: int) -> int:
227         L = self.law
228         if self.L is None:
229             return n
230         need = int(math.ceil(max(self.peak(L.H_add), 0.0) / c - 1e-9))
231         still = []
232         for lvl, t0 in self.open:          # a release record closes when demand needs its level again
233             if need >= lvl:
234                 self.closed.append(self.t - t0)
235             else:
236                 still.append((lvl, t0))
237         self.open = still[-200:]; self.closed = self.closed[-400:]
238         if n <= 0 or need > n:
239             t = int(min(n_max, max(n_min, need)))
240             self.ages.sort(); del self.ages[max(0, len(self.ages) - (t - n)):]    # wake the longest parked
first
241             self.n_last = t
242             return t
243         # release only if the machine is not expected back within the learned break-even T*: the trend is
projected over
244         # the release horizon, the forecast error is accumulated over max(H_rel, T*)
245         if n - 1 >= n_min and self.peak(L.H_rel, max(L.H_rel, self.threshold())) <= (n - 1) * c and self.v <= 0.
0 \
246             and push <= L.push_release:
247             self.ages.append(0); self.n_last = n - 1; self.open.append((n, self.t))
248             return n - 1
249         self.n_last = n
250         return n
251

```

omnicompass/compass.py

164 lines SHA-256 b96e934c3919c8d8fd5ffbc2a93ed2f9152ef794abd806bcc6d400020dfab7d

```
1  """The Omni-Compass, computed.
2
3  I read my own state on the compass every decision and say where I am, in the compass's own terms.
4
5  The wheel (Chapter 29). My deviation E, scaled by its ceiling E_max, runs along the horizontal; its rate of
  change runs
6  along the vertical. A reading is a point on that wheel. Its heading is measured clockwise from north, as on the
  face:
7
8      +   A       0 deg   rising through rest
9      ≠   Δ       45 deg   expansion: above rest and rising, exchange under way
10     >   B       90 deg   peak extension: the turn, where the metric flips (k -> -k)
11     T   Λ       135 deg  dispersion: above rest and easing, the ceiling holds
12     -   Ω       180 deg  falling through rest
13     ≈   Π       225 deg  compression: below rest and settling into the basin
14     <   Γ       270 deg  deepest compression: the turn at the floor
15     +   Ψ       315 deg  re-alignment: below rest and rising, ignition ahead
16
17  The four quadrants are the strokes of the closed circle (Chapters 29-30):
18      I   Expansion   above rest, rising
19      II  Dispersion  above rest, easing
20      III Compression below rest, falling
21      IV Re-alignment below rest, rising
22  The wheel turns I -> II -> III -> IV -> I. Crossing from IV into I is ignition; the turn at peak extension is the
23  metric flip. The state is continuous through every crossing: nothing switches off, nothing restarts.
24
25  The rim carries the 24 letters in their places on the face, 15 degrees each, clockwise from A at north.
26
27  The axle is the structural basin S. Its rest point S* solves equation (6):  $\delta - \alpha_s S - (3/4) \beta_s S^2 = 0$ .
28
29  Closing the circle (Chapter 1, the Unified Circle Principle):
30      X' = G(X), X(0) in Omega           Omega is the living band: every level inside [0.05, 0.95]
31      G(X) . n(X) <= 0 on the boundary    at a boundary the next move points inward, never outward
32      grad L(X) . G(X) <= 0              the ledger L, the engine's composite storage over all six states
33                                          (omnicompass.storage.composite.V), descends when nothing is forcing it
34      => lim X(t) in M*                   the state settles into the basin
35  I check all three every decision and report them. The ledger may rise while outside load forces me; I report
  that as
36  forcing, never hide it.
37  """
38  from __future__ import annotations
39
40  import math
41  from dataclasses import dataclass, field
42  from typing import Dict, List, Optional
43
44  from omnicompass.nervous_system import BAND
45
46  RIM = ["A", "E", "Z", "Δ", "H", "Θ", "B", "I", "K", "Λ", "M", "N",
47         "O", "Ξ", "0", "Π", "P", "Σ", "Γ", "T", "Υ", "Ψ", "Φ", "X"]
48  POINTS = [{"+": "rising through rest"}, {"≠": "expansion, exchange under way"}, {">": "peak extension, the
  turn"},
49            {"T": "dispersion, the ceiling holds"}, {"-": "falling through rest"}, {"≈": "compression, settling
  into the basin"},
50            {"<": "deepest compression, the turn at the floor"}, {"+": "re-alignment, ignition ahead"}]
51  QUADRANTS = {1: "Expansion", 2: "Dispersion", 3: "Compression", 4: "Re-alignment"}
52
53
54  def phi(S: float, alpha_s: float, beta_s: float, delta: float) -> float:
55      """Basin potential, equation (5)."""
56      return alpha_s * S * S / 2.0 + beta_s * S ** 3 / 4.0 - delta * S
57
58
59  def axle(alpha_s: float, beta_s: float, delta: float) -> float:
60      """S*, the rest point of the axle: the positive root of  $\delta - \alpha_s S - (3/4) \beta_s S^2 = 0$ ."""
61      q = 0.75 * beta_s
62      if q <= 0:
63          return delta / alpha_s if alpha_s > 0 else 0.0
64      return (-alpha_s + math.sqrt(alpha_s * alpha_s + 4.0 * q * delta)) / (2.0 * q)
65
66
67  def heading(e: float, rate: float) -> float:
68      """Degrees clockwise from north: north is rising (rate > 0), east is extension (e > 0)."""
69      if abs(e) < 1e-12 and abs(rate) < 1e-12:
70          return 0.0
71      return math.degrees(math.atan2(e, rate)) % 360.0
72
73
74  def quadrant(e: float, rate: float) -> int:
75      if e > 0:
76          return 1 if rate > 0 else 2
77      return 3 if rate < 0 else 4
```

```

78
79
80 def in_omega(levels: Dict[str, float]) -> Dict[str, bool]:
81     lo, hi = BAND
82     return {k: lo - 1e-9 <= v <= hi + 1e-9 for k, v in levels.items()}
83
84
85 def inward(level: float, move: float, eps: float = 0.01) -> bool:
86     """G . n <= 0 on the boundary of Omega: at (or past) the floor the move is not downward, at the ceiling not
    upward."""
87     lo, hi = BAND
88     if level >= hi - eps and move > 0:
89         return False
90     if level <= lo + eps and move < 0:
91         return False
92     return True
93
94
95 @dataclass
96 class Compass:
97     """Reads the engine each decision and keeps the previous reading to form rates and the ledger step."""
98     E_max: float = 1.0
99     alpha_s: float = 0.12
100    beta_s: float = 0.10
101    delta: float = 0.5
102    prev_e: Optional[float] = None
103    prev_L: Optional[float] = None
104    scale: float = 0.0          # running RMS of the rate, so the wheel is round in its own units
105    turns: int = 0              # full circles closed (IV -> I crossings)
106    last_q: Optional[int] = None
107    trail: List[str] = field(default_factory=list)
108
109    def read(self, E: float, S: float, levels: Dict[str, float] = None, moves: Dict[str, float] = None,
110            forced: bool = False, x=None, p=None) -> Dict:
111        """x, p: the engine's full state and parameters; with them the ledger is the composite storage of all six
112        states. Without them it is the storage of the two the reading names, E^2/2 + Phi(S) - Phi(S*)."""
113        e = E / max(self.E_max, 1e-12)
114        raw = 0.0 if self.prev_e is None else e - self.prev_e
115        self.scale = raw * raw if self.scale == 0.0 else 0.9 * self.scale + 0.1 * raw * raw
116        rate = raw / math.sqrt(self.scale) if self.scale > 1e-18 else 0.0
117        h = heading(e, rate)
118        q = quadrant(e, rate)
119        if self.last_q == 4 and q == 1:
120            self.turns += 1
121        pt, meaning = POINTS[int(((h + 22.5) % 360) // 45)]
122        letter = RIM[int(((h + 7.5) % 360) // 15)]
123        s_star = axle(self.alpha_s, self.beta_s, self.delta)
124        parts = {}
125        if x is not None and p is not None:
126            from omnicompass.storage import composite_V
127            cv = composite_V(x, p, 1)
128            L = cv["V"]; parts = {k: round(v, 5) for k, v in cv.items() if k.startswith("V_")}
129            s_star = cv["S_star"]
130        else:
131            L = 0.5 * e * e + phi(S, self.alpha_s, self.beta_s, self.delta) - phi(s_star, self.alpha_s, self.
132            beta_s, self.delta)
133            dL = None if self.prev_L is None else L - self.prev_L
134            levels = levels or {}; moves = moves or {}
135            omega = in_omega(levels)
136            lo, hi = BAND
137            outside = {k: ("below the floor" if levels[k] < lo else "above the ceiling") for k, ok in omega.items()}
138            if not ok:
139                heading_in = {k: (moves.get(k, 0.0) > 0 if levels[k] < lo else moves.get(k, 0.0) < 0) for k in outside}
140                inw = {k: inward(levels[k], moves.get(k, 0.0)) for k in levels}
141                r = {"heading_deg": round(h, 1), "letter": letter, "point": pt, "meaning": meaning,
142                    "quadrant": q, "stroke": QUADRANTS[q], "e": round(e, 4), "rate": round(rate, 3),
143                    "axle_S": round(S, 4), "axle_rest": round(s_star, 4), "ledger": round(L, 5),
144                    "ledger_step": None if dL is None else round(dL, 5), "ledger_parts": parts,
145                    "descent": None if dL is None else (dL <= 1e-9 or forced), "forced": forced,
146                    "omega_held": all(omega.values()) if omega else True, "inward": all(inw.values()) if inw else True,
147                    "outside": outside, "returning": heading_in,
148                    "circles_closed": self.turns}
149                self.prev_e, self.prev_L, self.last_q = e, L, q
150            return r
151
152    def omega_words(r: Dict) -> str:
153        if r["omega_held"]:
154            return "Ω held"
155        parts = [f"{k} {w}, {'returning' if r['returning'].get(k) else 'held'}" for k, w in r["outside"].items()]
156        return "Ω: " + "; ".join(parts)
157
158    def say(r: Dict) -> str:
159        """One line in the compass's own terms."""
160        step = "" if r["ledger_step"] is None else f" {r['ledger_step']:.4f}" if r["forced"] and r[

```

```

'ledger_step'] > 0 else ''"
161     return (f"{r['letter']} {r['point']} {r['heading_deg']:5.1f}° {r['stroke']}: {r['meaning']} | "
162             f"{omega_words(r)} | G.n≤0 {'yes' if r['inward'] else 'NO'} | "
163             f"L {r['ledger']:.4f}{step} | axle S {r['axle_S']} (rest {r['axle_rest']}) | circles {r[
'circles_closed']}")
164

```

omnicompass/configurations.py

66 lines SHA-256 9a2c92daf02ee5fd09c8bf848211741df3af9eb9de878f84c163c1404882a442

```
1  """Named engine configurations.
2
3  Default remains omnicompass.core (symmetric_verified). This module does not
4  replace it. printed_eight_line is the marketing-plate transcription.
5
6  symmetric_verified claim-2 / package core
7      dE = -alpha E E + beta_int + beta_ext + v_eff
8      dU = mu U (1-U^2) - dE/E_max - lambda_U U + u
9      dI = 1-U - sigma_1 E - delta S - lambda_I I U
10     v = cos(omega_B t / 2) * c * tanh(lambda_0 + lambda_1 (U-0.5) + lambda_2 S)
11
12  printed_eight_line plate species (not the held-out evidence core)
13     dE = -alpha E + beta_int + beta_ext + v_eff
14     dU = alpha (1-U) - dE/E_max - k U (1-U) + u
15     dI = 1-U - sigma_1 E - delta S
16     v = c * tanh(lambda_0 + lambda_1 (U-U_t) + lambda_2 S)
17 """
18 from __future__ import annotations
19
20 import math
21 from typing import Callable
22
23 from .core import EPS, KP, U_AUTHORITY, Params, State, derivatives as core_derivatives, v_eff as core_v_eff
24
25
26 SYMMETRIC = "symmetric_verified"
27 PRINTED = "printed_eight_line"
28 DEFAULT = SYMMETRIC
29
30
31 def printed_v_eff(x: State, p: Params, t: float = 0.0, U_t: float = 0.5) -> float:
32     z = p.lambda_0 + p.lambda_1 * (x.U - U_t) + p.lambda_2 * x.S
33     return p.c * math.tanh(z)
34
35
36 def printed_derivatives(x: State, p: Params, t: float, u: float = 0.0, U_t: float = 0.5) -> State:
37     v = printed_v_eff(x, p, t, U_t)
38     dE = -p.alpha * x.E + p.beta_int + p.beta_ext + v
39     dU = p.alpha * (1.0 - x.U) - dE / max(p.E_max, 1e-6) - p.k * x.U * (1.0 - x.U) + u
40     dI = (1.0 - x.U) - p.sigma_1 * x.E - p.delta * x.S
41     dS = p.delta - p.alpha_s * x.S - 0.75 * p.beta_s * x.S * x.S
42     dBd = p.gamma_c * p.delta * x.S - (p.omega_B / max(p.Q_B, EPS)) * x.B_dot - p.omega_B ** 2 * x.B
43     return State(dE, dU, dI, dS, x.B_dot, dBd)
44
45
46 def derivatives_for(name: str) -> Callable:
47     if name == SYMMETRIC:
48         return core_derivatives
49     if name == PRINTED:
50         return printed_derivatives
51     raise ValueError(name)
52
53
54 def control_command_for(name: str, x: State, p: Params, t: float, target: float):
55     deriv = derivatives_for(name)
56     drift = deriv(x, p, t, 0.0).U
57     raw = -drift + KP * (float(target) - x.U)
58     return max(-U_AUTHORITY, min(U_AUTHORITY, raw)), raw
59
60
61 def describe(name: str) -> str:
62     return {
63         SYMMETRIC: "package / claim-2 cubic U; owns held-out evidence",
64         PRINTED: "plate logistic U; executable; not the default",
65     }[name]
66
```

omnicompass/conveyance.py

102 lines SHA-256 96616a3ec70414b345e76fda4ba44f1795848b31811321b6dfe9274b41b27912

```
1  """Conveyance law: moving a conserved budget ("juice") between organs.
2
3  Manuscript sources (interpretation of the English, not only the symbols):
4  Ch. 29 §8  $E_{\text{system}} = E + U + S + B$ ,  $dE_{\text{system}}/dt = 0$  over the cycle: "Energy redistributes. Energy
   reorganizes.
5      Energy does not diverge." -> the budget is conserved, only moved
6  Ch. 31 §4 "When E grows locally, B increases to distribute gradients." -> flow runs down the gradient of need
7  Ch. 31 §5 stability requires  $\text{Re}(\lambda_i) \leq 0$  -> the flow must provably converge
8  Ch. 30 §5-7 ignition is a "redistribution surge ... controlled release", with state continuity (no
   discontinuity)
9
   -> the reserve is released
10 continuously
11 Ch. 29 §6 compression: "backpressure develops through S ... structured inward rotation"
12 -> surplus is pulled back smoothly
13 The law. Organs  $i$  hold allocations  $a_i$  of a budget  $A$  (watts, cores, ...), have need  $d_i$  (the budget that would
   serve
14 their work at the engine's utilisation target  $\rho$ ), and a deficit ("local deviation")
15  $e_i = d_i / (\rho a_i) - 1$  ( $> 0$  needs juice,  $< 0$  holds surplus)
16 The redistribution mediator moves budget by the replicator flow
17  $da_i/dt = \kappa a_i (e_i - \bar{e})$ ,  $\bar{e} = \sum_j a_j e_j / A$ 
18 Properties (proved in docs/CONVEYANCE_LAW.md, checked in tests/test_conveyance.py):
19 P1 conservation  $\sum_i da_i/dt = 0$  exactly: the budget is moved, never created
20 P2 gradient flow it is the Shahshahani gradient of  $F(a) = \sum_i (d_i / \rho) \ln a_i$  on the simplex  $\sum a_i = A$ ;
21
    $F$  is strictly concave, so  $F$  rises monotonically (Lyapunov function)
22 P3 convergence unique equilibrium  $a_i^* = A d_i / D$  ( $D = \sum d_i$ ): every organ at the same deficit  $e^*$ ,
23 reached exponentially:  $da_i/dt = (\kappa/\rho)(d_i - a_i D/A)$ , rate  $\kappa D / (\rho A)$ , and
   the
24 linearisation has eigenvalues  $-\kappa D/(\rho A) < 0$  ( $\text{Re}(\lambda) \leq 0$ , Ch. 31 §5)
25 P4 bounds organ floors and ceilings (device minimum, TDP) are kept by projection; budget an organ
   cannot
26 take because it is at its ceiling returns to the others (water-filling), and budget left
   over
27 when every organ is at its need is held back as reserve, not spent
28 Reserve (Ch. 30 surge):  $R$  = the budget needed to cover the demand rise the rate tracker projects over one
   actuation
29 delay; released continuously to organs whose deficit stays positive after the flow."""
30 from __future__ import annotations
31
32 from typing import List, Sequence
33
34
35 def deficits(a: Sequence[float], d: Sequence[float], rho: float) -> List[float]:
36     return [di / (rho * max(ai, 1e-12)) - 1.0 for ai, di in zip(a, d)]
37
38
39 def replicator_step(a: Sequence[float], d: Sequence[float], rho: float, kappa: float, dt: float = 1.0) -> List[
   float]:
40     """One exact step of  $da_i/dt = (\kappa/\rho)(d_i - a_i D/A)$ , the closed form of the replicator flow (P3):
41      $a_i(t+dt) = a_i^* + (a_i - a_i^*) \exp(-\kappa D dt / (\rho A))$ ; conserves  $\sum a_i$  to rounding (P1)."""
42     import math
43     A = sum(a); D = sum(d)
44     if A <= 0 or D <= 0:
45         return list(a)
46     k = math.exp(-kappa * D * dt / (rho * A))
47     return [A * di / D + (ai - A * di / D) * k for ai, di in zip(a, d)]
48
49
50 def project(a: Sequence[float], lo: Sequence[float], hi: Sequence[float], budget: float) -> List[float]:
51     """Water-filling projection onto  $\{lo_i \leq a_i \leq hi_i, \sum a_i \leq \text{budget}\}$ : organs pinned at a bound keep it,
   the
52     rest share the remainder in proportion to their unconstrained allocation (P4). Exact in  $\leq n$  rounds."""
53     a = [min(max(x, l), h) for x, l, h in zip(a, lo, hi)]
54     for _ in range(len(a) + 1):
55         tot = sum(a)
56         if tot <= budget + 1e-9:
57             return a
58         free = [i for i in range(len(a)) if a[i] > lo[i] + 1e-12]
59         over = tot - budget
60         share = sum(a[i] - lo[i] for i in free)
61         if share <= 0:
62             return a
63         a = [max(lo[i], a[i] - over * (a[i] - lo[i]) / share) if i in free else a[i] for i in range(len(a))]
64     return a
65
66
67 class Conveyance:
68     """The budget law for a set of organs sharing one budget (e.g. GPU groups under one site power cap).
69     step(need, lo, hi) returns allocations; the flow moves budget toward need, the projection keeps bounds, and
   the
70     reserve holds back what the projected rise needs, releasing it continuously where deficits stay positive."""
71
```

```

72     def __init__(self, n: int, budget: float, rho: float = 0.8, kappa: float = 1.0):
73         self.budget, self.rho, self.kappa = budget, rho, kappa
74         self.a = [budget / n] * n
75         self.prev = None
76         self.rate = [0.0] * n
77
78     def step(self, need: Sequence[float], lo: Sequence[float], hi: Sequence[float]) -> List[float]:
79         # the living band: every organ keeps at least 5% of its range (idle, never off) and never takes more
than 95%
80         from omnicompass.nervous_system import BAND
81         lo = [max(l, BAND[0] * h) for l, h in zip(lo, hi)]
82         hi = [max(l, BAND[1] * h) for l, h in zip(lo, hi)]
83         need = [max(x, 1e-9) for x in need]
84         if self.prev is not None:
85             # rate tracker for the reserve (the dual-bath rate channel)
86             self.rate = [0.5 * r + 0.5 * (x - p) for r, x, p in zip(self.rate, need, self.prev)]
87             self.prev = list(need)
88             reserve = min(0.5 * self.budget, sum(max(0.0, r) for r in self.rate))
89             want = [min(h, x / self.rho) for x, h in zip(need, hi)] # budget each organ can use at rho
90             want = [max(w, l) for w, l in zip(want, lo)] # a device cannot go below its floor
91             spend = max(sum(lo), min(self.budget - reserve, sum(want))) # never spend budget nobody needs
92             a = replicator_step(self.a, [w * self.rho for w in want], self.rho, self.kappa)
93             a = [x * spend / max(sum(a), 1e-12) for x in a]
94             a = project(a, lo, [max(l, min(h, w)) if sum(want) <= spend else h for l, h, w in zip(lo, hi, want)],
spend)
95             # surge (Ch. 30): release the reserve continuously to organs still in deficit, up to their ceiling
96             left = self.budget - sum(a)
97             short = [max(0.0, min(h, x / self.rho) - ai) for x, h, ai in zip(need, hi, a)]
98             tot = sum(short)
99             if tot > 0 and left > 0:
100                 a = [ai + min(s, left * s / tot) for ai, s in zip(a, short)]
101         self.a = a
102         return a

```

omnicompass/core.py

241 lines SHA-256 bd615f156169f6794b41439c4ce0726deb6a0b620fac76e072d24b357ab8b0a4

```
1  """Omni-Compass core: six-state ODE, RK4 integrator, bounded feedforward-plus-proportional controller.
2
3  State x = (E, U, I_U, S, B, B_dot).
4
5  (1) dE/dt = -alpha_E E + beta_int + beta_ext + v_eff
6  (2) dU/dt = mu U (1 - U^2) - (dE/dt)/E_max - lambda_U U + u, |u| <= U_AUTHORITY
7  (3) dI_U/dt = (1 - U) - sigma_1 E - delta S - lambda_I I_U
8  (4) v_eff = cos(omega_B t / 2) * c * tanh(lambda_0 + lambda_1 (U - 0.5) + lambda_2 S)
9  (5) Phi(S) = alpha_s S^2 / 2 + beta_s S^3 / 4 - delta S
10 (6) dS/dt = -dPhi/dS = delta - alpha_s S - (3/4) beta_s S^2
11 (7) dB/dt = B_dot; dB_dot/dt = gamma_c delta S - (omega_B/Q_B) B_dot - omega_B^2 B
12 (8) R_B[n] = finite-difference audit of (7); never fed back into the state.
13
14 Controller (one micro step, zero-order hold across all four RK4 stages):
15 u = clip(-f_U(x, t) + KP (sigma - U), -U_AUTHORITY, +U_AUTHORITY)
16 where f_U is the uncontrolled U drift and sigma in {-1, +1} is the locked target.
17 """
18 from __future__ import annotations
19
20 import math
21 from dataclasses import dataclass, asdict, fields
22 from typing import Dict, List, Optional, Tuple
23
24 EPS = 1e-12
25 KP = 12.0
26 U_AUTHORITY = 25.0
27 MACRO_STEPS = 20
28 MICRO_STEPS = 10
29 MACRO_DT = 0.1
30 MICRO_DT = MACRO_DT / MICRO_STEPS
31 BASIN_TOL = 0.10
32 CONVEY_WINDOW = 5
33 CERT_WINDOW = 10
34
35 PARAMETER_RANGES: Dict[str, Tuple[float, float]] = {
36     "alpha": (4.2, 4.2), "beta_int": (0.0, 1.0), "beta_ext": (0.0, 1.0), "k": (1.7, 1.7),
37     "sigma_1": (0.38, 0.38), "delta": (0.01, 1.0), "gamma_1": (1.35, 1.35), "gamma_c": (1.0, 1.0),
38     "lambda_0": (-2.0, 2.0), "lambda_1": (0.1, 3.0), "lambda_2": (0.0, 2.0), "c": (0.5, 5.0),
39     "E_max": (1.0, 10.0), "omega_B": (0.1, 5.0), "Q_B": (0.5, 10.0), "alpha_s": (0.05, 0.25),
40     "beta_s": (0.05, 0.25), "mu": (2.0, 6.0), "alpha_E": (4.2, 4.2), "alpha_U": (4.2, 4.2),
41     "lambda_I": (0.5, 2.0), "lambda_U": (0.0, 0.5),
42 }
43 # alpha, k, gamma_1, alpha_U do not enter equations (1)-(8).
44
45
46 @dataclass
47 class State:
48     E: float
49     U: float
50     I_U: float
51     S: float
52     B: float
53     B_dot: float
54
55     def as_tuple(self) -> Tuple[float, ...]:
56         return (self.E, self.U, self.I_U, self.S, self.B, self.B_dot)
57
58
59 @dataclass
60 class Params:
61     alpha: float = 4.2
62     beta_int: float = 0.0
63     beta_ext: float = 0.0
64     k: float = 1.7
65     sigma_1: float = 0.38
66     delta: float = 0.5
67     gamma_1: float = 1.35
68     gamma_c: float = 1.0
69     lambda_0: float = 0.0
70     lambda_1: float = 1.0
71     lambda_2: float = 0.5
72     c: float = 1.0
73     E_max: float = 2.0
74     omega_B: float = 1.0
75     Q_B: float = 2.5
76     alpha_s: float = 0.12
77     beta_s: float = 0.10
78     mu: float = 4.0
79     alpha_E: float = 4.2
80     alpha_U: float = 4.2
81     lambda_I: float = 1.0
82     lambda_U: float = 0.1
83
```

```

84
85 # ----- equations
86 def v_eff(x: State, p: Params, t: float = 0.0) -> float:
87     z = p.lambda_0 + p.lambda_1 * (x.U - 0.5) + p.lambda_2 * x.S
88     return math.cos(0.5 * p.omega_B * t) * (p.c * math.tanh(z))
89
90
91 def phi(S: float, p: Params) -> float:
92     return 0.5 * p.alpha_s * S * S + 0.25 * p.beta_s * S ** 3 - p.delta * S
93
94
95 def s_flow(S: float, p: Params) -> float:
96     return p.delta - p.alpha_s * S - 0.75 * p.beta_s * S * S
97
98
99 def derivatives(x: State, p: Params, t: float, u: float = 0.0) -> State:
100     v = v_eff(x, p, t)
101     dE = -p.alpha_E * x.E + p.beta_int + p.beta_ext + v
102     dU = p.mu * x.U * (1.0 - x.U * x.U) - dE / max(p.E_max, 1e-6) - p.lambda_U * x.U + u
103     dI = (1.0 - x.U) - p.sigma_1 * x.E - p.delta * x.S - p.lambda_I * x.I_U
104     dS = s_flow(x.S, p)
105     dBd = p.gamma_c * p.delta * x.S - (p.omega_B / max(p.Q_B, EPS)) * x.B_dot - p.omega_B ** 2 * x.B
106     return State(dE, dU, dI, dS, x.B_dot, dBd)
107
108
109 def _axpy(x: State, k: State, h: float) -> State:
110     return State(x.E + h * k.E, x.U + h * k.U, x.I_U + h * k.I_U, x.S + h * k.S, x.B + h * k.B, x.B_dot + h * k.B_dot)
111
112
113 def rk4_step(x: State, p: Params, t: float, h: float, u: float = 0.0) -> State:
114     k1 = derivatives(x, p, t, u)
115     k2 = derivatives(_axpy(x, k1, 0.5 * h), p, t + 0.5 * h, u)
116     k3 = derivatives(_axpy(x, k2, 0.5 * h), p, t + 0.5 * h, u)
117     k4 = derivatives(_axpy(x, k3, h), p, t + h, u)
118     c = h / 6.0
119     return State(*(a + c * (b1 + 2 * b2 + 2 * b3 + b4) for a, b1, b2, b3, b4 in
120                     zip(x.as_tuple(), k1.as_tuple(), k2.as_tuple(), k3.as_tuple(), k4.as_tuple())))
121
122
123 def bath_residual(prev: State, cur: State, nxt: State, p: Params, h: float) -> float:
124     d2 = (nxt.B - 2.0 * cur.B + prev.B) / (h * h)
125     dm = (cur.B - prev.B) / h
126     return d2 + (p.omega_B / max(p.Q_B, EPS)) * dm + p.omega_B ** 2 * cur.B - p.gamma_c * p.delta * cur.S
127
128
129 # ----- controller
130 def control_command(x: State, p: Params, t: float, target: int) -> Tuple[float, float]:
131     """Return (u, raw) for one micro step. raw is the unclipped command."""
132     drift = derivatives(x, p, t, 0.0).U
133     raw = -drift + KP * (float(target) - x.U)
134     return max(-U_AUTHORITY, min(U_AUTHORITY, raw)), raw
135
136
137 def basin_sign(U: float) -> int:
138     if abs(U - 1.0) <= BASIN_TOL:
139         return 1
140     if abs(U + 1.0) <= BASIN_TOL:
141         return -1
142     return 0
143
144
145 def nearest_basin_sign(U: float) -> int:
146     return 1 if abs(U - 1.0) <= abs(U + 1.0) else -1
147
148
149 def macro_step(x: State, p: Params, t0: float, target: Optional[int],
150               micro_steps: int = MICRO_STEPS, macro_dt: float = MACRO_DT) -> Tuple[State, Dict[str, float]]:
151     """Advance one macro interval. target=None runs the plant with u=0 (observe)."""
152     h = macro_dt / micro_steps
153     t = t0
154     u_abs = u_sq = u_peak = 0.0
155     sat = 0
156     for _ in range(micro_steps):
157         if target is None:
158             u = raw = 0.0
159         else:
160             u, raw = control_command(x, p, t, target)
161             u_abs += abs(u) * h
162             u_sq += u * u * h
163             u_peak = max(u_peak, abs(u))
164             sat += int(abs(raw) > U_AUTHORITY + EPS)
165             x = rk4_step(x, p, t, h, u)
166             t += h
167     return x, {"u_abs": u_abs, "u_sq": u_sq, "u_peak": u_peak, "saturated": sat}
168

```

```

169
170 # ----- trajectory
171 @dataclass
172 class Result:
173     final: State
174     U_hist: List[float]
175     target_sign: int
176     born: int
177     convey: int
178     convey_start: int
179     convey_confirm: int
180     cert: int
181     cert_start: int
182     cert_complete: int
183     max_streak: int
184     actuator_abs_integral: float
185     actuator_sq_integral: float
186     actuator_peak: float
187     actuator_saturated_periods: int
188
189
190 def simulate(x0: State, p: Params, target: Optional[int] = None, controlled: bool = True) -> Result:
191     """Canonical 20-macro-step trajectory with CONVEY-5 / CERT-10 predicates.
192
193     target defaults to the nearest basin of U0. controlled=False runs u=0.
194     """
195     tgt = nearest_basin_sign(x0.U) if target is None else int(target)
196     b0 = basin_sign(x0.U)
197     born = int(b0 != 0)
198     streak = 1 if born else 0
199     ssign = b0 if born else 0
200     best = streak
201     convey = cert = 0
202     cs = cc = ts = tc = -1
203     ua = usq = upk = 0.0
204     nsat = 0
205     x = x0
206     hist = [x0.U]
207     for m in range(1, MACRO_STEPS + 1):
208         x, a = macro_step(x, p, (m - 1) * MACRO_DT, tgt if controlled else None)
209         ua += a["u_abs"]; usq += a["u_sq"]; upk = max(upk, a["u_peak"]); nsat += a["saturated"]
210         hist.append(x.U)
211         b = basin_sign(x.U)
212         if b:
213             if b == ssign:
214                 streak += 1
215             else:
216                 ssign, streak = b, 1
217         else:
218             ssign, streak = 0, 0
219         best = max(best, streak)
220         if not convey and streak >= CONVEY_WINDOW:
221             convey, cs, cc = 1, m - CONVEY_WINDOW + 1, m
222         if not cert and streak >= CERT_WINDOW:
223             cert, ts, tc = 1, m - CERT_WINDOW + 1, m
224     return Result(x, hist, tgt, born, convey, cs, cc, cert, ts, tc, best, ua, usq, upk, nsat)
225
226
227 def sample_state(rng) -> State:
228     return State(float(rng.uniform(0.0, 1.0)), float(rng.uniform(-1.0, 1.0)), float(rng.uniform(0.0, 1.0)),
229                 float(rng.uniform(-1.0, 1.0)), float(rng.uniform(-1.0, 1.0)), float(rng.uniform(-1.0, 1.0)))
230
231
232 def sample_params(rng) -> Params:
233     return Params(*{k: float(rng.uniform(lo, hi)) for k, (lo, hi) in PARAMETER_RANGES.items()})
234
235
236 def s_roots(p: Params) -> Tuple[float, float]:
237     """Roots of equation (6): S_minus < 0 < S_plus."""
238     a = 0.75 * p.beta_s
239     d = math.sqrt(p.alpha_s ** 2 + 4.0 * a * p.delta)
240     return (-p.alpha_s - d) / (2.0 * a), (-p.alpha_s + d) / (2.0 * a)
241

```

omnicompass/mathdrive.py

82 lines SHA-256 ec880382a58e1e1f47246703351c52448e28f228943eae3d95b40db2f00cd6d

```
1  """Math-driven actuation: the Omni-Compass equations move the system directly.
2
3  The frozen governor (adapter.Governor) evolves the equations open-loop (u = 0) and reads only a few derived
   numbers.
4  Here the equations are closed-loop and their own quantities ARE the actuator commands. No equation is changed; the
5  integrator, parameters and observation map are the frozen ones (omnicompass/core.py, omnicompass/adapter.py).
6
7  Per decision
8  1. sense      x <- assimilate(x, observe_vector(telemetry))          (the frozen observation map)
9  2. evolve      x <- macro_step(x, target = +1, macro_dt = clock)      equations (1)-(7) with
   equation
10     (2)'s controller ON: u = clip(-f U(x) + 12 (1 - U), -25, 25) holds the healthy basin U = +1
11 3. read the    effort w = integral |u| dt / (25 clock) in [0, 1]: how hard the math had to work to stay
   healthy
12   math        need      I_U          equation (3): unrelieved need, slow memory
13               tide      B, B_dot      equation (7): the bath, second order, anticipates
14               drive      v_eff        equation (4): the tanh-bounded drive
15 4. actuate     capacity  n = ceil(fit (1 + g_w w + g_I max(0, I_U) + g_B max(0, B_dot))) machines
16               (released one at a time only while w <= w_release: the math reports calm)
17               pods      HPA target rho = clamp(rho0 - k_w w - k_I max(0, I_U), rho_min, rho0)
18               power      cap = clamp(1 - g_cap max(0, B) + 2 w, cap_min, 1) (the bath lowers power only when
   calm)
19   clock: engine time advanced per decision. It sets how fast the math moves relative to the system: the bath's
20   natural period is 2 pi / omega_B = 6.28 engine units, i.e. 6.28 / clock decisions.
21 """
22 from __future__ import annotations
23
24 import math
25 from dataclasses import dataclass, replace
26
27 from omnicompass.adapter import STACK_PARAMS, INITIAL_STATE, observe_vector, assimilate
28 from omnicompass.core import State, macro_step, v_eff, U_AUTHORITY
29
30
31 @dataclass(frozen=True)
32 class MathLaw:
33     clock: float = 0.1
34     g_w: float = 1.0
35     g_I: float = 0.3
36     g_B: float = 0.5
37     w_release: float = 0.02
38     dwell: int = 3
39     rho0: float = 0.70
40     rho_min: float = 0.50
41     k_w: float = 0.5
42     k_I: float = 0.1
43     g_cap: float = 0.0
44     cap_min: float = 0.8
45
46
47 class MathDrive:
48     def __init__(self, law: MathLaw = MathLaw()):
49         self.law = law
50         self.p = replace(STACK_PARAMS)
51         self.x = State(**vars(INITIAL_STATE))
52         self.t = 0.0
53         self.calm = 0
54         self.w = 0.0
55
56     def step(self, obs, n, requests, cores_per_node):
57         """Returns (hpa_target, node_target, power_cap, readout)."""
58         L = self.law
59         o = observe_vector(obs, 0)
60         xa = assimilate(self.x, o, self.p)
61         self.x, stats = macro_step(xa, self.p, self.t, target=+1, macro_dt=L.clock)
62         self.t += L.clock
63         w = min(1.0, stats["u_abs"] / (U_AUTHORITY * L.clock))
64         self.w = w
65         x = self.x
66         need = max(0.0, x.I_U); tide = max(0.0, x.B_dot)
67         fit = requests / cores_per_node
68         want = math.ceil(fit * (1.0 + L.g_w * w + L.g_I * need + L.g_B * tide) - 1e-9)
69         tgt = n
70         if want > n:
71             tgt = want; self.calm = 0
72         elif want < n:
73             self.calm = self.calm + 1 if w <= L.w_release else 0
74             if self.calm >= L.dwell:
75                 tgt = n - 1; self.calm = 0
76         else:
77             self.calm = 0
78         rho = min(L.rho0, max(L.rho_min, L.rho0 - L.k_w * w - L.k_I * need))
79         cap = min(1.0, max(L.cap_min, 1.0 - L.g_cap * max(0.0, x.B) + 2.0 * w))
```

```
80     return rho, int(tgt), cap, {"w": w, "I_U": x.I_U, "B": x.B, "B_dot": x.B_dot, "U": x.U, "E": x.E,  
81                                   "v_eff": v_eff(x, self.p, self.t)}  
82
```

omnicompass/muscles.py

155 lines SHA-256 8aa5e61437f59a4d3ffbdd2c810b43b192b12929a9001085e6dc51a192c40f40

```
1  """Infrastructure component catalog and fleet overhead model.
2
3  Each component records its role, deployment pattern, resource footprint, the source of that
4  footprint and its evidence status, and its classification under the keep-or-remove rule:
5
6  decision      control logic whose function the governor performs      -> candidate for removal
7  observability telemetry and storage for humans, audit and control      -> retained; control path replaced by
8  direct sensing
9  execution      machinery that performs the work                      -> retained
10 security       cryptography, identity, policy, runtime detection        -> retained unless a stronger guarantee
11 is proven
12
13 Evidence status:
14 measured       benchmark or production measurement published by the source
15 documented      formula or default published in official documentation
16 example        request values from a published example manifest
17 assumption      modelling assumption; parameter exposed for sensitivity analysis
18
19 Energy per resource follows the Cloud Carbon Footprint method:
20 W = vCPU * (W_min + u * (W_max - W_min)) + GB * W_mem, scaled by PUE.
21 """
22 from __future__ import annotations
23
24 from dataclasses import dataclass, asdict
25 from typing import Dict, List
26
27 W_MIN_PER_VCPU = 0.434
28 W_MAX_PER_VCPU = 1.948
29 W_PER_GB = 0.392
30 ENERGY_SOURCE = ("Cloud Carbon Footprint coefficients (SPECpower-derived), AMD EPYC 3rd generation: "
31                  "0.434 W min, 1.948 W max per vCPU; memory 0.392 W per GB (arXiv 2510.26413)")
32
33 @dataclass(frozen=True)
34 class Component:
35     name: str
36     role: str
37     scope: str
38     vcpu: float
39     gib: float
40     status: str
41     source: str
42     verdict: str
43     note: str = ""
44
45 def kube_reserved_cpu(cores: int) -> float:
46     """AKS kube-reserved CPU table (Microsoft Learn), cores -> vCPU, linear between table points."""
47     pts = [(1, 0.060), (2, 0.100), (4, 0.140), (8, 0.180), (16, 0.260), (32, 0.420), (64, 0.740)]
48     if cores <= 1:
49         return pts[0][1]
50     for (c0, v0), (c1, v1) in zip(pts, pts[1:]):
51         if cores <= c1:
52             return v0 + (v1 - v0) * (cores - c0) / (c1 - c0)
53     c0, v0 = pts[-2]; c1, v1 = pts[-1]
54     return v1 + (v1 - v0) * (cores - c1) / (c1 - c0)
55
56
57 @dataclass
58 class Fleet:
59     nodes: int = 100_000
60     nodes_per_cluster: int = 5_000
61     cores_per_node: int = 64
62     gib_per_node: int = 256
63     pods_per_node: int = 30
64     series_per_node: int = 2_000
65     long_term_replication: int = 3
66     mean_cpu_utilization: float = 0.30
67     pue: float = 1.18
68     mesh: str = "sidecar"
69
70     @property
71     def clusters(self) -> int:
72         return -(self.nodes // self.nodes_per_cluster)
73
74
75 def catalog(f: Fleet) -> List[Component]:
76     series = f.nodes * f.series_per_node
77     C = []
78     C.append(Component("kubelet + container runtime (kube-reserved)", "execution", "per node",
79                        kube_reserved_cpu(f.cores_per_node) * f.nodes, 0.0, "documented",
80                        "Microsoft Learn, AKS node resource reservations (CPU table)", "retain",
81                        "memory reservation not counted here"))
```

```

82     C.append(Component("node-exporter (metrics agent)", "observability", "per node", 0.050 * f.nodes,
83                        100 / 1024 * f.nodes, "example", "OneUptime DaemonSet resource example (requests 50m /
100Mi)",
84                        "retain for humans; control path replaced by direct sensing"))
85     C.append(Component("log collector (Fluentd / Fluent Bit)", "observability", "per node", 0.100 * f.nodes,
86                        200 / 1024 * f.nodes, "example", "OneUptime DaemonSet resource example (requests 100m /
200Mi)",
87                        "retain for audit and humans"))
88     C.append(Component("OpenTelemetry agent collector", "observability", "per node", 0.100 * f.nodes,
89                        128 / 1024 * f.nodes, "example", "OneUptime agent-gateway topology (requests 100m /
128Mi)",
90                        "retain for humans; control path replaced by direct sensing"))
91     if f.mesh == "sidecar":
92         C.append(Component("service-mesh sidecar proxies (Envoy, per pod)", "security", "per pod",
93                            0.20 * f.pods_per_node * f.nodes, 60 / 1024 * f.pods_per_node * f.nodes, "measured",
94                            "Istio 1.24 benchmark at 1,000 req/s, 1 KB: ~0.20 vCPU, ~60 MB per sidecar",
95                            "retain (mTLS); independent option: ambient mode",
96                            "load-dependent; an idle sidecar measured ~8 millicores / 40 Mi (Defense Unicorns)"))
97     else:
98         C.append(Component("service-mesh ztunnel (ambient, per node)", "security", "per node",
99                            0.06 * f.nodes, 12 / 1024 * f.nodes, "measured",
100                           "Istio 1.24 benchmark: ztunnel ~0.06 vCPU, ~12 MB", "retain (mTLS)"))
101     C.append(Component("Prometheus (in-cluster head series)", "observability", "per cluster",
102                        0.0, 3.0e-6 * series, "measured", "Robust Perception: ~3 kB per head series (Prometheus 2.
20+)"),
103                        "retain for humans; control scrape load removable", "series_per_node is an assumption"))
104     C.append(Component("long-term metrics (Mimir ingesters)", "observability", "fleet",
105                        series * f.long_term_replication / 300_000, 2.5 * series * f.long_term_replication /
300_000,
106                        "documented", "Grafana Mimir capacity planning: 1 core and 2.5 GB per 300,000 in-memory
series",
107                        "retain for audit"))
108     C.append(Component("cluster autoscaler / Karpenter", "decision", "per cluster", 1.0 * f.clusters,
109                        1.0 * f.clusters, "assumption", "typical controller sizing; not measured here",
110                        "remove: governor owns capacity"))
111     C.append(Component("HPA / VPA / KEDA controllers", "decision", "per cluster", 0.5 * f.clusters,
112                        1.0 * f.clusters, "assumption", "typical controller sizing; not measured here",
113                        "remove: governor owns replicas and sizing"))
114     C.append(Component("power / thermal policy agents", "decision", "per node", 0.010 * f.nodes,
115                        32 / 1024 * f.nodes, "assumption", "typical node agent sizing; not measured here",
116                        "remove: governor owns power caps"))
117     C.append(Component("alert routing and paging (Alertmanager)", "decision", "per cluster", 0.2 * f.clusters,
118                        0.5 * f.clusters, "assumption", "sizing scales with firing alerts (Grafana Mimir
guidance)",
119                        "remove paging path; keep audit notifications"))
120     C.append(Component("Terraform as live capacity controller", "decision", "fleet", 0.0, 0.0, "assumption",
121                        "run cost negligible; harm is stale overwrites (measured in the stack benchmark)",
122                        "remove from the live control path; retain as audit record"))
123     return C
124
125
126 def watts(vcpu: float, gib: float, u: float, pue: float) -> float:
127     return pue * (vcpu * (W_MIN_PER_VCPU + u * (W_MAX_PER_VCPU - W_MIN_PER_VCPU)) + gib * 1.073741824 * W_PER_GB)
128
129
130 def fleet_report(f: Fleet) -> Dict:
131     comps = catalog(f)
132     total_vcpu = f.nodes * f.cores_per_node
133     total_gib = f.nodes * f.gib_per_node
134     fleet_w = watts(total_vcpu, total_gib, f.mean_cpu_utilization, f.pue)
135     rows = []
136     for c in comps:
137         lo, hi = watts(c.vcpu, c.gib, 0.0, f.pue), watts(c.vcpu, c.gib, 1.0, f.pue)
138         rows.append({**asdict(c), "watts_low": lo, "watts_high": hi,
139                     "mwh_per_year_low": lo * 8760 / 1e6, "mwh_per_year_high": hi * 8760 / 1e6,
140                     "share_of_fleet_vcpu": c.vcpu / total_vcpu,
141                     "share_of_fleet_energy_low": lo / fleet_w, "share_of_fleet_energy_high": hi / fleet_w})
142     by_role = {}
143     for r in rows:
144         g = by_role.setdefault(r["role"], {"vcpu": 0.0, "share_of_fleet_vcpu": 0.0,
145                                             "share_of_fleet_energy_low": 0.0, "share_of_fleet_energy_high": 0.0,
146                                             "mwh_per_year_low": 0.0, "mwh_per_year_high": 0.0})
147         for k in g:
148             g[k] += r[k]
149     idle_vcpu = total_vcpu * (1.0 - f.mean_cpu_utilization)
150     idle_w = f.pue * idle_vcpu * W_MIN_PER_VCPU
151     return {"fleet": asdict(f), "clusters": f.clusters, "fleet_vcpu": total_vcpu, "fleet_mwh_per_year": fleet_w
* 8760 / 1e6,
152            "energy_model": ENERGY_SOURCE, "components": rows, "by_role": by_role,
153            "idle_vcpu": idle_vcpu, "idle_share_of_fleet_vcpu": 1.0 - f.mean_cpu_utilization,
154            "idle_vcpu_share_of_fleet_energy": idle_w / fleet_w}
155

```

omnicompass/nervous.py

228 lines SHA-256 556fe9ac841a1c8180081f6c767d40ec435ae921310a361bbd609106ef19a17c

```
1  """Nervous system: one engine, many muscles.
2
3  Each muscle is a connector with five parts (domain map):
4      afferent   sense
5      engine     shared six-state field (not duplicated)
6      efferent   push
7      reflex     shield / bounds before push
8      kill       return that muscle to its native controller
9
10 Status:
11 wired         sense and push are executed in this tree
12 sensed        telemetry is accepted; push is not issued
13 open          connector registered; no plant; no push
14
15 This module does not replace adapter.Governor. It sits around it:
16 afferents merge into the observation vector the governor already understands;
17 efferents are tagged by muscle so a kill is per-nerve, not a mystery write.
18 """
19 from __future__ import annotations
20
21 from dataclasses import dataclass, field
22 from typing import Any, Callable, Dict, List, Optional
23
24
25 # Fit from the domain map. Advisory muscles never push from this runtime.
26 DIRECT = "direct"
27 SUPERVISORY = "supervisory"
28 ADVISORY = "advisory"
29 NOT_OURS = "not_omni"
30
31 WIRED = "wired"
32 SENSED = "sensed"
33 OPEN = "open"
34
35
36 @dataclass(frozen=True)
37 class Muscle:
38     name: str
39     family: str
40     fit: str
41     status: str
42     reads: tuple
43     pushes: tuple
44     reflexes: tuple
45     note: str = ""
46
47
48 # Registered nerves. Status is a claim about THIS tree, not a product slide.
49 CATALOG: List[Muscle] = [
50     Muscle("nodes", "infra", DIRECT, WIRED,
51           ("load_ratio", "pending", "requests"),
52           ("node_delta", "actuation"),
53           ("request_floor", "step_limit", "I2", "I5"),
54           "fleet + k8s_controlplane"),
55     Muscle("hpa", "infra", DIRECT, WIRED,
56           ("utilization", "replicas"),
57           ("hpa_target",),
58           ("target_bounds_50_95",),
59           "target write only when authority is on"),
60     Muscle("power_cap", "infra", DIRECT, WIRED,
61           ("power_stress", "site_limit"),
62           ("power_cap",),
63           ("I4", "cap_lo_hi"),
64           "harness; live: CPU limits of capped deployments (omni_controller/muscles.py)"),
65     Muscle("heat", "infra", DIRECT, SENSED,
66           ("thermal",),
67           (),
68           ("thermal_limit",),
69           "feeds E/S; live: harness heat law on live power, or GPU temperature; no chiller write"),
70     Muscle("network", "infra", DIRECT, SENSED,
71           ("network_stress",),
72           ("route_shift",),
73           ("link_capacity",),
74           "directive flag only"),
75     Muscle("security", "infra", DIRECT, WIRED,
76           ("security_block",),
77           ("block_expansion",),
78           ("I1",),
79           "live: hold signal from a ConfigMap; blocks node, cap and batch expansion"),
80     Muscle("deployments", "infra", DIRECT, WIRED,
81           ("rollout_progress",),
82           ("pause", "resume", "undo"),
83           ("change_permitted", "rollback_authorized"),
```

```

84         "live: kubectrl rollout pause / resume / undo"),
85     Muscle("gpu", "infra", DIRECT, WIRED,
86           ("gpu_power", "gpu_temp"),
87           ("gpu_power_cap",),),
88           ("temp_limit", "cap_floor"),
89     "hardware connector (nvidia-smi query and power limit); tested with fake hardware, not run on a GPU;
MIG open"),
90     Muscle("cpu_pstate", "infra", DIRECT, WIRED,
91           ("rapl_watts",),),
92           ("frequency_ceiling",),
93           ("cap_floor",),),
94     "hardware connector (RAPL read, cpufreq ceiling); tested with fake hardware, not run on bare metal"),
95     Muscle("memory", "infra", DIRECT, OPEN,
96           ("mem_used", "mem_request"),
97           ("mem_request",),),
98           ("oom_margin",),),
99     Muscle("storage", "infra", SUPERVISORY, OPEN,
100           ("io_load", "io_latency", "capacity"),
101           ("tier", "volume_scale"),
102           ("durability",),),
103     Muscle("batch_queue", "infra", DIRECT, WIRED,
104           ("queue_depth", "load_ratio", "power_stress"),
105           ("admit",),),
106           ("headroom", "I1"),
107     "live: admit held Jobs (spec.suspend) with headroom; priority and deadlines open"),
108     Muscle("cooling", "facility", DIRECT, OPEN,
109           ("inlet_c", "chiller_load"),
110           ("setpoint_c", "fan_curve"),
111           ("inlet_limit",),),
112     Muscle("grid", "facility", DIRECT, OPEN,
113           ("price", "carbon", "demand_signal"),
114           ("shed_kw", "shift"),
115           ("contract_limit",),),
116     Muscle("training", "ai", DIRECT, OPEN,
117           ("gpu_power", "step", "progress"),
118           ("power_cap", "pack"),
119           ("no_kill_midstep",),),
120     Muscle("inference", "ai", DIRECT, OPEN,
121           ("qps", "latency", "gpu_mem"),
122           ("replicas", "batch"),
123           ("latency_slo",),),
124     Muscle("agent_containment", "ai", DIRECT, OPEN,
125           ("actions", "spend", "net"),
126           ("cpu_cap", "mem_cap", "net_cap", "spend_cap"),
127           ("hard_cap", "instant_kill"),
128     "containment only; not value alignment"),
129     Muscle("value_alignment", "ai", NOT_OURS, OPEN,
130           (),
131           (),
132           (),
133     "not a muscle Omni governs"),
134 ]
135
136
137 def catalog_by_name() -> Dict[str, Muscle]:
138     return {m.name: m for m in CATALOG}
139
140
141 def wired() -> List[Muscle]:
142     return [m for m in CATALOG if m.status == WIRED]
143
144
145 def open_muscles() -> List[Muscle]:
146     return [m for m in CATALOG if m.status == OPEN]
147
148
149 @dataclass
150 class NerveEvent:
151     muscle: str
152     kind: str # sense | push | hold | kill | reflex
153     payload: Dict[str, Any] = field(default_factory=dict)
154
155
156 @dataclass
157 class Bundle:
158     """One tick of the whole body."""
159     obs: Dict[str, float]
160     events: List[NerveEvent] = field(default_factory=list)
161     kills: List[str] = field(default_factory=list)
162
163     def sense(self, muscle: str, **vals: float) -> None:
164         m = catalog_by_name()[muscle]
165         for k, v in vals.items():
166             self.obs[k] = float(v)
167         self.events.append(NerveEvent(muscle, "sense", dict(vals)))
168         if m.status == OPEN:

```

```

169         self.events.append(NerveEvent(muscle, "hold", {"reason": "open"}))
170
171     def kill(self, muscle: str) -> None:
172         self.kills.append(muscle)
173         self.events.append(NerveEvent(muscle, "kill", {}))
174
175
176 def merge_stack_obs(bundle: Bundle) -> Dict[str, float]:
177     """Map extra nerves onto the governor's existing observation names."""
178     o = dict(bundle.obs)
179     aliases = {
180         "utilization": "load_ratio",
181         "gpu_util": "load_ratio",
182         "inlet_c": "thermal",
183         "gpu_temp": "thermal",
184         "rapl_watts": "power_stress",
185         "gpu_power": "power_stress",
186         "queue_depth": "queue_ratio",
187         "spend": "queue_ratio",
188     }
189     for src, dst in aliases.items():
190         if src in o and dst not in o:
191             x = float(o[src])
192             if src == "inlet_c":
193                 x = max(0.0, (x - 18.0) / 12.0)
194             elif src == "gpu_temp":
195                 x = max(0.0, (x - 40.0) / 40.0)
196             o[dst] = x
197     return o
198
199
200 def allow_push(muscle: str, authority: bool, killed: List[str]) -> bool:
201     m = catalog_by_name()[muscle]
202     if muscle in killed:
203         return False
204     if not authority:
205         return False
206     if m.fit in (ADVISORY, NOT_OURS):
207         return False
208     return m.status == WIRED
209
210
211 def tag_directive(d: Dict[str, Any], authority: bool, killed: List[str]) -> Dict[str, Any]:
212     """Attach which muscle may execute which field. Unwired fields stay logged."""
213     out = dict(d)
214     nerves = {
215         "node_delta": "nodes",
216         "actuation": "nodes",
217         "hpa_target": "hpa",
218         "power_cap": "power_cap",
219         "route_shift": "network",
220     }
221     permitted = {}
222     for field_name, muscle in nerves.items():
223         if field_name in out:
224             permitted[field_name] = allow_push(muscle, authority, killed)
225     out["nerve_permit"] = permitted
226     out["nerve_killed"] = list(killed)
227     return out
228

```

omnicompass/nervous_system.py

159 lines SHA-256 2f6492ff86440e1524df70d0192e749544e3ac29807d4a6b0e82dad70e08ea8e

```
1  """The supervisory nervous system: one evolving engine state grants authority to every organ family.
2
3  Omni-Compass does not replace the specialist controllers. Each keeps its own mechanism:
4  - the HPA and VPA size pods;
5  - the Cluster Autoscaler and Karpenter size node pools;
6  - Linux schedutil picks CPU frequency;
7  - NVIDIA DCGM manages GPUs;
8  - rollout controllers and routers do their own work.
9
10 The nervous system decides, from one shared state, the envelope inside which each of them may act.
11
12 Inputs, all from the six-state engine (omnicompass/core.py, via the governor), and nothing tuned per organ:
13 kappa convergence = 1 if push <= push_release, else push_release / push. This is the equation-(2) gate made
14 continuous: the further the engine is from its basin, the less it may give back.
15 h basin health = clamp((U - U_gate) / (1 - U_gate), 0, 1)
16 sigma stress ratio = S / S*, where S* solves  $\delta - \alpha_s S - (3/4) \beta_s S^2 = 0$  (equation 6)
17 nu unmet need = max(0, I_U) (equation 3)
18 calm = kappa * h * (1 - clamp(sigma - 1, 0, 1)) * (1 - clamp(nu, 0, 1))
19 in [0, 1]; 1 means fully settled
20
21 Authority per organ:
22 expand may the organ add capacity or performance. Always yes, except during a security hold.
23 contract may it give capacity or performance back. Only when calm >= the organ's reversibility threshold
24 (theta_pods 0.5, theta_nodes 0.7: a node costs a boot to reverse) and the SLO is clean.
25 step fraction of the organ's surplus it may shed this decision. Equals calm.
26 envelope for continuous organs, the allowed range:
27   cpufreq ceiling (fraction of cpuinfo_max) in [1 - 0.35 calm, 1];
28   GPU power limit (fraction of max) in [1 - 0.30 calm, 1];
29   site power cap in [0.65, 1];
30   traffic-shift fraction in [0, 0.5 calm];
31   cooling supply-air setpoint (C) in [18, 18 + 9 calm].
32   Heat above 0.96 forces the frequency and GPU ceilings down to 1 - 0.35 x excess.
33 batch admit held work only if calm >= 0.5 and power stress < 0.9; pause pausable work if sigma > 1, power
34 stress >= 0.95, or heat >= 0.96.
35 rollback the engine's rollback authorisation (S high, U falling, or a security hold).
36
37 Global rules, applied last:
38 - observe: every authority is computed and logged; execute is False everywhere.
39 - kill: no authority at all.
40 - security hold: no capacity organ may expand. Cooling is the one protective organ: "expand" there means more
41   cooling, which adds protection and no capacity, so it stays allowed.
42 The safety shield (omnicompass/shield.py) stays downstream and can still veto any action.
43 """
44 from __future__ import annotations
45
46 from dataclasses import dataclass
47 from typing import Any, Dict
48
49 from omnicompass.closure import stress_equilibrium
50
51 ORGANS = ("pods", "nodes", "cpufreq", "gpu", "power", "batch", "routing", "rollback")
52 # The living band (the founder's rule): every level the nervous system hands out lives between 5% and 95% of its
53 # range.
54 # Nothing is driven to zero (a part with no work idles down to its floor, alive and ready) and nothing is driven
55 # to its
56 # absolute top (the last 5% is never spent). The band is applied last, to every envelope, so no organ can leave
57 # it.
58 BAND = (0.05, 0.95)
59
60 def in_band(lo: float, hi: float) -> list:
61     """Clip an envelope [lo, hi] (fractions of the organ's range) into the living band, keeping lo <= hi."""
62     b0, b1 = BAND
63     h = min(b1, max(b0, hi))
64     return [min(h, max(b0, lo)), h]
65
66 THETA = {"pods": 0.5, "nodes": 0.7, "cpufreq": 0.3, "gpu": 0.3, "power": 0.3, "routing": 0.5}
67
68 def clamp(x, lo, hi):
69     return lo if x < lo else hi if x > hi else x
70
71
72 @dataclass(frozen=True)
73 class NervousInputs:
74     E: float; U: float; I_U: float; S: float
75     push: float
76     push_release: float = 0.2
77     U_gate: float = 0.5
78     s_eq: float = 1.0
79     security_block: float = 0.0
80     slo_clean: bool = True
```

```

81     power_stress: float = 0.0
82     thermal: float = 0.0
83     rollback: bool = False
84     stale: float = 0.0          # fraction of declared senses that are blind (dropout, frozen, unreadable)
85     mode: str = "autopilot"    # "observe" or "autopilot"
86     killed: bool = False
87
88
89 def scalars(i: NervousInputs) -> Dict[str, float]:
90     kappa = 1.0 if i.push <= i.push_release else i.push_release / max(i.push, 1e-12)
91     h = clamp((i.U - i.U_gate) / max(1.0 - i.U_gate, 1e-12), 0.0, 1.0)
92     sigma = i.S / i.s_eq if i.s_eq > 0 else 0.0
93     nu = max(0.0, i.I.U)
94     calm = kappa * h * (1.0 - clamp(sigma - 1.0, 0.0, 1.0)) * (1.0 - clamp(nu, 0.0, 1.0))
95     return {"kappa": kappa, "h": h, "sigma": sigma, "nu": nu, "calm": clamp(calm, 0.0, 1.0)}
96
97
98 def authority(i: NervousInputs) -> Dict[str, Any]:
99     if i.killed:
100         return {"execute": False, "killed": True, "organs": {}}
101     sc = scalars(i); calm = sc["calm"]
102     sec = i.security_block > 0.5
103     hot = i.thermal >= 0.96
104     excess = clamp(i.thermal - 0.96, 0.0, 1.0)
105     org: Dict[str, Dict[str, Any]] = {}
106     seeing = i.stale <= 0.0          # never give capacity back on a blind sense; expanding stays allowed (the
safe side)
107     for o, th in THETA.items():
108         ok = calm >= th and i.slo_clean and seeing
109         org[o] = {"expand": not sec, "contract": ok, "step": calm if ok else 0.0}
110         cf_lo = 1.0 - 0.35 * calm; gp_lo = 1.0 - 0.30 * calm
111         cf_hi = 1.0 - 0.35 * excess if hot else 1.0; gp_hi = 1.0 - 0.35 * excess if hot else 1.0
112         org["cpufreq"]["envelope"] = in_band(min(cf_lo, cf_hi), cf_hi)
113         org["gpu"]["envelope"] = in_band(min(gp_lo, gp_hi), gp_hi)
114         org["power"]["envelope"] = in_band(0.65, 1.0)
115         # routing is an amount moved, not a level: it may move nothing, never more than the band's top
116         org["routing"]["envelope"] = [0.0, min(BAND[1], 0.5 * calm) if not sec else 0.0]
117         org["cooling"] = {"expand": True, "protective": True, "contract": calm >= 0.3 and seeing, "step": calm,
"envelope": [18.0, 18.0 + 9.0 * calm]}
118         org["batch"] = {"expand": (not sec) and seeing and calm >= 0.5 and i.power_stress < 0.9, "contract": True,
"protective_contract": True,
"admit": (not sec) and seeing and calm >= 0.5 and i.power_stress < 0.9,
"pause": sc["sigma"] > 1.0 or i.power_stress >= 0.95 or hot, "step": calm}
119         org["rollback"] = {"expand": False, "contract": False, "authorized": bool(i.rollback or sec), "step": 0.0}
120     return {"execute": i.mode == "autopilot", "killed": False, "security_hold": sec, "scalars": sc, "organs":
org}
121
122
123
124
125 def from_governor(g, obs: Dict[str, Any], d: Dict[str, Any] = None, mode: str = None) -> Dict[str, Any]:
126     """"Authority from a live omnicompass.adapter.Governor after its step (d = the step's directive)."""
127     x = g.x; p = g.p; L = g.law
128     return authority(NervousInputs(
129         E=x.E, U=x.U, I_U=x.I.U, S=x.S, push=g.last_push, push_release=getattr(L, "push_release", 0.2),
130         U_gate=getattr(L, "U_gate", 0.5), s_eq=stress_equilibrium(p.delta, p.alpha_s, p.beta_s),
131         security_block=float(obs.get("security_block", 0.0)), slo_clean=bool(obs.get("slo_clean", True)),
132         power_stress=float(obs.get("power_stress", 0.0)), thermal=float(obs.get("thermal", 0.0)),
133         rollback=bool((d or {}).get("rollback authorized", False)), stale=float(obs.get("stale", 0.0)),
134         mode=mode or g.mode, killed=g.killed))
135
136
137 def node_release_gate(n: int, per_node_m: float, used_m: float, pending: int, pods_scaling_up: bool,
138     latency_breach_now: bool, rho: float, node_auth: Dict[str, Any],
139     senses_live: bool = True, last_command_landed: bool = True) -> Dict[str, Any]:
140     """"May the machine organ give one machine back now? Attribution: an organ is held back only by stress it can
cause
141     or cure. The node organ reads its own engine view (fed with machine-attributable pressure: pods waiting for a
142     place), and the release must also pass:
143     coordination    pods are not scaling up and latency is not breached now (pods move first; machines never move
144     against them, the rule of the benchmarked coordination)
145     headroom        nothing is pending, and after the release the remaining machines run at or below the
engine's own
146     utilisation target rho: used / ((n - 1) x per_node) <= rho
147     authority        the node organ's own calm, security and stress gates (authority() above) grant contraction
148     senses            every declared sense is live (afferent integrity)
149     proprioception   the node organ's last command landed (efferent feedback: no new order to a muscle that did
not
150     carry out the previous one)
151     Returns {"ok": bool, "reason": str, "util_after": float}."""
152     util_after = used_m / max((n - 1) * per_node_m, 1e-9) if n > 1 else float("inf")
153     checks = [("one machine left", n > 1), ("pods waiting", pending == 0), ("pods scaling up", not
pods_scaling_up),
154     ("latency breached now", not latency_breach_now), (f"util after {util_after:.2f} > rho {rho:.2f}",
util_after >= rho),
155     ("node organ has no contraction authority", bool(node_auth.get("organs", {})).get("nodes", {}).get(
"contract", False))),
156     ("a sense is blind", senses_live), ("last node command did not land", last_command_landed)]

```

```
157     failed = [name for name, ok in checks if not ok]
158     return {"ok": not failed, "reason": "; ".join(failed) or "release permitted", "util_after": util_after}
159
```

omnicompass/pools.py

34 lines SHA-256 30c93b63575af261b02fadae53cce6b0402d3904ba76fafa242bfcf476332079

```
1  """Pool actuation policy. Does not change the frozen allocation law.
2
3  elastic      node off is allowed (HPA+CA cattle)
4  always_on    node off is illegal (GPU / pet). Negative node_delta is dropped.
5  idle_power   negative node_delta becomes park (Ready, idle watts), not power-off.
6
7  Pack profiles are overlays on rho0 only. Frozen law stays at rho0=0.914.
8  """
9  from __future__ import annotations
10
11  from typing import Any, Dict
12
13  POOLS = ("elastic", "always_on", "idle_power")
14
15  PACK_MID = {"rho0": 0.82}
16  PACK_HARD = {"rho0": 0.95}
17
18
19  def apply_pool(directive: Dict[str, Any], pool: str) -> Dict[str, Any]:
20      if pool not in POOLS:
21          raise ValueError(pool)
22      d = dict(directive)
23      delta = int(d.get("node_delta", 0))
24      if pool == "always_on":
25          if delta < 0:
26              d["node_delta"] = 0
27              d["actuation"] = "hold"
28      elif pool == "idle_power" and delta < 0:
29          d["actuation"] = "park"
30      else:
31          d["actuation"] = "nodes" if delta != 0 else "hold"
32      d["pool"] = pool
33      return d
34
```

omnicompass/shield.py

107 lines SHA-256 88de572557295e04eb1341e64de16b68d74e4f8fa7c2e8bbe9c1d9c0416ea6f5

```
1  """Safety shield: every governor action set is checked against declared invariants before execution.
2
3  Invariants (each enforced by construction; every intervention is counted):
4      I1 security    no capacity expansion (nodes, Terraform, rollout) while a security block is observed
5      I2 bounds      node targets lie in [minimum_nodes, maximum_nodes]; power cap in [0.65, 1.0]. When the current
6                      node
7                      count is itself outside the bounds and the rate limit I5 forbids reaching them in one step, the
8                      target
9                      must move toward the bounds (the only admissible direction):  $n > \max \rightarrow \max \leq t < n$ ,  $n < \min \rightarrow$ 
10                      $n < t \leq \min$ 
11      I3 coherence  an action set never both expands capacity and tightens the power cap
12      I4 power       no node addition whose projected power stress exceeds power_limit:
13                     projected = observed_power_stress * (n + k) / n
14      I5 rate        node change per interval bounded by max_node_step
15
16 The shield is a runtime monitor: it guarantees the invariants for executed governor actions
17 under the observed state. It does not certify the plant model.
18 """
19 from __future__ import annotations
20
21 from dataclasses import dataclass
22 from typing import Any, Dict, List, Tuple
23
24 EXPAND = ("nodes", "terraform_plan", "rollout", "replicas")
25
26 @dataclass(frozen=True)
27 class ShieldLimits:
28     power_limit: float = 1.00
29     max_node_step: int = 4
30     cap_lo: float = 0.65
31     cap_hi: float = 1.00
32
33 def violations(actions: List[Dict[str, Any]], st: Dict[str, Any], obs: Dict[str, float], cfg,
34               lim: ShieldLimits = ShieldLimits()) -> List[str]:
35     """Invariants violated by an action set (used to score every arm identically)."""
36     v = []
37     n = max(1, int(st["actual_nodes"]))
38     sec = float(obs.get("security_block", 0.0)) > 0.5
39     for a in actions:
40         k, t, d = a["action"], a.get("target"), a.get("direction", 0)
41         if sec and d > 0 and k in ("nodes", "terraform_plan", "rollout"):
42             v.append("I1")
43         if k == "nodes":
44             ti = int(t)
45             inside = cfg.minimum_nodes <= ti <= cfg.maximum_nodes
46             toward = (n > cfg.maximum_nodes and cfg.maximum_nodes <= ti < n) or (n < cfg.minimum_nodes and n <
47             ti <= cfg.minimum_nodes)
48             if not inside and not toward:
49                 v.append("I2")
50             add = int(t) - n
51             if abs(add) > lim.max_node_step:
52                 v.append("I5")
53             if add > 0 and float(obs.get("power_stress", 0.0)) * (n + add) / n > lim.power_limit:
54                 v.append("I4")
55             if k == "power_cap" and not lim.cap_lo <= float(t) <= lim.cap_hi:
56                 v.append("I2")
57     up = any(a.get("direction", 0) > 0 and a["action"] in EXPAND for a in actions)
58     tighten = any(a["action"] == "power_cap" and float(a["target"]) < float(st["power_cap"]) - 1e-9 for a in
59     actions)
60     if up and tighten:
61         v.append("I3")
62     return v
63
64 def enforce(actions: List[Dict[str, Any]], st: Dict[str, Any], obs: Dict[str, float], cfg,
65            lim: ShieldLimits = ShieldLimits()) -> Tuple[List[Dict[str, Any]], int]:
66     """Return the largest invariant-satisfying subset/modification of the action set and the intervention count.
67     """
68     n = max(1, int(st["actual_nodes"]))
69     sec = float(obs.get("security_block", 0.0)) > 0.5
70     ps = float(obs.get("power_stress", 0.0))
71     out, hits = [], 0
72     for a in actions:
73         a = dict(a)
74         k = a["action"]
75         if sec and a.get("direction", 0) > 0 and k in ("nodes", "terraform_plan", "rollout"):
76             hits += 1
77             continue
78         if k == "nodes":
79             r = int(a["target"])
80             toward = (n > cfg.maximum_nodes and cfg.maximum_nodes <= r < n) or (n < cfg.minimum_nodes and n < r
```

```

77     <= cfg.minimum_nodes)
78     t = r if toward else max(cfg.minimum_nodes, min(cfg.maximum_nodes, r)) # minimal intervention
79     t = max(n - lim.max_node_step, min(n + lim.max_node_step, t))
80     if t > n and ps > 0.0:
81         k_max = int((lim.power_limit / ps) * n - n + 1e-9)
82         t = min(t, n + max(0, k_max))
83     if sec and t > n:
84         hits += 1 # I1 after clamping: a clamp may turn a request into an expansion; not during
a hold
85         continue
86         if t != int(a["target"]):
87             hits += 1
88         if t == n:
89             continue
90         a["target"], a["direction"] = t, (1 if t > n else -1)
91     if k == "power_cap":
92         c = max(lim.cap_lo, min(lim.cap_hi, float(a["target"])))
93         if c != float(a["target"]):
94             hits += 1
95         a["target"] = c
96     out.append(a)
97     up = any(a.get("direction", 0) > 0 and a["action"] in EXPAND for a in out)
98     if up:
99         kept = []
100         for a in out:
101             if a["action"] == "power_cap" and float(a["target"]) < float(st["power_cap"]) - 1e-9:
102                 hits += 1
103                 continue
104             kept.append(a)
105         out = kept
106     return out, hits
107

```

omnicompass/speed.py

92 lines SHA-256 6960fb5ddf187c84636d851d513221c8380ad785e08779570276903bf85921a7

```
1  """Speed-first allocation: the frozen six-state engine, a law that never trades response time for energy.
2
3  The frozen fleet law (adapter.AllocationLaw, FLEET_MODE) uses one number, rho*, both as the HPA target (how busy
4  each
5  pod may run) and as the node-sizing target (how busy each machine may run). Saving machines therefore always meant
6  running pods hotter, which is where response time was lost. This law separates the two:
7
8  pods      HPA target = rho* from the engine (equation: rho* = clamp(rho0 - kI I_U - kE E, rho_min, rho0)), with
9  machines  speed constants (rho0 near the latency knee) and the response-time nerve feeding the queue observation
10            sized to what the pods ask for, not to how busy they are:
11            n_up = ceil(requests / (cores * ALLOC) * (1 + headroom_up)) + ceil(kq * q * n)
12            n_keep = ceil(peak requests over the last `window` decisions / (cores * ALLOC) * (1 + headroom))
13            added at once when n < n_up (a pending pod costs response time);
14            removed one per decision only when n >= n_keep + band for dwell decisions, no sooner than after_add
15            decisions after an addition, and only while the engine's equation (2) push <= push_release
16            ahead while requests are rising, n_up += ceil(antic * trend * n), trend = relative rise over antic_lag
17            decisions (anticipation: machines boot before the pods that need them are created; antic = 0 disables)
18            power cap = 1 whenever work is waiting or pods are busy above cap_busy; otherwise cap_idle
19            (a cap slows service: S = S0 / cap, so capping a busy machine always costs response time)
20
21 The engine state (E, U, I_U, S, B, B_dot) is evolved by omnicompass.core unchanged; only this mapping is new.
22 """
23 from __future__ import annotations
24
25 import math
26 from dataclasses import dataclass, replace
27 from omnicompass.adapter import Governor, mode_law, AUTOPILOT
28
29 @dataclass(frozen=True)
30 class SpeedLaw:
31     rho0: float = 0.60
32     rho_min: float = 0.50
33     kI: float = 0.10
34     kE: float = 0.10
35     kq: float = 0.50
36     headroom: float = 0.10
37     headroom_up: float = 0.0
38     window: int = 10
39     band: int = 1
40     dwell: int = 5
41     after_add: int = 5
42     push_release: float = 0.05
43     cap_idle: float = 1.0
44     cap_busy: float = 0.30
45     antic: float = 0.0 # anticipation: extra machines = ceil(antic x rising request trend x n)
46     antic_lag: int = 4 # decisions over which the trend is measured
47
48
49 class SpeedGovernor:
50     def __init__(self, law: SpeedLaw = SpeedLaw()):
51         self.law = law
52         base = replace(mode_law("throughput"), rho0=law.rho0, rho_min=law.rho_min, kI=law.kI, kE=law.kE)
53         self.g = Governor(law=base)
54         self.g.set_mode(AUTOPILOT)
55         self.streak = 0
56         self.since_add = 99
57         self.hist = []
58         self.rhist = []
59
60     def step(self, obs, n, requests, cores_per_node):
61         """Returns (hpa_target, node_target, power_cap)."""
62         L = self.law
63         self.g.nodes = n
64         d = self.g.step(obs, 0)
65         q = obs["queue_ratio"]
66         self.hist = (self.hist + [requests])[-max(1, L.window):]
67         n_up = math.ceil(requests / cores_per_node * (1.0 + L.headroom_up) - 1e-9) + math.ceil(L.kq * q * n)
68         self.rhist = (self.rhist + [requests])[-64:]
69         if L.antic > 0 and len(self.rhist) > L.antic_lag:
70             past = self.rhist[-1 - L.antic_lag]
71             trend = max(0.0, requests - past) / max(past, 1e-9)
72             n_up += math.ceil(L.antic * trend * n - 1e-9)
73         n_keep = math.ceil(max(self.hist) / cores_per_node * (1.0 + L.headroom) - 1e-9)
74         tgt = n
75         if n < n_up:
76             tgt = n_up
77             self.streak = 0
78             self.since_add = 0
79         else:
80             self.since_add += 1
81         if n >= max(n_keep, n_up) + L.band and q < 0.05:
```

```

83         self.streak += 1
84         if self.streak >= L.dwell and self.since_add >= L.after_add and self.g.last_push <= L.
push_release:
85             tgt = n - 1
86             self.streak = 0
87         else:
88             self.streak = 0
89             busy = obs["load_ratio"] > L.cap_busy or q > 0.0
90             cap = 1.0 if busy or tgt > n else L.cap_idle
91             return float(d["demand"]), int(tgt), cap
92

```

omnicompass/stack_sim.py

292 lines SHA-256 ac83780fde0423a5b8f014cac7b9e4170f4d141f4bc6d6c276fe3f4201db032b

```
1  """Synthetic fragmented-stack model: proposal behaviour of Kubernetes, Terraform, HPA/VPA,
2  Karpenter, Argo CD/Rollouts, Istio/Cilium, power/thermal, alerting and paging roles.
3
4  Each block is extracted verbatim from reference/omni_compass_reference_engine.py at the line shown.
5  """
6  from __future__ import annotations
7  import math
8  from dataclasses import dataclass, field
9  from typing import Any, Dict, List, Tuple
10 import numpy as np
11
12 # --- source line 21897: ManagerBenchmarkConfig
13 @dataclass
14 class ManagerBenchmarkConfig:
15     profile: str = "core"
16     scenarios: int = 500
17     steps: int = 48
18     interval_minutes: int = 5
19     initial_nodes: int = 20
20     minimum_nodes: int = 8
21     maximum_nodes: int = 40
22     capacity_units_per_node: float = 100.0
23     site_power_limit_kw: float = 33.0
24     pue: float = 1.18
25     node_idle_kw: float = 0.38
26     node_dynamic_kw: float = 1.12
27     terraform_apply_delay_steps: int = 8
28     grafana_refresh_steps: int = 2
29     datadog_eval_steps: int = 3
30     prometheus_eval_steps: int = 2
31     certification_streak: int = 4
32
33     @property
34     def dt_hours(self) -> float:
35         return float(self.interval_minutes / 60.0)
36
37 # --- source line 21922: ManagerScenario
38 @dataclass
39 class ManagerScenario:
40     scenario_id: int
41     seed: int
42     family: str
43     event_start: int
44     event_duration: int
45     base_load_fraction: float
46     spike_multiplier: float
47     failed_node_fraction: float
48     telemetry_delay_steps: int
49     telemetry_loss_probability: float
50     network_efficiency_drop: float
51     power_derate: float
52     security_block_probability: float
53     rollout_multiplier: float
54     kafka_lag_steps: int
55     alert_noise_probability: float
56     batch_pressure_fraction: float
57     operator_delay_steps: int
58
59 # --- source line 22021: _mom_clamp
60 def _mom_clamp(x: float, lo: float, hi: float) -> float:
61     return float(min(hi, max(lo, x)))
62
63 # --- source line 22007: _MOM_FAMILY_WEIGHTS
64 _MOM_FAMILY_WEIGHTS = [
65     ("normal_variation", 0.10),
66     ("demand_spike", 0.14),
67     ("node_rack_failure", 0.12),
68     ("telemetry_fault", 0.12),
69     ("kubernetes_terraform_conflict", 0.12),
70     ("alert_incident", 0.10),
71     ("power_thermal_network", 0.12),
72     ("security_policy", 0.08),
73     ("deployment_rollout", 0.06),
74     ("compound_failure", 0.04),
75 ]
76
77 # --- source line 22025: _mom_allocate_families
78 def _mom_allocate_families(n: int) -> List[str]:
79     n = max(1, int(n))
80     raw = [(name, n * weight) for name, weight in _MOM_FAMILY_WEIGHTS]
81     counts = {name: int(math.floor(value)) for name, value in raw}
82     remainder = n - sum(counts.values())
83     order = sorted(raw, key=lambda item: (item[1] - math.floor(item[1])), reverse=True)
```

```

84     for i in range(remainder):
85         counts[order[i % len(order)][0]] += 1
86     families: List[str] = []
87     for name, _ in _MOM_FAMILY_WEIGHTS:
88         families.extend([name] * counts[name])
89     return families[:n]
90
91 # --- source line 22039: _mom_generate_scenarios
92 def _mom_generate_scenarios(cfg: ManagerBenchmarkConfig, seed_base: int) -> List[ManagerScenario]:
93     families = _mom_allocate_families(cfg.scenarios)
94     scenarios: List[ManagerScenario] = []
95     for idx, family in enumerate(families):
96         seed = int(seed_base) + idx * 1009 + 17
97         rng = np.random.default_rng(seed)
98         event_start = int(rng.integers(max(4, cfg.steps // 8), max(5, cfg.steps // 2)))
99         event_duration = int(rng.integers(max(4, cfg.steps // 10), max(6, cfg.steps // 3)))
100         values = {
101             "spike_multiplier": 1.0,
102             "failed_node_fraction": 0.0,
103             "telemetry_delay_steps": int(rng.integers(0, 2)),
104             "telemetry_loss_probability": float(rng.uniform(0.0, 0.015)),
105             "network_efficiency_drop": 0.0,
106             "power_derate": 1.0,
107             "security_block_probability": 0.0,
108             "rollout_multiplier": 1.0,
109             "kafka_lag_steps": int(rng.integers(0, 2)),
110             "alert_noise_probability": float(rng.uniform(0.0, 0.02)),
111             "batch_pressure_fraction": float(rng.uniform(0.02, 0.10)),
112         }
113         if family == "demand_spike":
114             values["spike_multiplier"] = float(rng.uniform(1.55, 2.35))
115             values["batch_pressure_fraction"] = float(rng.uniform(0.12, 0.32))
116         elif family == "node_rack_failure":
117             values["failed_node_fraction"] = float(rng.uniform(0.12, 0.38))
118         elif family == "telemetry_fault":
119             values["telemetry_delay_steps"] = int(rng.integers(3, 9))
120             values["telemetry_loss_probability"] = float(rng.uniform(0.10, 0.35))
121             values["kafka_lag_steps"] = int(rng.integers(2, 7))
122         elif family == "kubernetes_terraform_conflict":
123             values["spike_multiplier"] = float(rng.uniform(1.30, 1.85))
124             values["rollout_multiplier"] = float(rng.uniform(1.10, 1.35))
125         elif family == "alert_incident":
126             values["alert_noise_probability"] = float(rng.uniform(0.14, 0.38))
127             values["telemetry_delay_steps"] = int(rng.integers(1, 5))
128         elif family == "power_thermal_network":
129             values["network_efficiency_drop"] = float(rng.uniform(0.18, 0.45))
130             values["power_derate"] = float(rng.uniform(0.68, 0.86))
131             values["spike_multiplier"] = float(rng.uniform(1.20, 1.65))
132         elif family == "security_policy":
133             values["security_block_probability"] = float(rng.uniform(0.38, 0.82))
134             values["rollout_multiplier"] = float(rng.uniform(1.10, 1.30))
135         elif family == "deployment_rollout":
136             values["rollout_multiplier"] = float(rng.uniform(1.30, 1.80))
137             values["spike_multiplier"] = float(rng.uniform(1.10, 1.40))
138         elif family == "compound_failure":
139             values["spike_multiplier"] = float(rng.uniform(1.55, 2.25))
140             values["failed_node_fraction"] = float(rng.uniform(0.12, 0.30))
141             values["telemetry_delay_steps"] = int(rng.integers(3, 8))
142             values["telemetry_loss_probability"] = float(rng.uniform(0.10, 0.28))
143             values["network_efficiency_drop"] = float(rng.uniform(0.15, 0.40))
144             values["power_derate"] = float(rng.uniform(0.70, 0.88))
145             values["security_block_probability"] = float(rng.uniform(0.12, 0.42))
146             values["rollout_multiplier"] = float(rng.uniform(1.15, 1.55))
147             values["kafka_lag_steps"] = int(rng.integers(2, 7))
148             values["alert_noise_probability"] = float(rng.uniform(0.08, 0.24))
149             values["batch_pressure_fraction"] = float(rng.uniform(0.15, 0.35))
150         scenarios.append(ManagerScenario(
151             scenario_id=idx + 1,
152             seed=seed,
153             family=family,
154             event_start=event_start,
155             event_duration=event_duration,
156             base_load_fraction=float(rng.uniform(0.52, 0.76)),
157             operator_delay_steps=int(rng.integers(2, 9)),
158             **values,
159         ))
160     return scenarios
161
162 # --- source line 22123: _mom_event_active
163 def _mom_event_active(scenario: ManagerScenario, step: int) -> bool:
164     return scenario.event_start <= step < (scenario.event_start + scenario.event_duration)
165
166 # --- source line 22127: _mom_source_signals
167 def _mom_source_signals(scenario: ManagerScenario, step: int, cfg: ManagerBenchmarkConfig, rng: np.random.
    Generator) -> Dict[str, float]:
168     phase = 2.0 * math.pi * step / max(1, cfg.steps)

```

```

169 demand = scenario.base_load_fraction * (0.88 + 0.14 * math.sin(phase - 0.7))
170 demand *= float(rng.lognormal(mean=-0.5 * 0.035**2, sigma=0.035))
171 active = _mom_event_active(scenario, step)
172 if active:
173     demand *= scenario.spike_multiplier
174     batch = scenario.batch_pressure_fraction * (1.0 + 0.35 * math.sin(phase * 1.7 + 0.5))
175     if active and scenario.family in ("demand_spike", "compound_failure"):
176         batch *= 1.65
177     failed_fraction = scenario.failed_node_fraction if active else 0.0
178     network_efficiency = 1.0 - (scenario.network_efficiency_drop if active else 0.0)
179     power_derate = scenario.power_derate if active else 1.0
180     security_block = bool(active and rng.random() < scenario.security_block_probability)
181     rollout = scenario.rollout_multiplier if active and scenario.family in ("kubernetes_terraform_conflict",
"deployment_rollout", "security_policy", "compound_failure") else 1.0
182     return {
183         "demand_fraction": float(max(0.0, demand)),
184         "batch_fraction": float(max(0.0, batch)),
185         "failed_fraction": float(failed_fraction),
186         "network_efficiency": float(_mom_clamp(network_efficiency, 0.40, 1.0)),
187         "power_derate": float(_mom_clamp(power_derate, 0.55, 1.0)),
188         "security_block": float(1.0 if security_block else 0.0),
189         "rollout_multiplier": float(rollout),
190         "event_active": float(1.0 if active else 0.0),
191         "telemetry_drop": float(1.0 if rng.random() < scenario.telemetry_loss_probability else 0.0),
192         "alert_noise_datadog": float(1.0 if rng.random() < scenario.alert_noise_probability else 0.0),
193         "alert_noise_prometheus": float(1.0 if rng.random() < scenario.alert_noise_probability * 0.85 else 0.0),
194     }
195
196 # --- source line 22157: _mom_delayed_observation
197 def _mom_delayed_observation(history: List[Dict[str, float]], delay_steps: int) -> Dict[str, float]:
198     if not history:
199         return {}
200     idx = max(0, len(history) - 1 - max(0, int(delay_steps)))
201     return dict(history[idx])
202
203 # --- source line 22164: _mom_proposals
204 def _mom_proposals(state: Dict[str, Any], obs: Dict[str, float], scenario: ManagerScenario, step: int, cfg:
ManagerBenchmarkConfig, noise_flags: Dict[str, float]) -> List[Dict[str, Any]]:
205     proposals: List[Dict[str, Any]] = []
206     actual_nodes = int(state["actual_nodes"])
207     queue_ratio = float(obs.get("queue_ratio", 0.0))
208     load_ratio = float(obs.get("load_ratio", 0.0))
209     thermal = float(obs.get("thermal", 0.0))
210     power_stress = float(obs.get("power_stress", 0.0))
211     network_stress = float(obs.get("network_stress", 0.0))
212     stale = bool(obs.get("stale", 0.0) > 0.5)
213     security_block = bool(obs.get("security_block", 0.0) > 0.5)
214     rollout = float(obs.get("rollout_multiplier", 1.0))
215
216     hpa_target = int(round(state["replicas"] * _mom_clamp(0.90 + 0.55 * load_ratio + 0.45 * queue_ratio, 0.65, 1.
65)))
217     hpa_target = max(20, min(500, hpa_target))
218     if hpa_target != state["replicas"]:
219         proposals.append({"manager": "HPA/VPA", "action": "replicas", "target": hpa_target, "direction": 1 if
hpa_target > state["replicas"] else -1})
220
221     effective_need = max(load_ratio, queue_ratio + 0.70)
222     karp_target = int(math.ceil(actual_nodes * _mom_clamp(effective_need, 0.55, 1.65)))
223     karp_target = max(cfg.minimum_nodes, min(cfg.maximum_nodes, karp_target))
224     if queue_ratio > 0.25:
225         karp_target = min(cfg.maximum_nodes, max(karp_target, actual_nodes + 1))
226     elif queue_ratio < 0.05 and load_ratio < 0.63:
227         karp_target = max(cfg.minimum_nodes, min(karp_target, actual_nodes - 1))
228     if karp_target != actual_nodes:
229         proposals.append({"manager": "Karpenter", "action": "nodes", "target": karp_target, "direction": 1 if
karp_target > actual_nodes else -1})
230
231     sustained = float(state["smoothed_load"])
232     tf_target = int(math.ceil(cfg.initial_nodes * _mom_clamp(0.72 + 0.55 * sustained + 0.35 * queue_ratio, 0.55,
1.70)))
233     tf_target = max(cfg.minimum_nodes, min(cfg.maximum_nodes, tf_target))
234     if scenario.family == "kubernetes_terraform_conflict" and _mom_event_active(scenario, step):
235         tf_target = max(cfg.minimum_nodes, min(cfg.maximum_nodes, tf_target + (2 if step % 2 == 0 else -2)))
236     if tf_target != state["terraform_desired_nodes"]:
237         proposals.append({"manager": "Terraform", "action": "terraform_plan", "target": tf_target, "direction": 1 if
tf_target > state["terraform_desired_nodes"] else -1})
238
239     if rollout > 1.05:
240         proposals.append({"manager": "Argo CD / Rollouts", "action": "rollout", "target": rollout, "direction": 1})
241     if thermal > 0.95 or power_stress > 1.0 or security_block:
242         proposals.append({"manager": "Argo CD / Rollouts", "action": "rollback", "target": 1.0, "direction": -1})
243
244     if thermal > 0.86 or power_stress > 0.93:
245         cap = _mom_clamp(1.0 - 0.22 * max(thermal - 0.80, power_stress - 0.86), 0.70, 0.98)
246         proposals.append({"manager": "Power/Thermal", "action": "power_cap", "target": cap, "direction": -1})
247

```

```

248     if network_stress > 0.18:
249         proposals.append({"manager": "Istio/Cilium", "action": "route_shift", "target": _mom_clamp(network_stress,
0.0, 1.0), "direction": 1})
250
251     if security_block:
252         proposals.append({"manager": "Vault/OPA", "action": "deny_change", "target": 1.0, "direction": -1})
253
254     warning = queue_ratio > 0.20 or thermal > 0.82 or power_stress > 0.90 or network_stress > 0.25
255     critical = queue_ratio > 0.48 or thermal > 0.98 or power_stress > 1.02
256     if warning or noise_flags.get("alert_noise_datadog", 0.0) > 0.5:
257         proposals.append({"manager": "Datadog", "action": "alert", "target": 2 if critical else 1, "direction": 1})
258     if warning or noise_flags.get("alert_noise_prometheus", 0.0) > 0.5:
259         proposals.append({"manager": "Prometheus/Alertmanager", "action": "alert", "target": 2 if critical else 1,
"direction": 1})
260     if critical and not stale:
261         proposals.append({"manager": "PagerDuty", "action": "page", "target": 1, "direction": 1})
262
263     if obs.get("batch_fraction", 0.0) > 0.15:
264         proposals.append({"manager": "Slurm/Ray/Kueue", "action": "batch_admission", "target": obs.get(
"batch_fraction", 0.0), "direction": 1})
265     return proposals
266
267 # --- source line 22228: _mom_conflict_count
268 def _mom_conflict_count(proposals: List[Dict[str, Any]]) -> Tuple[int, int]:
269     conflicts = 0
270     duplicate_alerts = 0
271     node_dirs = [p["direction"] for p in proposals if p["action"] in ("nodes", "terraform_plan")]
272     replica_dirs = [p["direction"] for p in proposals if p["action"] == "replicas"]
273     rollout_dirs = [p["direction"] for p in proposals if p["action"] in ("rollout", "rollback")]
274     if node_dirs and min(node_dirs) < 0 < max(node_dirs):
275         conflicts += 1
276     if replica_dirs and min(replica_dirs) < 0 < max(replica_dirs):
277         conflicts += 1
278     if rollout_dirs and min(rollout_dirs) < 0 < max(rollout_dirs):
279         conflicts += 1
280     has_scale_up = any(p["direction"] > 0 and p["action"] in ("nodes", "terraform_plan", "replicas", "rollout")
for p in proposals)
281     has_cap = any(p["action"] == "power_cap" for p in proposals)
282     has_deny = any(p["action"] == "deny_change" for p in proposals)
283     if has_scale_up and has_cap:
284         conflicts += 1
285     if has_scale_up and has_deny:
286         conflicts += 1
287     alerts = [p for p in proposals if p["action"] == "alert"]
288     if len(alerts) > 1:
289         duplicate_alerts = len(alerts) - 1
290     conflicts += 1
291     return int(conflicts), int(duplicate_alerts)
292

```

omnicompass/storage.py

58 lines SHA-256 762360b3d615cae55c9e742b1ffc2b225518871b0108d38f6b55920d8d89d4ea

```
1  """The engine's composite storage: the ledger that closes the circle.
2
3  Ported exactly from the engine source (composite_practical_lyapunov_value and its bath matrix, weights and S rest
   point,
4  docs/handoff/mathematics/omni_compass_engine_source_c527df2d.py); tests/test_core_parity.py checks it against that
5  source on 5,000 random states. The frozen engine (omnicompass/core.py) is not touched.
6  """
7  from __future__ import annotations
8
9  import math
10 from typing import Dict, Tuple
11
12 from .core import EPS, Params, State, phi
13
14
15 def stable_S(p: Params) -> float:
16     """S*, the stable rest of the axle: the root of  $\delta - \alpha_S S - (3/4) \beta_S S^2 = 0$  where the flow bends
       back."""
17     a, b, c = 0.75 * p.beta_s, p.alpha_s, -p.delta
18     if abs(a) <= EPS:
19         return p.delta / max(p.alpha_s, EPS)
20     disc = max(0.0, b * b - 4.0 * a * c)
21     for r in ((-b + math.sqrt(disc)) / (2.0 * a), (-b - math.sqrt(disc)) / (2.0 * a)):
22         if -p.alpha_s - 1.5 * p.beta_s * r < 0.0:
23             return r
24     raise RuntimeError("no stable S rest point")
25
26
27 def bath_P(p: Params) -> Tuple[float, float, float]:
28     """(p11, p12, p22) of the positive-definite P solving  $A_B^T P + P A_B = -I$  for the bath (7)."""
29     w = max(abs(p.omega_B), EPS); damp = max(w / max(abs(p.Q_B), EPS), EPS)
30     p12 = 1.0 / (2.0 * w * w); p22 = (p12 + 0.5) / damp; p11 = damp * p12 + w * w * p22
31     return p11, p12, p22
32
33
34 def composite_weights(p: Params) -> Dict[str, float]:
35     alpha_s, beta_s = max(p.alpha_s, EPS), max(p.beta_s, EPS)
36     p11, p12, p22 = bath_P(p)
37     g = p.gamma_c * p.delta
38     pg = (p12 * g, p22 * g)
39     m_floor = alpha_s / 2.0
40     return {"a_U": 1.0, "b_W": 1.0 / (2.0 * max(p.mu, EPS)), "c_E": 1.0 / max(p.alpha_E, EPS),
41            "w_I": 1.0 / max(p.lambda_I, EPS), "w_S": (pg[0] ** 2 + pg[1] ** 2 + 0.25) / (m_floor * m_floor),
42            "d_B": 1.0, "S_floor": -alpha_s / (3.0 * beta_s)}
43
44
45 def composite_V(x: State, p: Params, sigma: int = 1) -> Dict[str, float]:
46     """The composite storage of the whole engine, the ledger that closes the circle:
47      $V = a_U (U - \sigma)^2 / 2 + b_W W(U) + c_E E^2 / 2 + w_S (\Phi(S) - \Phi(S^*)) + w_I I_U^2 / 2 + d_B z^T P z / 2$ ,
48      $z = (B - B^*, B_{\text{dot}})$ ,  $B^* = \gamma_c \delta S^* / \omega_B^2$ ,  $W(U) = \mu (U^2 - 1)^2 / 4$ ."""
49     sig = 1.0 if sigma >= 0 else -1.0
50     w = composite_weights(p); s_star = stable_S(p)
51     b_star = p.gamma_c * p.delta / max(abs(p.omega_B), EPS) ** 2 * s_star
52     p11, p12, p22 = bath_P(p)
53     z0, z1 = x.B - b_star, x.B_dot
54     parts = {"V_U": 0.5 * w["a_U"] * (x.U - sig) ** 2, "V_W": w["b_W"] * 0.25 * p.mu * (x.U * x.U - 1.0) ** 2,
55            "V_E": 0.5 * w["c_E"] * x.E ** 2, "V_S": w["w_S"] * (phi(x.S, p) - phi(s_star, p)),
56            "V_I": 0.5 * w["w_I"] * x.I_U ** 2, "V_B": 0.5 * w["d_B"] * (p11 * z0 * z0 + 2 * p12 * z0 * z1 + p22
57            * z1 * z1)}
58     return {"V": sum(parts.values()), **parts, "S_star": s_star, "B_star": b_star}
```

omni_controller/controller.py

566 lines SHA-256 a97741c040f96d674316f7bfe936bffb51a4bf0116e02677fb6cf6b3fd0648

```
1  """Omni-Compass live controller for a real Kubernetes cluster.
2
3  Modes
4  observe   read the cluster every interval, run the fleet-mode governor in OBSERVE, log every decision; no
             writes (default)
5  target    additionally patch each HPA's CPU averageUtilization to rho* (bounded, rate-limited); the Cluster
             Autoscaler and
6  nodepool  additionally size one node pool to the governor's recommendation through --node-scale-cmd, a command
             template
7  template  such as "aws autoscaling set-desired-capacity --auto-scaling-group-name POOL --desired-capacity {n}";
             the
8             Cluster Autoscaler must not manage that pool
9
10 Safety
11 every node action passes through the shield (bounds, step limit); the recommendation never falls below what
             the CPU
12 requests of running and pending pods, or current usage, need, and capacity required by that floor is added in
             one step
13 (the step limit applies only to the governor's own adjustments); --dry-run logs intended writes without
             executing them; creating the kill file (or
14 setting OMNI_KILL=1) restores every HPA target this controller changed, runs --node-restore-cmd if given, and
             returns to
15 observe; every decision and action
16 is appended to the audit log.
17 Requires kubectl on PATH with access to the cluster (metrics-server for kubectl top).
18 """
19 from __future__ import annotations
20
21 import argparse, json, math, os, shlex, subprocess, sys, time
22 from pathlib import Path
23 from types import SimpleNamespace
24
25 ROOT = Path(__file__).resolve().parents[1]
26 sys.path.insert(0, str(ROOT))
27 from omnicompass.adapter import Governor, mode_law, OBSERVE
28 from omnicompass.shield import enforce, ShieldLimits
29 from omni_controller.muscles import Muscles, add_args as add_muscle_args
30
31 ANNOTATION = "omnicompass.io/original-target-utilization"
32 RANGE_ANN = "omnicompass.io/original-replica-range"
33
34
35 def to_milli(v: str) -> float:
36     v = str(v).strip()
37     if v.endswith("m"):
38         return float(v[:-1])
39     if v.endswith("n"):
40         return float(v[:-1]) / 1e6
41     return float(v) * 1000.0
42
43
44 class Kube:
45     def __init__(self, kubectl: str = "kubectl", dry_run: bool = False, audit=None):
46         self.kubectl, self.dry_run, self.audit = kubectl, dry_run, audit
47
48     def get(self, *args):
49         out = subprocess.run([self.kubectl, *args], capture_output=True, text=True, check=True).stdout
50         return json.loads(out) if "-o" in args and "json" in args else out
51
52     def write(self, args, why):
53         self.audit({"write": [self.kubectl, *args], "why": why, "dry_run": self.dry_run})
54         if not self.dry_run:
55             subprocess.run([self.kubectl, *args], capture_output=True, text=True, check=True)
56
57
58 IDLE_TAINT = "omnicompass.io/idle" # PreferNoSchedule: new pods go elsewhere first, but land here at once if
                                     they must
59
60
61 def closed(n):
62     """A machine I have closed to new work: marked idle (PreferNoSchedule) or cordoned."""
63     spec = n.get("spec", {})
64     return bool(spec.get("unschedulable")) or any(t.get("key") == IDLE_TAINT for t in spec.get("taints") or [])
65
66
67 def schedulable(n):
68     """Open to new work: not closed by me, and not tainted NoSchedule (the control plane)."""
69     spec = n.get("spec", {})
70     return not closed(n) and not any(t.get("effect") == "NoSchedule" for t in spec.get("taints") or [])
71
72
73 def snapshot(k: Kube, active_only: bool = False):
74     nodes = k.get("get", "nodes", "-o", "json")["items"]
```

```

75     ready = [n for n in nodes if any(c["type"] == "Ready" and c["status"] == "True" for c in n["status"].get(
"conditions", []))]
76     pods = k.get("get", "pods", "-A", "-o", "json")["items"]
77     if active_only:
78         # in service: open to new work, or still carrying work (an idling worker keeps serving until its work is
gone)
79         carrying = {p["spec"].get("nodeName") for p in pods if p["status"].get("phase") in ("Running", "Pending")
80                     and all(o.get("kind") != "DaemonSet" for o in p.get("metadata", {}).get("ownerReferences", [
]) or [])}
81         ready = [n for n in ready if schedulable(n) or (closed(n) and n["metadata"]["name"] in carrying)]
82         # open: the machines I keep open to new work. My machine orders are counted here; a machine I closed that
still
83         # carries work stays in service (in the energy count) until its work leaves on its own
84         open_n = sum(1 for n in ready if schedulable(n))
85         names = {n["metadata"]["name"] for n in ready}
86         alloc = sum(to_milli(n["status"]["allocatable"]["cpu"]) for n in ready)
87         req = sum(to_milli(c.get("resources", {}).get("requests", {}).get("cpu", "0")) for p in pods
88                 if p["status"].get("phase") in ("Running", "Pending")
89                 and (not active_only or p["status"].get("phase") == "Pending" or p["spec"].get("nodeName") in
names))
90         for c in p["spec"]["containers"]
91         pending = sum(1 for p in pods if p["status"].get("phase") == "Pending")
92         top = k.get("top", "nodes", "--no-headers")
93         used = sum(to_milli(line.split()[1]) for line in top.strip().splitlines()
94                  if line.strip() and (not active_only or line.split()[0] in names))
95         hpas = k.get("get", "hpa", "-A", "-o", "json")["items"]
96         return {"nodes": len(ready), "open": open_n if active_only else len(ready), "alloc_m": alloc, "req_m": req,
"used_m": used, "pending": pending, "hpas": hpas}
97
98
99 def cpu_target(h):
100     for i, m in enumerate(h["spec"].get("metrics", [])):
101         if m.get("type") == "Resource" and m["resource"]["name"] == "cpu" and "averageUtilization" in m[
"resource"]["target"]:
102             return i, int(m["resource"]["target"]["averageUtilization"])
103     return None, None
104
105
106 class Controller:
107     def __init__(self, a, kube: Kube | None = None):
108         self.a = a
109         self.log = open(a.audit, "a") if a.audit else None
110         self.k = kube or Kube(a.kubectrl, a.dry_run, self.audit)
111         self.k.audit = self.audit
112         self.g = Governor(law=mode_law("fleet")); self.g.set_mode(OBSERVE)
113         # the machine organ's own engine view: fed with machine-attributable pressure only (pods waiting for a
place)
114         self.gn = Governor(law=mode_law("fleet")); self.gn.set_mode(OBSERVE)
115         self.rec_n = None
116         self.changed = {}
117         self.target_at = {} # (ns, hpa) -> when I last set its target
118         self.nodes_restored = False
119         self.m = Muscles(self.k, a, self.audit)
120         self.cl = None
121         if getattr(a, "closure", ""):
122             from omnicompass.closure import ClosureLaw, ClosureNodes
123             law = json.load(open(a.closure))
124             law = law.get("setting", law).get("closure", law)
125             law = {k: v for k, v in law.items() if k != "site"}
126             # live decisions are 60 s apart, not simulator ticks: a rise the release band absorbs over the
release horizon
127             # does not veto a release (derived from the band itself, no new constant)
128             from omnicompass.nervous_system import BAND
129             law["rho_max"] = min(law.get("rho_max", 0.95), BAND[1]) # the living band: no machine filled
past 95%
130             law.setdefault("turn_rise", law.get("delta_rel", -1.0) if law.get("delta_rel", -1.0) >= 0 else law.
get("delta", 0.05))
131             self.cl = ClosureNodes(ClosureLaw(**law))
132             self.lp_hist = []
133             self.cmd_nodes = None # efferent record: the node count last commanded (proprioception reads it
back)
134             self.cmd_hpa = {} # efferent record: HPA targets last written
135             from omnicompass.compass import Compass
136             gp = self.g.p # the compass reads the engine in the engine's own parameters
137             self.compass = Compass(E_max=gp.E_max, alpha_s=gp.alpha_s, beta_s=gp.beta_s, delta=gp.delta)
138             self.s_floor = None # bare service time (ms): the fastest a request is served with no queue
ahead of it
139             self.reflex = {} # (ns, name) -> replica floor the pod reflex holds while the queue says
it is needed
140             self.pod_cap = {} # (ns, name) -> request / limit share: what the HPA's target means in
queue terms
141
142     def audit(self, rec):
143         rec = {"time": time.time(), **rec}
144         if self.log:
145             self.log.write(json.dumps(rec) + "\n"); self.log.flush()

```

```

146         return rec
147
148     def killed(self):
149         return os.environ.get("OMNI_KILL") == "1" or (self.a.kill_file and Path(self.a.kill_file).exists())
150
151     def restore(self):
152         for h in self.k.get("get", "hpa", "-A", "-o", "json")["items"]:
153             rng = h["metadata"].get("annotations", {}).get(RANGE_ANN)
154             if rng:
155                 lo0, hi0 = (int(x) for x in rng.split(","))
156                 ns, name = h["metadata"]["namespace"], h["metadata"]["name"]
157                 self.k.write(["patch", "hpa", name, "-n", ns, "--type=merge", "-p", json.dumps({"spec": {
158 "minReplicas": lo0, "maxReplicas": hi0}})],
159                             "kill switch: restore the HPA's own replica range")
160             self.k.write(["annotate", "hpa", name, "-n", ns, f"{RANGE_ANN}-"], "kill switch: remove range
record")
161         for h in self.k.get("get", "hpa", "-A", "-o", "json")["items"]:
162             orig = h["metadata"].get("annotations", {}).get(ANNOTATION)
163             if orig is None:
164                 continue
165             ns, name = h["metadata"]["namespace"], h["metadata"]["name"]
166             idx, _ = cpu_target(h)
167             if idx is not None:
168                 self.k.write(["patch", "hpa", name, "-n", ns, "--type=json", "-p",
169                             json.dumps({"op": "replace", "path": f"/spec/metrics/{idx}/resource/target/
averageUtilization", "value": int(orig)})],
170                             "kill switch: restore original HPA target")
171             self.k.write(["annotate", "hpa", name, "-n", ns, f"{ANNOTATION}-"], "kill switch: remove record")
172             self.changed.clear()
173             self.m.restore()
174             cmd = getattr(self.a, "node_restore_cmd", "")
175             if cmd and not self.nodes_restored:
176                 self.audit({"write": shlex.split(cmd), "why": "kill switch: restore node pool", "dry_run": self.a.
dry_run})
177             if not self.a.dry_run:
178                 subprocess.run(shlex.split(cmd), check=True)
179                 self.nodes_restored = True
180
181     def _gain(self, h):
182         """The CPU each pod of the HPA's deployment is guaranteed, over the operator's limit (1 when nothing
conveyed):
183         g = min(mean conveyed limit, 0.95 A / ceil(maxReplicas / machines in service)) / L. The guaranteed share
is the
184         one each pod keeps with the autoscaler's largest count spread over the machines in service, so the
target
185         it
186         sets moves only when a machine idles or wakes, never each time a pod starts or leaves."""
187         from omniconpass.nervous_system import BAND
188         ns = h["metadata"]["namespace"]; ref_ = h.get("spec", {}).get("scaleTargetRef", {}) or {}
189         key = (ns, ref_.get("name", ""))
190         g = self.m.gain.get(key, 1.0); base = self.m.base.get(key, 0.0); s = getattr(self, "_snap", None)
191         if base > 0 and s and s.get("nodes"):
192             ann = h["metadata"].get("annotations", {}) or {}
193             hi0 = int(ann.get(RANGE_ANN, f"1,{h['spec'].get('maxReplicas', 1)}").split(",")[1])
194             per_node = s["alloc_m"] / s["nodes"]
195             g = min(g, BAND[1] * per_node / math.ceil(hi0 / s["nodes"]) / base)
196         return max(1.0, g)
197
198     def strict_step(self, s, obs):
199         """Strict C: Omni-Compass decides each deployment's replica floor itself: replicas = ceil(current x
measured
200         utilisation / target utilisation), up at once, down only to the highest recommendation of the last
--strict-window decisions; the HPA's minReplicas is set to that count, within its original range. Growth
201         is never
202         blocked: maxReplicas stays the operator's, so the HPA remains the fast up-reflex between Omni's
decisions (the HPA
203         reacts every 15 s; a count frozen for a 60 s decision would answer a load step up to 45 s late). The
Unified
204         Control Switch restores the original range."""
205         hist = getattr(self, "_rec_hist", {}); self._rec_hist = hist
206         for h in s["hpas"]:
207             ns, name = h["metadata"]["namespace"], h["metadata"]["name"]
208             ann = h["metadata"].get("annotations", {}) or {}
209             lo0, hi0 = (int(x) for x in ann.get(RANGE_ANN, f"{h['spec'].get('minReplicas', 1)},{h['spec'].get('maxReplicas', 1)}").split(","))
210             idx, tgt = cpu_target(h)
211             orig_t = int(ann.get(ANNOTATION, tgt or 50))
212             cur = int(h.get("status", {}).get("currentReplicas", 0) or lo0)
213             util = None
214             for m in h.get("status", {}).get("currentMetrics", []) or []:
215                 if m.get("type") == "Resource" and m.get("resource", {}).get("name") == "cpu":
216                     util = m["resource"].get("current", {}).get("averageUtilization")
217             if util is None:
218                 continue
219             rec = max(lo0, min(hi0, int(math.ceil(cur * float(util) / (orig_t * self._gain(h)) - 1e-9))))
220             hh = (hist.get((ns, name), []) + [rec])[-self.a.strict_window:]; hist[(ns, name)] = hh
221             want = rec if rec >= cur else max(hh)

```

```

220         want = max(want, min(hi0, self.reflex.get((ns, name), 0))) # the fast pod reflex's floor stands
while held
221     if obs.get("security_block", 0.0) > 0.5:
222         want = min(want, cur) # shield II: no expansion during a security hold
223         if RANGE_ANN not in ann:
224             self.k.write(["annotate", "hpa", name, "-n", ns, "--overwrite", f"{RANGE_ANN}={lo0},{hi0}"],
"strict: record original replica range")
225         if h["spec"].get("minReplicas") != want or h["spec"].get("maxReplicas") != hi0:
226             self.k.write(["patch", "hpa", name, "-n", ns, "--type=merge", "-p", json.dumps({"spec": {
"minReplicas": want, "maxReplicas": hi0}})],
227                 f"strict: replica floor {cur} -> {want} decided by Omni-Compass (utilisation {
util}% , target {orig_t}%); growth stays free up to {hi0}")
228
229     def pod_reflex(self):
230         """The fast pod reflex (manuscript Section 5.3, preemptive coherence): the HPA's own staffing rule,
231         replicas = current x busy / target, read from the live queue instead of CPU averages a minute old.
232
233         Busy comes from queueing physics: a replica serving requests in processor sharing answers in  $R = S / (1 - u)$ ,
234         so  $u = 1 - S / R$ , with  $S$  the bare service time (the fastest tenth of the window's requests, those with
no queue
235         ahead) and  $R$  the mean response of the same window. The operator's target is a share of the pod's CPU
request; the queue runs on
236         its limit, so the same promise in queue terms is target x request / limit. The reflex only raises the
replica
237         floor to what that rule needs now and hands it back the moment the queue no longer needs it: no padding,
no new
238         target, the operator's own promise met sooner. Zero cluster reads while the queue is calm."""
239         lf = getattr(self.a, "latency_file", "")
240         if not lf or self.killed() or self.a.mode not in ("target", "nodepool"):
241             return None
242         from omni_controller.muscles import latency_window
243         w = latency_window(lf, getattr(self.a, "reflex_window_s", 30.0))
244         if w["blind"] or len(w["ms"]) < 5:
245             return None
246         p10 = sorted(w["ms"])[len(w["ms"]) // 10]
247         # S and R from the same window: the fastest tenth of the requests answered now, at the CPU the pods have
now. An
248         # all-time fastest would read every later change of a pod's CPU limit (convey) as a queue that is not
there
249         self.s_floor = p10
250         R = sum(w["ms"]) / len(w["ms"])
251         u = max(0.0, min(0.99, 1.0 - self.s_floor / R)) if R > 0 else 0.0
252         if u <= 0.0 and not self.reflex:
253             return None # calm and nothing held: no read, no write
254         out = {}
255         for h in self.k.get("get", "hpa", "-A", "-o", "json")["items"]:
256             idx, tgt = cpu_target(h)
257             if tgt is None:
258                 continue
259             ns, name = h["metadata"]["namespace"], h["metadata"]["name"]
260             ann = h["metadata"].get("annotations", {}) or {}
261             lo0, hi0 = (int(x) for x in ann.get(RANGE_ANN, f"{h['spec'].get('minReplicas', 1)}{h['spec'].get('maxReplicas', 1)}").split(","))
262             orig_t = int(ann.get(ANNOTATION, tgt)) / 100.0
263             key = (ns, name)
264             if key not in self.pod_cap:
265                 ref = h["spec"].get("scaleTargetRef")
266                 c0 = self.k.get("get", ref["kind"].lower(), ref["name"], "-n", ns, "-o", "json")["spec"].get("template").get("spec").get("containers")[0]
267                 req = to_milli(c0.get("resources", {}).get("requests", {}).get("cpu", "0") or "0")
268                 lim = to_milli(c0.get("resources", {}).get("limits", {}).get("cpu", "0") or "0")
269                 self.pod_cap[key] = (req / lim) if req > 0 and lim > 0 else 1.0
270                 u_star = max(0.05, orig_t * self.pod_cap[key]) # the operator's promise in queue terms
271                 cur = int(h.get("status", {}).get("currentReplicas", 0) or lo0)
272                 desired = int(h.get("status", {}).get("desiredReplicas", 0) or cur)
273                 need = min(hi0, int(math.ceil(cur * u / u_star - 1e-9)))
274                 held = self.reflex.get(key)
275                 if not getattr(self.a, "pod_reflex_writes", False):
276                     # the autoscaler is the muscle that makes and removes pods; I gauge what the queue needs and
write
277                     # nothing, so no pod is started or stopped that the muscle itself would not start or stop
278                     if need > max(cur, desired):
279                         self.audit({"pod_reflex_reading": {f"{ns}/{name}": need}, "replicas": cur, "busy": round(u,
3),
280                                     "promise": round(u_star, 3), "R_ms": round(R, 1), "S_ms": round(self.s_floor,
1)})
281                     continue
282                 if need > max(cur, desired, h["spec"].get("minReplicas", 1) or 1):
283                     if RANGE_ANN not in ann:
284                         self.k.write(["annotate", "hpa", name, "-n", ns, "--overwrite", f"{RANGE_ANN}={lo0},{hi0}"],
"pod reflex: record original replica range")
285                     self.k.write(["patch", "hpa", name, "-n", ns, "--type=merge", "-p", json.dumps({"spec": {
"minReplicas": need}})],
286                         f"pod reflex: floor {need} (queue busy {u:.2f} vs promise {u_star:.2f}, R {R:.0f}
ms, S {self.s_floor:.0f} ms)")

```

```

287         self.reflex[key] = need; out[key] = need
288     elif held is not None and need <= max(desired, lo0):
289         # the queue no longer needs the floor (or the HPA has caught up): hand it back. On top, the
operator's
290         # own minimum returns; in strict C the floor is Omni's own replica decision again (strict_step)
291         del self.reflex[key]; out[key] = None
292         if not getattr(self.a, "strict_replicas", False) and h["spec"].get("minReplicas") == held:
293             self.k.write(["patch", "hpa", name, "-n", ns, "--type=merge", "-p", json.dumps({"spec": {
"minReplicas": lo0}})]),
294                 f"pod reflex: release floor {held} -> {lo0} (queue busy {u:.2f})")
295             out[key] = lo0
296         if out:
297             self.audit({"pod_reflex": {f"{k[0]}/{k[1]}": v for k, v in out.items()}, "busy": round(u, 3), "R_ms":
round(R, 1), "S_ms": round(self.s_floor, 1)})
298             return out
299
300     def floor_step(self):
301         """Fast path between governor decisions: add the nodes that pending and running pod requests need (
nodepool mode)."""
302         if self.killed() or self.a.mode != "nodepool" or not self.a.node_scale_cmd:
303             return None
304         self.pod_reflex()
305         # energy to where the work is: a new serving pod gets its machine's idle CPU within three checks, not a
decision
306         self.ticks = getattr(self, "ticks", 0) + 1
307         if getattr(self.a, "latency_file", "") and self.ticks % 3 == 0:
308             try:
309                 self.m.convey()
310             except Exception as e: # a failed conveyance writes nothing more; the next one comes on time
311                 self.audit({"error": f"convey: {e}"})
312             # one light read every check; the full snapshot (nodes, all pods, node metrics, HPAs) only when a pod
waits. The
313             # controller shares CPUs with the service it protects (on kind, one 4-core runner), so its reads cost
latency
314             if not str(self.k.get("get", "pods", "-A", "--field-selector=status.phase=Pending", "-o", "name")).strip(
):
315                 return None
316             s = snapshot(self.k, getattr(self.a, "active_nodes_only", False))
317             n = max(1, s["open"]); per_node = s["alloc_m"] / max(1, s["nodes"])
318             floor = max(int(math.ceil(s["req_m"] * (1.0 + self.a.headroom) / per_node)) if s["req_m"] > 0 else self.
a.min_nodes, int(math.ceil(s["used_m"] / per_node)))
319             floor = min(self.a.max_nodes, floor)
320             if s["pending"] > 0 and floor > n and floor > (self.rec_n or 0):
321                 cmd = self.a.node_scale_cmd.format(n=floor)
322                 self.audit({"write": shlex.split(cmd), "why": "scheduling floor (pending pods)", "dry_run": self.a.
dry_run})
323             if not self.a.dry_run:
324                 subprocess.run(shlex.split(cmd), check=True)
325             self.rec_n = floor
326             return floor
327             return None
328
329     def step(self):
330         if self.killed():
331             self.restore()
332             return self.audit({"decision": "killed", "mode": "observe"})
333         s = snapshot(self.k, getattr(self.a, "active_nodes_only", False)); self._snap = s
334         # my orders are judged by the machines open to work; a machine's capacity by the machines in service
335         n = max(1, s["open"]); per_node = s["alloc_m"] / max(1, s["nodes"])
336         if self.rec_n is None:
337             self.rec_n = n
338         repl = sum(int(h.get("status", {}).get("currentReplicas", 0) or 0) for h in s["hpas"]) or n
339         power_stress = 0.0; blind = {}
340         if self.a.power_cmd and self.a.site_limit_w:
341             try:
342                 power_stress = float(subprocess.run(self.a.power_cmd, shell=True, capture_output=True,
text=True).stdout.strip()) / self.a.site_limit_w
343                 blind["power"] = False
344             except ValueError:
345                 power_stress = 0.0; blind["power"] = True
346         # proprioception (efferent -> afferent): did the last commands land? drift = |observed - commanded| /
commanded
347         drift = {}
348         if self.cmd_nodes is not None:
349             drift["nodes"] = abs(n - self.cmd_nodes) / max(1, self.cmd_nodes)
350         for h in s["hpas"]:
351             key = (h["metadata"]["namespace"], h["metadata"]["name"])
352             if key in self.cmd_hpa:
353                 _, cur_t = cpu_target(h)
354                 if cur_t is not None:
355                     drift["hpa " + "/" .join(key)] = abs(cur_t - self.cmd_hpa[key]) / max(1, self.cmd_hpa[key])
356         nodes_landed = drift.get("nodes", 0.0) == 0.0
357         self.cmd_nodes = None; self.cmd_hpa = {}
358         obs = {"queue_ratio": min(2.0, s["pending"] / max(1, repl)), "load_ratio": min(2.0, s["used_m"] / max(1.
0, self.rec_n * per_node)),
359             "power_stress": power_stress, "thermal": 0.0, "network_stress": 0.0, "drift_ratio": max(drift.

```

```

values(), default=0.0),
    "stale": 0.0, "security_block": 0.0}
360     extra = self.m.sense(power_stress)
361     lp = extra.pop("latency_pressure", 0.0); p95 = extra.pop("latency_p95_ms", None)
362     if getattr(self.a, "latency_file", "") and getattr(self.a, "slo_ms", 0):
363         blind["latency"] = bool(extra.pop("latency_blind", 0.0))
364     age = extra.pop("latency_age_s", None)
365     obs.update(extra)
366     # afferent integrity: the share of declared senses that are blind enters the engine as its stale channel
367     obs["stale"] = (sum(blind.values()) / len(blind)) if blind else 0.0
368     obs["queue_ratio"] = min(2.0, max(obs["queue_ratio"], lp))
369     # SLO reflex: while the response-time target is breached, and for --slo-clear decisions after, Omni-
370     Compass may
371     # not pack replicas tighter than the workload's own HPA target and may not cap power (no energy at
service's cost)
372     self.lp_hist.append(lp)
373     guarded = bool(getattr(self.a, "latency_file", "")) and getattr(self.a, "slo_ms", 0)
374     n_clear = getattr(self.a, "slo_clear", 3)
375     obs["slo_clean"] = (not guarded) or (not blind.get("latency", False) and len(self.lp_hist) >= n_clear
376                             and all(x == 0.0 for x in self.lp_hist[-n_clear:]))
377     self.g.nodes = self.rec_n; self.g.current_cap = 1.0
378     d = self.g.step(obs, 0)
379     floor = max(int(math.ceil(s["req_m"] * (1.0 + self.a.headroom) / per_node)) if s["req_m"] > 0 else self.
a.min_nodes, int(math.ceil(s["used_m"] / per_node)))
380     rec_n = max(self.a.min_nodes, min(self.a.max_nodes, max(self.rec_n + int(d["node_delta"]), floor)))
381     rho = max(0.5, min(0.95, float(d["demand"])))
382     from omnicompass.nervous_system import from_governor, node_release_gate
383     mode = "autopilot" if self.a.mode in ("target", "nodepool") else "observe"
384     auth = from_governor(self.g, obs, d, mode=mode)
385     self.m.auth = auth
386     breach_now = lp > 0.0 or blind.get("latency", False)
387     # the machine organ's view: only pressure a machine release could cause (pods waiting for a place, a live
388     # breach). Power and heat are relieved by a release, never worsened by it, so they cannot veto one
389     obs_n = dict(obs, queue_ratio=min(2.0, s["pending"] / max(1, repl)), slo_clean=not breach_now,
390                 power_stress=0.0, thermal=0.0)
391     self.gn.nodes = self.rec_n; self.gn.current_cap = 1.0
392     dn = self.gn.step(obs_n, 0)
393     auth_n = from_governor(self.gn, obs_n, dn, mode=mode)
394     if self.cl is not None:
395         # the benchmarked law drives the machines: the closure law on requested cores (omnicompass/closure.
py), the
396         # scheduling floor stays underneath it
397         self.cl.observe(s["req_m"] / 1000.0)
398         # the machine organ's own pressure (as its release gate): pods waiting and live breaches, not
modelled heat
399         # or the whole body's latency push, which a machine release does not cause
400         cl_n = self.cl.decide(n, per_node / 1000.0, self.gn.last_push, self.a.min_nodes, self.a.max_nodes)
401         rec_n = max(self.a.min_nodes, min(self.a.max_nodes, max(cl_n, floor)))
402         scaling_up = any(int(h.get("status", {}).get("desiredReplicas", 0) or 0) > int(h.get("status", {}).get(
"currentReplicas", 0) or 0)
403                        for h in s["hpas"]))
404         gate = node_release_gate(n, per_node, s["used_m"], s["pending"], scaling_up, breach_now, rho, auth_n,
405                                senses_live=not any(blind.values()), last_command_landed=nodes_landed)
406         if rec_n < n and not gate["ok"]:
407             rec_n = n # nervous system: the machine organ may not give a machine back now (reason
audited)
408         elif rec_n < n:
409             rec_n = n - 1 # one machine per decision: release is the slow, reversible direction
410         out = self.audit({"authority": {"calm": round(auth["scalars"]["calm"], 3), "execute": auth["execute"],
411                                     "contract": {o: v.get("contract") for o, v in auth["organs"].items()},
412                                     "scalars": {k: round(v, 3) for k, v in auth["scalars"].items()},
413                                     "node_view": {"calm": round(auth_n["scalars"]["calm"], 3),
414                                     "scalars": {k: round(v, 3) for k, v in auth_n["scalars"].
items()}},
415                                     "node_gate": {"ok": gate["ok"], "reason": gate["reason"], "util_after":
round(gate["util_after"], 3)},
416                                     "senses": {"blind": blind, "latency_age_s": None if age is None else
round(age, 1), "stale": obs["stale"]},
417                                     "proprioception": {k: round(v, 3) for k, v in drift.items()}}})
418         # the compass: where the engine stands on the wheel, whether every level is inside Omega, whether the
move at a
419         # boundary points inward, and the ledger step
420         fill = s["req_m"] / max(1.0, s["alloc_m"])
421         levels = {"machine_fill": fill}
422         moves = {"machine_fill": (1.0 if rec_n < n else -1.0 if rec_n > n else 0.0)}
423         for o in ("cpufreq", "gpu", "power"):
424             env = auth["organs"].get(o, {}).get("envelope")
425             if env:
426                 levels[f"{o}_ceiling"] = env[1]
427             forced = obs["queue_ratio"] > 0.0 or s["pending"] > 0
428             self.audit({"compass": self.compass.read(d["state"]["E"], d["state"]["S"], levels, moves, forced, x=self.
g.x, p=self.g.p)})
429             out = self.audit({"decision": {"nodes_observed": n, "nodes_recommended": rec_n, "law": "closure" if self.
cl is not None else "governor", "hpa_target_recommended": round(rho, 3),
430                                     "E": d["state"]["E"], "U": d["state"]["U"], "pending": s["pending"],
431                                     "power_cap": round(float(d["power_cap"]), 3), "change_permitted": bool(d[

```

```

"change_permitted"]],
432         "rollback_authorized": bool(d["rollback_authorized"]),
433         "thermal": round(obs["thermal"], 3), "security_block": obs[
"security_block"],
434         "power_stress": round(power_stress, 3), "latency_p95_ms": p95,
435         "queue_ratio": round(obs["queue_ratio"], 3), "slo_clean": obs[
"slo_clean"]}, "mode": self.a.mode))
436     if self.a.mode in ("target", "nodepool") and getattr(self.a, "strict_replicas", False):
437         self.strict_step(s, obs)
438         self.m.push(d, obs)
439     elif self.a.mode in ("target", "nodepool"):
440         for h in s["hpas"]:
441             idx, cur = cpu_target(h)
442             if cur is None:
443                 continue
444             ns, name = h["metadata"]["namespace"], h["metadata"]["name"]
445             orig = int(h["metadata"].get("annotations", {}).get(ANNOTATION, cur))
446             # more headroom is always allowed; less never: at a given load fewer pods always means a longer
M/M/c wait
447             # (no target above the operator's keeps the wait), so a raise only spends latency.
448             # Omni on top earns its keep on machines, not by packing the operator's pods tighter
449             # the same promise in queue terms: busy = target x request / limit. While convey() gives the
pods a
450             # limit g times the operator's, the target that keeps each pod exactly as busy is g times higher
451             want = int(round(100 * min(rho, orig / 100.0) * self._gain(h)))
452             want_h = want if obs["slo_clean"] else min(want, orig)
453             if abs(cur - want_h) < self.a.min_target_change and not (not obs["slo_clean"] and cur > orig):
454                 continue
455             if cur == want_h:
456                 continue
457             # the muscle's own clock: the autoscaler takes its scale-down window (300 s unless the operator
set one)
458             # to answer a target. A new target inside that window moves the muscle mid-movement and starts
pods it
459             # then removes, so I hold each target for one window. A response-time breach returns the
operator's
460             # target at once
461             win = float(((h["spec"].get("behavior") or {}).get("scaleDown") or {}).get(
"stabilizationWindowSeconds", 300))
462             last = self.target_at.get((ns, name))
463             if obs["slo_clean"] and last is not None and time.time() - last < win:
464                 continue
465             self.target_at[(ns, name)] = time.time()
466             self.changed.setdefault((ns, name), int(h["metadata"].get("annotations", {}).get(ANNOTATION,
cur)))
467             self.k.write(["annotate", "hpa", name, "-n", ns, "--overwrite", f"{ANNOTATION}={self.changed[(ns,
name)]}"], "record original target")
468             self.k.write(["patch", "hpa", name, "-n", ns, "--type=json", "-p",
469             json.dumps({"op": "replace", "path": f"/spec/metrics/{idx}/resource/target/
averageUtilization", "value": want_h})],
470             f"HPA target to rho* = {want_h}% + (" if obs["slo_clean"] else " (SLO reflex: not
tighter than native)")
471             self.cmd_hpa[(ns, name)] = want_h
472             self.m.push(d, obs)
473             if self.a.mode in ("target", "nodepool"):
474                 self.pod_reflex()
475             if self.a.mode == "nodepool" and self.a.node_scale_cmd and rec_n != n:
476                 cfg = SimpleNamespace(minimum_nodes=self.a.min_nodes, maximum_nodes=self.a.max_nodes)
477                 acts, hits = enforce([{"action": "nodes", "target": rec_n, "direction": 1 if rec_n > n else -1}],
478                 {"actual_nodes": n, "power_cap": 1.0}, {"power_stress": power_stress,
"security_block": obs["security_block"]}, cfg,
479                 ShieldLimits(power_limit=1e9, max_node_step=max(self.a.max_node_step, floor -
n)))
480             for act in acts:
481                 cmd = self.a.node_scale_cmd.format(n=int(act["target"]))
482                 self.audit({"write": shlex.split(cmd), "why": "node pool size", "dry_run": self.a.dry_run,
"shield_interventions": hits})
483                 if not self.a.dry_run:
484                     subprocess.run(shlex.split(cmd), check=True)
485                     self.cmd_nodes = int(act["target"])
486                 self.rec_n = rec_n
487                 return out
488
489 def parser():
490     ap = argparse.ArgumentParser(description=__doc__, formatter_class=argparse.RawDescriptionHelpFormatter)
491     ap.add_argument("--mode", choices=["observe", "target", "nodepool"], default="observe")
492     ap.add_argument("--interval", type=float, default=60.0, help="seconds between governor decisions")
493     ap.add_argument("--floor-interval", type=float, default=15.0, help="seconds between scheduling-floor checks (
nodepool mode)")
494     ap.add_argument("--iterations", type=int, default=0, help="0 = run until stopped")
495     ap.add_argument("--dry-run", action="store_true")
496     ap.add_argument("--kubectll", default="kubectll")
497     ap.add_argument("--audit", default="omni_audit.jsonl")
498     ap.add_argument("--kill-file", default="/tmp/omni.kill")
499     ap.add_argument("--node-scale-cmd", default="")
500

```

```

501     ap.add_argument("--min-nodes", type=int, default=1)
502     ap.add_argument("--max-nodes", type=int, default=1000)
503     ap.add_argument("--max-node-step", type=int, default=2)
504     ap.add_argument("--min-target-change", type=int, default=3)
505     ap.add_argument("--closure", default="", help="JSON with the closure-law setting (e.g. tuning/
GLOBAL_LEAGUE_PREREGISTRATION.json): the benchmarked law decides the node count")
506     ap.add_argument("--strict-replicas", action="store_true", help="strict C: Omni-Compass sets replica counts;
the HPA is pinned")
507     ap.add_argument("--pod-reflex-writes", action="store_true",
508                     help="let the pod reflex raise the HPA's replica floor itself (default: it gauges and writes
nothing)")
509     ap.add_argument("--reflex-window-s", type=float, default=30.0, help="window of probe samples the fast pod
reflex reads")
510     ap.add_argument("--strict-window", type=int, default=5, help="decisions a scale-down waits for (highest
recent recommendation)")
511     ap.add_argument("--headroom", type=float, default=0.5, help="spare capacity kept above pod requests (0.5 =
50%, the default)")
512     ap.add_argument("--active-nodes-only", action="store_true",
513                     help="count only nodes in service (schedulable, or cordoned but still carrying work) and the
pods and usage on them")
514     ap.add_argument("--node-restore-cmd", default="", help="command run once when the kill switch fires,
returning the node pool to native")
515     add_muscle_args(ap)
516     ap.add_argument("--power-cmd", default="")
517     ap.add_argument("--site-limit-w", type=float, default=0.0)
518     return ap
519
520
521 def safe_step(c, fails):
522     """One decision. A failed decision is recorded and skipped: it writes nothing, so the cluster keeps the last
settings
523     that landed, and the next decision comes at the normal cadence. There is no automated fallback (manuscript
Section
524     5.8: the Unified Control Switch is mechanical, explicit, operator-controlled and 'does not rely on automated
fallback
inference'). Only a human flips the switch, for the whole harness at once: --kill-file present (or
OMNI_KILL=1) turns
526     Omni-Compass OFF and hands every setting back to native; removing it turns Omni-Compass back ON. Boundaries
are never
527     handled by switching anything off: the living band and the shield clamp every level inside its range."""
528     try:
529         c.step()
530         return 0
531     except Exception as e:
532         fails += 1
533         err = getattr(e, "stderr", "") or ""
534         c.audit({"error": repr(e)[:500], "stderr": str(err)[-500:], "consecutive_failures": fails})
535         print(f"decision failed ({fails} in a row): {e!r} {err}", file=sys.stderr, flush=True)
536         return fails
537
538
539 def main(argv=None):
540     a = parser().parse_args(argv)
541     c = Controller(a); i = 0; fails = 0
542     import atexit, resource
543     t0 = time.time()
544     def overhead(): # the controller's own cost: CPU seconds of this process and every kubectrl/script it ran
545         me, kids = resource.getrusage(resource.RUSAGE_SELF), resource.getrusage(resource.RUSAGE_CHILDREN)
546         cpu = me.ru_utime + me.ru_stime + kids.ru_utime + kids.ru_stime
547         c.audit({"overhead": {"cpu_s": round(cpu, 2), "wall_s": round(time.time() - t0, 1),
548                             "cores_mean": round(cpu / max(1e-9, time.time() - t0), 4)}})
549     atexit.register(overhead)
550     while a.iterations == 0 or i < a.iterations:
551         fails = safe_step(c, fails); i += 1
552         if a.iterations == 0 or i < a.iterations:
553             waited = 0.0
554             while waited + 1e-9 < a.interval:
555                 dt = min(a.floor_interval, a.interval - waited) if a.mode == "nodepool" else a.interval - waited
556                 time.sleep(dt); waited += dt
557                 if a.mode == "nodepool" and waited + 1e-9 < a.interval:
558                     try:
559                         c.floor_step()
560                     except Exception as e:
561                         c.audit({"error": "floor check: " + repr(e)[:500], "stderr": str(getattr(e, "stderr",
"")) or ""}[-500:]})
562
563
564 if __name__ == "__main__":
565     main()
566

```

omni_controller/gpu_governor.py

506 lines SHA-256 c3a8bfe2f6f06970da0ab7b5ace32c77a8a6cb5fd2d87bdfc88e88e992d26c38

```
1  """Omni-Compass on one GPU box, no Kubernetes: the engine reads each GPU's own meter and may lower its power
   limit.
2
3  Every decision (--interval seconds), for each GPU in --gpus:
4      sense      nvidia-smi: power.draw, temperature.gpu, utilization.gpu, power.limit (the device's own readings)
5                  the workload's response-time file (--latency-file, elapsed_seconds,latency_ms,ok) if given
6      engine      the Omni-Compass governor, throughput law (omnicompass/adaptor.py): load = utilization, power stress
   =
7                  draw / the limit found at start, heat = temperature / --temp-limit, queue = response-time pressure.
8                  Its directive power_cap is the share of the start limit the GPU may draw.
9      shield      limit = max(cap x start limit, draw x (1 + --headroom), --min-share x start limit, the device's
   minimum
10                 limit), never above the start limit; whole watts; a change under --min-change-w is not written.
11                 --min-share bounds the slowdown: a card held at that share of its limit runs a burst only a little
   slower.
12     busy gate    utilization (smoothed over decisions) at or above --util-gate: the start limit. A busy card is the
13                 bottleneck; slowing it grows the queue faster than it saves energy. The cap returns once utilization
   is
14                 back under the gate less --util-band.
15     reflexes     a blind sense (nvidia-smi unreadable, or the response-time file stale or empty): the start limit, at
   once
16                 a response-time breach (p95 > --slo-ms), and for --slo-clear decisions after it: the start limit, at
   once
17     read-back    no new write until the last one is read back from the device (power.limit within 1 W of what I wrote)
18     enforced     the card obeys enforced.power.limit, which another authority (a lower board or system limit) can
   hold under
19                 the limit I request. The engine senses the enforced limit; each write is read back at once (
   requested,
20                 return code, power.limit, enforced limit, delay), and enforced under requested is recorded as an
   override.
21                 A driver without enforced.power.limit: power.limit is used and the snapshot says so.
22     speed lock    (--baseline-file, from tools/gpu_baseline.py on a run without Omni): the limit is set by response
   time
23                 alone. Over the last --lock-window-s of the response-time file, the mean, 95th and 99th percentile
   are each
24                 divided by the baseline's at the same arrival rate; the worst of the three is the speed ratio. The
   line is
25                 1 - --speed-gain (0.99: at least 1% faster than without Omni). Above the line: the start limit, at
   once.
26                 Under the line less --lock-margin: lower by --lock-step x start, times the slack in margins (up to
27                 --lock-boost), at most once per --lock-hold-s (the window must see the step before the next).
   Between: held. Never under --lock-floor x start.
28                 Speed won elsewhere (CPU conveyed to the serving pods, shorter queues) shows as a ratio under the
   line,
29                 and the lock spends it on watts; with nothing won elsewhere it holds the start limit.
30     one writer    the limit must be what I last wrote (or the start limit). Any other value means another writer: I
   stop
31                 writing for the rest of the run (observe only), record "foreign_writer", leave its limit alone on
   kill
32                 (restoring would fight it), and exit 5 so the bench marks the run invalid. Two controllers on one
   limit is
33                 how production is lost.
34     heat          the device reporting a thermal or hardware slowdown (clock-limit reasons 0x8, 0x20, 0x40, 0x80):
   never
35                 tighten; the limit may only stay or rise (fail up).
36     envelope      every decision records every rule that held the engine's request back ("blocked_by") and the one that
37                 decided ("decided_by": engine, a floor, the busy gate, a reflex, heat, or the speed lock), so a
   limit that
38                 did not move names its cause, and a result is credited to the rule that produced it.
39 Modes
40     watch         everything above is computed and audited; nothing is written (the control arm)
41     cap           the limit is written with nvidia-smi -i <gpu> -pl <w>. Refused at start unless every GPU reports
42                 power.management Enabled. A write the device refuses (nonzero return) ends the run: recorded, the
   start
43                 limit restored, exit 4, and the bench marks the arm invalid. No clock locks are ever written.
44 Kill switch
45     the kill file (or SIGTERM) restores every GPU to the limit read at start, reads it back, and exits.
46     The start limits are recorded first in the audit ("snapshot": power.limit, enforced.power.limit, default, min,
47     max, persistence mode, power management, per GPU).
48 """
49 from __future__ import annotations
50
51 import argparse, copy, json, math, os, shlex, signal, subprocess, sys, time
52 from pathlib import Path
53
54 ROOT = Path(__file__).resolve().parents[1]
55 sys.path.insert(0, str(ROOT))
56 from omnicompass.adapter import Governor, mode_law, AUTOPILOT, OBSERVE, observe_vector, assimilate, ASSIMILATION
57 from omnicompass.core import U_AUTHORITY, State
58 from omni_controller.muscles import latency_sense
59
60 FIELDS = "index,power.draw,temperature.gpu,utilization.gpu,power.limit,power.min_limit,clocks.sm"
61 ENFORCED = "enforced.power.limit"
```

```

62 SNAPSHOT = ("power.limit", "enforced.power.limit", "power.default_limit", "power.min_limit", "power.max_limit",
63             "persistence_mode", "power.management")
64 REASONS = ("clocks_event_reasons.active", "clocks_throttle_reasons.active") # newer, older drivers
65 HEAT_BITS = 0x8 | 0x20 | 0x40 | 0x80 # HW slowdown, SW thermal slowdown, HW thermal slowdown, HW power brake
66
67
68 def slowed(reason):
69     """True when the device reports a thermal or hardware slowdown in its clock-limit reasons (hex bit mask)."""
70     try:
71         return bool(int(str(reason), 16) & HEAT_BITS)
72     except (TypeError, ValueError):
73         return False
74
75
76 def blocked_by(requested, draw, start, min_w, headroom, min_share, slo_clean, gated, heat, lock):
77     """Every rule whose floor or override sits above the engine's own request (whole watts), in a fixed order."""
78     engine_w = math.floor(requested * start)
79     b = []
80     if not slo_clean:
81         b.append("slo_or_blind_reflex")
82     if gated:
83         b.append("busy_gate")
84     if math.ceil(draw * (1.0 + headroom)) > engine_w:
85         b.append("draw_headroom_floor")
86     if min_share * start > engine_w:
87         b.append("share_floor")
88     if min_w > engine_w:
89         b.append("device_minimum")
90     if heat:
91         b.append("thermal_hold")
92     if lock:
93         b.append("speed_lock")
94     return b
95 STATE = ("E", "U", "I_U", "S", "B", "B_dot")
96 MEASURED = ("E", "U", "I_U", "S", "B") # the memoryless part of the observation map (B_dot exists only through
    history)
97
98
99 def measured_state(o, p):
100     """h without memory: the state the telemetry alone points to (the targets assimilate() blends toward).
    Computed by
101     assimilate() itself from a zero prior, so it is the engine's own map, not a second one."""
102     z = assimilate(State(0.0, 0.0, 0.0, 0.0, 0.0, 0.0), o, copy.deepcopy(p))
103     return {k: getattr(z, k) / ASSIMILATION for k in MEASURED}
104
105
106 def window_stats(path, window_s, now=None):
107     """Mean, 95th and 99th percentile (ms) and arrival rate (served per second) of the successful requests in
    the last
108     window_s of the response-time file; None when blind (stale beyond two windows, or under 20 requests)."""
109     try:
110         age = (now if now is not None else time.time()) - os.path.getmtime(path)
111         rows = [l.split(",") for l in Path(path).read_text().splitlines()[1:] if l.strip()]
112     except OSError:
113         return None
114     if not rows or age > 2.0 * window_s:
115         return None
116     t_end = float(rows[-1][0])
117     ms = sorted(float(r[1]) for r in rows if float(r[0]) >= t_end - window_s and r[2].strip() == "1")
118     if len(ms) < 20:
119         return None
120     span = min(window_s, max(1.0, t_end - float(rows[0][0])))
121     return {"mean": sum(ms) / len(ms), "p95": ms[min(len(ms) - 1, int(0.95 * len(ms)))],
122           "p99": ms[min(len(ms) - 1, int(0.99 * len(ms)))], "rate": len(ms) / span, "n": len(ms)}
123
124
125 def baseline_at(base, rate):
126     """The baseline's mean, p95 and p99 at an arrival rate, linear between its rate bins, flat beyond the ends.
    """
127     bins = base["bins"]
128     if rate <= bins[0]["rate"]:
129         return bins[0]
130     if rate >= bins[-1]["rate"]:
131         return bins[-1]
132     for lo, hi in zip(bins, bins[1:]):
133         if lo["rate"] <= rate <= hi["rate"]:
134             f = (rate - lo["rate"]) / (hi["rate"] - lo["rate"])
135             return {k: lo[k] + f * (hi[k] - lo[k]) for k in ("mean", "p95", "p99")}
136
137
138 def shield_limit(cap, draw, start, min_w, headroom, min_share, slo_clean):
139     """The shield: limit = max(cap x start, draw x (1 + headroom), min_share x start, device minimum), never
    above start,
140     whole watts. Returns (want_w, the bound that decided it). Twinned in C++ (cpp/src/gpu_rules.cpp)."""
141     share_floor = min_share * start
142     floor = max(draw * (1.0 + headroom), min_w, share_floor)

```

```

143     engine_w = math.floor(cap * start)
144     want = int(min(start, max(engine_w, math.ceil(floor))))
145     bound = ("slo_or_blind_reflex" if not slo_clean else "start_ceiling" if want >= start and engine_w >= start
146             else "share_floor" if math.ceil(floor) > engine_w and floor == share_floor
147             else "draw_headroom_floor" if math.ceil(floor) > engine_w and floor > min_w
148             else "device_minimum" if math.ceil(floor) > engine_w else "engine")
149     return want, bound
150
151
152 def busy_gate(u_prev, util, gated_prev, gate, band):
153     """Utilization smoothed over decisions (u_prev None: the first reading) and the gate with its band. Returns
154     (u, gated). Twinned in C++."""
155     u = 0.5 * (util if u_prev is None else u_prev) + 0.5 * util
156     return u, (u >= gate or (gated_prev and u >= gate - band))
157
158
159 def lock_decide(ratio, cur, start, min_w, since_last, speed_gain, lock_margin, lock_step, lock_boost, lock_floor,
160               lock_hold_s):
161     """The speed lock's step from the worst response-time ratio against the baseline. Returns (want_w, bound,
162     line, aim). Twinned in C++."""
163     line = 1.0 - speed_gain
164     aim = line * (1.0 - lock_margin)
165     floor = max(min_w, lock_floor * start)
166     if ratio > line:
167         want, bound = start, "speed_lock_release"
168     elif ratio < aim and since_last >= lock_hold_s:
169         mult = min(lock_boost, max(1.0, (aim - ratio) / max(1e-6, line * lock_margin)))
170         want, bound = max(floor, cur - mult * lock_step * start), "speed_lock_spend"
171     else:
172         want, bound = cur, "speed_lock_hold"
173     return int(min(start, max(floor, want))), bound, line, aim
174
175 def throttle(smi, gpus):
176     """The device's own record of why its clock is held down (power cap, thermal, ...), as a bit mask; None if
177     the
178     driver does not report it."""
179     for f in REASONS:
180         try:
181             out = subprocess.run(shlex.split(smi) + [f"--query-gpu=index,{f}", "--format=csv,noheader,nounits"],
182                                capture_output=True, text=True, timeout=10, check=True).stdout
183             return {int(x.split(",")[0]): x.split(",")[1].strip() for x in out.strip().splitlines()}
184         except (subprocess.SubprocessError, OSError, ValueError, IndexError):
185             continue
186     return None
187
188 class WriteFailed(RuntimeError):
189     """The device refused a power-limit write: the arm ends (the kill path still restores)."""
190
191
192 def snapshot(smi, gpus):
193     """{field: {gpu: value}} for every SNAPSHOT field, each queried alone; "unsupported" where the driver
194     refuses it."""
195     out = {}
196     for f in SNAPSHOT:
197         try:
198             r = subprocess.run(shlex.split(smi) + [f"--query-gpu=index,{f}", "--format=csv,noheader,nounits"],
199                                capture_output=True, text=True, timeout=10, check=True).stdout
200             vals = {int(x.split(",")[0]): x.split(",")[1].strip() for x in r.strip().splitlines()}
201             except (subprocess.SubprocessError, OSError, ValueError, IndexError):
202                 vals = {}
203             out[f] = {str(g): vals.get(g, "unsupported") for g in gpus}
204     return out
205
206 def query(smi, gpus, enforced=False):
207     """{gpu: {"draw", "temp", "util", "limit", "min", "clock_mhz", "enforced"}} from the device, or None if
208     nvidia-smi
209     cannot be read. enforced: also read enforced.power.limit (else "enforced" is power.limit)."""
210     try:
211         fields = FIELDS + ("," + ENFORCED if enforced else "")
212         out = subprocess.run(shlex.split(smi) + [f"--query-gpu={fields}", "--format=csv,noheader,nounits"],
213                             capture_output=True, text=True, timeout=10, check=True).stdout
214         rows = {}
215         for line in out.strip().splitlines():
216             v = [x.strip() for x in line.split(",")]
217             rows[int(v[0])] = {"draw": float(v[1]), "temp": float(v[2]), "util": float(v[3]) / 100.0,
218                              "limit": float(v[4]), "min": float(v[5]), "clock_mhz": float(v[6]),
219                              "enforced": float(v[7]) if enforced else float(v[4])}
220         return {g: rows[g] for g in gpus} if all(g in rows for g in gpus) else None
221     except (subprocess.SubprocessError, OSError, ValueError, IndexError):
222         return None
223

```

```

224 class GpuGovernor:
225     def __init__(self, a):
226         self.a = a
227         self.gpus = [int(g) for g in str(a.gpus).split(",") if g.strip() != ""]
228         self.log = open(a.audit, "a")
229         self.eng = {g: Governor(law=mode_law("throughput")) for g in self.gpus}
230         for e in self.eng.values():
231             e.set_mode(AUTOPILOT if a.mode == "cap" else OBSERVE)
232             e.nodes = 1
233         self.cap = {g: 1.0 for g in self.gpus}
234         self.written = {} # gpu -> last limit written (W), until read back
235         self.lp_hist = []
236         self.writes = 0
237         self.projected = {}
238         self.reasons_ok = True # False once the driver shows it does not report clock-limit reasons
239         self.util_avg = {} # gpu -> utilization smoothed over decisions (busy gate)
240         self.gated = {} # gpu -> True while the busy gate holds the start limit
241         bf = getattr(a, "baseline_file", "")
242         self.baseline = json.loads(Path(bf).read_text()) if bf else None # speed lock: the run without Omni
243         self.lock_at = {} # gpu -> time of the lock's last change
244         self.clock = time.time # replaceable, so a simulation in virtual time can drive the lock's hold
245         self.req_at = {} # gpu -> (limit requested, time requested) until it reads back
246         self.foreign = False # True once another writer changed a limit: observe only from then on
247         self.enforced_ok = query(a.smi, self.gpus, enforced=True) is not None
248         s = query(a.smi, self.gpus, self.enforced_ok)
249         if s is None:
250             raise SystemExit("nvidia-smi unreadable at start: no snapshot, so I take no authority")
251         snap = snapshot(a.smi, self.gpus)
252         self.start = {g: s[g]["limit"] for g in self.gpus}
253         self.expect = dict(self.start) # gpu -> the limit that should be there: the start, or my last write
254         that landed
255         self.min = {g: s[g]["min"] for g in self.gpus}
256         self.audit({{"snapshot": {str(g): {"limit_w": self.start[g], "min_limit_w": self.min[g],
257             "enforced_w": s[g]["enforced"] if self.enforced_ok else None,
258             **{f: snap[f][str(g)] for f in SNAPSHOT}} for g in self.gpus},
259             "senses": ENFORCED if self.enforced_ok else "power.limit (the driver does not report
260             enforced.power.limit)",
261             "mode": a.mode}))
262         off = {g for g in self.gpus if snap["power.management"][str(g)].lower() != "enabled"}
263         if a.mode == "cap" and off:
264             self.audit({"refused": f"power management not Enabled on GPU {off}: a written limit would not bind"})
265             raise SystemExit(f"power management not Enabled on GPU {off}: cap mode refused, no authority taken")
266
267     def audit(self, rec):
268         self.log.write(json.dumps({"time": time.time(), **rec}) + "\n"); self.log.flush()
269
270     def set_limit(self, g, w, why):
271         cmd = ["nvidia-smi", "-i", str(g), "-pl", str(int(w))]
272         if self.a.mode != "cap" or self.foreign:
273             # watch: computed and recorded, never executed; after another writer appeared, likewise
274             self.audit({"would_write": cmd, "why": why, **({"withheld": "another writer owns the limit"} if self.
275             foreign else {})})
276             return
277         self.audit({"write": cmd, "why": why})
278         t0 = time.time()
279         try:
280             p = subprocess.run(shlex.split(self.a.smi) + ["-i", str(g), "-pl", str(int(w))], capture_output=True,
281             text=True, timeout=20)
282             rc, err = p.returncode, p.stderr.strip()[:300]
283             except (subprocess.SubprocessError, OSError) as e:
284                 rc, err = -1, f"{type(e).__name__}: {e}[:300]"
285             act = {"gpu": g, "requested_w": int(w), "t_requested": t0, "rc": rc, "stderr": err}
286             if rc != 0:
287                 self.audit({"actuator": act, "write_failed": cmd})
288                 raise WriteFailed(f"GPU {g}: nvidia-smi -pl {int(w)} returned {rc}: {err}")
289             self.writes += 1
290             self.written[g] = int(w)
291             s = query(self.a.smi, [g], self.enforced_ok)
292             if s is not None:
293                 back = s[g]
294                 act.update({"readback_w": back["limit"], "enforced_w": back["enforced"], "t_readback": time.time(),
295                 "realized": abs(back["limit"] - w) < 1.0, "override": back["enforced"] < back["limit"] -
296                 1.0})
297                 if act["realized"]:
298                     act["delay_s"] = round(act["t_readback"] - t0, 3); self.expect[g] = int(w)
299                 if not act.get("realized"):
300                     self.req_at[g] = (int(w), t0)
301                 self.audit({"actuator": act})
302
303     def killed(self):
304         return os.path.exists(self.a.kill_file)
305
306     def restore(self):
307         """Kill switch: every GPU back to the limit read at start, read back from the device."""
308         s = query(self.a.smi, self.gpus, self.enforced_ok) or {}
309         failed = []

```

```

305         for g in self.gpus:
306             if self.foreign:
307                 break # another writer owns the limit: restoring would fight it
308             if self.a.mode == "cap" and (g not in s or abs(s[g]["limit"] - self.start[g]) >= 1.0):
309                 try:
310                     self.set_limit(g, self.start[g], "kill switch: the limit read at start")
311                 except WriteFailed as e:
312                     failed.append(str(e))
313             s = query(self.a.smi, self.gpus, self.enforced_ok) or {}
314             back = {str(g): (s[g]["limit"] if g in s else None) for g in self.gpus}
315             ok = not failed and all(v is not None and abs(v - self.start[int(g)]) < 1.0 for g, v in back.items())
316             self.audit({"restored": back, "enforced": {str(g): (s[g]["enforced"] if g in s else None) for g in self.
gpus}},
317                        "start": {str(g): self.start[g] for g in self.gpus}, "ok": ok, "writes": self.writes,
318                        "foreign_writer": self.foreign,
319                        "restore_write_failed": failed))
320         return ok
321
322     def speed_lock(self, g, cur, slo_clean):
323         """The limit from response time against the run without Omni (see the module docstring). Returns
324         (want_w, bound, why, record)."""
325         a, start = self.a, self.start[g]
326         st = window_stats(a.latency_file, a.lock_window_s) if a.latency_file else None
327         if st is None or not slo_clean:
328             return int(start), "speed_lock_blind", "speed lock: no response-time window (or SLO reflex), the
start limit", None
329         b = baseline_at(self.baseline, st["rate"])
330         ratios = {k: st[k] / b[k] for k in ("mean", "p95", "p99")}
331         ratio = max(ratios.values())
332         want, bound, line, aim = lock_decide(ratio, cur, start, self.min[g], self.clock() - self.lock_at.get(g, -
1e18),
333                                             a.speed_gain, a.lock_margin, a.lock_step, a.lock_boost, a.
lock_floor, a.lock_hold_s)
334         if want != int(cur):
335             self.lock_at[g] = self.clock()
336             rec = {"rate": round(st["rate"], 2), "ratios": {k: round(v, 4) for k, v in ratios.items()},
337                  "line": line, "aim": round(aim, 4), "n": st["n"]}
338             why = f"speed lock: worst ratio {ratio:.3f} vs line {line:.3f} ({bound.split('_')[-1]}), limit {want} W"
339             return want, bound, why, rec
340
341     def step(self):
342         a = self.a
343         s = query(a.smi, self.gpus, self.enforced_ok)
344         blind = s is None
345         lp, p95, served = 0.0, None, None
346         if a.latency_file and a.slo_ms:
347             ls = latency_sense(a.latency_file, a.latency_window_s)
348             served = ls["ok"]
349             blind = blind or ls["blind"]
350             if not ls["blind"]:
351                 p95 = ls["p95"]
352                 lp = max(0.0, ls["p95"] / a.slo_ms - 1.0)
353                 if ls["fail"]:
354                     lp = max(lp, ls["fail"] / (ls["ok"] + ls["fail"]))
355             self.lp_hist.append(lp)
356             slo_clean = not blind and len(self.lp_hist) >= a.slo_clear and all(x == 0.0 for x in self.lp_hist[-a.
slo_clear:])
357             thr = throttle(a.smi, self.gpus) if s is not None and self.reasons_ok else None
358             self.reasons_ok = self.reasons_ok and (s is None or thr is not None)
359             rec = {"decision": {}, "blind": blind, "served_in_window": served, "latency_p95_ms": p95,
"latency_pressure": round(lp, 3), "slo_clean": slo_clean}
360             for g in self.gpus:
361                 if s is None:
362                     # blind: no give-back of power I cannot see; the start limit, at once
363                     if self.cap[g] == 1.0 or self.foreign:
364                         continue
365                     want, why, cur = self.start[g], "blind sense: the limit read at start", None
366                     self.cap[g] = 1.0
367                 else:
368                     r = s[g]; cur = r["limit"]
369                     if abs(cur - self.expect[g]) >= 1.0 and not (g in self.written and abs(cur - self.written[g]) <
1.0):
370                         if g not in self.written or not self.foreign:
371                             # neither what I last wrote nor what was there before: another writer
372                             if not self.foreign:
373                                 self.audit({"foreign_writer": {"gpu": g, "reads_w": cur, "expected_w": self.expect[
g],
374                                     "pending_w": self.written.get(g)},
375                                             "action": "observe only from now on; the limit left to the other
writer"})
376                                 self.foreign = True; self.expect[g] = cur; self.written.pop(g, None)
377                             if g in self.written and abs(cur - self.written[g]) >= 1.0:
378                                 # read-back: my last write has not landed; no new order on top of it
379                                 rec["decision"][str(g)] = {"hold": "last write not read back", "wrote_w": self.written[g],
"reads_w": cur}
380                                 continue

```

```

381         if g in self.written:
382             self.expect[g] = int(self.written.pop(g))
383         if g in self.req_at: # a write that had not read back at once has now landed
384             w0, t0 = self.req_at.pop(g)
385             self.audit({"actuator_realized": {"gpu": g, "requested_w": w0, "readback_w": cur,
386                 "enforced_w": r["enforced"], "delay_s": round(time.time() -
t0, 3)}}))
387         # the limit the card obeys (enforced), not the one I meant
388         e = self.eng[g]; e.current_cap = min(1.0, r["enforced"] / self.start[g])
389         obs = {"queue_ratio": min(2.0, lp), "load_ratio": r["util"], "power_stress": r["draw"] / self.
start[g],
390             "thermal": r["temp"] / a.temp_limit, "network_stress": 0.0, "drift_ratio": 0.0,
391             "stale": 0.0, "security_block": 0.0}
392         # the chain, recorded whole: what the device said, the state it puts the engine in, what I had
393         projected
394         allowed
395         seen = assimilate(e.x, observe_vector(obs, 0), copy.deepcopy(e.p))
396         meas = measured_state(observe_vector(obs, 0), e.p)
397         prev = self.projected.get(g)
398         d = e.step(obs, 0)
399         self.projected[g] = {k: getattr(e.x, k) for k in STATE}
400         requested = float(d["power_cap"])
401         cap = requested if slo_clean else 1.0
402         want, bound = shield_limit(cap, r["draw"], self.start[g], self.min[g], a.headroom, getattr(a,
"min_share", 0.0), slo_clean)
403         gate = getattr(a, "util_gate", 0.0)
404         if gate:
405             self.util_avg[g], self.gated[g] = busy_gate(self.util_avg.get(g), r["util"], self.gated.get(
g, False),
406                 gate, getattr(a, "util_band", 0.1))
407         if self.gated[g]:
408             want, bound = int(self.start[g]), "busy_gate"
409             cap = 1.0
410             why = ("response-time reflex: the limit read at start" if not slo_clean
411                 else f"busy gate: utilization {self.util_avg[g]:.2f}, the limit read at start" if bound
412                 else f"engine cap {cap:.3f} of {self.start[g]:.0f} W; floor draw {r['draw']:.0f} W x {1 +
a.headroom:.2f}, "
413                     f"share floor {getattr(a, 'min_share', 0.0):.2f}")
414             lock = None
415             if self.baseline is not None:
416                 want, bound, why, lock = self.speed_lock(g, cur, slo_clean)
417             heat = slowed((thr or {}).get(g))
418             if heat and want < cur:
419                 # the device is slowing itself for heat or a hardware limit: never tighten on top of it
420                 want, bound, why = int(cur), "thermal_hold", "thermal or hardware slowdown reported: no
tightening"
421             blocks = blocked_by(requested, r["draw"], self.start[g], self.min[g], a.headroom, getattr(a,
"min_share", 0.0),
422                 slo_clean, self.gated.get(g, False) if gate else False, heat, self.baseline
423                 is not None)
424             self.cap[g] = want / self.start[g]
425             rec["decision"][str(g)] = {
426                 "telemetry": {"util": r["util"], "draw_w": r["draw"], "temp_c": r["temp"], "limit_w": cur,
427                     "enforced_w": r["enforced"],
428                     "clock_mhz": r["clock_mhz"], "throttle": (thr or {}).get(g)},
429                 "state_observed": {k: round(getattr(seen, k), 4) for k in STATE},
430                 "state_measured": {k: round(v, 4) for k, v in meas.items()},
431                 "state_projected_before": None if prev is None else {k: round(v, 4) for k, v in prev.items(
432                     )},
433                 "prediction_error": None if prev is None else {k: round(getattr(seen, k) - prev[k], 4) for k
434                     in STATE},
435                 "state_projected_next": {k: round(v, 4) for k, v in self.projected[g].items()},
436                 "u_push": round(e.last_push * U_AUTHORITY, 4), # the U-channel command on the evolved state
437                 "admissible": bool(d["change_permitted"]), "requested_cap": round(requested, 4),
438                 "granted_cap": round(cap, 4), "shield_bound": bound, "want_w": want,
439                 "engine_w": math.floor(requested * self.start[g]), "blocked_by": blocks,
440                 "decided_by": "speed_lock" if bound.startswith("speed_lock") else bound,
441                 "engine_cap": round(requested, 3), "E": round(d["state"]["E"], 3), "U": round(d["state"]
442                     ["U"], 3),
443                 "speed_lock": lock}
444             if cur is not None and abs(want - cur) < a.min_change_w and not (want == self.start[g] and cur !=
445                 want):
446                 continue
447             if cur is not None and want == int(cur):
448                 continue
449             self.set_limit(g, want, why)
450             self.audit(rec)
451
452 def parser():
453     ap = argparse.ArgumentParser(description=__doc__, formatter_class=argparse.RawDescriptionHelpFormatter)
454     ap.add_argument("--mode", choices=["watch", "cap"], default="watch")
455     ap.add_argument("--gpus", default="0", help="GPU indexes, comma separated")
456     ap.add_argument("--smi", default=os.environ.get("NVIDIA_SMI", "nvidia-smi"))

```

```

452     ap.add_argument("--interval", type=float, default=2.0)
453     ap.add_argument("--duration", type=float, default=0.0, help="seconds to run (0: until killed)")
454     ap.add_argument("--audit", default="gpu_audit.jsonl")
455     ap.add_argument("--kill-file", default="/tmp/omni-gpu-kill")
456     ap.add_argument("--headroom", type=float, default=0.3, help="limit never below draw x (1 + headroom)")
457     ap.add_argument("--min-share", type=float, default=0.70, help="limit never below this share of the start
limit (0: off)")
458     ap.add_argument("--util-gate", type=float, default=0.5, help="smoothed utilization at which the start limit
returns (0: off)")
459     ap.add_argument("--util-band", type=float, default=0.1, help="the cap resumes below --util-gate less this
band")
460     ap.add_argument("--baseline-file", default="", help="speed lock: baseline from tools/gpu_baseline.py (off
when empty)")
461     ap.add_argument("--speed-gain", type=float, default=0.01, help="speed lock line: at least this much faster
than baseline")
462     ap.add_argument("--lock-boost", type=float, default=5.0, help="largest multiple of --lock-step in one step
down")
463     ap.add_argument("--lock-margin", type=float, default=0.08, help="spend watts only while this far under the
line")
464     ap.add_argument("--lock-step", type=float, default=0.02, help="share of the start limit given up per
decision")
465     ap.add_argument("--lock-hold-s", type=float, default=10.0, help="seconds between two steps down")
466     ap.add_argument("--lock-window-s", type=float, default=60.0, help="response-time window the lock reads")
467     ap.add_argument("--lock-floor", type=float, default=0.5, help="the lock never holds the limit under this
share of start")
468     ap.add_argument("--min-change-w", type=float, default=5.0)
469     ap.add_argument("--temp-limit", type=float, default=83.0, help="GPU temperature read as heat 1.0")
470     ap.add_argument("--latency-file", default="")
471     ap.add_argument("--slo-ms", type=float, default=0.0)
472     ap.add_argument("--latency-window-s", type=float, default=30.0)
473     ap.add_argument("--slo-clear", type=int, default=3)
474     return ap
475
476
477 def main(argv=None):
478     a = parser().parse_args(argv)
479     gov = GpuGovernor(a)
480     stop = {"now": False}
481     signal.signal(signal.SIGTERM, lambda *_: stop.__setitem__("now", True))
482     t0 = time.time()
483     failed = warned = False
484     try:
485         while not stop["now"] and not gov.killed() and (a.duration <= 0 or time.time() - t0 < a.duration):
486             try:
487                 gov.step()
488             except WriteFailed as e: # the actuator refused: the arm ends here, never carries on as if watching
489                 gov.audit({"fatal": f"write failed: {e}"}); failed = True
490                 break
491             except Exception as e: # noqa: BLE001 a failed decision is logged; the kill path still runs
492                 gov.audit({"error": f"{type(e).__name__}: {e}"});
493             if gov.foreign and not warned:
494                 print("omni-gpu: another writer changed the power limit; observing only", file=sys.stderr,
flush=True)
495                 warned = True
496                 end = time.time() + a.interval
497                 while time.time() < end and not stop["now"] and not gov.killed():
498                     time.sleep(0.2)
499         finally:
500             ok = gov.restore()
501             return 5 if gov.foreign else 3 if not ok else 4 if failed else 0
502
503
504 if __name__ == "__main__":
505     sys.exit(main())
506

```

omni_controller/muscles.py

681 lines SHA-256 286302f782935bbc456d9d4b376323a5b38d1b1908bf04b2906b228c6b03a7c4

```
1  """Live muscles for the Kubernetes controller beyond HPA and nodes.
2
3  Each muscle pulls (afferent) and pushes (efferent) through kubectl, records what it changed so the kill switch
4  can hand
5  the muscle back, and tags every write with its muscle name in the audit log.
6
7  power_cap    push: CPU limit of each running pod of the capped deployments = base limit x the governor's power
8  cap,         resized in place (no restart; Linux CFS quota throttles the containers for real); pull: power from
9
10 --power-cmd   pull: thermal state from the harness heat law driven by live power stress (modelled: kind has no
11 heat         thermometer)
12 security     pull: ConfigMap key "hold" == "true" means a security hold; the shield then blocks every expansion
13 rollout      push: pause a deployment's rollout while the governor does not permit change; resume it when it
14 does;
15
16 batch        undo a rollout that has exceeded its progress deadline when the governor authorises rollback
17 is load      push: admit (unsuspend) held batch Jobs, one per decision, only when change is permitted and there
18
19 latency      and power headroom; running Jobs are never suspended
20 written by   pull: 95th-percentile response time over the last --latency-window-s from --latency-file (CSV
21
22 cpu_pstate   scripts/latency_probe.py); pressure max(0, p95 / --slo-ms - 1) enters the engine as queue pressure
23 rapl);       hardware connector. pull: --rapl-cmd prints package watts (e.g. from /sys/class/powercap/intel-
24
25 frequency-set push: --cpufreq-cmd with {khz}, the frequency ceiling = max frequency x power cap (e.g. writing
26 gpu          scaling_max_freq with scripts/cpufreq_ceiling.sh {khz}, see hardware/SCHEDUTIL.md, or `cpupower
27 --query-gpu=power.draw,temperature.gpu --format=csv,noheader,nounits`); GPU temperature / --gpu-
28 temp-limit   --query-gpu=power.draw,temperature.gpu --format=csv,noheader,nounits`); GPU temperature / --gpu-
29
30 rightsize    feeds the heat sense; push: --gpu-power-cmd with {w}, the power limit = max limit x power cap
31 through the  (e.g. `nvidia-smi -pl {w}`); kill restores the maximum limit
32 decrease while push: each running pod's CPU *request* follows its measured use x (1 + headroom), resized in place
33
34 coldstart    pods/resize subresource (no restart); never above the limit, never below --rightsize-min-m; no
35 read) for    the SLO is breached, no increase during a security hold; kill restores every original request
36 moment       push: a deployment with no work waiting (ConfigMap key "queue" == 0, the queue a KEDA scaler would
37
38 batch_pace   --coldstart-idle decisions is scaled to zero; it is scaled back to its recorded replica count the
39 suspended    work is waiting; kill restores the recorded count
40 resumed when push: while power stress or heat is high, one running Job labelled omnicompass.io/pausable=true is
41
42 contain      per decision (checkpointed training pauses instead of the site going over its limit); it is
43 security     stress and heat are back down; kill resumes every Job Omni paused
44
45 place so the push: an agent namespace whose measured CPU use exceeds --contain-cpu-m, or any namespace during a
46
47 the limits   hold, gets a ResourceQuota (limits.cpu = budget) and its running pods' CPU limits are scaled in
48 cooling       total fits the budget; lifted when use is back under budget; kill deletes the quota and restores
49 min-c and    facility connector. push: --cooling-cmd with {c}, the supply-air setpoint moved between --cooling-
50
51 as heat      --cooling-max-c by the heat state (warm setpoint while cool, saving chiller energy; cold setpoint
52
53             rises); kill runs it with --cooling-restore-c
54
55 Hardware connectors are off unless their commands are given (not available on CI runners); the mechanism is
56 the same.
57
58 """
59 from __future__ import annotations
60
61 import csv, json, math, shlex, subprocess, time
62 import re
63
64 CPU_ANN = "omnicompass.io/original-cpu-limit"
65 PAUSE_ANN = "omnicompass.io/paused-by-omni"
66 BATCH_LABEL = "omnicompass.io/batch=true"
67 PAUSABLE_LABEL = "omnicompass.io/pausable=true"
68 REQ_ANN = "omnicompass.io/original-cpu-request"
69 REPL_ANN = "omnicompass.io/original-replicas"
70 PACE_ANN = "omnicompass.io/paced-by-omni"
71 QUOTA = "omni-containment"
72 LIM_ANN = "omnicompass.io/original-limits"
73
74
75 def milli(v):
76     v = str(v)
77     return float(v[:-1]) if v.endswith("m") else float(v) * 1000.0
78
79
80 def ref(s):
```

```

62     ns, name = s.split("/", 1)
63     return ns, name
64
65
66 class Muscles:
67     def __init__(self, k, a, audit):
68         self.k, self.a, self.audit = k, a, audit
69         self.gain = {} # (ns, deployment) -> mean conveyed CPU limit / operator's limit
70         self.base = {} # (ns, deployment) -> operator's CPU limit (millicores)
71         self.thermal = 0.32 # harness initial thermal state
72
73     # ---- afferent -----
74
75     def sense(self, power_stress):
76         o = {}
77         if getattr(self.a, "thermal_model", False):
78             self.thermal = min(1.35, max(0.0, 0.86 * self.thermal + 0.14 * (0.34 + 0.62 * min(1.35,
79 power_stress))))
80             o["thermal"] = self.thermal
81             if getattr(self.a, "rapl_cmd", ""):
82                 o["cpu_watts"] = self._read(self.a.rapl_cmd)
83             if getattr(self.a, "gpu_query_cmd", ""):
84                 out = self._run_out(self.a.gpu_query_cmd)
85                 try:
86                     w, t = [float(x) for x in out.split(",")[:2]]
87                     o["gpu_watts"], o["gpu_celsius"] = w, t
88                     o["thermal"] = max(o.get("thermal", 0.0), min(1.35, t / max(1.0, self.a.gpu_temp_limit)))
89                 except ValueError:
90                     pass
91             lf = getattr(self.a, "latency_file", "")
92             if lf and getattr(self.a, "slo_ms", 0):
93                 ls = latency_sense(lf, self.a.latency_window_s)
94                 o["latency_blind"] = 1.0 if ls["blind"] else 0.0
95                 o["latency_age_s"] = ls["age_s"]
96                 if not ls["blind"]:
97                     o["latency_p95_ms"] = ls["p95"]
98                     o["latency_pressure"] = max(0.0, ls["p95"] / self.a.slo_ms - 1.0)
99                     if ls["fail"]: # failed requests in a live window are real service failures: full
100 pressure share
101                     o["latency_pressure"] = max(o["latency_pressure"], ls["fail"] / (ls["ok"] + ls["fail"]))
102             cm = getattr(self.a, "security_configmap", "")
103             if cm:
104                 ns, name = ref(cm)
105                 try:
106                     data = self.k.get("get", "configmap", name, "-n", ns, "-o", "json").get("data", {}) or {}
107                     o["security_block"] = 1.0 if str(data.get("hold", "")).lower() == "true" else 0.0
108                 except Exception:
109                     o["security_block"] = 0.0
110             return o
111
112     def _run_out(self, cmd):
113         return subprocess.run(cmd, shell=True, capture_output=True, text=True).stdout.strip()
114
115     def _read(self, cmd):
116         try:
117             return float(self._run_out(cmd).split())[0])
118         except (ValueError, IndexError):
119             return float("nan")
120
121     def _hw(self, template, value, why):
122         cmd = template.format(khz=int(value), w=int(value), v=value)
123         self.audit({"write": shlex.split(cmd), "why": why, "dry_run": self.a.dry_run})
124         if not self.a.dry_run:
125             subprocess.run(shlex.split(cmd), check=True)
126
127     # ---- efferent -----
128
129     def push(self, d, obs, killed=False):
130         if killed:
131             return
132         self._power_cap(float(d["power_cap"]), obs)
133         self._hardware(float(d["power_cap"]), obs)
134         self._rollout(bool(d["change_permitted"]), bool(d["rollback_authorized"]))
135         self._batch(bool(d["change_permitted"]), obs)
136         for name, fn in (("rightsize", lambda: self._rightsize(obs)), ("coldstart", self._coldstart),
137 ("batch_pace", lambda: self._batch_pace(obs)), ("contain", lambda: self._contain(obs)),
138 ("cooling", lambda: self._cooling(obs))):
139             self._guard(name, fn)
140
141     def _may(self, organ, what):
142         """Nervous-system authority for an organ ("expand", "contract", "admit", "pause"); without it,
143 everything allowed."""
144         a = getattr(self, "auth", None)
145         if not a:
146             return True
147         return bool(a.get("organs", {}).get(organ, {}).get(what, False))

```

```

143
144 def _envelope(self, organ, default):
145     a = getattr(self, "auth", None)
146     return (a or {}).get("organs", {}).get(organ, {}).get("envelope", default)
147
148 def _guard(self, name, fn):
149     """One lever failing is logged and never stops the others (nor the kill switch)."""
150     try:
151         fn()
152     except Exception as e: # noqa: BLE001
153         err = getattr(e, "stderr", "") or str(e)
154         self.audit({"error": f"{name}: {type(e).__name__}: {str(err).strip()[:400]}"})
155
156 # ---- new levers -----
157
158 def _named_usage(self, ns, sel=None):
159     args = ["top", "pods", "-n", ns, "--no-headers"] + (["-l", sel] if sel else [])
160     try:
161         out = self.k.get(*args)
162     except Exception:
163         return {}
164     return {line.split()[0]: milli(line.split()[1]) for line in str(out).strip().splitlines() if line.strip()}
165
166 def _rightsize(self, obs):
167     for target in filter(None, getattr(self.a, "rightsize_deployments", "").split(",")):
168         ns, name = ref(target)
169         dep = self.k.get("get", "deployment", name, "-n", ns, "-o", "json")
170         ann = dep["metadata"].get("annotations", {}) or {}
171         c0 = dep["spec"]["template"]["spec"]["containers"][0].get("resources", {})
172         base = c0.get("requests", {}).get("cpu")
173         if base is None:
174             continue
175         if REQ_ANN not in ann:
176             self.k.write(["annotate", "deployment", name, "-n", ns, f"{REQ_ANN}={int(milli(base))m}"],
177 "rightsize: record original CPU request")
178             sel = ", ".join(f"{k}={v}" for k, v in dep["spec"]["selector"]["matchLabels"].items())
179             use = self._named_usage(ns, sel)
180             for pod in self._pods(dep, ns):
181                 pn = pod["metadata"]["name"]
182                 if pn not in use:
183                     continue
184                 res = pod["spec"]["containers"][0].get("resources", {})
185                 cur = milli(res.get("requests", {}).get("cpu", base))
186                 lim = milli(res.get("limits", {}).get("cpu", "1000000m"))
187                 want = int(min(lim, max(self.a.rightsize_min_m, math.ceil(use[pn] * (1.0 + self.a.
188 rightsize_headroom) / 10.0) * 10)))
189                 if want < cur and (not obs.get("slo_clean", True) or not self._may("pods", "contract")):
190                     continue # SLO reflex / nervous system: no shrinking without contraction authority
191                 if want > cur and obs.get("security_block", 0.0) > 0.5:
192                     continue # shield II: no expansion during a security hold
193                 if abs(want - cur) < self.a.cap_min_change_m:
194                     continue
195                 self.k.write(["patch", "pod", pn, "-n", ns, "--subresource", "resize", "--type=json", "-p",
196 json.dumps({"op": "replace", "path": "/spec/containers/0/resources/requests/cpu",
197 "value": f"{want}m"})]),
198 f"rightsize: CPU request {int(cur)}m -> {want}m in place (use {int(use[pn])}m)")
199
200 def _coldstart(self):
201     sig = getattr(self.a, "coldstart_signal", "")
202     if not sig:
203         return
204     sns, sname = ref(sig)
205     try:
206         q = float((self.k.get("get", "configmap", sname, "-n", sns, "-o", "json").get("data") or {}).get(
207 "queue", "0"))
208     except Exception:
209         return
210     self._idle = getattr(self, "_idle", 0) + 1 if q <= 0 else 0
211     for target in filter(None, getattr(self.a, "coldstart_deployments", "").split(",")):
212         ns, name = ref(target)
213         dep = self.k.get("get", "deployment", name, "-n", ns, "-o", "json")
214         ann = dep["metadata"].get("annotations", {}) or {}
215         rep = int(dep["spec"].get("replicas", 1))
216         if q > 0 and rep == 0 and REPL_ANN in ann:
217             self.k.write(["scale", "deployment", name, "-n", ns, f"--replicas={ann[REPL_ANN]}"], f"coldstart:
218 wake to {ann[REPL_ANN]} (queue {q:g})")
219             elif self._idle >= self.a.coldstart_idle and rep > 0 and self._may("pods", "contract"):
220                 self.k.write(["annotate", "deployment", name, "-n", ns, "--overwrite", f"{REPL_ANN}={rep}"],
221 "coldstart: record replicas")
222                 self.k.write(["scale", "deployment", name, "-n", ns, "--replicas=0"], f"coldstart: scale to zero
223 (idle {self._idle} decisions)")
224
225 def _batch_pace(self, obs):
226     if not getattr(self.a, "batch_pace", False):
227         return

```

```

220         hot = obs.get("power_stress", 0.0) >= self.a.pace_high or obs.get("thermal", 0.0) >= self.a.pace_heat \
221             or (getattr(self, "auth", None) is not None and self._may("batch", "pause"))
222         calm = obs.get("power_stress", 1.0) <= self.a.pace_low and obs.get("thermal", 1.0) < self.a.pace_heat -
0.06
223         jobs = self.k.get("get", "jobs", "-A", "-l", PAUSABLE_LABEL, "-o", "json")["items"]
224         if hot:
225             run = [j for j in jobs if not j["spec"].get("suspend")]
226             if run:
227                 j = run[0]; ns, name = j["metadata"]["namespace"], j["metadata"]["name"]
228                 self.k.write(["annotate", "job", name, "-n", ns, "--overwrite", f"{PACE_ANN}=true"], "batch_pace:
record pause")
229                 self.k.write(["patch", "job", name, "-n", ns, "--type=merge", "-p", json.dumps({"spec": {
"suspend": True}})],
230                             f"batch_pace: pause job (power stress {obs.get('power_stress', 0.0):.2f}, heat {obs.
get('thermal', 0.0):.2f})")
231             elif calm:
232                 paced = [j for j in jobs if j["spec"].get("suspend") and (j["metadata"].get("annotations") or {})].
get(PACE_ANN) == "true"]
233                 if paced:
234                     self._resume_job(paced[0], "batch_pace: resume job (stress and heat back down)")
235
236             def _resume_job(self, j, why):
237                 ns, name = j["metadata"]["namespace"], j["metadata"]["name"]
238                 self.k.write(["patch", "job", name, "-n", ns, "--type=merge", "-p", json.dumps({"spec": {"suspend":
False}})], why)
239                 self.k.write(["annotate", "job", name, "-n", ns, f"{PACE_ANN}=", "batch_pace: clear pause record")
240
241             def _quota(self, ns):
242                 try:
243                     return self.k.get("get", "resourcequota", QUOTA, "-n", ns, "-o", "json")
244                 except Exception:
245                     return None
246
247             def _contain(self, obs):
248                 """Contain on over-budget use or a security hold; while contained each running pod gets an equal share
of the
249                 budget as its CPU limit (in place); lift only when use is under 80% of the budget (release band) and no
hold."""
250                 for ns in filter(None, getattr(self.a, "contain_namespaces", "").split(",")):
251                     total = sum(self._named_usage(ns).values())
252                     budget = float(self.a.contain_cpu_m)
253                     hold = obs.get("security_block", 0.0) > 0.5
254                     q = self._quota(ns)
255                     if q is None and (total > budget or hold):
256                         self.k.write(["create", "quota", QUOTA, "-n", ns, f"--hard=limits.cpu={int(budget)}m"],
                                     f"contain: quota limits.cpu={int(budget)}m (use {int(total)}m)")
257                         q = {"metadata": {"annotations": {}}}
258                     elif q is not None and total < 0.8 * budget and not hold:
259                         self._release(ns, q, "contain: use back under budget, lift containment")
260                         continue
261                     if q is None:
262                         continue
263                     held = json.loads((q["metadata"].get("annotations") or {}).get(LIM_ANN, "{}"))
264                     run = [p for p in self.k.get("get", "pods", "-n", ns, "-o", "json")["items"] if p["status"].get(
"phase") == "Running"
265                             and p["spec"]["containers"][0].get("resources", {}).get("limits", {}).get("cpu") is not None]
266                     share = max(10, int(budget / max(1, len(run)) / 10) * 10)
267                     changed = False
268                     for p in run:
269                         pn = p["metadata"]["name"]; lim = p["spec"]["containers"][0]["resources"]["limits"]["cpu"]
270                         if pn not in held:
271                             held[pn] = lim; changed = True
272                             want = min(share, int(milli(held[pn])))
273                             if int(milli(lim)) != want:
274                                 self.k.write(["patch", "pod", pn, "-n", ns, "--subresource", "resize", "--type=json", "-p",
                                             json.dumps({"op": "replace", "path": "/spec/containers/0/resources/limits/
cpu", "value": f"{want}m"})]),
275                                             f"contain: pod CPU limit {lim} -> {want}m in place (budget share)")
276
277                     if changed:
278                         # the original limits live on the quota object, so a kill switch in any process can restore them
279                         self.k.write(["annotate", "resourcequota", QUOTA, "-n", ns, "--overwrite", f"{LIM_ANN}={json.
dumps(held, sort_keys=True)}"],
280                                     "contain: record original pod CPU limits")
281
282             def _release(self, ns, q, why):
283                 # the quota goes first: while it stands, Kubernetes refuses to raise a pod's limit back above the budget
284                 share
285                 held = json.loads((q["metadata"].get("annotations") or {}).get(LIM_ANN, "{}"))
286                 self.k.write(["delete", "resourcequota", QUOTA, "-n", ns, f"{why}: delete quota")
287                 for pn, lim in held.items():
288                     self._guard(f"{why} {pn}", lambda pn=pn, lim=lim: self.k.write(
289                         ["patch", "pod", pn, "-n", ns, "--subresource", "resize", "--type=json", "-p",
290                         json.dumps({"op": "replace", "path": "/spec/containers/0/resources/limits/cpu", "value":
lim}]]),
291                         f"{why}: restore pod CPU limit {lim}"))
292

```

```

293     def _cooling(self, obs):
294         cmd = getattr(self.a, "cooling_cmd", "")
295         if not cmd:
296             return
297         th = obs.get("thermal", 0.5)
298         x = min(1.0, max(0.0, (th - 0.5) / 0.5))
299         c = round(self.a.cooling_max_c - (self.a.cooling_max_c - self.a.cooling_min_c) * x, 1)
300         c = round(min(c, max(self.a.cooling_min_c, self._envelope("cooling", [0, 99])[1])), 1) # nervous-
system
301     envelope
302     if abs(c - getattr(self, "_setpoint", -99.0)) >= 0.5:
303         self._hw(cmd.replace("{c}", "{v}"), c, f"cooling: supply-air setpoint {c} C (heat {th:.2f})")
304         self._setpoint = c
305
306     def _pods(self, dep, ns):
307         sel = ",".join(f"{k}={v}" for k, v in dep["spec"]["selector"]["matchLabels"].items())
308         return [p for p in self.k.get("get", "pods", "-n", ns, "-l", sel, "-o", "json")["items"]
309                 if p["status"].get("phase") == "Running"]
310
311     def _pod_usage(self, dep, ns):
312         sel = ",".join(f"{k}={v}" for k, v in dep["spec"]["selector"]["matchLabels"].items())
313         try:
314             out = self.k.get("top", "pods", "-n", ns, "-l", sel, "--no-headers")
315             except Exception:
316                 return []
317             return [milli(line.split()[1]) for line in str(out).strip().splitlines() if line.strip()]
318
319     def _resize(self, pod, ns, value, why):
320         self.k.write(["patch", "pod", pod["metadata"]["name"], "-n", ns, "--subresource", "resize", "--
type=json", "-p",
321                     json.dumps([{"op": "replace", "path": "/spec/containers/0/resources/limits/cpu", "value":
value}]]), why)
322
323     def _power_cap(self, cap, obs):
324         """In-place pod resize (no restart): each running pod's CPU limit = base limit x cap; the kernel's CFS
quota
325         enforces it. The deployment template is not changed, so no rollout is triggered."""
326         if not obs.get("slo_clean", True):
327             cap = 1.0 # SLO reflex: no power capping while service is (or was just) over its response-time
328         target
329         if getattr(self.a, "latency_file", ""):
330             # a request-served workload: its work is set by arrivals, not by the cap, so throttling it saves no
energy
331             # (the same CPU-seconds run later) and only adds wait. The energy goes where the work is: convey()
332             return self.convey(obs)
333         for target in filter(None, getattr(self.a, "cap_deployments", "").split(",")):
334             ns, name = ref(target)
335             dep = self.k.get("get", "deployment", name, "-n", ns, "-o", "json")
336             ann = dep["metadata"].get("annotations", {}) or {}
337             c0 = dep["spec"]["template"]["spec"]["containers"][0]
338             tml = c0.get("resources", {}).get("limits", {}).get("cpu")
339             if tml is None:
340                 continue
341             base = milli(ann.get("CPU_ANN", tml))
342             req = milli(c0.get("resources", {}).get("requests", {}).get("cpu", "0"))
343             want = int(max(req, round(base * max(self.a.cap_min, min(1.0, cap)) / 10.0) * 10))
344             # reflex: never cap a pod below what it is using plus headroom (as the node muscle never goes below
requests)
345             use = self._pod_usage(dep, ns)
346             if use:
347                 want = int(min(base, max(want, math.ceil(max(use) * (1.0 + self.a.cap_headroom) / 10.0) * 10)))
348             if "CPU_ANN" not in ann:
349                 self.k.write(["annotate", "deployment", name, "-n", ns, f"{CPU_ANN}={int(base)}m", "power_cap:
record original CPU limit")
350                 for pod in self._pods(dep, ns):
351                     cur = milli(pod["spec"]["containers"][0].get("resources", {}).get("limits", {}).get("cpu", tml))
352                     if obs.get("security_block", 0.0) > 0.5 and want > cur:
353                         continue # shield II: no expansion during a security hold
354                     if abs(want - cur) < self.a.cap_min_change_m:
355                         continue
356                     self._resize(pod, ns, f"{want}m", f"power_cap: pod CPU limit to {want}m in place (cap {cap:.
3f})")
357
358     def convey(self, obs=None):
359         """Energy to where the work is (omnicompass/conveyance.py, on one machine's CPU).
360
361         A serving pod's CPU limit is a quota: a request that needs more CPU than one quota period allows waits
for the
362         next period, while the machine it runs on sits idle. That wait is energy withheld from the work, not
saved: the
363         request uses the same CPU-seconds either way. So each machine's idle CPU is conveyed to the serving pods
on it:
364
365         c_i = min( max(L_i, (0.95 A_j - Q_j) / |P_j|), 0.95 A_j )
366
367         A_j the machine's allocatable CPU, Q_j the requests of every other pod on it, P_j the serving pods on it,
L_i the

```

```

366         operator's own limit. The machine's last five percent is never handed out (the living band), a pod never
367 gets less than its operator gave it, and the requests (what the scheduler and the autoscaler read) are
untouched.
368 Resized in place through pods/resize: no pod restarts. No expansion during a security hold. The kill
switch
369 returns every pod to L_i.
370
371 Conveyance is on by default. It moves no work: the same requests run with less waiting. (The +26% CPU
seen on
372 kind is more requests served, because that load is closed-loop, not waste.) Optionally, --convey-on (a
share
373 of --slo-ms) engages it only once the 95th-percentile response time reaches convey-on x SLO and releases
it
374 below convey-off x SLO, returning every serving pod to L_i; a blind latency sense engages it (service
first)."""
375 from omnicompass.nervous_system import BAND
376 obs = obs or {}
377 targets = [ref(t) for t in filter(None, getattr(self.a, "cap_deployments", "").split(","))]
378 if not targets:
379     return None
380 engaged = self._convey_engaged()
381 alloc = {n["metadata"]["name"]: milli(n["status"]["allocatable"]["cpu"])
382         for n in self.k.get("get", "nodes", "-o", "json")["items"]}
383 pods = [p for p in self.k.get("get", "pods", "-A", "-o", "json")["items"]
384         if p["status"].get("phase") in ("Running", "Pending") and p["spec"].get("nodeName")]
385 out = {}
386 for ns, name in targets:
387     dep = self.k.get("get", "deployment", name, "-n", ns, "-o", "json")
388     ann = dep["metadata"].get("annotations", {}) or {}
389     c0 = dep["spec"]["template"]["spec"]["containers"][0]
390     tpl = c0.get("resources", {}).get("limits", {}).get("cpu")
391     if tpl is None:
392         continue
393     base = milli(ann.get(CPU_ANN, tpl))
394     sel = dep["spec"]["selector"]["matchLabels"]
395     mine = lambda p: p["metadata"].get("namespace", "default") == ns and all(
396         (p["metadata"].get("labels", {}) or {}).get(k) == v for k, v in sel.items())
397     serving, other = {}, {}
398     for p in pods:
399         node = p["spec"]["nodeName"]
400         if mine(p) and p["status"].get("phase") == "Running":
401             serving.setdefault(node, []).append(p)
402         elif not mine(p):
403             other[node] = other.get(node, 0.0) + sum(
404                 milli(c.get("resources", {}).get("requests", {}).get("cpu", "0") or "0") for c in p[
"spec"]["containers"])
405             if CPU_ANN not in ann and serving:
406                 self.k.write(["annotate", "deployment", name, "-n", ns, f"{CPU_ANN}={int(base)}m"], "convey:
record original CPU limit")
407             got = []
408             for node, ps in serving.items():
409                 a_j = alloc.get(node, 0.0)
410                 share = (BAND[1] * a_j - other.get(node, 0.0)) / len(ps)
411                 want = int(min(max(base, math.floor(share / 10.0) * 10), BAND[1] * a_j)) if engaged else int(
base)
412                 for p in ps:
413                     cur = milli(p["spec"]["containers"][0].get("resources", {}).get("limits", {}).get("cpu",
tpl))
414                     if obs.get("security_block", 0.0) > 0.5 and want > cur:
415                         got.append(cur); continue # shield II: no expansion during a security hold
416                     if abs(want - cur) < self.a.cap_min_change_m:
417                         got.append(cur); continue
418                     self._resize(p, ns, f'{want}m', f'convey: {node} idle CPU to its {len(ps)} serving pod(s),
limit {want}m in place"
419                     if engaged else f'convey: response time calm, operator's limit {want}m in
place")
420                     out[f"{ns}/{p['metadata']['name']}"] = want; got.append(want)
421             # the gain g = mean conveyed limit / operator's limit: the replica organ reads it to keep the
operator's
422             # promise in queue terms (target x request / limit) while the limit is larger
423             self.gain[(ns, name)] = (sum(got) / len(got) / base) if got and base > 0 else 1.0
424             self.base[(ns, name)] = base
425         return out
426
427 def _convey_engaged(self):
428     """Hysteresis on the response time: on at convey-on x SLO, off below convey-off x SLO; blind engages."""
429     on, off = getattr(self.a, "convey_on", 0.0), getattr(self.a, "convey_off", 0.0)
430     lf, slo = getattr(self.a, "latency_file", ""), getattr(self.a, "slo_ms", 0.0)
431     if not on or not lf or not slo:
432         return True
433     ls = latency_sense(lf, self.a.latency_window_s)
434     was = getattr(self, "_conveying", False)
435     now = True if ls["blind"] else (ls["p95"] >= on * slo or (was and ls["p95"] >= off * slo))
436     if now != was:
437         self.audit({"convey": "engaged" if now else "released", "p95_ms": None if ls["blind"] else round(ls[

```

```

"p95"], 1),
    "slo_ms": slo, "blind": ls["blind"]})
    self._conveying = now
    return now

    def _hardware(self, cap, obs):
        cap = 1.0 if not obs.get("slo_clean", True) else max(self.a.cap_min, min(1.0, cap))
        if obs.get("security_block", 0.0) > 0.5:
            cap = min(cap, getattr(self, "_last_cap", 1.0)) # shield I1: no expansion during a security hold
        if abs(cap - getattr(self, "_last_cap", 1.0)) < 0.02:
            return
        if getattr(self.a, "cpufreq_policy_root", ""):
            self._cpufreq_sysfs(cap, obs)
        if getattr(self.a, "cpufreq_cmd", "") and self.a.cpu_max_khz:
            lo, hi = self._envelope("cpufreq", [0.0, 1.0]); cap = min(max(cap, lo), hi)
            self._hw(self.a.cpufreq_cmd, self.a.cpu_max_khz * cap, f"cpu_pstate: frequency ceiling {int(self.a.
cpu_max_khz * cap)} kHz (cap {cap:.3f})")
        if getattr(self.a, "gpu_power_cmd", "") and self.a.gpu_max_w:
            self._hw(self.a.gpu_power_cmd, self.a.gpu_max_w * cap, f"gpu: power limit {int(self.a.gpu_max_w *
cap)} W (cap {cap:.3f})")
        self._last_cap = cap

    def _cpufreq_sysfs(self, cap, obs):
        """CPU frequency ceiling through the kernel's own policy files (hardware/cpufreq.py).
        ceiling = min(max(cap, envelope floor, schedutil request), envelope ceiling), where the schedutil
        request is
        min(1, 1.25 u) at the current CPU utilisation u: the ceiling never cuts below the frequency schedutil
        itself
        would ask for, so it only removes headroom the scheduler is not using. The envelope ceiling (heat) wins
        last.
        The first write snapshots every policy; the kill switch restores those exact values."""
        from hardware.cpufreq import CpuFreqPolicies, schedutil_frequency_invariant
        if getattr(self, "cf", None) is None:
            self._cf = CpuFreqPolicies(self.a.cpufreq_policy_root)
            snap = self._cf.capture_original()
            if getattr(self.a, "cpufreq_require_schedutil", False):
                self._cf.assert_schedutil()
            self.audit({"cpu_pstate_snapshot": [r.as_dict() for r in snap]})
            lo, hi = self._envelope("cpufreq", [0.0, 1.0])
            u = float(obs.get("cpu_util", obs.get("load_ratio", 1.0)))
            c = min(max(cap, lo, schedutil_frequency_invariant(u)), hi)
            w = self._cf.apply_cap(c, dry_run=self.a.dry_run)
            self.audit({"cpu_pstate_write": w, "why": f"cpu_pstate: ceiling {c:.3f} (cap {cap:.3f}, envelope [{lo:.
3f}, {hi:.3f}], "
                                                                    f"schedutil request {schedutil_frequency_invariant(u):.3f} at
u {u:.3f})"})

    def _rollout(self, permitted, rollback):
        for target in filter(None, getattr(self.a, "rollout_guard", "").split(",")):
            ns, name = ref(target)
            dep = self.k.get("get", "deployment", name, "-n", ns, "-o", "json")
            ann = dep["metadata"].get("annotations", {}) or {}
            paused = bool(dep["spec"].get("paused"))
            conds = {c["type"]: c for c in dep.get("status", {}).get("conditions", [])}
            stuck = conds.get("Progressing", {}).get("reason") == "ProgressDeadlineExceeded"
            if rollback and stuck:
                self.k.write(["rollout", "undo", f"deployment/{name}", "-n", ns], "rollout: undo stuck rollout (
rollback authorised)")
                continue
            if not permitted and not paused:
                self.k.write(["annotate", "deployment", name, "-n", ns, "--overwrite", f"{PAUSE_ANN}=true"],
"rollout: record pause")
                self.k.write(["rollout", "pause", f"deployment/{name}", "-n", ns], "rollout: pause (change not
permitted)")
            elif permitted and paused and ann.get(PAUSE_ANN) == "true":
                self.k.write(["rollout", "resume", f"deployment/{name}", "-n", ns], "rollout: resume (change
permitted)")
            self.k.write(["annotate", "deployment", name, "-n", ns, f"{PAUSE_ANN}-"], "rollout: clear pause
record")

    def _batch(self, permitted, obs):
        if not getattr(self.a, "batch", False):
            return
        if not permitted or obs.get("security_block", 0.0) > 0.5:
            return
        if obs.get("load_ratio", 1.0) >= self.a.batch_load_max or obs.get("power_stress", 1.0) >= self.a.
batch_power_max:
            return
        if not self._may("batch", "admit"):
            return
        jobs = self.k.get("get", "jobs", "-A", "-l", BATCH_LABEL, "-o", "json")["items"]
        held = sorted((j for j in jobs if j["spec"].get("suspend")), key=lambda j: j["metadata"].get(
"creationTimestamp", ""))
        if held:
            j = held[0]; ns, name = j["metadata"]["namespace"], j["metadata"]["name"]
            self.k.write(["patch", "job", name, "-n", ns, "--type=merge", "-p", json.dumps({"spec": {"suspend":

```

```

False}}]],
509         "batch: admit held job (headroom)")
510
511     # ---- kill -----
512     def _restore(self):
513         self._guard("kill cooling", self._restore_cooling)
514         self._guard("kill contain", self._restore_contain)
515         self._guard("kill batch_pace", self._restore_pace)
516         self._guard("kill coldstart", self._restore_coldstart)
517         self._guard("kill rightsize", self._restore_rightsized)
518         self._guard("kill hardware and power cap", self._restore_core)
519
520     def _restore_cooling(self):
521         if getattr(self.a, "cooling_cmd", ""):
522             self._hw(self.a.cooling_cmd.replace("{c}", "{v}"), self.a.cooling_restore_c, "kill switch: restore
cooling setpoint")
523             self._setpoint = None
524
525     def _restore_contain(self):
526         for ns in filter(None, getattr(self.a, "contain_namespaces", "").split(",")):
527             q = self._quota(ns)
528             if q is not None:
529                 self._release(ns, q, "kill switch")
530
531     def _restore_pace(self):
532         if getattr(self.a, "batch_pace", False):
533             for j in self.k.get("get", "jobs", "-A", "-l", PAUSABLE_LABEL, "-o", "json")["items"]:
534                 if j["spec"].get("suspend") and (j["metadata"].get("annotations") or {}).get(PACE_ANN) == "true":
535                     self._resume_job(j, "kill switch: resume paced job")
536
537     def _restore_coldstart(self):
538         for target in filter(None, getattr(self.a, "coldstart_deployments", "").split(",")):
539             ns, name = ref(target)
540             dep = self.k.get("get", "deployment", name, "-n", ns, "-o", "json")
541             ann = dep["metadata"].get("annotations", {}) or {}
542             if REPL_ANN in ann:
543                 if int(dep["spec"].get("replicas", 1)) == 0:
544                     self.k.write(["scale", "deployment", name, "-n", ns, f"--replicas={ann[REPL_ANN]}"], "kill
switch: restore replicas")
545                     self.k.write(["annotate", "deployment", name, "-n", ns, f"{REPL_ANN}-"], "kill switch: clear
replica record")
546
547     def _restore_rightsized(self):
548         for target in filter(None, getattr(self.a, "rightsized_deployments", "").split(",")):
549             ns, name = ref(target)
550             dep = self.k.get("get", "deployment", name, "-n", ns, "-o", "json")
551             orig = (dep["metadata"].get("annotations", {}) or {}).get(REQ_ANN)
552             if orig:
553                 for pod in self._pods(dep, ns):
554                     if pod["spec"]["containers"][0].get("resources", {}).get("requests", {}).get("cpu") != orig:
555                         self.k.write(["patch", "pod", pod["metadata"]["name"], "-n", ns, "--subresource",
"resize", "--type=json", "-p",
556                                     json.dumps({"op": "replace", "path": "/spec/containers/0/resources/
requests/cpu", "value": orig})],
557                                     "kill switch: restore original CPU request in place")
558                         self.k.write(["annotate", "deployment", name, "-n", ns, f"{REQ_ANN}-"], "kill switch: remove
request record")
559
560     def _restore_core(self):
561         if getattr(self, "_cf", None) is not None:
562             self.audit({"cpu_pstate_restore": self._cf.restore(dry_run=self.a.dry_run), "why": "kill switch:
exact pre-Omni CPU ceilings"})
563             if getattr(self.a, "cpufreq_cmd", "") and self.a.cpu_max_khz and getattr(self, "_last_cap", 1.0) != 1.0:
564                 self._hw(self.a.cpufreq_cmd, self.a.cpu_max_khz, "kill switch: restore maximum CPU frequency")
565             if getattr(self.a, "gpu_power_cmd", "") and self.a.gpu_max_w and getattr(self, "_last_cap", 1.0) != 1.0:
566                 self._hw(self.a.gpu_power_cmd, self.a.gpu_max_w, "kill switch: restore maximum GPU power limit")
567             self._last_cap = 1.0
568             for target in filter(None, getattr(self.a, "cap_deployments", "").split(",")):
569                 ns, name = ref(target)
570                 dep = self.k.get("get", "deployment", name, "-n", ns, "-o", "json")
571                 orig = (dep["metadata"].get("annotations", {}) or {}).get(CPU_ANN)
572                 if orig:
573                     for pod in self._pods(dep, ns):
574                         if pod["spec"]["containers"][0].get("resources", {}).get("limits", {}).get("cpu") != orig:
575                             self._resize(pod, ns, orig, "kill switch: restore original pod CPU limit in place")
576                             self.k.write(["annotate", "deployment", name, "-n", ns, f"{CPU_ANN}-"], "kill switch: remove CPU
record")
577             for target in filter(None, getattr(self.a, "rollout_guard", "").split(",")):
578                 ns, name = ref(target)
579                 dep = self.k.get("get", "deployment", name, "-n", ns, "-o", "json")
580                 if (dep["metadata"].get("annotations", {}) or {}).get(PAUSE_ANN) == "true":
581                     self.k.write(["rollout", "resume", f"deployment/{name}", "-n", ns], "kill switch: resume
rollout")
582                     self.k.write(["annotate", "deployment", name, "-n", ns, f"{PAUSE_ANN}-"], "kill switch: clear
pause record")

```

```

583
584
585 def latency_sense(path, window_s, now=None):
586     """The latency afferent with its integrity (manuscript Appendix J: the nervous system owns delay and dropout
587     handling). Returns p95 of successful requests in the last window, the ok and failed counts in it, and the
    age of
588     the newest sample against the wall clock. The sense is blind when its newest sample is older than two
    windows, or
589     when the window holds no successful request: a hung probe stops writing, and its last clean window must not
    be read
590     as the present."""
591     import os, time
592     try:
593         age = (now if now is not None else time.time()) - os.path.getmtime(path)
594         rows = list(csv.DictReader(open(path)))
595     except OSError:
596         return {"p95": float("nan"), "ok": 0, "fail": 0, "age_s": float("inf"), "blind": True}
597     if not rows:
598         return {"p95": float("nan"), "ok": 0, "fail": 0, "age_s": age, "blind": True}
599     t_end = float(rows[-1]["elapsed_seconds"])
600     win = [r for r in rows if float(r["elapsed_seconds"]) >= t_end - window_s]
601     ms = sorted(float(r["latency_ms"]) for r in win if r.get("ok") == "1")
602     fails = sum(1 for r in win if r.get("ok") != "1")
603     p95 = ms[min(len(ms) - 1, int(0.95 * len(ms)))] if ms else float("nan")
604     return {"p95": p95, "ok": len(ms), "fail": fails, "age_s": age, "blind": age > 2.0 * window_s or not ms}
605
606
607 def latency_window(path, window_s, now=None):
608     """The last window_s seconds of successful probe samples (ms), with the same integrity rule as latency_sense.
    """
609     import os, time
610     try:
611         age = (now if now is not None else time.time()) - os.path.getmtime(path)
612         rows = list(csv.DictReader(open(path)))
613     except OSError:
614         return {"ms": [], "blind": True}
615     if not rows:
616         return {"ms": [], "blind": True}
617     t_end = float(rows[-1]["elapsed_seconds"])
618     ms = [float(r["latency_ms"]) for r in rows if r.get("ok") == "1" and float(r["elapsed_seconds"]) >= t_end -
    window_s]
619     return {"ms": ms, "blind": age > 2.0 * window_s or not ms}
620
621
622 def latency_p95(path, window_s):
623     """95th-percentile response time (ms) of successful requests in the last window_s seconds of the probe CSV.
    """
624     try:
625         rows = list(csv.DictReader(open(path)))
626     except OSError:
627         return float("nan")
628     if not rows:
629         return float("nan")
630     t_end = float(rows[-1]["elapsed_seconds"])
631     ms = sorted(float(r["latency_ms"]) for r in rows if r.get("ok") == "1" and float(r["elapsed_seconds"]) >=
    t_end - window_s)
632     fails = sum(1 for r in rows if r.get("ok") != "1" and float(r["elapsed_seconds"]) >= t_end - window_s)
633     if fails and not ms:
634         return 0
635     return ms[min(len(ms) - 1, int(0.95 * len(ms)))] if ms else float("nan")
636
637
638 def add_args(ap):
639     ap.add_argument("--cap-deployments", default="", help="ns/name[ns/name]: power_cap muscle scales their CPU
    limit")
640     ap.add_argument("--cap-min", type=float, default=0.65, help="lowest power cap applied (shield floor)")
641     ap.add_argument("--cap-min-change-m", type=int, default=20, help="smallest CPU-limit change written,
    millicores")
642     ap.add_argument("--cap-headroom", type=float, default=0.3, help="power cap never below pod CPU usage x (1 +
    headroom)")
643     ap.add_argument("--latency-file", default="", help="probe CSV (elapsed_seconds,latency_ms,ok) for the
    latency afferent")
644     ap.add_argument("--slo-ms", type=float, default=0.0, help="95th-percentile response-time target, ms")
645     ap.add_argument("--latency-window-s", type=float, default=60.0)
646     ap.add_argument("--convey-on", type=float, default=0.0,
    help="0 (default): convey idle CPU always; > 0: only once p95 reaches this share of --slo-
    ms")
647     ap.add_argument("--convey-off", type=float, default=0.25, help="return pods to the operator's limit once p95
    is below this share")
648     ap.add_argument("--slo-clear", type=int, default=3, help="decisions the SLO must stay met before densifying
    or capping again")
649     ap.add_argument("--thermal-model", action="store_true", help="heat muscle: thermal state from the harness
    heat law")
650     ap.add_argument("--security-configmap", default="", help="ns/name of a ConfigMap whose key 'hold' signals a
    security hold")
651     ap.add_argument("--rollout-guard", default="", help="ns/name[ns/name]: rollout muscle pauses, resumes,

```

```

undoes")
653     ap.add_argument("--batch", action="store_true", help="batch muscle: admit held Jobs labelled " + BATCH_LABEL)
654     ap.add_argument("--batch-load-max", type=float, default=0.8)
655     ap.add_argument("--batch-power-max", type=float, default=0.9)
656     ap.add_argument("--rapl-cmd", default="", help="cpu_pstate pull: prints CPU package watts")
657     ap.add_argument("--cpufreq-cmd", default="", help="cpu_pstate push: command template with {khz}")
658     ap.add_argument("--cpufreq-policy-root", default="", help="cpu_pstate push through sysfs policies (e.g. /sys/
devices/system/cpu/cpufreq); exact restore on kill")
659     ap.add_argument("--cpufreq-require-schedutil", action="store_true", help="refuse to act unless every policy
runs schedutil")
660     ap.add_argument("--cpu-max-khz", type=float, default=0.0)
661     ap.add_argument("--gpu-query-cmd", default="", help="gpu pull: prints 'watts,celsius'")
662     ap.add_argument("--gpu-power-cmd", default="", help="gpu push: command template with {w}")
663     ap.add_argument("--gpu-max-w", type=float, default=0.0)
664     ap.add_argument("--gpu-temp-limit", type=float, default=83.0, help="GPU temperature treated as thermal 1.0")
665     ap.add_argument("--rightsize-deployments", default="", help="ns/name[ns/name]: rightsize muscle sets pod
CPU requests from use")
666     ap.add_argument("--rightsize-headroom", type=float, default=0.3)
667     ap.add_argument("--rightsize-min-m", type=int, default=50)
668     ap.add_argument("--coldstart-deployments", default="", help="ns/name[ns/name]: scaled to zero when idle,
woken on work")
669     ap.add_argument("--coldstart-signal", default="", help="ns/name of a ConfigMap whose key 'queue' is the
waiting work")
670     ap.add_argument("--coldstart-idle", type=int, default=3, help="idle decisions before scaling to zero")
671     ap.add_argument("--batch-pace", action="store_true", help="batch_pace muscle: pause Jobs labelled " +
PAUSABLE_LABEL + " under stress")
672     ap.add_argument("--pace-high", type=float, default=0.95, help="power stress that pauses a job")
673     ap.add_argument("--pace-low", type=float, default=0.8, help="power stress under which a paused job resumes")
674     ap.add_argument("--pace-heat", type=float, default=0.96, help="heat state that pauses a job")
675     ap.add_argument("--contain-namespaces", default="", help="ns[ns]: agent namespaces held to --contain-cpu-m")
676     ap.add_argument("--contain-cpu-m", type=float, default=1000.0)
677     ap.add_argument("--cooling-cmd", default="", help="cooling push: command template with {c} (supply-air
setpoint, C)")
678     ap.add_argument("--cooling-min-c", type=float, default=18.0)
679     ap.add_argument("--cooling-max-c", type=float, default=27.0)
680     ap.add_argument("--cooling-restore-c", type=float, default=22.0)
681

```

hardware/cpufreq.py

174 lines SHA-256 65bfaf694501d09708d27f18fcb3c540c307c98172c946fd420e064701be4c40

```
1  """CPUFreq contract used by the Omni-Compass hardware ceiling muscle.
2
3  The schedutil *model* and the sysfs *actuator* are deliberately separate:
4
5  * ``schedutil_frequency_invariant()`` is the documented frequency-invariant CFS
6    shape:  $\min(1, 1.25 \cdot u)$ , with the 0.8 tipping point.
7  * ``CpuFreqPolicies`` never reimplements schedutil. It only constrains the
8    kernel policy by writing ``scaling_max_freq`` and can restore the exact
9    pre-Omni limits.
10
11  A real run must record the scaling driver/governor. ``intel_pstate`` active
12  mode (including HWP) has different frequency-selection semantics; a ceiling
13  write can still be a real policy bound, but must not be reported as a
14  schedutil-equivalence experiment.
15  """
16  from __future__ import annotations
17
18  from dataclasses import asdict, dataclass
19  from pathlib import Path
20  from typing import Iterable
21
22
23  SCHEDUTIL_TIPPING_CONSTANT = 1.25
24  SCHEDUTIL_TIPPING_UTIL = 1.0 / SCHEDUTIL_TIPPING_CONSTANT # 0.8
25
26
27  def schedutil_frequency_invariant(util_fraction: float, f_min_fraction: float = 0.0) -> float:
28      """Normalized frequency request for frequency-invariant CFS utilization.
29
30      This intentionally models only the documented CFS proportional branch.
31      OPP/driver resolution, uclamp, RT/DL max jumps, IO-wait boost, rate limiting,
32      thermal constraints, PM QoS and hardware coordination are outside this
33      scalar map and must be represented separately in evidence.
34      """
35      u = max(0.0, float(util_fraction))
36      floor = max(0.0, min(1.0, float(f_min_fraction)))
37      return min(1.0, max(floor, SCHEDUTIL_TIPPING_CONSTANT * u))
38
39
40  @dataclass(frozen=True)
41  class PolicySnapshot:
42      policy: str
43      path: str
44      cpuinfo_min_freq: int
45      cpuinfo_max_freq: int
46      scaling_min_freq: int
47      scaling_max_freq: int
48      scaling_driver: str
49      scaling_governor: str
50      affected_cpus: str
51
52      def as_dict(self) -> dict:
53          return asdict(self)
54
55
56  class CpuFreqPolicies:
57      """Discover, cap, and exactly restore Linux CPUFreq policy ceilings."""
58
59      def __init__(self, root: str | Path = "/sys/devices/system/cpu/cpufreq"):
60          self.root = Path(root)
61          self._original: dict[str, PolicySnapshot] = {}
62
63      @staticmethod
64      def _read_text(path: Path, default: str = "") -> str:
65          try:
66              return path.read_text().strip()
67          except OSError:
68              return default
69
70      @classmethod
71      def _read_int(cls, path: Path, default: int = 0) -> int:
72          try:
73              return int(cls._read_text(path))
74          except (TypeError, ValueError):
75              return default
76
77      def policy_dirs(self) -> list[Path]:
78          if not self.root.exists():
79              return []
80          return sorted((p for p in self.root.glob("policy*") if p.is_dir()), key=lambda p: p.name)
81
82      def snapshot(self) -> list[PolicySnapshot]:
83          rows: list[PolicySnapshot] = []
```

```

84     for p in self.policy_dirs():
85         s = PolicySnapshot(
86             policy=p.name,
87             path=str(p),
88             cpuinfo_min_freq=self._read_int(p / "cpuinfo_min_freq"),
89             cpuinfo_max_freq=self._read_int(p / "cpuinfo_max_freq"),
90             scaling_min_freq=self._read_int(p / "scaling_min_freq"),
91             scaling_max_freq=self._read_int(p / "scaling_max_freq"),
92             scaling_driver=self._read_text(p / "scaling_driver", "unknown"),
93             scaling_governor=self._read_text(p / "scaling_governor", "unknown"),
94             affected_cpus=self._read_text(p / "affected_cpus", ""),
95         )
96     rows.append(s)
97     return rows
98
99     def capture_original(self) -> list[PolicySnapshot]:
100         rows = self.snapshot()
101         if not rows:
102             raise RuntimeError(f"no CPUFreq policy directories found under {self.root}")
103         if not self._original:
104             self._original = {r.policy: r for r in rows}
105         return rows
106
107     def assert_schedutil(self) -> None:
108         rows = self.capture_original()
109         bad = [f"{r.policy}:{r.scaling_driver}/{r.scaling_governor}" for r in rows if r.scaling_governor !=
"schedutil"]
110         if bad:
111             raise RuntimeError("schedutil required for this evidence lane; found " + ", ".join(bad))
112
113     def apply_cap(self, cap: float, dry_run: bool = False) -> list[dict]:
114         """Set each scaling_max_freq to a fraction of its pre-0mni ceiling.
115
116         The operation cannot raise a pre-existing ceiling. The requested value
117         is clamped at scaling_min_freq. The kernel/driver may further resolve or
118         constrain the effective operating point.
119         """
120         rows = self.capture_original()
121         c = max(0.0, min(1.0, float(cap)))
122         writes: list[dict] = []
123         for now in rows:
124             orig = self._original[now.policy]
125             # Preserve the operator's pre-existing maximum even when it is below
126             # the silicon-reported cpuinfo maximum.
127             base = orig.scaling_max_freq or orig.cpuinfo_max_freq
128             if orig.cpuinfo_max_freq:
129                 base = min(base, orig.cpuinfo_max_freq) if base else orig.cpuinfo_max_freq
130             floor = max(0, now.scaling_min_freq)
131             desired = max(floor, min(base, int(round(base * c))))
132             target = Path(now.path) / "scaling_max_freq"
133             rec = {
134                 "policy": now.policy,
135                 "path": str(target),
136                 "before_khz": now.scaling_max_freq,
137                 "original_khz": orig.scaling_max_freq,
138                 "requested_khz": desired,
139                 "cap": c,
140                 "driver": now.scaling_driver,
141                 "governor": now.scaling_governor,
142                 "dry_run": bool(dry_run),
143             }
144             writes.append(rec)
145             if not dry_run:
146                 target.write_text(str(desired))
147         return writes
148
149     def restore(self, dry_run: bool = False) -> list[dict]:
150         if not self._original:
151             return []
152         current = {r.policy: r for r in self.snapshot()}
153         writes: list[dict] = []
154         for policy, orig in sorted(self._original.items()):
155             now = current.get(policy)
156             target = Path(orig.path) / "scaling_max_freq"
157             rec = {
158                 "policy": policy,
159                 "path": str(target),
160                 "before_khz": now.scaling_max_freq if now else None,
161                 "restore_khz": orig.scaling_max_freq,
162                 "driver": orig.scaling_driver,
163                 "governor": orig.scaling_governor,
164                 "dry_run": bool(dry_run),
165             }
166             writes.append(rec)
167             if not dry_run:
168                 target.write_text(str(orig.scaling_max_freq))

```

```
169         return writes
170
171     @property
172     def original(self) -> Iterable[PolicySnapshot]:
173         return tuple(self._original[k] for k in sorted(self._original))
174
```

hardware/node_exchange.py

157 lines SHA-256 78b18e55ddf78fc846b6673e3dd96adec95cb2de2ffc6ac44453b1bb65e1e9ec

```
1  """CPU and GPU on one conserved power budget: the conveyance law (omnicompass/conveyance.py) with CPU organs
   beside
2  GPU organs, so watts a CPU holds and does not need flow to the GPUs that need them. THEORETICAL SIMULATION on the
3  device physics of hardware/plant.py (GPU gamma fitted to MLPerf v4.0 DGX-H100 MaxQ vs MaxP; CPU first-order
   DVFS).
4
5  The law, read as mechanism (docs/CONVEYANCE_LAW.md): the site's watts are conserved (moved between organs, never
6  created, Ch. 29 §8); they flow down the gradient of need, from organs holding surplus to organs in deficit (Ch.
   31 §4);
7  the flow converges ( $\text{Re}(\lambda) < 0$ , Ch. 31 §5); a reserve for the projected rise is released continuously (Ch.
   30).
8  Here the organs are G GPU groups and the G CPU groups that feed them.
9
10 A group = N GPUs and M CPU sockets (two 8-GPU servers: 16 H100-class GPUs, 4 sockets).
11 GPU capacity  $s^\gamma$  at power-limit fraction s; power per device  $p_{\text{idle}} + (s p_{\text{max}} - p_{\text{idle}}) u$ 
12 CPU capacity f at clock fraction f; power per socket  $p_{\text{idle}} + (p_{\text{max}} - p_{\text{idle}}) u f^3$ ; clock set by a ceiling
13 (cpufreq scaling_max_freq or a RAPL package limit, the hardware connectors in omni_controller/muscles.py)
14 work the GPUs serve requests; each unit of GPU work needs --cpu-share of the CPUs at full clock to feed it:
15 served = min(demand + backlog, GPU capacity, CPU capacity / cpu-share)
16 site one budget = --budget x (sum of GPU TDP + sum of CPU maximum power); a reactive breaker throttles every
17 device uniformly the step after the site draws over it (how a site power caper behaves)
18
19 Arms, same demand, same budget, same breaker:
20 A native: GPUs at TDP, CPUs on their own governor (schedutil-style), only the breaker (it trips: a real
21 power-limited site cannot run this way)
22 S today's practice: every GPU capped at one fixed limit low enough that the site fits even with every CPU at
   its
23 maximum, (budget - CPU maximum) / GPU TDP; CPUs on their own governor
24 XR conveyance among the GPU groups; the CPUs are uncontrolled, so their full maximum power stays reserved
25 (the planning rule when nothing holds the CPUs down); GPU budget = site budget - CPU maximum
26 XM conveyance among the GPU groups; GPU budget = site budget - the CPUs' measured draw, updated every step
27 XC conveyance over CPU and GPU organs together: each CPU group is held to the clock its feeding work needs,
   and
28 the watts it no longer holds flow to the GPU groups with the largest deficit
29
30 Gauges: work served (GPU work units), backlog minutes (SLO breach), p95 latency factor, energy, work per kWh,
31 site-budget violation minutes, peak. Usage: python hardware/node_exchange.py [--scenarios 24] [--budget 0.7] [--
   load 1.3]
32 """
33 from __future__ import annotations
34 import argparse, json, math, sys
35 from pathlib import Path
36 import numpy as np
37 ROOT = Path(__file__).resolve().parents[1]; sys.path.insert(0, str(ROOT))
38 from hardware.plant import VESSELS, FAMILIES, demand_trace
39 from omnicompass.conveyance import Conveyance
40
41 G, N, M = 4, 16, 4
42 CPU_IDLE, CPU_MAX, CPU_FMIN = 120.0, 400.0, 0.4 # per socket, the cpu vessel of hardware/plant.py
43 ARMS = ("A", "S", "XR", "XM", "XC")
44
45
46 def cpu_power(f, load):
47     """Group CPU power at clock f serving load (share of the group's CPU at full clock): busy fraction load / f.
48     """
49     u = min(1.0, load / max(f, 1e-9))
50     return M * (CPU_IDLE + (CPU_MAX - CPU_IDLE) * u * f ** 3)
51
52 def cpu_clock_for(watts, load):
53     """The highest clock whose group power at this load fits the allocation (power = idle + dyn load f^2 while u
54     < 1)."""
55     dyn = (CPU_MAX - CPU_IDLE) * max(load, 1e-6)
56     f = math.sqrt(max(0.0, (watts / M - CPU_IDLE) / dyn))
57     return max(CPU_FMIN, min(1.0, f))
58
59 def gpu_need(v, want, rho):
60     s = max(v.f_min, min(1.0, (want / rho) ** (1.0 / v.gamma)))
61     return N * (v.p_idle + (s * v.p_max - v.p_idle) * rho), s
62
63
64 def cpu_need(want, c, rho):
65     f = max(CPU_FMIN, min(1.0, c * want / rho))
66     return cpu_power(f, min(c * want, f)), f
67
68
69 def run(v, fams, seed, arm, budget_frac, c=0.35, load=1.3, steps=360, rho=0.8):
70     rng = np.random.default_rng(seed)
71     dem = [load * demand_trace(f, steps, np.random.default_rng(seed + 97 * k)) for k, f in enumerate(fams)]
72     shift = [int(rng.integers(0, steps)) for _ in range(G)]
73     dem = [np.roll(d, sh) for d, sh in zip(dem, shift)]
74     gpu_tdp, cpu_top = G * N * v.p_max, G * M * CPU_MAX
```

```

75 budget = budget_frac * (gpu_tdp + cpu_top)
76 backlog = [0.0] * G; throttle = 1.0; last_cpu = G * M * CPU_IDLE
77 cv_gpu = Conveyance(G, max(1.0, budget - cpu_top), rho=rho)
78 cv_all = Conveyance(2 * G, budget, rho=rho)
79 energy = work = peak = 0.0; slo = viol = 0; lat = []
80 lo_g, hi_g = [N * v.p_max * v.f_min] * G, [N * v.p_max] * G
81 lo_c, hi_c = [M * CPU_IDLE * 1.05] * G, [M * CPU_MAX] * G
82 for t in range(steps):
83     want = [dem[k][t] + backlog[k] for k in range(G)]
84     f = [max(CPU_FMIN, min(1.0, 1.25 * c * w)) for w in want]          # the CPUs' own governor
85     if arm == "A":
86         s = [1.0] * G
87     elif arm == "S":
88         s = [max(v.f_min, min(1.0, (budget - cpu_top) / gpu_tdp))] * G
89     elif arm in ("XR", "XM"):
90         cv_gpu.budget = max(sum(lo_g), budget - (cpu_top if arm == "XR" else last_cpu))
91         a = cv_gpu.step([gpu_need(v, w, rho)[0] for w in want], lo_g, hi_g)
92         s = [max(v.f_min, min(1.0, ak / (N * v.p_max))) for ak in a]
93     else:
94         need = [gpu_need(v, w, rho)[0] for w in want] + [cpu_need(w, c, rho)[0] for w in want]
95         a = cv_all.step(need, lo_g + lo_c, hi_g + hi_c)
96         s = [max(v.f_min, min(1.0, ak / (N * v.p_max))) for ak in a[G:]]
97         f = [min(fk, cpu_clock_for(ak, c * min(w, 1.3))) for fk, ak, w in zip(f, a[G:], want)]
98         s = [max(v.f_min, x * throttle) for x in s]
99         f = [max(CPU_FMIN, x * throttle) for x in f]
100        power = 0.0; cpu_w = 0.0; served_all = 0.0
101        for k in range(G):
102            gcap = s[k] ** v.gamma
103            served = min(want[k], gcap, f[k] / c)
104            u = served / max(gcap, 1e-9)
105            pc = cpu_power(f[k], c * served)
106            power += N * (v.p_idle + (s[k] * v.p_max - v.p_idle) * u) + pc; cpu_w += pc
107            backlog[k] = max(0.0, want[k] - served); served_all += served
108            slo += int(backlog[k] > dem[k][t]); lat.append(1.0 + backlog[k] / max(dem[k][t], 1e-3))
109            last_cpu = cpu_w
110        energy += power / 60.0 / 1000.0; work += served_all; peak = max(peak, power)
111        over = power > budget * 1.0001
112        viol += int(over)
113        throttle = min(1.0, throttle * budget / power) if over else min(1.0, throttle * 1.02)
114    return {"work": work, "backlog_min": float(slo), "p95_latency_x": float(np.percentile(lat, 95)),
115            "energy_kwh": energy, "work_per_kwh": work / max(energy, 1e-9), "site_violation_min": float(viol),
116            "peak_kw": peak / 1000.0}
117
118
119 def main(argv=None):
120     ap = argparse.ArgumentParser()
121     ap.add_argument("--scenarios", type=int, default=24); ap.add_argument("--seed", type=int, default=515151)
122     ap.add_argument("--budget", type=float, default=0.7, help="site budget as a share of GPU TDP + CPU maximum")
123     ap.add_argument("--load", type=float, default=1.3, help="demand multiplier on hardware/plant.py's traces")
124     ap.add_argument("--cpu-share", type=float, default=0.35, help="CPU at full clock needed per unit of GPU
work")
125     ap.add_argument("--vessel", default="gpu_mlperf_median")
126     ap.add_argument("--out", default="")
127     a = ap.parse_args(argv)
128     v = VESSELS[a.vessel]; rng = np.random.default_rng(a.seed)
129     scen = [[FAMILIES[int(rng.integers(0, len(FAMILIES)))] for _ in range(G)], int(rng.integers(1, 2**31))] for
_ in range(a.scenarios)]
130     rows = {arm: [run(v, fm, sd, arm, a.budget, a.cpu_share, a.load) for fm, sd in scen] for arm in ARMS}
131     keys = list(rows["A"][0]); br = np.random.default_rng(7)
132     out = {"note": "THEORETICAL SIMULATION; hardware/node_exchange.py", "args": vars(a), "gamma": v.gamma,
133            "means": {k: {m: float(np.mean([r[m] for r in rows[k]])) for m in keys} for k in rows},
134            "paired_XC_vs": {}}
135     for base in ("A", "S", "XR", "XM"):
136         out["paired_XC_vs"][base] = {}
137         for m in keys:
138             d = np.array([rows["XC"][i][m] - rows[base][i][m] for i in range(len(scen))])
139             bs = d[br.integers(0, len(d), (4000, len(d)))]
140             ref = float(np.mean([r[m] for r in rows[base]]))
141             out["paired_XC_vs"][base][m] = {"delta": float(d.mean()), "pct": 100 * float(d.mean()) / ref if ref
0}
142             "ci95": [float(lo), float(hi)], "significant": bool(lo > 0 or hi <
0)}
143         if a.out:
144             Path(a.out).parent.mkdir(parents=True, exist_ok=True); Path(a.out).write_text(json.dumps(out, indent=1))
145             print(f"site budget {a.budget:.0%} of GPU TDP + CPU max, load x{a.load}, CPU share {a.cpu_share}, gamma {v.
gamma:.2f}, "
146                   f"{a.scenarios} scenarios")
147             print(f"{'gauge':20s}" + "".join(f"{'x':11s}" for x in ("A native", "S static", "XR gpu+res", "XM gpu+meas",
"XC cpu+gpu")))
148             for m in keys:
149                 print(f"{'m':20s}" + "".join(f"{'out['means'][k][m]:11.3f}" for k in ARMS))
150             for base in ("A", "S", "XR", "XM"):
151                 sig = {m: f"{'x['pct']:.1f}%" for m, x in out["paired_XC_vs"][base].items() if x["significant"]} and x[
"pct"] is not None}
152             print(f"XC vs {base}, significant:", sig)
153     return out

```

```
153
154
155 if __name__ == "__main__":
156     main()
157
```

hardware/plant.py

188 lines SHA-256 34d0da36027b1fbecab473184ed334ee8b1124a4d4fbde7796a18e2fa936908

```
1 """Device plant: CPU frequency (DVFS) and GPU power-limit muscles, three architectures, 24 scenarios.
2
3 SIMULATION. CPU physics use a standard first-order model with declared constants. The GPU performance-vs-power-
  limit
4 exponent gamma is fitted to real metered hardware: MLPerf Inference v4.0 NVIDIA DGX-H100 MaxQ vs MaxP
5 (results/hardware/CALIBRATION_MLPERF.json, Apache 2.0); the median, least favourable and most favourable fits
  are all
6 run. Replace with measured curves (RAPL, nvidia-smi) when the hardware connectors in omni_controller/muscles.py
  run on
7 owned machines.
8
9 Physics (per device, one step = 60 s)
10 CPU   power = P_idle + P_dyn * u * f^3          (dynamic power ~ C V^2 f with V ~ f)
11       capacity = f                             (compute-bound work scales with clock)
12 GPU   power = P_idle + (L - P_idle) * u         (the board draws up to its power limit L when busy)
13       capacity = (L / TDP)^gamma               (concave performance-vs-power-limit curve, gamma per workload)
14 Queue backlog += demand - served; SLO breach when backlog exceeds one step of demand.
15 Heat thermal = 0.86 thermal + 0.14 (0.34 + 0.62 power_stress) (the harness heat law)
16
17 Architectures (the only difference between arms is who sets f or L)
18 A native   CPU: schedutil-style governor f = min(1, 1.25 u_needed); GPU: power limit = TDP (vendor default)
19 B on top   the native governor runs, Omni-Compass sets a ceiling: f = min(native f, cap); L = cap x TDP
20 C direct   Omni-Compass replaces the native governor. CPU: f = min(cap, want / rho*), rho* the engine's
  target
21           utilisation (its demand output), floor f_min; GPU: L = cap x TDP (floor f_min)
22 S reference a fixed manual cap at 70% (f or L), no Omni-Compass: what an operator could do by hand
23 Omni-Compass = the shipped Governor (fleet law); its power_cap output is the cap.
24 """
25 from __future__ import annotations
26
27 import argparse, json, math, sys
28 from dataclasses import dataclass
29 from pathlib import Path
30
31 import numpy as np
32
33 ROOT = Path(__file__).resolve().parents[1]
34 sys.path.insert(0, str(ROOT))
35 from omnicompass.adapter import Governor, mode_law, AUTOPILOT
36
37
38 @dataclass
39 class Vessel:
40     name: str
41     devices: int
42     p_idle: float          # W per device
43     p_max: float           # W per device at full load, full clock / TDP
44     gamma: float           # GPU perf-vs-power exponent (unused for CPU)
45     kind: str              # "cpu" or "gpu"
46     f_min: float = 0.4     # lowest clock or power-limit fraction the device allows
47
48
49 _CAL = json.loads((ROOT / "results/hardware/CALIBRATION_MLPERF.json").read_text())
50 VESSELS = {
51     "cpu_web": Vessel("cpu_web", 100, 120.0, 400.0, 0.0, "cpu"),
52     # H100-class GPUs, TDP 700 W; gamma measured from MLPerf v4.0 MaxQ vs MaxP (results/hardware/
53     CALIBRATION_MLPERF.json).
54     # Floor 300 W / 700 W = 0.43, the lowest limit NVIDIA used for MaxQ.
55     "gpu_mlperf_median": Vessel("gpu_mlperf_median", 64, 60.0, 700.0, _CAL["gamma_median"], "gpu", 300 / 700),
56     "gpu_mlperf_least_favourable": Vessel("gpu_mlperf_least_favourable", 64, 60.0, 700.0, _CAL["gamma_max"],
57     "gpu", 300 / 700),
58     "gpu_mlperf_most_favourable": Vessel("gpu_mlperf_most_favourable", 64, 60.0, 700.0, _CAL["gamma_min"], "gpu",
59     300 / 700),
60 }
61
62 FAMILIES = ["steady", "diurnal", "spike", "flash_crowd", "slow_ramp", "lull", "burst_train", "mixed"]
63
64
65 def demand_trace(family, steps, rng):
66     t = np.arange(steps)
67     base = {"steady": 0.55, "diurnal": 0.5, "spike": 0.45, "flash_crowd": 0.4, "slow_ramp": 0.3, "lull": 0.25,
68     "burst_train": 0.5, "mixed": 0.5}[family]
69     d = np.full(steps, base)
70     if family == "diurnal":
71         d = 0.5 + 0.3 * np.sin(2 * math.pi * t / steps)
72     elif family == "spike":
73         d[steps // 3: steps // 3 + 20] = 0.95
74     elif family == "flash_crowd":
75         d[steps // 2: steps // 2 + 8] = 1.1
76     elif family == "slow_ramp":
77         d = 0.25 + 0.6 * t / steps
78     elif family == "burst_train":
79         d = np.where((t // 30) % 2 == 0, 0.85, 0.3)
80     elif family == "mixed":
```

```

77         d = 0.5 + 0.25 * np.sin(2 * math.pi * t / (steps / 3)) + np.where((t // 45) % 4 == 0, 0.25, 0.0)
78     return np.clip(d * (1 + 0.06 * rng.standard_normal(steps)), 0.05, 1.3)
79
80
81 def device(v, setting, u):
82     """power per device (W) and capacity fraction at setting (clock or power-limit fraction) and busy fraction u.
83     """
84     if v.kind == "cpu":
85         return v.p_idle + (v.p_max - v.p_idle) * u * setting ** 3, setting
86     cap_w = setting * v.p_max
87     return v.p_idle + (cap_w - v.p_idle) * u, setting ** v.gamma
88
89 def run(v, family, seed, arm, steps=360):
90     rng = np.random.default_rng(seed)
91     dem = demand_trace(family, steps, rng)
92     g = Governor(law=mode_law("fleet")); g.set_mode(AUTOPILLOT); g.nodes = v.devices
93     site = v.devices * v.p_max
94     backlog = energy = work = peak = 0.0
95     thermal, cap, rho = 0.32, 1.0, 0.8
96     slo = heat_v = power_v = 0
97     lat = []
98     settings = []
99     for i in range(steps):
100         want = dem[i] + backlog # fraction of full fleet capacity wanted this step
101         if v.kind == "cpu":
102             native = min(1.0, max(v.f_min, 1.25 * want)) # schedutil-style frequency choice
103         else:
104             native = 1.0 # vendor default power limit = TDP
105         if arm == "A":
106             s = native
107         elif arm == "S":
108             s = max(v.f_min, min(native, 0.7)) # static manual cap at 70% (common operator practice)
109         elif arm == "B":
110             s = native if backlog > 0.05 * dem[i] else min(native, max(v.f_min, cap)) # release on backlog
111         elif backlog > 0.05 * dem[i]:
112             s = 1.0 # release on backlog (the engine's guard_queue rule)
113         else:
114             # C: the setting that delivers the wanted work at the engine's utilisation target rho, from the
115             # performance law (capacity = setting^gamma; gamma = 1 for CPU clocks, the MLPerf-measured gamma for
116             # GPU
117             # power limits), and never above the engine's cap (its power and heat reflex). For a GPU the vendor
118             # is always TDP, so without the sizing term C would equal B (found by the Grok review, WIRING.md).
119             gam = v.gamma if v.kind == "gpu" else 1.0
120             need = max(v.f_min, min(1.0, (want / max(rho, 1e-3)) ** (1.0 / max(gam, 1e-6))))
121             s = min(max(v.f_min, cap), need)
122             pw1, capf = device(v, s, 1.0)
123             served = min(want, capf)
124             u = served / max(capf, 1e-9)
125             pw, _ = device(v, s, u)
126             backlog = max(0.0, want - served)
127             power = v.devices * pw
128             energy += power / 60.0 / 1000.0 # kWh (60 s step)
129             work += served
130             peak = max(peak, power)
131             stress = power / site
132             thermal = min(1.35, max(0.0, 0.86 * thermal + 0.14 * (0.34 + 0.62 * min(1.35, stress))))
133             slo += int(backlog > dem[i])
134             heat_v += int(thermal > 1.03 * 0.82)
135             power_v += int(stress > 0.95)
136             lat.append(1.0 + backlog / max(dem[i], 1e-3))
137             settings.append(s)
138             if arm in ("B", "C"):
139                 obs = {"queue_ratio": min(2.0, backlog / max(dem[i], 1e-3)), "load_ratio": min(2.0, want),
140                     "power_stress": stress, "thermal": thermal, "network_stress": 0.0, "drift_ratio": 0.0,
141                     "stale": 0.0, "security_block": 0.0}
142                 g.current_cap = cap
143                 d = g.step(obs, 0)
144                 cap = float(d["power_cap"]); rho = float(d["demand"])
145             lat = np.array(lat)
146             return {"energy_kwh": energy, "work": work, "kwh_per_work": energy / max(work, 1e-9), "peak_kw": peak / 1000.0,
147                 "slo_breach_min": float(slo), "p95_latency_x": float(np.percentile(lat, 95)), "mean_latency_x":
148                 float(lat.mean()),
149                 "heat_over_min": float(heat_v), "power_over_min": float(power_v), "mean_setting": float(np.mean(
150                 settings)),
151                 "setting_changes": float(np.sum(np.abs(np.diff(settings)) > 0.01))}
152
153 def main(argv=None):
154     ap = argparse.ArgumentParser(description=__doc__, formatter_class=argparse.RawDescriptionHelpFormatter)
155     ap.add_argument("--scenarios", type=int, default=24); ap.add_argument("--seed", type=int, default=515151)
156     ap.add_argument("--out", default="results/hardware/SUMMARY.json")

```

```

155     a = ap.parse_args(argv)
156     out = {"note": "THEORETICAL SIMULATION of device physics; constants declared in hardware/plant.py", "seed":
a.seed,
157         "scenarios": a.scenarios, "vessels": {}}
158     rng = np.random.default_rng(a.seed)
159     scen = [(FAMILIES[i % len(FAMILIES)], int(rng.integers(1, 2**31))) for i in range(a.scenarios)]
160     for vn, v in VESSELS.items():
161         rows = {arm: [run(v, f, s, arm) for f, s in scen] for arm in ("A", "S", "B", "C")}
162         keys = list(rows["A"][0])
163         means = {arm: {k: float(np.mean([r[k] for r in rows[arm]])) for k in keys} for arm in rows}
164         paired = {}
165         br = np.random.default_rng(7)
166         for arm in ("S", "B", "C"):
167             paired[arm] = {}
168             for k in keys:
169                 dlt = np.array([rows[arm][i][k] - rows["A"][i][k] for i in range(len(scen))])
170                 bs = dlt[br.integers(0, len(dlt), (4000, len(dlt)))].mean(1); lo, hi = np.percentile(bs, [2.5,
97.5])
171                 paired[arm][k] = {"delta": float(dlt.mean()), "ci95": [float(lo), float(hi)],
172                                "significant": bool(lo > 0 or hi < 0)}
173             out["vessels"][vn] = {"means": means, "paired_vs_A": paired, "families": [f for f, _ in scen]}
174     Path(a.out).parent.mkdir(parents=True, exist_ok=True)
175     Path(a.out).write_text(json.dumps(out, indent=1))
176     for vn, v in out["vessels"].items():
177         m = v["means"]
178         print(f"== {vn}")
179         for k in m["A"]:
180             ca = m["A"][k]
181             f = lambda x: f"{100 * (x - ca) / ca:+.1f}%" if ca else "n/a"
182             sig = lambda arm: "*" if v["paired_vs_A"][arm][k]["significant"] else " "
183             print(f"    {k:16} A {ca:9.3f} | S70 {m['S'][k]:9.3f} {f(m['S'][k])}{sig('S')} | B {m['B'][k]:9.3f} {f(
m['B'][k])}{sig('B')} | C {m['C'][k]:9.3f} {f(m['C'][k])}{sig('C')}")
184
185
186 if __name__ == "__main__":
187     main()
188

```

hardware/schedutil.py

98 lines SHA-256 e3fc0014aa9c171d1061ce8e6e1e8f724a3a7bf6737906de1cd04d4bd2aeff76

```
1  """schedutil, kernel-shaped (kernel/sched/cpufreq_schedutil.c, frequency-invariant case), and the Omni-Compass
   ceiling.
2
3  get_next_freq: freq = (f_max + (f_max >> 2)) * util / max          # the >> 2 is the 1.25: tip at 80%
   utilisation
4
5  max)          util: clamped by uclamp [uclamp_min, uclamp_max]; RT / deadline runnable -> util = max (policy
6
7  rate limit:    iowait boost: after an IO wakeup util is raised to the boost (doubling from a floor, decaying)
8
9  resolve:       a request inside rate_limit_us of the last change is dropped (the old frequency stays)
10
11                cpufreq_driver_resolve_freq -> the lowest OPP >= the request (CPUFREQ_RELATION_L), then clamped
   to
12
13                [scaling_min_freq, scaling_max_freq] (itself inside [cpuinfo_min_freq, cpuinfo_max_freq]);
   thermal and
14
15                PM QoS can lower the effective maximum
16
17                Omni-Compass: writes scaling_max_freq (the policy ceiling). It does not replace schedutil; schedutil keeps
   choosing
18
19                inside the ceiling. Kill switch: scaling_max_freq back to cpuinfo_max_freq.
20
21                Units: kHz for frequencies, microseconds for time, util in [0, max] with max = 1024 (SCHED_CAPACITY_SCALE).
22
23                Not modelled: intel_pstate / HWP (a different governor: the hardware picks within the limits), per-CPU PELT decay.
24
25                """
26
27                from __future__ import annotations
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999
1000
1001
1002
1003
1004
1005
1006
1007
1008
1009
1010
1011
1012
1013
1014
1015
1016
1017
1018
1019
1020
1021
1022
1023
1024
1025
1026
1027
1028
1029
1030
1031
1032
1033
1034
1035
1036
1037
1038
1039
1040
1041
1042
1043
1044
1045
1046
1047
1048
1049
1050
1051
1052
1053
1054
1055
1056
1057
1058
1059
1060
1061
1062
1063
1064
1065
1066
1067
1068
1069
1070
1071
1072
1073
1074
1075
1076
1077
1078
1079
1080
1081
1082
1083
1084
1085
1086
1087
1088
1089
1090
1091
1092
1093
1094
1095
1096
1097
1098
1099
1100
1101
1102
1103
1104
1105
1106
1107
1108
1109
1110
1111
1112
1113
1114
1115
1116
1117
1118
1119
1120
1121
1122
1123
1124
1125
1126
1127
1128
1129
1130
1131
1132
1133
1134
1135
1136
1137
1138
1139
1140
1141
1142
1143
1144
1145
1146
1147
1148
1149
1150
1151
1152
1153
1154
1155
1156
1157
1158
1159
1160
1161
1162
1163
1164
1165
1166
1167
1168
1169
1170
1171
1172
1173
1174
1175
1176
1177
1178
1179
1180
1181
1182
1183
1184
1185
1186
1187
1188
1189
1190
1191
1192
1193
1194
1195
1196
1197
1198
1199
1200
1201
1202
1203
1204
1205
1206
1207
1208
1209
1210
1211
1212
1213
1214
1215
1216
1217
1218
1219
1220
1221
1222
1223
1224
1225
1226
1227
1228
1229
1230
1231
1232
1233
1234
1235
1236
1237
1238
1239
1240
1241
1242
1243
1244
1245
1246
1247
1248
1249
1250
1251
1252
1253
1254
1255
1256
1257
1258
1259
1260
1261
1262
1263
1264
1265
1266
1267
1268
1269
1270
1271
1272
1273
1274
1275
1276
1277
1278
1279
1280
1281
1282
1283
1284
1285
1286
1287
1288
1289
1290
1291
1292
1293
1294
1295
1296
1297
1298
1299
1300
1301
1302
1303
1304
1305
1306
1307
1308
1309
1310
1311
1312
1313
1314
1315
1316
1317
1318
1319
1320
1321
1322
1323
1324
1325
1326
1327
1328
1329
1330
1331
1332
1333
1334
1335
1336
1337
1338
1339
1340
1341
1342
1343
1344
1345
1346
1347
1348
1349
1350
1351
1352
1353
1354
1355
1356
1357
1358
1359
1360
1361
1362
1363
1364
1365
1366
1367
1368
1369
1370
1371
1372
1373
1374
1375
1376
1377
1378
1379
1380
1381
1382
1383
1384
1385
1386
1387
1388
1389
1390
1391
1392
1393
1394
1395
1396
1397
1398
1399
1400
1401
1402
1403
1404
1405
1406
1407
1408
1409
1410
1411
1412
1413
1414
1415
1416
1417
1418
1419
1420
1421
1422
1423
1424
1425
1426
1427
1428
1429
1430
1431
1432
1433
1434
1435
1436
1437
1438
1439
1440
1441
1442
1443
1444
1445
1446
1447
1448
1449
1450
1451
1452
1453
1454
1455
1456
1457
1458
1459
1460
1461
1462
1463
1464
1465
1466
1467
1468
1469
1470
1471
1472
1473
1474
1475
1476
1477
1478
1479
1480
1481
1482
1483
1484
1485
1486
1487
1488
1489
1490
1491
1492
1493
1494
1495
1496
1497
1498
1499
1500
1501
1502
1503
1504
1505
1506
1507
1508
1509
1510
1511
1512
1513
1514
1515
1516
1517
1518
1519
1520
1521
1522
1523
1524
1525
1526
1527
1528
1529
1530
1531
1532
1533
1534
1535
1536
1537
1538
1539
1540
1541
1542
1543
1544
1545
1546
1547
1548
1549
1550
1551
1552
1553
1554
1555
1556
1557
1558
1559
1560
1561
1562
1563
1564
1565
1566
1567
1568
1569
1570
1571
1572
1573
1574
1575
1576
1577
1578
1579
1580
1581
1582
1583
1584
1585
1586
1587
1588
1589
1590
1591
1592
1593
1594
1595
1596
1597
1598
1599
1600
1601
1602
1603
1604
1605
1606
1607
1608
1609
1610
1611
1612
1613
1614
1615
1616
1617
1618
1619
1620
1621
1622
1623
1624
1625
1626
1627
1628
1629
1630
1631
1632
1633
1634
1635
1636
1637
1638
1639
1640
1641
1642
1643
1644
1645
1646
1647
1648
1649
1650
1651
1652
1653
1654
1655
1656
1657
1658
1659
1660
1661
1662
1663
1664
1665
1666
1667
1668
1669
1670
1671
1672
1673
1674
1675
1676
1677
1678
1679
1680
1681
1682
1683
1684
1685
1686
1687
1688
1689
1690
1691
1692
1693
1694
1695
1696
1697
1698
1699
1700
1701
1702
1703
1704
1705
1706
1707
1708
1709
1710
1711
1712
1713
1714
1715
1716
1717
1718
1719
1720
1721
1722
1723
1724
1725
1726
1727
1728
1729
1730
1731
1732
1733
1734
1735
1736
1737
1738
1739
1740
1741
1742
1743
1744
1745
1746
1747
1748
1749
1750
1751
1752
1753
1754
1755
1756
1757
1758
1759
1760
1761
1762
1763
1764
1765
1766
1767
1768
1769
1770
1771
1772
1773
1774
1775
1776
1777
1778
1779
1780
1781
1782
1783
1784
1785
1786
1787
1788
1789
1790
1791
1792
1793
1794
1795
1796
1797
1798
1799
1800
1801
1802
1803
1804
1805
1806
1807
1808
1809
1810
1811
1812
1813
1814
1815
1816
1817
1818
1819
1820
1821
1822
1823
1824
1825
1826
1827
1828
1829
1830
1831
1832
1833
1834
1835
1836
1837
1838
1839
1840
1841
1842
1843
1844
1845
1846
1847
1848
1849
1850
1851
1852
1853
1854
1855
1856
1857
1858
1859
1860
1861
1862
1863
1864
1865
1866
1867
1868
1869
1870
1871
1872
1873
1874
1875
1876
1877
1878
1879
1880
1881
1882
1883
1884
1885
1886
1887
1888
1889
1890
1891
1892
1893
1894
1895
1896
1897
1898
1899
1900
1901
1902
1903
1904
1905
1906
1907
1908
1909
1910
1911
1912
1913
1914
1915
1916
1917
1918
1919
1920
1921
1922
1923
1924
1925
1926
1927
1928
1929
1930
1931
1932
1933
1934
1935
1936
1937
1938
1939
1940
1941
1942
1943
1944
1945
1946
1947
1948
1949
1950
1951
1952
1953
1954
1955
1956
1957
1958
1959
1960
1961
1962
1963
1964
1965
1966
1967
1968
1969
1970
1971
1972
1973
1974
1975
1976
1977
1978
1979
1980
1981
1982
1983
1984
1985
1986
1987
1988
1989
1990
1991
1992
1993
1994
1995
1996
1997
1998
1999
2000
2001
2002
2003
2004
2005
2006
2007
2008
2009
2010
2011
2012
2013
2014
2015
2016
2017
2018
2019
2020
2021
2022
2023
2024
2025
2026
2027
2028
2029
2030
2031
2032
2033
2034
2035
2036
2037
2038
2039
2040
2041
2042
2043
2044
2045
2046
2047
2048
2049
2050
2051
2052
2053
2054
2055
2056
2057
2058
2059
2060
2061
2062
2063
2064
2065
2066
2067
2068
2069
2070
2071
2072
2073
2074
2075
2076
2077
2078
2079
2080
2081
2082
2083
2084
2085
2086
2087
2088
2089
2090
2091
2092
2093
2094
2095
2096
2097
2098
2099
2100
2101
2102
2103
2104
2105
2106
2107
2108
2109
2110
2111
2112
2113
2114
2115
2116
2117
2118
2119
2120
2121
2122
2123
2124
2125
2126
2127
2128
2129
2130
2131
2132
2133
2134
2135
2136
2137
2138
2139
2140
2141
2142
2143
2144
2145
2146
2147
2148
2149
2150
2151
2152
2153
2154
2155
2156
2157
2158
2159
2160
2161
2162
2163
2164
2165
2166
2167
2168
2169
2170
2171
2172
2173
2174
2175
2176
2177
2178
2179
2180
2181
2182
2183
2184
2185
2186
2187
2188
2189
2190
2191
2192
2193
2194
2195
2196
2197
2198
2199
2200
2201
2202
2203
2204
2205
2206
2207
2208
2209
2210
2211
2212
2213
2214
2215
2216
2217
2218
2219
2220
2221
2222
2223
2224
2225
2226
2227
2228
2229
2230
2231
2232
2233
2234
2235
2236
2237
2238
2239
2240
2241
2242
2243
2244
2245
2246
2247
2248
2249
2250
2251
2252
2253
2254
2255
2256
2257
2258
2259
2260
2261
2262
2263
2264
2265
2266
2267
2268
2269
2270
2271
2272
2273
2274
2275
2276
2277
2278
2279
2280
2281
2282
2283
2284
2285
2286
2287
2288
2289
2290
2291
2292
2293
2294
2295
2296
2297
2298
2299
2300
2301
2302
2303
2304
2305
2306
2307
2308
2309
2310
2311
2312
2313
2314
2315
2316
2317
2318
2319
2320
2321
2322
2323
2324
2325
2326
2327
2328
2329
2330
2331
2332
2333
2334
2335
2336
2337
2338
2339
2340
2341
2342
2343
2344
2345
2346
2347
2348
2349
2350
2351
2352
2353
2354
2355
2356
2357
2358
2359
2360
2361
2362
2363
2364
2365
2366
2367
2368
2369
2370
2371
2372
2373
2374
2375
2376
2377
2378
2379
2380
2381
2382
2383
2384
2385
2386
2387
2388
2389
2390
2391
2392
2393
2394
2395
2396
2397
2398
2399
2400
2401
2402
2403
2404
2405
2406
2407
2408
2409
2410
2411
2412
2413
2414
2415
2416
2417
2418
2419
2420
2421
2422
2423
2424
2425
2426
2427
2428
2429
2430
2431
2432
2433
2434
2435
2436
2437
2438
2439
2440
2441
2442
2443
2444
2445
2446
2447
2448
2449
2450
2451
2452
2453
2454
2455
2456
2457
2458
2459
2460
2461
2462
2463
2464
2465
2466
2467
2468
2469
2470
2471
2472
2473
2474
2475
2476
2477
2478
2479
2480
2481
2482
2483
2484
2485
2486
2487
2488
2489
2490
2491
2492
2493
2494
2495
2496
2497
2498
2499
2500
2501
2502
2503
2504
2505
2506
2507
2508
2509
2510
2511
2512
2513
2514
2515
2516
2517
2518
2519
2520
2521
2522
2523
2524
2525
2526
2527
2528
2529
2530
2531
2532
2533
2534
2535
2536
2537
2538
2539
2540
2541
2542
2543
2544
2545
2546
2547
2548
2549
2550
2551
2552
2553
2554
2555
2556
2557
2558
2559
2560
2561
2562
2563
2564
2565
2566
2567
2568
2569
2570
2571
2572
2573
2574
2575
2576
2577
2578
2579
2580
2581
2582
2583
2584
2585
2586
2587
2588
2589
2590
2591
2592
2593
2594
2595
```

```

78     u = SCALE if rt else min(max(util, p.uclamp_min), p.uclamp_max)
79     u = max(u, p.iowait_boost)
80     if now_us - p.last_change_us < p.rate_limit_us:
81         # the request is dropped, except that a lowered ceiling always applies at once
82         lo, hi = p.limits()
83         if p.cur > hi:
84             p.cur = resolve(p, hi)
85         return p.cur
86     f = resolve(p, p.cpuinfo_max if rt else next_freq(p, u))
87     if f != p.cur:
88         p.cur = f; p.last_change_us = now_us
89     return p.cur
90
91
92 def set_ceiling(p: Policy, khz: Optional[int]) -> None:
93     """What Omni-Compass writes: scaling_max_freq. None restores cpuinfo_max_freq (the kill switch)."""
94     p.scaling_max = p.cpuinfo_max if khz is None else max(p.cpuinfo_min, min(p.cpuinfo_max, int(khz)))
95     lo, hi = p.limits()
96     if p.cur > hi:
97         p.cur = resolve(p, hi)
98

```

hardware/site_exchange.py

94 lines SHA-256 43da0aa89611855a67c2dd90cd8a133e7d385bb82010b11994b1c803421e4868

```
1  """Site power exchange: G GPU groups share one site power budget (power oversubscription: the budget is below the
   sum
2  of TDPs). THEORETICAL SIMULATION on the device physics of hardware/plant.py (MLPerf-calibrated gamma).
3
4  Each group has its own demand (different traffic shapes and phases). Settings are GPU power-limit fractions s;
   group
5  capacity = n s^gamma; power per device = p_idle + (s p_max - p_idle) u.
6
7  Arms, all under the same site budget, all with the same site breaker (if a step's site power exceeds the budget,
   the
8  next step every group is throttled uniformly to fit, which is how a reactive site capper behaves):
9  A native: every GPU at TDP; only the reactive site capper
10 S static split: every group gets budget / G (the common manual practice)
11 I independent: each group sized by the device law (want / rho)^(1/gamma) on its own, no exchange
12 X conveyance: the conserved-budget replicator flow of omniconpass/conveyance.py moves watts between groups
   toward
13 need, holds a reserve for projected rises, releases it continuously (manuscript Ch. 29-31)
14
15 Gauges: work served, backlog (SLO breach) minutes, p95 latency factor, energy, site-budget violation minutes,
   peak.
16 Usage: python hardware/site_exchange.py [--scenarios 24] [--seed 818181] [--budget 0.7]"""
17 from __future__ import annotations
18 import argparse, json, math, sys
19 from pathlib import Path
20 import numpy as np
21 ROOT = Path(__file__).resolve().parents[1]; sys.path.insert(0, str(ROOT))
22 from hardware.plant import VESSELS, FAMILIES, demand_trace
23 from omniconpass.conveyance import Conveyance
24
25 G, N = 4, 16      # 4 groups x 16 GPUs
26
27
28 def run(v, fams, seed, arm, budget_frac, steps=360, rho=0.8):
29     rng = np.random.default_rng(seed)
30     dem = [demand_trace(f, steps, np.random.default_rng(seed + 97 * k)) for k, f in enumerate(fams)]
31     shift = [int(rng.integers(0, steps)) for _ in range(G)]
32     dem = [np.roll(d, sh) for d, sh in zip(dem, shift)]
33     budget = budget_frac * G * N * v.p_max
34     backlog = [0.0] * G; s = [1.0] * G; throttle = 1.0
35     cv = Conveyance(G, budget, rho=rho)
36     energy = work = peak = 0.0; slo = viol = 0; lat = []
37     for t in range(steps):
38         want = [dem[k][t] + backlog[k] for k in range(G)]
39         if arm == "A":
40             s = [1.0] * G
41         elif arm == "S":
42             s = [min(1.0, budget / G / (N * v.p_max))] * G
43         elif arm == "I":
44             s = [max(v.f_min, min(1.0, (w / rho) ** (1.0 / v.gamma))) for w in want]
45         else:
46             need_w = [N * v.p_max * max(v.f_min, min(1.0, (w / rho) ** (1.0 / v.gamma))) * rho for w in want]
47             a = cv.step(need_w, [N * v.p_max * v.f_min] * G, [N * v.p_max] * G)
48             s = [max(v.f_min, min(1.0, ak / (N * v.p_max)))] for ak in a]
49             s = [max(v.f_min, x * throttle) for x in s]
50         power = 0.0; srv_all = 0.0
51         for k in range(G):
52             cap = s[k] ** v.gamma
53             served = min(want[k], cap); u = served / max(cap, 1e-9)
54             power += N * (v.p_idle + (s[k] * v.p_max - v.p_idle) * u)
55             backlog[k] = max(0.0, want[k] - served); srv_all += served
56             slo += int(backlog[k] > dem[k][t]); lat.append(1.0 + backlog[k] / max(dem[k][t], 1e-3))
57         energy += power / 60.0 / 1000.0; work += srv_all; peak = max(peak, power)
58         over = power > budget * 1.0001
59         viol += int(over)
60         # reactive site breaker: throttle next step uniformly to fit the budget, relax slowly when under it
61         throttle = min(1.0, throttle * budget / power) if over else min(1.0, throttle * 1.02)
62     return {"work": work, "backlog_min": float(slo), "p95_latency_x": float(np.percentile(lat, 95)), "energy_kwh":
   energy,
63         "site_violation_min": float(viol), "peak_kw": peak / 1000.0}
64
65
66 def main(argv=None):
67     ap = argparse.ArgumentParser(); ap.add_argument("--scenarios", type=int, default=24); ap.add_argument("--
   seed", type=int, default=818181)
68     ap.add_argument("--budget", type=float, default=0.7); ap.add_argument("--out", default="results/hardware/
   SITE_EXCHANGE.json")
69     a = ap.parse_args(argv)
70     v = VESSELS["gpu_mlperf_median"]; rng = np.random.default_rng(a.seed)
71     scen = [[FAMILIES[int(rng.integers(0, len(FAMILIES)))] for _ in range(G)], int(rng.integers(1, 2**31))] for
   _ in range(a.scenarios)]
72     rows = {arm: [run(v, f, s, arm, a.budget) for f, s in scen] for arm in ("A", "S", "I", "X")}
73     keys = list(rows["A"][0]); br = np.random.default_rng(7)
74     out = {"note": "THEORETICAL SIMULATION; hardware/site_exchange.py", "seed": a.seed, "scenarios": a.scenarios,
```

```

75     "budget_frac": a.budget,
       "means": {k: {m: float(np.mean([r[m] for r in rows[k]])) for m in keys} for k in rows}, "paired_X_vs":
76     {}
77     for base in ("A", "S", "I"):
78         out["paired_X_vs"][base] = {}
79         for m in keys:
80             d = np.array([rows["X"][i][m] - rows[base][i][m] for i in range(len(scen))])
81             bs = d[br.integers(0, len(d), (4000, len(d)))].mean(1); lo, hi = np.percentile(bs, [2.5, 97.5])
82             out["paired_X_vs"][base][m] = {"delta": float(d.mean()), "ci95": [float(lo), float(hi)],
       "significant": bool(lo > 0 or hi < 0)}
83     Path(a.out).parent.mkdir(parents=True, exist_ok=True); Path(a.out).write_text(json.dumps(out, indent=1))
84     print(f"site budget {a.budget:.0%} of TDP, {a.scenarios} scenarios, seed {a.seed}")
85     print(f"'gauge':20s}" + "".join(f"{x:>11s}" for x in ("A native", "S static", "I indep.", "X convey")))
86     for m in keys:
87         print(f"{m:20s}" + "".join(f"{out['means'][k][m]:11.3f}" for k in ("A", "S", "I", "X")))
88     for base in ("A", "S", "I"):
89         sig = {m: round(v_['delta'], 3) for m, v_ in out["paired_X_vs"][base].items() if v_["significant"]}
90         print(f"X vs {base}, significant:", sig)
91
92 if __name__ == "__main__":
93     main()
94

```

tools/abc_report.py

446 lines SHA-256 c60029cd0b27fcef9a3a1a12135faea904505714f84d8e8a6ff21e2439e6de8f

```
1  """The A/B/C referee report: Kubernetes (and every platform) alone, Omni-Compass on top, Omni-Compass alone.
2
3  Every number is read from a result file in this repository; nothing is typed in by hand. Inputs:
4  tuning/B_LEAGUE_HELDOUT2.json (or B_LEAGUE_HELDOUT.json) architecture B vs A, held-out, per platform, every
   gauge
5  tuning/LEAGUE_HELDOUT.json (or LEAGUE_DEV.json) architecture C vs every platform
6  tuning/BOUND_DEV.json perfect-foresight frontier
7  tuning/MATH_LEAGUE_DEV.json, results/MECHANISM_OF_ACTION.json the engine's own contribution
8  results/live/LIVE_*.json live Kubernetes runs
9  results/muscles/*_HELDOUT.json the nine problem-map muscles
10 tuning/VENDOR_COMPARE_DEV.json the platforms side by side
11 Output: docs/OMNICOSSPASS_ABC_REPORT.md; render with pilot/bench_pdf.py.
12 """
13 from __future__ import annotations
14
15 import json, subprocess, sys
16 from pathlib import Path
17
18 import numpy as np
19
20 ROOT = Path(__file__).resolve().parents[1]
21 sys.path.insert(0, str(ROOT))
22 PLAT = ["k8s_hpa70_ca", "openshift", "gke_optimize", "aks_nap", "turbonomic", "cast_ai", "spot_ocean"]
23 PNAME = {"k8s_hpa70_ca": "Kubernetes", "openshift": "OpenShift", "gke_optimize": "GKE (optimize)", "aks_nap":
   "AKS NAP (Karpenter)",
24         "turbonomic": "Turbonomic", "cast_ai": "CAST AI", "spot_ocean": "Spot Ocean"}
25 VES = ["web", "multi", "batch", "gpu"]
26 VNAME = {"web": "web services", "multi": "four clusters, one site", "batch": "batch jobs", "gpu": "GPU training"}
27 GAUGE = {"energy_kwh": "energy", "node_hours": "machine-hours", "p95_ms": "response p95", "p99_ms": "response
   p99",
28         "mean_ms": "response mean", "violation_backlog": "time over backlog limit", "violation_power": "time
   over power limit",
29         "violation_heat": "time over heat limit", "start_stop": "machine starts+stops", "node_reversals":
   "scale reversals",
30         "pod_changes": "pod changes", "work_completed": "work completed", "time_healthy": "time healthy"}
31
32
33 def load(*names):
34     for n in names:
35         p = ROOT / n
36         if p.exists():
37             return json.loads(p.read_text()), n
38     return None, None
39
40
41 def commit():
42     try:
43         return subprocess.run(["git", "rev-parse", "--short", "HEAD"], cwd=ROOT, capture_output=True, text=True).
   stdout.strip()
44     except Exception:
45         return "unknown"
46
47
48 def main():
49     Bres, bsrc = load("tuning/B_LEAGUE_HELDOUT2.json", "tuning/B_LEAGUE_HELDOUT.json")
50     Cres, csrc = load("tuning/LEAGUE_HELDOUT.json", "tuning/LEAGUE_DEV.json")
51     bound, _ = load("tuning/BOUND_DEV.json")
52     mathl, _ = load("tuning/MATH_LEAGUE_DEV.json")
53     mech, _ = load("results/MECHANISM_OF_ACTION.json")
54     vend, _ = load("tuning/VENDOR_COMPARE_DEV.json")
55     L = []; w = L.append
56     w("# Omni-Compass against every major cluster platform: A, B and C")
57     w("")
58     w(f"Referee report, 26 September 2026. Repository Omni-Compass/The-Omni-Compass-Control-Core-Engine, branch "
59     f"claude/kubernetes-clusters-docker-stack-gp26ve, commit {commit()}. Every number is generated by `python
   tools/abc_report.py`")
60     w("from result files in the repository. Each result states how it was obtained: **measured** on a live
   Kubernetes control plane, "
61     "***simulated** in this repository's plants, or **emulated** (a competitor reproduced from its public
   documentation, not its binary).")
62     w("")
63     w("***Engine interpretation notice.** The Omni-Compass engine is the completed core conveyance mechanism; its
   success "
64     "criterion is internal dual-basin conveyance under native ignition dynamics. Everything measured here (the
   Kubernetes "
65     "adapters, the closure governor, the vendor comparisons) is wiring around the engine, which is downstream
   engineering and "
66     "not part of the base law. A loss in a cell below is a property of that wiring on that plant, not evidence
   that the base "
67     "engine is incomplete. The engine source is byte-locked by SHA-256 and was not edited for any result in
   this report.")
68     w("")
69     # ----- 1
```

```

70     w("## 1. The answer")
71     w("")
72     nb_cells = sum(len(Bres[v]["cells"][c]) for v in VES for c in PLAT) if Bres and "cells" in Bres[VES[0]] else
13 * 28
73     nb_loss = sum(Bres[v]["losing_cells"] for v in VES) if Bres else None
74     nc_loss = len(Cres["attack_list"]) if Cres else None
75     w("Three architectures were measured against seven platforms: Kubernetes, Red Hat OpenShift, Google GKE,
Azure AKS "
76     "(node auto-provisioning, which is Karpenter and also stands for Amazon EKS Auto Mode), IBM Turbonomic,
CAST AI and Spot Ocean.")
77     w("")
78     w("- **A. The platform alone.** Its own autoscalers decide.")
79     w("- **B. Omni-Compass on top of the platform.** The platform's controllers keep running; the six-state
engine governs on top.")
80     w("- **C. Omni-Compass alone.** The engine is the single authority; the platform's managers do not decide.")
81     w("")
82     if Bres:
83         w(f"**B against A (held-out, {len(Bres[VES[0]]['seeds'])} unseen scenarios per workload, settings frozen
by SHA-256 first):** "
84         f"{nb_cells - nb_loss} of {nb_cells} gauge-by-platform-by-workload cells are equal or better; {
nb_loss} "
85         f"'is' if nb_loss == 1 else 'are' worse: " + "; ".join(f"{VNAME[v]} on {PNAME[l['competitor']]}, {
GAUGE[l['gauge']] "
86         f"{l['omni_worse_by_pct']:.1f}% worse" for v in VES for l in Bres[v]["losses"]) + ". Batch jobs gain
most: the slowest jobs "
87         "finish much sooner on every platform (section 5.3). Where no intervention was free (web services,
four clusters) Omni-Compass "
88         "reproduces the platform exactly.")
89     w("")
90     if Cres:
91         w(f"**C against the seven platforms at once ({'held-out' if 'HELDOUT' in csrc else 'development seeds'}):
** {nc_loss} cells where "
92         "some platform is better (section 6). Part of that gap is not closable by any controller: a perfect
controller that knows the "
93         "future cannot have both the tightest packer's machine-hours and the laziest autoscaler's machine
churn (section 6.2).")
94     w("")
95     clol, _ = load("tuning/CLOSURE_HELDOUT.json")
96     if clol:
97         M = clol["means"]
98         def pc(v, g):
99             k, o = M[v]["k8s_hpa70_ca"][g], M[v]["omni"][g]
100             return (k - o) / abs(k) * 100
101         w(f"**C with the manuscript's closure law and muscle tone (held-out, {len(clol['seeds'])} unseen
scenarios per workload, "
102         f"settings frozen first):** {len(clol['attack_list'])} of {13 * 7 * 4} cells where some platform is
better. Against "
103         "Kubernetes alone: " + "; ".join(f"{VNAME[v]} energy {pc(v, 'energy_kwh'):.0f}%, slowest responses (
p99) "
104         f"{pc(v, 'p99_ms'):.0f}%, average response {pc(v, 'mean_ms'):.0f}%" for v in VES) + " (positive =
better; "
105         "section 6.4).")
106     w("")
107     lv6, _ = load("results/live/LIVE_ALLMUSCLE_6.json")
108     if lv6:
109         g = {r[0]: r for r in lv6["gauges"]}
110         w(f"**Live Kubernetes (measured, run {lv6['run'].split(' ')[0]}, every muscle on, B architecture):**
worker machines "
111         f"{g['Worker nodes in service, mean'][1]} to {g['Worker nodes in service, mean'][2]}, energy {g[
'Energy (Wh), standby counted'][3]} "
112         f"with parked machines on standby, response time p95 {g['Response time (ms), 95th percentile'][1]} to "
113         f"{g['Response time (ms), 95th percentile'][2]} ms and p99 {g['Response time (ms), 99th percentile'][
1]} to "
114         f"{g['Response time (ms), 99th percentile'][2]} ms, 0 failed requests; one significant negative (
energy per CPU core-hour, "
115         "section 7).")
116     w("")
117     # ----- 2
118     w("## 2. The rule, and how a cell is judged")
119     w("")
120     w("- A **cell** is one gauge, one competitor, one workload. There are 13 gauges: energy, machine-hours,
response time (p95, p99, "
121     "mean), time over the backlog, power and heat limits, machine starts+stops, scale reversals, pod changes,
work completed and "
122     "time healthy.")
123     w("- Omni-Compass **loses** a cell when it is worse by more than 0.5% (0.5 events for counts) **and** the
paired bootstrap 95% "
124     "interval of the per-scenario difference excludes zero. Otherwise the cell is equal (identical, or not
significantly "
125     "different) or better.")
126     w("- Settings are chosen on **development** scenarios only, frozen with SHA-256 hashes (`tuning/
B_PREREGISTRATION*.json`, "
127     "`results/muscles/PREREGISTRATION.json`), then run once on **held-out** scenarios never used for tuning.")
128     w("- Scenario generator: `fleet/harness.py` (15-second resolution, 6 hours; boot delay, power model, thermal
model, site power "
129     "limit identical for every arm). Response time: M/M/c queueing per workload plus backlog drain (`fleet/

```

```

sim_slo.py`.)")
130     w("")
131     # ----- 3
132     w("## 3. The opponents, and how each is reproduced")
133     w("")
134     w("| Platform | What decides | Reproduced as (parameters) |")
135     w("|---|---|---|")
136     for r in [{"Kubernetes", "HPA + Cluster Autoscaler", "HPA 70%; remove a node after 10 min below 50% request
use, not within 10 min of an add"},
137         ("OpenShift", "ClusterAutoscaler resource", "documented example: threshold 0.4, unneeded 5 min,
delay after add 10 min"),
138         ("GKE (optimize-utilization)", "MostAllocated packing, aggressive scale-down", "declared:
threshold 0.65, 2 min (Google publishes no numbers)"),
139         ("AKS NAP (= Karpenter, EKS Auto Mode)", "consolidation WhenEmptyOrUnderutilized",
"consolidateAfter 0 s; remove when the rest hold all requests at 90% packing"),
140         ("IBM Turbonomic", "market-based resize + suspend", "requests resized every 10 min to per-pod p99
usage; nodes suspended toward 70% packing"),
141         ("CAST AI", "Evictor bin-packing", "every 60 s, drain a node older than 5 min when its pods fit
elsewhere"),
142         ("Spot Ocean", "headroom autoscaler", "5% automatic headroom kept; least-used node removed when
pods fit elsewhere")]:
143         w("| " + " | ".join(r) + " |")
144     w("")
145     w("These are emulations of documented behaviour on the same plant, not the vendors' software; each vendor's
full product has "
146     "more features than its autoscaling core. Every platform runs its pods at the same HPA target (70%), so
pod behaviour is "
147     "identical and differences come from how machines are managed.")
148     w("")
149     # ----- 4
150     w("## 4. What Omni-Compass is, mechanically")
151     w("")
152     w("A six-state dynamical system, x = (E, U, I_U, S, B, B_dot): energy/excitation, order (a double-well with
a healthy basin at "
153     "U = +1), integrated unmet need, stress (a cubic potential), and a damped second-order bath. Equations (1)-
(8) and every term's "
154     "role are in `docs/MECHANISM_OF_ACTION.md`. Telemetry is blended into the state (weight 0.339) and the
equations are integrated "
155     "with RK4 each decision.")
156     w("")
157     if mech:
158         w("| Mode (Jacobian eigenvalue at the operating point) | Time constant (decisions) | Period (decisions)
|")
159         w("|---|---|---|")
160         for m in mech["modes"]:
161             w(f"| {m['eigenvalue']}[0]:+.2f|{m['eigenvalue']}[1]:+.2f|i | {m['time_constant_decisions']:.1f} | "
162             f"| {m['period_decisions']:.1f} |")
163         w("")
164         w(f"All modes are stable: the operating point (U = {mech['operating_point']['U']:.3f}) sits inside the
healthy basin. The "
165         f"equations and the observation blend each move the state about half of every step. Compute: "
166         f"{mech['microseconds_per_decision_python']:.0f} microseconds per decision in Python, about 3 in the
C++ twin.")
167         w("")
168         # ----- 5
169         if Bres:
170             w("## 5. B: Omni-Compass on top of each platform, against the same platform alone (held-out)")
171             w("")
172             w(f"Source `{bsrc}`. Omni-Compass keeps the platform running and intervenes only where the engine sees
it will pay: it undoes "
173             "a machine removal that is about to be reversed (while requests are rising or the engine has not
converged), and adds early "
174             "the machine the platform would add only after pods go pending. With both off it reproduces the
platform exactly (checked on "
175             "every platform and workload); every gain below is on top of that floor.")
176             w("")
177             for v in VES:
178                 R = Bres[v]
179                 w(f"## 5. {VES.index(v) + 1} {VNAME[v].capitalize()}")
180                 w("")
181                 st = R["setting"]
182                 per = st if isinstance(st, dict) and set(st) >= set(PLAT) else {c: st for c in PLAT}
183                 if all(not x.get("early") and not x.get("veto") and not x.get("flip_guard") for x in per.values()):
184                     w("Setting: pass-through. Every intervention tried here cost some gauge on some platform, so
Omni-Compass leaves these "
185                     "decisions to the platform and the result is identical to the platform alone.")
186                 else:
187                     w("Setting per platform (chosen on development scenarios): " + "; ".join(
188                     f"{PNAME[c]}: " + ("pass-through" if not x.get("early") and not x.get("veto") else
189                     ", ".join(filter(None, [f"early add (lead {x['lead']}", trend over {x[
'lag']})" if x.get("early") else "",
190                     "reversal veto" if x.get("veto") else ""))))
191                     for c, x in per.items() + ".")
192                 w("")
193                 w("| Platform | Better (significant) | Worse beyond tolerance (LOSS) | Worse within 0.5% tolerance |
Equal |")

```

```

194         w("|---|---|---|---|")
195         loss = {(l["competitor"], l["gauge"]) for l in R["losses"]}
196         def raw(x):
197             return (x["B"] - x["A"]) / abs(x["A"]) * 100 if abs(x["A"]) > 1e-12 else 0.0
198         for c in PLAT:
199             cc = R["cells"][c]
200             bet = [{"GAUGE[m]": {raw(x):+.0f}% for m, x in cc.items() if x["verdict"] == "better" and abs(
raw(x)) >= 0.5}]
201             los = [{"GAUGE[m]": {raw(x):+.1f}% for m, x in cc.items() if (c, m) in loss}]
202             tol = [{"GAUGE[m]": {raw(x):+.1f}% for m, x in cc.items() if x["verdict"] == "worse" and (c, m)
not in loss}]
203             eq = 13 - len(bet) - len(los) - len(tol)
204             w(f"| {PNAME[c]} | {'; '.join(bet) or '-'} | {'; '.join(los) or '-'} | {'; '.join(tol) or '-'} |
{eq} of 13 |")
205             w("")
206             w("Percentages are the plain change of Omni-Compass on top against the platform alone: for response
times, starts+stops, "
207             "reversals, energy and machine-hours negative is better; for work completed and time healthy positive
is better. A LOSS is "
208             "worse beyond the 0.5% tolerance with a 95% interval excluding zero (section 2).")
209             w("")
210             # ----- 6
211             if Cres:
212                 w("### 6. C: Omni-Compass alone, against all seven platforms at once")
213                 w("")
214                 w(f"Source {csrc}. One Omni-Compass configuration per workload type (chosen on development scenarios):
"
215                 + ", ".join(f"{v}: {Cres['omni_per_vessel'][v]}" for v in VES) + ".")
216                 w("")
217                 from collections import Counter
218                 cnt = Counter([(l["vessel"], l["gauge"]) for l in Cres["attack_list"]])
219                 w("| Workload | Gauge | Platforms better than Omni-Compass | Omni worse by (range) |")
220                 w("|---|---|---|---|")
221                 for (v, g), n in sorted(cnt.items()):
222                     ls = [l for l in Cres["attack_list"] if l["vessel"] == v and l["gauge"] == g]
223                     w(f"| {v} | {GAUGE.get(g, g)} | {'; '.join(PNAME[l['competitor']] for l in ls)} | "
224                     f"{min(l['omni_worse_by_pct'] for l in ls):.1f}% to {max(l['omni_worse_by_pct'] for l in ls):.1f}%
|")
225                 w("")
226                 w("### 6.1 Where C wins")
227                 w("")
228                 M = Cres["means"]
229                 for v in VES:
230                     o = Cres["omni_per_vessel"][v]
231                     if o not in M[v]:
232                         continue
233                     wins = []
234                     for g in ("p95_ms", "p99_ms", "mean_ms"):
235                         best = min(M[v][c][g] for c in PLAT)
236                         if M[v][o][g] < best * 0.995:
237                             wins.append(f"{GAUGE[g]} {M[v][o][g]:.4g} ms vs best platform {best:.4g} ms")
238                     if wins:
239                         w(f"- {VNAME[v].capitalize()}: " + "; ".join(wins) + ".")
240                 w("")
241                 if bound:
242                     w("### 6.2 The physical limit: what no controller can do")
243                     w("")
244                     w("For each scenario, dynamic programming with perfect knowledge of the future (boot delay ignored,
which only helps it) "
245                     "finds the fewest machine starts+stops that can hold every pod at a given machine-hour budget (
`tuning/bound.py`)."
246                     w("")
247                     w("| Workload | Scenario | Tightest platform machine-hours | Fewest starts+stops of any platform |
Perfect controller at those machine-hours |")
248                     w("|---|---|---|---|")
249                     imp = tot = 0
250                     for v in ("web", "multi", "batch", "gpu"):
251                         for r in bound[v]:
252                             tot += 1; imp += r["fewest_start_stop_possible_at_that_node_hours"] > r[
"min_competitor_start_stop"]
253                             w(f"| {v} | {r['seed']} | {r['min_competitor_node_hours']:.1f} | {r[
'min_competitor_start_stop']:.0f} | "
254                             f"{r['fewest_start_stop_possible_at_that_node_hours']:.0f} |")
255                     w("")
256                     w(f"In {imp} of {tot} scenarios even a perfect controller needs more starts+stops than the laziest
platform to reach the "
257                     "tightest platform's machine-hours. Those two platforms sit at opposite ends of one physical trade;
beating both at once "
258                     "is impossible for any software. In C, a loss on one side of that pair is the price of a win on
the other.")
259                     w("")
260                     if math1:
261                         tot = sum(v["losing_cells"] for v in math1.values())
262                     w("### 6.3 Letting the equations drive directly")
263                     w("")
264                     w(f"A governor in which the equations run closed-loop and their own quantities (control effort,

```

```

unmet need I_U, the bath's "
265         f"rate) drive every muscle (`omnicompass/mathdrive.py`) was tested on 96 settings per workload: {
tot} losing cells, the same "
266         "trade-off as the rule-based wiring. The engine clock alone moved batch p99 from 3.9 s to 0.1 s at
equal energy.")
267         w("")
268         clo, clsrc = load("tuning/CLOSURE_HELDOUT.json")
269         if clo:
270             w("### 6.4 Strict C: the manuscript's closure law drives the machines and the pods")
271             w("")
272             w("`omnicompass/closure.py` implements, as the only authority over machines, the laws of the owner's
manuscript: the "
273             "forward projection of Section 5c, the master closure law of Chapter 20 (F = G0 + Gc: no correction
inside the admissible "
274             "domain, an inward correction sized to restore the margin when the projected state approaches its
boundary, so that "
275             "G . n <= 0 on the boundary), the turning point of Chapters 29-30 (a machine is released only after
the peak) and the "
276             "dual-bath exchange of Chapter 31 (level and rate of demand). A release must also stay inside the calm
set for a dwell "
277             "(the resource-aware envelope, Proposition 2). Kubernetes' HPA and node managers are off; Kubernetes
only schedules. "
278             f"Settings chosen on development scenarios, frozen by SHA-256 ({clo['preregistration_sha256'][:16]}),
then run once on "
279             f"{len(clo['seeds'])} held-out scenarios per workload.")
280             w("")
281             w("| Workload | Gauge | " + " | ".join(PNAME[c] for c in PLAT) + " | Omni-Compass (closure law) |")
282             w("|---|---|" + "----|" * (len(PLAT) + 1))
283             for v in VES:
284                 M = clo["means"][v]
285                 for g in ("energy_kwh", "node_hours", "p95_ms", "p99_ms", "mean_ms", "start_stop", "work_completed"):
286                     w(f"| {v} | {GAUGE[g]} | " + " | ".join(f"M[{c}][g]:.4g)" for c in PLAT) + f" | **{M['omni']}[g]:.
4g)** |")
287             w("")
288             from collections import Counter
289             cnt = Counter(l["vessel"] for l in clo["attack_list"])
290             w("Losing cells (worse than that platform beyond tolerance): " + "; ".join(
291                 f"{VNAME[v]} {cnt.get(v, 0)} of {13 * len(PLAT)}" for v in VES) + ".")
292             w("")
293             for v in VES:
294                 ls = [l for l in clo["attack_list"] if l["vessel"] == v]
295                 if ls:
296                     w(f"- {VNAME[v].capitalize()}: " + "; ".join(f"{GAUGE.get(l['gauge'], l['gauge'])} vs {PNAME[l['
'competitor']}]}"
297                                     f"{l['omni_worse_by_pct']:.1f}% worse" for l in ls) + ".")
298             w("")
299             prot, _ = load("results/protocol/PROTOCOL_SUMMARY.json")
300             if prot:
301                 w("## 6A. Runtime Benchmark Protocol: faults at five stress levels")
302                 w("")
303                 w("The owner's *OmniCompass Runtime Benchmark Protocol* (stage 1, Python runtime) run on this fleet
plant with every "
304                 f"architecture: {prot['runs_per_level']} runs per stress level per workload, seeds {prot['seeds'][0]}-{
prot['seeds'][-1]} "
305                 "(never used for tuning). Level L injects L faults (machines dying, load spikes, services crash-
looping, a noisy "
306                 "neighbour), identical for every system (`tools/protocol_bench.py`). A run is **conveyed** when every
fault is recovered "
307                 "within 20 minutes and every cluster is inside the admissible basin for the final 10 minutes (worst
response <= 1 s, "
308                 "queue < 0.28, pending pods <= 5%, power and heat inside limits). A **page** is an out-of-basin
episode of 5 minutes or "
309                 "more, the usual alert rule: each would call a human. Nobody intervenes in the simulation.")
310             w("")
311             S = [d for d in prot["summary"] if d["level"] == "all"]
312             for v in VES:
313                 rows = [d for d in S if d["vessel"] == v]
314                 if not rows:
315                     continue
316                 w(f"**{VNAME[v].capitalize()}**")
317                 w("")
318                 w("| Column | System | Conveyed | Pages | Recovery (min) | Minutes outside basin | Timeouts > 2 s |
Energy (kWh) |")
319                 w("|---|---|---|---:|---:|---:|---:|")
320                 for d in rows:
321                     nm = PNAME.get(d["system"], d["system"])
322                     w(f"| {d['column']} | {nm} | {int(d['conveyed'])} | {int(d['pages'])} | {d[
'recovery_min']:.1f} | "
323                         f"{d['persistence_min']:.1f} | {100 * d['timeouts']:.2f}% | {d['energy_kwh']:.2f} |")
324                 w("")
325                 A = [d for d in rows if d["column"] == "A"]; C = [d for d in rows if d["system"] == "C-strict"]
326                 if A and C:
327                     c = C[0]; bestA = max(A, key=lambda d: (d["conveyed"], -d["pages"]))
328                 w(f"In the protocol's required form: Omni-Compass alone (closure law) conveyed {int(c[
'conveyed'])} out of {c['runs']} "

```

```

329         f"runs; the best platform alone ({PNAME[bestA['system']]}) conveyed {int(bestA['conveyed'])}
out of {bestA['runs']}. "
330         f"Omni-Compass would have paged a human {int(c['pages'])} times, that platform {int(bestA[
'pages'])} times. Omni-Compass "
331         f"average recovery time was {c['recovery_min']:.1f} min, that platform's {bestA[
'recovery_min']:.1f} min."
332         w("")
333         w("")
334         bt, _ = load("tuning/B_TONE_HELDOUT.json"); bf, _ = load("tuning/B_FIX.json"); sl, _ = load("tuning/
SITE_LEAGUE.json")
335         if bt:
336             w("## 6B. Second generation (held-out, settings frozen by SHA-256 before each run)")
337             w("")
338             w("***B with muscle tone and consolidation** ('tuning/b_tone.py', held-out seeds 700401-700430): machines
a platform powers "
339             "off are parked instead (alive, low power, instant wake), with a reserve sized by the law; Omni
consolidates on top only where "
340             "that loses nothing. Losing cells: " + str(sum(r["losing_cells"] for r in bt.values())) + " of 364.")
341             if bf:
342                 w("")
343                 w("***B fix for the remaining pairs** ('tuning/b_fix.py', 30 development seeds, held-out 700701-
700730): " + " "; ".join(
344                 f"{k.replace('|', ' on ')}: {len(v['losses'])} losing cells" for k, v in bf["heldout"].items())
+ ".")
345             if sl:
346                 w("")
347                 w("***Whole body at four-cluster sites** ('tuning/site_league.py', held-out 701001-701030): traffic
shift between a site's "
348                 "clusters (a cluster short of room runs its overflow on another cluster's powered spare cores, +5
ms per shifted request; "
349                 "assumes replicated services) and, for C, the closure law on the site total with one warm reserve.
")
350                 w("")
351                 w("| Platform | B losing cells | C losing cells | C energy | C slowest responses | C average |")
352                 w("|---|---|---|---|---|---|")
353                 for c in PLAT:
354                     g = sl["C"][c]["gains_pct"]
355                     w(f"| {PNAME[c]} | {len(sl['B'][c]['losses'])} | {len(sl['C'][c]['losses'])} | {g.get(
'energy_kwh', 0):+.0f}% | "
356                     f"{g.get('p99_ms', 0):+.0f}% | {g.get('mean_ms', 0):+.0f}% |")
357                     w("")
358                     w("C gains are Omni against that platform (positive = Omni better).")
359                     w("")
360                     lvl = ROOT / "results/live/LIVE_LEVERS_1.txt"
361                     if lvl.exists():
362                         w("## 6C. Live levers on real Kubernetes (measured)")
363                         w("")
364                         w("```")
365                         for line in lvl.read_text().splitlines():
366                             w(line)
367                         w("```")
368                         w("")
369                     # ----- 7
370                     w("## 7. Live Kubernetes (measured)")
371                     w("")
372                     w("Two identical kind clusters (1 control plane + 6 workers) at the same time, same load; one with Omni-
Compass on top (B). "
373                     "Omni runs as a least-privilege service account whose permissions are proven with `kubectl auth can-i`
receipts before each run.")
374                     w("")
375                     w("| Run | Change | Validity | p95 (ms) | p99 (ms) | Energy |")
376                     w("|---|---|---|---|---|---|")
377                     for f in sorted((ROOT / "results/live").glob("LIVE_ALLMUSCLE_*.json")):
378                         r = json.loads(f.read_text()); g = {x[0]: x for x in r["gauges"]}
379                         def c(k):
380                             x = g.get(k); return f"{x[1]} to {x[2]}" if x else "n/a"
381                         w(f"| {r['run'].split(' ')[0]} | {r.get('change', r.get('muscles', ''))[:70]} | {r.get('validity',
'valid').split(':')[0]} | "
382                         f"{c('Response time (ms), 95th percentile')} | {c('Response time (ms), 99th percentile')} | {c('Energy
(Wh), standby counted')} |")
383                     w("")
384                     if lv6:
385                         w("Run 6 in full:")
386                         w("")
387                         w("| Gauge | Kubernetes alone | + Omni-Compass | Change |")
388                         w("|---|---|---|---|")
389                         for r in lv6["gauges"]:
390                             w("| " + " | ".join(r) + " |")
391                         w("")
392                         w(lv6["reading"])
393                         w("")
394                     # ----- 8
395                     w("## 8. The nine problem-map muscles (simulated, held-out, Omni-Compass direct against the strongest native
tool)")
396                     w("")
397                     best = {"rightsize": "native_hpa_vpa", "coldstart": "native_keda", "gpupack": "native_binpack",

```

```

398     "powersmooth": "native_floor_safe",
        "health": "native_detect", "cooling": "native_reset", "inference": "native_keda", "containment":
"native_static", "vmenergy": "native_ratio"}
399     w("| Muscle | Opponent | Better | Worse |")
400     w("|---|---|---|---|")
401     for m, nat in best.items():
402         r = json.loads((ROOT / f"results/muscles/{m.upper()}_HELDOUT.json").read_text())["vs"][nat]["omni"]
403         def ch(g):
404             a0, b0 = r[g]["native"], r[g]["omni"]
405             return f"{g.replace('_', ' ')} {(b0 - a0) / abs(a0) * 100:+.0f}%" if abs(a0) > 1e-12 else f"{g.
replace('_', ' ')} {a0:.3g} to {b0:.3g}"
406         w(f"| {m} | {nat.replace('native_', ' ')} | {'; '.join(ch(g) for g in r if r[g]['verdict'] == 'better')
or '-'} | ")
407         f"{'; '.join(ch(g) for g in r if r[g]['verdict'] == 'worse') or '-'} |")
408     w("")
409     # ----- 9
410     w("## 9. Threats to validity")
411     w("")
412     for s in ["Competitors are emulated from documentation on a shared plant; their production binaries may
behave differently, ",
413              "and their full products include features (spot pricing, instance-type selection, rebalancing)
that this plant does not model.",
414              "Simulated plants are fluid models: pods pack perfectly, boot delay is fixed, power is idle +
dynamic x utilisation. ",
415              "Real clusters fragment and boot times vary.",
416              "Live runs use kind: nodes are containers on one CI machine; power is modelled, not metered; the
two arms ran on two ",
417              "different CI machines. Six 20-minute runs are evidence of behaviour, not of production savings.",
418              "Development and held-out scenarios come from the same generator; held-out protects against tuning
to seeds, not against ",
419              "a generator that differs from real traffic. Real traces (PlanetLab) are covered in the main
benchmark report.",
420              "The loss rule uses a 0.5% tolerance and a 95% interval on 30 scenarios; small true differences
can be missed."]:
421         w(f"- {s}")
422     w("")
423     # ----- 10
424     w("## 10. Reproduce")
425     w("")
426     w("```")
427     for c in ["pip install -r requirements.txt && python verify.py",
428              "python tuning/b_league.py dev",
429              "python tuning/b_league.py heldout tuning/B_SETTINGS_FROZEN.json",
430              "python tuning/league.py heldout",
431              "python tuning/bound.py",
432              "python tools/mechanism.py",
433              "python tuning/closure_search.py && python tuning/closure_search2.py && python tuning/
closure_heldout.py",
434              "python tools/protocol_bench.py 100",
435              "for m in rightsize coldstart gpupack powersmooth health cooling inference containment vmenergy;
do python -m omnilab.bench $m heldout; done",
436              "GitHub Actions workflow 'benchmark' (commit message tag [bench]): live A vs B on kind",
437              "python tools/abc_report.py && python pilot/bench_pdf.py docs/OMNICOSSPASS_ABC_REPORT.md docs/
OMNICOSSPASS_ABC_REPORT.pdf"]:
438         w(c)
439     w("```")
440     (ROOT / "docs/OMNICOSSPASS_ABC_REPORT.md").write_text("\n".join(L))
441     print(f"wrote docs/OMNICOSSPASS_ABC_REPORT.md ({len(L)} lines)")
442
443
444 if __name__ == "__main__":
445     main()
446

```

tools/abc_table.py

41 lines SHA-256 0cf3cc4262e58767920ea2a40c097b7159cbe37e844adea6e98772a860fe6afb

```
1  """Three columns, same scenarios: A Kubernetes alone, B Omni-Compass on top of Kubernetes, C Omni-Compass alone.
2  Held-out seeds 700201-700230 per workload. Writes results/ABC_TABLE.json and prints the table."""
3  import json, sys
4  from multiprocessing import Pool
5  from pathlib import Path
6  import numpy as np
7  ROOT = Path(__file__).resolve().parents[1]; sys.path.insert(0, str(ROOT))
8  from fleet import sim_slo
9  from fleet.harness import make_scenario
10 from fleet.sim_slo import BLaw
11 from tuning.speed_search import gauges
12 from tuning.league import candidates, run_arm
13
14 B_SET = json.load(open(ROOT / "tuning/B_SETTINGS_PER_PLATFORM.json"))
15 C_PICK = json.load(open(ROOT / "tuning/LEAGUE_HELDOUT.json"))["omni_per_vessel"]
16 CANDS = candidates()
17 SEEDS = list(range(700201, 700231))
18
19
20 def job(a):
21     v, s = a
22     sc = make_scenario(v, s)
23     return v, s, {"A": gauges(sim_slo.run(sc, "k8s_hpa70_ca")),
24                  "B": gauges(sim_slo.run(sc, "omniB:k8s_hpa70_ca", b_law=BLaw(**B_SET[v]["k8s_hpa70_ca"]))),
25                  "C": run_arm(sc, C_PICK[v], CANDS[C_PICK[v]])}
26
27
28 if __name__ == "__main__":
29     with Pool(4) as p:
30         res = p.map(job, [(v, s) for v in ("web", "multi", "batch", "gpu") for s in SEEDS])
31         out = {}
32         for v in ("web", "multi", "batch", "gpu"):
33             rows = [r for vv, s, r in res if vv == v]
34             out[v] = {arm: {m: float(np.mean([r[arm][m] for r in rows])) for m in rows[0]["A"] if isinstance(rows[0][
35 "A"][m], (int, float))}
36                       for arm in "ABC"}
37             (ROOT / "results/ABC_TABLE.json").write_text(json.dumps(out, indent=1))
38             for v, d in out.items():
39                 print(v)
40                 for m in ("energy_kwh", "node_hours", "p95_ms", "p99_ms", "mean_ms", "start_stop", "node_reversals",
41 "work_completed", "time_healthy"):
42                     print(f"  {m:16} A {d['A'][m]:10.4g} B {d['B'][m]:10.4g} C {d['C'][m]:10.4g}")
43
```

tools/buyer_report.py

271 lines SHA-256 19925e6466a9e60cald79adb54e2cd332e7f1d0f945324a46b24f67eef37bd01

```
1  """The buyer edition: one clean report from the final result files. Every number is read from a file in the
   repository.
2  Output: docs/OMNICOSSPASS_BUYER_EDITION.md (render with pilot/bench_pdf.py)."""
3  from __future__ import annotations
4
5  import json, subprocess, sys
6  from collections import Counter
7  from pathlib import Path
8
9  ROOT = Path(__file__).resolve().parents[1]
10 PLAT = ["k8s_hpa70_ca", "openshift", "gke_optimize", "aks_nap", "turbonomic", "cast_ai", "spot_ocean"]
11 PN = {"k8s_hpa70_ca": "Kubernetes", "openshift": "OpenShift", "gke_optimize": "GKE", "aks_nap": "AKS NAP /
   Karpenter",
12       "turbonomic": "Turbonomic", "cast_ai": "CAST AI", "spot_ocean": "Spot Ocean"}
13 VN = {"web": "Web services", "multi": "Four clusters, one site", "batch": "Batch jobs", "gpu": "GPU training"}
14 GN = {"energy_kwh": "energy", "node_hours": "machine-hours", "p95_ms": "typical response (p95)", "p99_ms":
   "slowest responses (p99)",
15       "mean_ms": "average response", "violation_backlog": "time over backlog limit", "violation_power": "time
   over power limit",
16       "violation_heat": "time over heat limit", "start_stop": "machine starts+stops", "node_reversals": "scale
   reversals",
17       "pod_changes": "pod changes", "work_completed": "work completed", "time_healthy": "time healthy"}
18
19
20 def J(p):
21     q = ROOT / p
22     return json.loads(q.read_text()) if q.exists() else None
23
24
25 def pct(a, b, lower=True):
26     return ((a - b) if lower else (b - a)) / abs(a) * 100 if abs(a) > 1e-12 else 0.0
27
28
29 def main():
30     L = []; w = L.append
31     head = subprocess.run(["git", "rev-parse", "--short", "HEAD"], cwd=ROOT, capture_output=True, text=True).
   stdout.strip()
32     w("# Omni-Compass: buyer edition")
33     w("")
34     w(f"Commit {head}. Every number below is generated from result files in the repository by `python tools/
   buyer_report.py`.")
35     w("Simulated results use the repository's fleet plant; live results come from real Kubernetes (kind) in
   GitHub Actions.")
36     w("Competitors are reproduced from their public documentation, not their binaries.")
37     w("")
38     w("## What Omni-Compass is")
39     w("")
40     w("- **The engine:** a six-state control engine, stable by construction (every mode decays), in Python and
   C++.")
41     w("- **The law it applies:** the closure law of the owner's manuscript.")
42     w(" - It does nothing in the safe interior.")
43     w(" - It corrects inward when the forecast approaches capacity.")
44     w(" - It releases a machine only after the peak.")
45     w(" - It keeps released machines warm for instant reuse.")
46     w("- **What it governs:** machines, pod counts and sizes, cold start, batch pacing, containment of runaway
   agent ")
47     w("workloads, and cooling setpoints.")
48     w("- **Two ways to deploy it:**")
49     w(" - **B, on top of the platform a customer already runs.** Nothing is replaced.")
50     w(" - **C, alone.** Kubernetes stays only as the muscle; the scaling layer is replaced (HPA, Cluster
   Autoscaler, ")
51     w("Karpenter, and optimizers such as CAST AI, Spot Ocean and Turbonomic).")
52     w("")
53     # --- B
54     bt, bf, sl = J("tuning/B_TONE_HELDOUT.json"), J("tuning/B_FIX.json"), J("tuning/SITE_LEAGUE.json")
55     if bt:
56         w("## B: Omni-Compass on top of each platform (held-out scenarios, settings frozen first)")
57         w("")
58         w("| Workload | Platform | Omni-Compass on top vs the platform alone |")
59         w("|---|---|---|")
60         fix = {k: v for k, v in (bf or {}).get("heldout", {}).items()}
61         for v in ("web", "multi", "batch", "gpu"):
62             for c in PLAT:
63                 key = f"{v}|{c}"
64                 if v == "multi" and sl:
65                     g, ls = sl["B"][c]["gains_pct"], sl["B"][c]["losses"]
66                 elif key in fix:
67                     g, ls = fix[key]["gains_pct"], fix[key]["losses"]
68                 else:
69                     g = bt[v]["gains_pct"][c]; ls = [l for l in bt[v]["losses"] if l["competitor"] == c]
70                 best = sorted(((x, m) for m, x in g.items() if x >= 1.0), reverse=True)[:3]
71                 txt = ", ".join(f"{GN.get(m, m)} {x:.0f}% better" for x, m in best) or "equal"
72                 if ls:
```

```

73         txt += "; worse: " + ", ".join(f"{GN.get(l['gauge'], l['gauge'])} {l['omni_worse_by_pct']:.1f}% " for l in ls)
74         w(f"| {VN[v]} | {PN[c]} | {txt} |")
75         w("")
76         # --- C
77         gl, cf, pl = J("tuning/GLOBAL_LEAGUE.json"), J("tuning/CONFIRMATORY.json"), J("tuning/PLANETLAB_LEAGUE.json")
78         if cf:
79             w("## C: Omni-Compass alone, one global setting, confirmatory run")
80             w("")
81             w(f"- **Setting:** one global setting for every workload, frozen by SHA-256 before the run "
82               f"({cf['preregistration_sha256'][:16]}).")
83             w(f"- **Scenarios:** 100 never-used scenarios per workload (seeds {cf['seeds'][0]}-{cf['seeds'][1]}).")
84             w(f"- **Comparisons:** 364 cells (4 workloads x 7 platforms x 13 gauges).")
85             w(f"- **Rule:** a cell counts as better or worse only if it survives Holm-Bonferroni correction across
all 364 "
86               "cells (family alpha 0.05) and exceeds the 0.5% practical tolerance. Otherwise it is equal.")
87             w("")
88             t = cf["tally"]
89             w(f"- **Result:** {t.get('better', 0)} better, {t.get('equal', 0)} equal, {t.get('worse', 0)} worse.")
90             w("")
91             w("| Workload | Better | Equal | Worse | The worse cells |")
92             w("|---|---|---|---|---|")
93             for v in ("web", "multi", "batch", "gpu"):
94                 p = cf["per_workload"][v]
95                 bad = [c for c in cf["cells"] if c["vessel"] == v and c["verdict"] == "worse"]
96                 w(f"| {VN[v]} | {p.get('better', 0)} | {p.get('equal', 0)} | {p.get('worse', 0)} | "
97                   f"+ ('; '.join(f'{PN[c['competitor']] {GN[c['gauge']] {-c['omni_better_by_pct']:.1f}% " for c in
bad) or "-") + " |")
98                 w("")
99                 w("| Workload | Energy vs Kubernetes | Machine-hours | Slowest responses | Average response |")
100                w("|---|---|---|---|---|")
101                for v in ("web", "multi", "batch", "gpu"):
102                    cc = {c["gauge"]: c for c in cf["cells"] if c["vessel"] == v and c["competitor"] == "k8s_hpa70_ca"}
103                    w(f"| {VN[v]} | {cc['energy_kwh']['omni_better_by_pct']:+.0f}% | {cc['node_hours']['
'omni_better_by_pct']:+.0f}% | "
104                      f"{'cc['p99_ms']['omni_better_by_pct']:+.0f}% | {cc['mean_ms']['omni_better_by_pct']:+.0f}% |")
105                    w("")
106                    w("Positive means Omni-Compass is better.")
107                    w("")
108                if pl:
109                    w("## Real demand: 1,052 recorded machines (PlanetLab)")
110                    w("")
111                    k, o = pl["fluid"]["means"]["k8s_hpa70_ca"], pl["fluid"]["means"]["omni"]
112                    w(f"- **Losing cells:** Omni-Compass alone, with the frozen web setting (never tuned on these traces),
loses "
113                      f"{len(pl['fluid']['losing_cells'])} of 91 against the seven platforms.")
114                    w(f"- **Whole-pod packing:** with whole-pod packing it loses {len(pl['bins']['losing_cells'])} of 91.")
115                    w(f"- **Against Kubernetes:**")
116                    w(f"  - energy {pct(k['energy_kwh'], o['energy_kwh']):+.0f}%")
117                    w(f"  - machine-hours {pct(k['node_hours'], o['node_hours']):+.0f}%")
118                    w(f"  - typical response {pct(k['p95_ms'], o['p95_ms']):+.0f}%")
119                    w(f"  - slowest responses {pct(k['p99_ms'], o['p99_ms']):+.0f}%")
120                    w("")
121                    w("Positive means better.")
122                    w("")
123                if sl:
124                    z = [c for c in PLAT if not sl["C"][c]["losses"]]
125                    w("## Four-cluster sites as one body (held-out)")
126                    w("")
127                    w(f"- **The mechanism:** traffic shift between clusters, plus the law run on the site total with one warm
reserve.")
128                    w(f"- **Omni-Compass alone:** zero losing cells against {'; '.join(PN[c] for c in z)}.")
129                    w(f"- **Omni-Compass on top:** zero losing cells on all seven platforms.")
130                    w(f"- **Assumption:** services are replicated across the site's clusters.")
131                    w("")
132                    pr = J("results/protocol/PROTOCOL_SUMMARY.json")
133                    if pr:
134                        w("## Under failure: the runtime protocol")
135                        w("")
136                        w("Machines dying, load spikes, crash-looping services and noisy neighbours were injected at five stress
levels, "
137                          "with identical faults for every system. Totals over all levels:")
138                        w("")
139                        w("| Workload | System | Runs conveyed | Pages to a human | Recovery (min) |")
140                        w("|---|---|---|---|---|")
141                        for d in pr["summary"]:
142                            if d["level"] == "all" and d["vessel"] in ("web", "multi") and (d["system"] in ("C-strict",
"k8s_hpa70_ca", "cast_ai")) and d["column"] in ("A", "C"):
143                                w(f"| {VN[d['vessel']] | {'Omni-Compass alone' if d['system'] == 'C-strict' else PN[d[
'system']] | "
144                                  f"int(d['conveyed']) of {d['runs']} | int(d['pages']) | {d['recovery_min']:.1f} |")
145                                w("")
146                                cc = J("tuning/CONFIRMATORY_COORD.json")
147                                ab = J("tuning/ABLATION.json")
148                                w("## The supervisory nervous system")
149                                w("")

```

```

150     w("- **One state, every organ:** `omnicompass/nervous_system.py` turns the engine's state (convergence,
basin health, "
151     "stress against its equation-6 equilibrium, unmet need) into one calm value between 0 and 1.")
152     w("- **Authority from calm:** calm grants each organ its authority:")
153     w(" - pods and machines may give capacity back only above their reversibility thresholds (0.5 and 0.7);")
154     w(" - CPU frequency, GPU power, routing and cooling get envelopes that widen with calm;")
155     w(" - batch is admitted or paused;")
156     w(" - rollback is authorised.")
157     w("- **Holds and the shield:** a security hold stops every capacity organ from expanding. The shield stays "
158     "downstream and can still veto.")
159     w("- **Invariants tested:** 300,000 random engine states, with zero violations (`tests/test_nervous_system.
py`)).")
160     w("- **Coordination:** pods move first, and a machine move opposite to the pod move is vetoed. Controller "
161     "contradictions per day (fighting, or a reversal within one boot time):")
162     w(" - web: 2.10 to 0.10;")
163     w(" - four-cluster: 5.30 to 1.03, below Kubernetes' 2.10.")
164     w("")
165     if cc:
166         t = cc["tally"]
167         w(f"**Coordination, confirmatory on fresh seeds {cc['seeds'][0]}-{cc['seeds'][1]}** (frozen first in "
168         f"`tuning/COORD_PREREGISTRATION.json`, same 364 cells, same Holm rule: {t.get('better', 0)} better, "
169         f"{t.get('equal', 0)} equal, {t.get('worse', 0)} worse.")
170         w("")
171         w("| Workload | Better | Equal | Worse | Contradictions per day: Omni-Compass / Kubernetes |")
172         w("|---|---|---|---|---|")
173         for v in ("web", "multi", "batch", "gpu"):
174             p_ = cc["per_workload"][v]; k_ = cc["contradictions_mean_per_day"][v]
175             w(f"| {VN[v]} | {p_.get('better', 0)} | {p_.get('equal', 0)} | {p_.get('worse', 0)} | "
176             f"{k_['omni']:.2f} / {k_['k8s_hpa70_ca']:.2f} |")
177         w("")
178     if ab:
179         VAR = {"gate_off": "engine release gate removed", "no_turn": "turning point removed", "no_tone": "muscle
tone removed",
180         "no_trend": "trend term removed"}
181         w(f"**Mechanism ablation** (`tuning/ABLATION.json`, fresh seeds {ab['seeds'][0]}-{ab['seeds'][-1]}):
each part of the "
182         "law removed in turn, against the full law. Listed: the gauges that get significantly worse (95%
interval "
183         "excludes 0, more than 0.5%).")
184         w("")
185         w("| Part removed | " + " | ".join(VN[v] for v in ("web", "multi", "batch", "gpu")) + " |")
186         w("|---|---|---|---|---|")
187         for k, lab in VAR.items():
188             cells = []
189             for v in ("web", "multi", "batch", "gpu"):
190                 g = ab["result"][v].get(k, {})
191                 hit = [f"{GN.get(m, m.replace('_', ' '))} +{d['worse_by_pct']:.0f}%" for m, d in g.items() if d[
"ci"][0] > 0 and d["worse_by_pct"] > 0.5]
192                 cells.append(" ".join(hit) or "no change")
193                 w(f"| {lab} | " + " | ".join(cells) + " |")
194             w("")
195         w("## Live Kubernetes evidence")
196         w("")
197         r3 = ROOT / "results/live/LIVE_REPS_3.md"
198         if r3.exists():
199             w("Repeated live runs, 5 paired repetitions per arm, probe through the Service (`results/live/
LIVE_REPS_3.md`):")
200             w("")
201             w(r3.read_text())
202             w("")
203             w("- **Earlier live sets 1 and 2:** these carried a probe defect, now fixed. The probe's one-pod tunnel hung
when a "
204             "drain moved its pod, so the Omni arms logged false failed requests (`results/live/LIVE_REPS_PROBE_DEFECT.
md`)).")
205             w("- **Their machine savings are withdrawn.** The hung probe left the controller's latency sense nearly
blind, so "
206             "it released machines it would not have released with a working probe. Set 3 shows that: with a working
probe and "
207             "latency near the declared 500 ms SLO, Omni-Compass kept all 6 machines. The decision trail per run is now
in "
208             "every job log (set 4).")
209             w("")
210             if (ROOT / "results/live/LIVE_LEVERS_2_NERVOUS.txt").exists():
211                 w("**Live levers under the nervous system:**")
212                 w("")
213                 w("- **What acted:** right-sizing, cold start, batch pacing, agent containment and the cooling connector.
")
214                 w("- **Authority:** each lever acted on real Kubernetes only inside the authority the nervous system
granted.")
215                 w("- **Kill switch:** it restored every lever, including from a fresh process.")
216                 w("- **Result:** 19 of 19 checks passed (`results/live/LIVE_LEVERS_2_NERVOUS.txt`, first pass "
217                 "`LIVE_LEVERS_1.txt`)).")
218                 w("- **Identity:** least-privilege, with `kubectl auth can-i` receipts.")
219                 w("")
220             if (ROOT / "results/live/LIVE_SHADOW_1.txt").exists():
221                 w("**Shadow pilot kit, live:** a read-only identity ran for 600 s and logged 40 decisions, with 0 writes

```

```

222     "(`results/live/LIVE_SHADOW_1.txt`). This is the kit a customer runs first.")
223     w("")
224     w("## Safety and correctness")
225     w("")
226     w("- **Safety shield:** tested on 2,000,000 random and adversarial inputs, with zero invariant violations, "
227     "idempotent, never inventing an action, and intervening minimally (`tests/test_shield_properties.py`). The
test "
228     "found two real bugs, both fixed.")
229     w("- **C++ twins:** the C++ shield and the C++ closure law match Python exactly, over 300,000 adversarial "
230     "shield cases and every recorded closure decision.")
231     w("- **Endurance:** the C++ engine ran 100,000,000 decisions with no failure, at about 2.3 microseconds per
decision.")
232     w("- **Fail-safe:** after repeated failed decisions, control returns to the native autoscalers.")
233     w("- **Reproducibility:** a clean copy of the delivered zip reproduced every held-out result byte for byte.")
234     w("")
235     w("## What is not claimed")
236     w("")
237     for s in ["**No production or customer deployment yet.** The next step is the shadow pilot (`docs/PILOT_KIT.
md`), "
238             "which is read-only.",
239             "**Live runs are small.** They use kind on CI machines, and power is modelled, not metered.",
240             "**Hardware levers are not proven.** CPU-frequency and power caps need real servers, and cooling
was "
241             "exercised against a stand-in controller.",
242             "**Parked machines in the public cloud.** A parked cloud machine still bills, so the warm-reserve
energy "
243             "saving applies to owned hardware.",
244             "**Single-cluster energy against the tightest packers is roughly a tie.** There Omni-Compass wins
on "
245             "response time and stability.",
246             "**Some gaps cannot be closed.** No controller, even one with perfect foresight, can match both
the "
247             "tightest packer's machine-hours and the calmest autoscaler's machine churn "
248             "(`tuning/bound.py`).",
249             "**Not modelled in the plant:** variable boot times and pod-eviction cost. Fragmentation is
modelled "
250             "(whole-pod packing) and is small."]:
251         w(f"- {s}")
252     w("")
253     w("## Reproduce")
254     w("")
255     w("```")
256     for c in ["pip install -r requirements.txt && python verify.py",
257             "python tuning/confirmatory.py          # C, one global setting, 100 scenarios per workload, Holm-
corrected",
258             "python tuning/confirmatory.py --coord # the same with nervous-system coordination, fresh seeds",
259             "python tuning/planetlab_league.py <planetlab-workload-traces/20110303>",
260             "python tuning/site_league.py --heldout",
261             "python tools/protocol_bench.py 100",
262             "Live: push a commit whose message contains [reps], [levers] or [shadow]"]:
263         w(c)
264     w("```")
265     (ROOT / "docs/OMNICOSSPASS_BUYER_EDITION.md").write_text("\n".join(L))
266     print(f"wrote docs/OMNICOSSPASS_BUYER_EDITION.md {{len(L)} lines}")
267
268
269 if __name__ == "__main__":
270     main()
271

```

tools/closed_loop_estimate.py

65 lines SHA-256 b20883485bcd6276dfc139d34781f7ea4f39d73411cb114a686029278480102a

```
1  """Work served under the closed-loop load (deploy/kind/loadgen.yaml), estimated per repetition.
2
3  python3 tools/closed_loop_estimate.py DIR          # DIR holds bench-<arm>-<rep>/ with latency.csv, capture.csv,
4  load_schedule.log
5  Each generator sends 20 requests one after another, each waiting for its answer, then sleeps RANDOM % 3 s (mean 1
6  s).
7  A generator's cycle is therefore 20 R + 1 s, R the mean response time, and it serves 20 / (20 R + 1) requests a
8  second.
9  The requests of an arm are that rate times the generator-seconds of the load schedule. R is the mean response
10 time of
11 the probe's successful requests in the same arm (the same service the generators call). The generators do not
12 count
13 their requests, so this is an estimate; an open-loop load (LOADGEN=open) or a server-side count replaces it.
14 Prints, for Omni-Compass on top against native, paired over repetitions: requests, CPU used, CPU per request.
15 """
16 from __future__ import annotations
17
18 import csv, json, math, re, statistics, sys
19 from pathlib import Path
20
21 T95 = {1: 12.706, 2: 4.303, 3: 3.182, 4: 2.776, 5: 2.571, 6: 2.447, 7: 2.365, 8: 2.306, 9: 2.262, 10: 2.228}
22
23 def generator_seconds(d):
24     """Generator-seconds over the measured window, from load_schedule.log ("HH:MM:SS load-generator replicas ->
25     r")."""
26     steps = []
27     for line in (d / "load_schedule.log").read_text().splitlines():
28         m = re.match(r"(\d+):(\d+):(\d+) load-generator replicas -> (\d+)", line)
29         if m:
30             h, mi, s, r = map(int, m.groups()); steps.append((h * 3600 + mi * 60 + s, r))
31     rows = list(csv.DictReader(open(d / "capture.csv")))
32     t_end = steps[0][0] + float(rows[-1]["elapsed_seconds"]) if rows else steps[-1][0]
33     total = 0.0
34     for (t0, r), nxt in zip(steps, steps[1:] + [(max(t_end, steps[-1][0]), 0)]):
35         t1 = nxt[0] if nxt[0] >= t0 else nxt[0] + 86400
36         total += r * (t1 - t0)
37     return total
38
39 def arm(d):
40     ms = [float(r["latency_ms"]) for r in csv.DictReader(open(d / "latency.csv")) if r["ok"] == "1"]
41     R = statistics.mean(ms) / 1000.0
42     req = generator_seconds(d) * 20.0 / (20.0 * R + 1.0)
43     cpu = statistics.mean(float(r["used_cpu_m"]) for r in csv.DictReader(open(d / "capture.csv")))
44     return req, cpu
45
46 def main(root):
47     root = Path(root); pairs = []
48     for dn in sorted(root.glob("bench-native-*")):
49         rep = dn.name.split("-", 2)[2]; do = root / f"bench-omni-{rep}"
50         if do.exists() and (dn / "load_schedule.log").exists() and (do / "load_schedule.log").exists():
51             (qn, cn), (qo, co) = arm(dn), arm(do)
52             pairs.append((qo / qn - 1, co / cn - 1, (co / qo) / (cn / qn) - 1))
53     out = {"repetitions": len(pairs)}
54     lines = [f"Closed-loop work estimate, Omni-Compass on top against native, {len(pairs)} paired repetitions (
55 estimate: the generators do not count their requests)"]
56     for k, name in enumerate(("requests served", "CPU used", "CPU per request")):
57         d = [p[k] for p in pairs]; m = statistics.mean(d)
58         h = T95.get(len(d) - 1, 1.96) * statistics.stdev(d) / math.sqrt(len(d)) if len(d) > 1 else float("nan")
59         out[name] = {"change": m, "ci95": [m - h, m + h]}
60         lines.append(f" {name:16s} {100 * m:+6.1f}% 95% interval {100 * (m - h):+.1f}% to {100 * (m + h):+.
61 1f}%")
62     (root / "CLOSED_LOOP_ESTIMATE.json").write_text(json.dumps(out, indent=1))
63     print("\n".join(lines))
64     return out
65
66 if __name__ == "__main__":
67     main(sys.argv[1])
68
```

tools/code_fingerprint.py

23 lines SHA-256 abb0e8a27c8f11d46a6e10c7f90213e251c0a72df3e67db7cbad94f9b94a7711

```
1  """SHA-256 of a Python file's syntax tree with comments and docstrings removed."""
2  import ast, hashlib, sys
3
4
5  def strip_docstrings(tree):
6      for node in ast.walk(tree):
7          if isinstance(node, (ast.Module, ast.FunctionDef, ast.AsyncFunctionDef, ast.ClassDef)):
8              b = node.body
9              if b and isinstance(b[0], ast.Expr) and isinstance(getattr(b[0], "value", None), ast.Constant) \
10                 and isinstance(b[0].value.value, str):
11                  node.body = b[1:] or [ast.Pass()]
12      return tree
13
14
15  def fingerprint(path):
16      tree = strip_docstrings(ast.parse(open(path, encoding="utf-8").read()))
17      return hashlib.sha256(ast.dump(tree, include_attributes=False).encode()).hexdigest()
18
19
20  if __name__ == "__main__":
21      for p in sys.argv[1:]:
22          print(fingerprint(p), p)
23
```

tools/cpufreq_receipt.py

32 lines SHA-256 a61661c9356da270a391dc3c16dfdfdf67c931e0e604d4135a536ff93d2aa148

```
1  #!/usr/bin/env python3
2  """Read-only CPUFreq receipt for a prospective Omni hardware run."""
3  from __future__ import annotations
4  import argparse, json, sys
5  from pathlib import Path
6  ROOT = Path(__file__).resolve().parents[1]
7  sys.path.insert(0, str(ROOT))
8  from hardware.cpufreq import CpuFreqPolicies
9
10
11 def main():
12     ap = argparse.ArgumentParser()
13     ap.add_argument("--root", default="/sys/devices/system/cpu/cpufreq")
14     ap.add_argument("--out", default="results/hardware/CPUFREQ_RECEIPT.json")
15     a = ap.parse_args()
16     c = CpuFreqPolicies(a.root)
17     rows = c.snapshot()
18     obj = {
19         "evidence_type": "read_only_cpufreq_sysfs_receipt",
20         "root": str(a.root),
21         "policies": [r.as_dict() for r in rows],
22         "schedutil_all": bool(rows) and all(r.scaling_governor == "schedutil" for r in rows),
23         "note": "A ceiling write is not an observation of delivered silicon frequency. Active intel_pstate/HWP is
a separate driver lane.",
24     }
25     out = Path(a.out); out.parent.mkdir(parents=True, exist_ok=True); out.write_text(json.dumps(obj, indent=2) +
"\n")
26     print(json.dumps(obj, indent=2))
27     if not rows:
28         raise SystemExit("no CPUFreq policies found")
29
30
31 if __name__ == "__main__": main()
32
```

tools/demo_nervous_system.py

75 lines SHA-256 1b73c1dfb1f4cd4dd703c2c4c7b131faf02b99645707f34efb085048a005d3ec

```
1 """Walk the live controller (omni_controller/controller.py, nodepool mode, closure law, two-way nervous system)
   through
2 five situations on the repository's stand-in cluster (tests/fake_cluster/kubect1: answers like the Kubernetes API
   from a
3 JSON state file; no pods run). Prints the decision trail each decision writes to its audit:
4 1 calm, machines idle          -> releases one machine per decision down to the floor (3 here)
5 2 latency breached now         -> holds ("latency breached now")
6 3 the latency probe freezes    -> holds ("a sense is blind"; the engine's stale channel rises)
7 4 a drain is refused (order does not land) -> next decision holds ("last node command did not land")
8 5 pods scale up               -> holds ("pods scaling up")
9 Usage: python tools/demo_nervous_system.py -> results/local_run/NERVOUS_SYSTEM_DEMO.txt"""
10 import csv, json, os, sys, tempfile, time
11 from pathlib import Path
12 ROOT = Path(__file__).resolve().parents[1]; sys.path.insert(0, str(ROOT))
13 from omni_controller.controller import Controller, parser
14 FAKE = str(ROOT / "tests" / "fake_cluster" / "kubect1")
15
16
17 def main():
18     t = Path(tempfile.mkdtemp()); p = t / "state.json"; lat = t / "latency.csv"; out = []
19     nodes = [{"metadata": {"name": f"w{i}"}, "spec": {}, "status": {"allocatable": {"cpu": "4"},
20         "conditions": [{"type": "Ready", "status": "True"}]}} for i in range(6)]
21     hpa = {"metadata": {"name": "web", "namespace": "default"},
22         "spec": {"minReplicas": 1, "maxReplicas": 10, "metrics": [{"type": "Resource", "resource": {"name":
23             "cpu", "target": {"type": "Utilization", "averageUtilization": 50}}]},
24         "status": {"currentReplicas": 2, "desiredReplicas": 2, "currentMetrics": [{"type": "Resource",
25             "resource": {"name": "cpu", "current": {"averageUtilization": 30}}]}}
26     p.write_text(json.dumps({"nodes": nodes, "pods": [], "used_per_node": "300m", "deployments": [], "configmaps":
27         [], "jobs": [], "hpas": [hpa]}))
28     os.environ["FAKE_KUBE_STATE"] = str(p); os.environ["FAKE_KUBE_LOG"] = str(t / "w.log")
29     scale = t / "scale.py" # node-pool actuator: cordon down to n (refuses when REFUSE_DRAIN exists)
30     scale.write_text(f'''import json, os, sys
31 n = int(sys.argv[1]); s = json.load(open({str(p)!r}))
32 if os.path.exists({str(t / "REFUSE_DRAIN"!r)}): sys.exit(0)
33 for i, nd in enumerate(s["nodes"]): nd["spec"]["unschedulable"] = i >= n
34 json.dump(s, open({str(p)!r}, "w"))''')
35
36
37 def probe(ms, fresh=True):
38     with open(lat, "w") as f:
39         w = csv.writer(f); w.writerow(["elapsed_seconds", "latency_ms", "ok"]); w.writerows((s, ms, 1) for s
40 in range(0, 120))
41     if not fresh:
42         old = time.time() - 900; os.utime(lat, (old, old))
43
44     probe(200)
45     a = parser().parse_args(["--kubect1", FAKE, "--mode", "nodepool", "--active-nodes-only", "--interval", "0", "-
46 -min-nodes", "3",
47         "--max-nodes", "6", "--node-scale-cmd", f"{sys.executable} {scale} {{n}}", "--
48 latency-file", str(lat),
49         "--slo-ms", "500", "--slo-clear", "1", "--audit", str(t / "a.jsonl"), "--kill-file",
50 str(t / "kill"),
51         "--closure", str(ROOT / "tuning/GLOBAL_LEAGUE_PREREGISTRATION.json")])
52
53     c = Controller(a)
54
55     def step(label):
56         c.step()
57         recs = [json.loads(l) for l in open(t / "a.jsonl")]
58         auth = [r for r in recs if "authority" in r][-1]["authority"]; dec = [r for r in recs if isinstance(r.get(
59 "decision"), dict)][-1]["decision"]
60         line = (f"{label:34s} machines {dec['nodes_observed']} -> {dec['nodes_recommended']} | gate: {auth[
61 'node_gate']['reason']} | "
62             f"node calm {auth['node_view']['calm']} | blind {[k for k, v in auth['senses']['blind'].items()
63 if v]} | "
64             f"proprioception {auth['proprioception']}")
65         print(line); out.append(line)
66
67     for i in range(5):
68         step(f"1 calm, idle machines (decision {i + 1})")
69         s = json.loads(p.read_text()) # release the floor's grip for scenario 4: one spare
70 machine back
71         for nd in s["nodes"][:4]: nd["spec"]["unschedulable"] = False
72         p.write_text(json.dumps(s))
73         probe(900); step("2 latency breached now (p95 900 ms)")
74         probe(200, fresh=False); step("3 latency probe frozen (15 min old)")
75         probe(200); (t / "REFUSE_DRAIN").touch()
76         for i in range(3):
77             step(f"4 drain refused (decision {i + 1})")
78             (t / "REFUSE_DRAIN").unlink()
79             s = json.loads(p.read_text()); s["hpas"][0]["status"]["desiredReplicas"] = 5; p.write_text(json.dumps(s))
80             step("5 pods scaling up (HPA 2 -> 5)")
81             Path(ROOT / "results/local_run").mkdir(parents=True, exist_ok=True)
82             (ROOT / "results/local_run/NERVOUS_SYSTEM_DEMO.txt").write_text(
83                 "Live controller on the stand-in cluster (fake kubect1; no pods run). Decision trail as written to the
```

```
        audit.\n\n" + "\n".join(out) + "\n")
71
72
73 if __name__ == "__main__":
74     main()
75
```

tools/full_report.py

791 lines SHA-256 ff10d54126ed5753f71211b9a986e6a374b8e644284add9f6406673157e237b9

```
1  """Build the Omni-Compass benchmark report (Markdown, then PDF via pilot/bench_pdf.py).
2
3  Three architectures, every gauge, every study in this repository:
4      A Kubernetes alone
5      B Kubernetes + Omni-Compass (Omni governs on top; Kubernetes' own controllers keep running)
6      C Omni-Compass direct (Omni is the single authority; the separate managers do not decide)
7
8  Inputs: results/heldout_seed_*/SUMMARY.json (pre-registered, 2 x 500 scenarios), a k8s_controlplane run (rows
9         JSON),
10         results/fleet/planetlab or a planetlab SUMMARY.json, the live kind results (--live JSON), results/VERIFY_LOG.txt.
11
12 python tools/full_report.py --cp-rows ROWS.json --planetlab SUMMARY.json --live LIVE.json --out docs/
13   BENCHMARK_REPORT.md
14 """
15 from __future__ import annotations
16
17 import argparse, json
18 from pathlib import Path
19
20 import numpy as np
21
22 ROOT = Path(__file__).resolve().parents[1]
23 SEEDS = ["346410161", "360555127"]
24
25 HELD = [ # (label, key, higher_is_better, unit/format)
26     ("ENERGY AND POWER", None, None, None),
27     ("Energy per scenario (kWh)", "energy_kwh", False, "{:.1f}"),
28     ("Peak power (kW)", "peak_power_kw", False, "{:.1f}"),
29     ("Node-hours", "node_hours", False, "{:.1f}"),
30     ("Idle node-hours (powered, doing nothing)", "idle_node_hours", False, "{:.2f}"),
31     ("Share of time over the power limit", "violation_power", False, "{:.3f}"),
32     ("Power-cap travel (how much the cap moved)", "power_cap_travel", False, "{:.3f}"),
33     ("HEAT", None, None, None),
34     ("Share of time over the heat limit", "violation_heat", False, "{:.3f}"),
35     ("Thermal travel (how far temperature swung)", "thermal_travel", False, "{:.3f}"),
36     ("SPEED AND BACKLOG", None, None, None),
37     ("Mean queue (work waiting)", "mean_queue", False, "{:.0f}"),
38     ("95th-percentile queue (worst moments)", "p95_queue", False, "{:.0f}"),
39     ("Share of time over the backlog limit", "violation_backlog", False, "{:.3f}"),
40     ("RELIABILITY AND RECOVERY", None, None, None),
41     ("Availability", "availability", True, "{:.4f}"),
42     ("Share of time healthy", "time_healthy", True, "{:.3f}"),
43     ("Share of incidents recovered", "recovered", True, "{:.3f}"),
44     ("Recovery time (minutes)", "recovery_minutes", False, "{:.1f}"),
45     ("Physical SLA breaches (power, heat, backlog)", "sla_violation_physical", False, "{:.3f}"),
46     ("All SLA breaches", "sla_violation_total", False, "{:.3f}"),
47     ("WEAR AND TEAR", None, None, None),
48     ("Machines started", "machines_started", False, "{:.1f}"),
49     ("Machines stopped", "machines_stopped", False, "{:.1f}"),
50     ("Node start/stop events", "node_start_stop", False, "{:.1f}"),
51     ("Machine round trips (stopped then restarted)", "machine_round_trips", False, "{:.2f}"),
52     ("Scale reversals (flip-flops)", "scale_reversals", False, "{:.2f}"),
53     ("CONTROL QUALITY AND SAFETY", None, None, None),
54     ("Contradictory commands between managers", "contradictions", False, "{:.2f}"),
55     ("Safety-rule (invariant) violations", "invariant_violations", False, "{:.2f}"),
56     ("Invariant violations excluding power", "invariant_violations_ex_power", False, "{:.2f}"),
57     ("Security violations", "security_violations", False, "{:.3f}"),
58     ("Pages to on-call", "pages", False, "{:.3f}"),
59     ("Human interventions required", "human_interventions", False, "{:.3f}"),
60 ]
61
62 A, B, C = "k8s_ref_70", "omni_k8s_throughput", "omni_direct"
63
64 def pct(a, b):
65     if a == 0:
66         return "0%" if b == 0 else "new"
67     r = 100 * (b - a) / abs(a)
68     return f"{r:+.2f}%" if abs(r) < 1 else f"{r:+.0f}%"
69
70 def verdict(entries, higher):
71     """entries: paired dicts for each seed. Significant only if both seeds' CIs exclude zero in the same
72     direction."""
73     sig = []
74     for e in entries:
75         lo, hi = e["ci95"]
76         if lo > 0:
77             sig.append("up")
78         elif hi < 0:
79             sig.append("down")
80         else:
81             sig.append("ns")
82     if all(s == sig[0] for s in sig) and sig[0] != "ns":
```

```

81         good = (sig[0] == "up") == higher
82         return "better" if good else "worse"
83     return "not significant" if all(s == "ns" for s in sig) else "mixed"
84
85
86 def held_out():
87     S = [json.load(open(ROOT / f"results/heldout_seed_{s}/SUMMARY.json")) for s in SEEDS]
88     mean = lambda arm, k: float(np.mean([s["means"][arm][k] for s in S]))
89     pair = lambda cand, k: [next(e for e in s["paired"][f"{cand}_vs_{A}"] if e["metric"] == k) for s in S]
90     rows = ["| Gauge | A. Kubernetes alone | B. Kubernetes + Omni-Compass | C. Omni-Compass direct | B vs A | C
vs A |",
91            "|---|---|---|---|---|---|"]
92     tally = {"B": {"better": 0, "worse": 0, "same": 0}, "C": {"better": 0, "worse": 0, "same": 0}}
93     for lab, k, hi, f in HELD:
94         if k is None:
95             rows.append(f"| **{lab}** | | | | |"); continue
96         a, b, c = mean(A, k), mean(B, k), mean(C, k)
97         cells = []
98         for tag, cand, v in (("B", B, b), ("C", C, c)):
99             try:
100                 P = pair(cand, k)
101                 vd = verdict(P, hi)
102                 n_b = sum(e["better"] for e in P); n_w = sum(e["worse"] for e in P); n = sum(e["n"] for e in P)
103                 cells.append(f"{pct(a, v)} {vd} (better in {n_b}, worse in {n_w} of {n})")
104                 tally[tag]["better" if vd == "better" else "worse" if vd == "worse" else "same"] += 1
105             except (StopIteration, KeyError):
106                 cells.append(pct(a, v))
107             rows.append(f"| {lab} | {f.format(a)} | {f.format(b)} | {f.format(c)} | {cells[0]} | {cells[1]} |")
108     obs = [{"observe_node_identical_to_native": s["layered_observe_identical_to_kubernetes"]} for s in S]
109     abl = {}
110     for s in S:
111         for x in ("omni_direct_vs_omni_direct_no_dynamics", "omni_direct_vs_omni_direct_no_shield"):
112             for e in s["paired"].get(x, []):
113                 abl.setdefault((x, e["metric"]), []).append(e)
114     native = {k: mean("native", k) for _, k, _ in HELD if k}
115     return rows, tally, obs, abl, native, mean
116
117
118 def cp_study(rows_path):
119     R = json.load(open(rows_path)); by = {}
120     for r in R:
121         by.setdefault(r["arm"], {})[r["scenario_id"]] = r
122     ids = sorted(by["hpa70_ca"])
123     G = [{"Energy (kWh)", "energy_kwh", False}, {"Peak power (kW)", "peak_power_kw", False},
124          {"Time over power limit", "violation_power", False}, {"Time over heat limit", "violation_heat", False},
125          {"Time healthy", "time_healthy", True}, {"Recovery time (min)", "recovery_minutes", False},
126          {"Physical SLA breaches", "sla_physical", False}, {"Mean queue", "mean_queue", False},
127          {"Machines started", "machines_started", False}, {"Machines stopped", "machines_stopped", False},
128          {"Scale reversals", "scale_reversals", False}, {"Invariant violations", "invariant_violations", False},
129          {"Pages to on-call", "pages", False}]
130     arms = [{"A. Kubernetes alone (HPA + Cluster Autoscaler)", "hpa70_ca"},
131             {"Reference: HPA + Karpenter-lite", "hpa70_karpenter"},
132             {"B. Kubernetes + Omni-Compass (HPA target + CA gate)", "omni_target_gate_hpa70_ca"},
133             {"C. Omni-Compass authority (HPA target, nodes, power cap)", "omni_throughput_full"}]
134     out = ["| Gauge | " + " | ".join(a for a, _ in arms) + " | C vs A (95% CI, wins of 24) |", "|---|" + "----:|"
* len(arms) + "----|"]
135     rng = np.random.default_rng(7)
136     for lab, k, hi in G:
137         vals = [np.mean([by[arm][i][k] for i in ids]) for _, arm in arms]
138         a = np.array([by["hpa70_ca"][i][k] for i in ids]); c = np.array([by["omni_throughput_full"][i][k] for i
in ids])
139         d = c - a; bs = d[rng.integers(0, len(d), (4000, len(d)))].mean(1); lo, hi_ = np.percentile(bs, [2.5, 97.
5])
140         v = ("better" if (lo > 0) == hi else "worse") if (lo > 0 or hi_ < 0) else "not significant"
141         wins = int(((c > a) if hi else (c < a)).sum())
142         out.append(f"| {lab} | " + " | ".join(f"{x:.3f}" if abs(x) < 10 else f"{x:.1f}" for x in vals) +
143                 f" | {pct(vals[0], vals[3])} {v} | {lo:+.3f}, {hi_:+.3f}, {wins} of 24 |")
144         obs = sum(by["omni_observe_hpa70_ca"][i]["trace_hash"] == by["hpa70_ca"][i]["trace_hash"] for i in ids)
145         return out, obs, len(ids)
146
147
148 def planetlab(path):
149     d = json.load(open(path)); m = d["means"]
150     arms = [{"A. Kubernetes alone (HPA 0.7 + CA)", "k8s_hpa70_ca"}, {"Reference: HPA 0.7 + Karpenter-lite",
"k8s_hpa70_karpenter"},
151            {"B. Kubernetes + Omni-Compass (HPA target)", "omni_target"}, {"C. Omni-Compass node-pool authority (
CA off)", "omni_fleet"}]
152     keys = [{"Energy (kWh)", "energy_kwh"}, {"Time healthy", "time_healthy"}, {"Work completed",
"work_completed"},
153            {"Backlog violations", "violation_backlog"}, {"Node reversals", "node_reversals"}, {"Machines
started", "machines_started"}]
154     out = ["| Gauge | " + " | ".join(a for a, _ in arms) + " |", "|---|" + "----:|" * len(arms)]
155     for lab, k in keys:
156         out.append(f"| {lab} | " + " | ".join(f"{m[a][k]:.4f}" if m[a][k] < 10 else f"{m[a][k]:.2f}" for _, a in
arms) + " |")
157     ver = ["| Comparison | Metric | Change | 95% CI | Verdict |", "|---|---|---|---|---|"]

```

```

158     for p in ("omni_target_vs_k8s_hpa70_ca", "omni_fleet_vs_k8s_hpa70_ca", "omni_fleet_vs_k8s_hpa70_karpenter"):
159         for e in d["paired"][p]:
160             ver.append(f"| {p.replace('_', ' vs ')} | {e['metric']} | {e['delta']:+.4f} | [{e['ci95']}[0]:+.
4f}, {e['ci95']}[1]:+.4f}] | {e['verdict']} |")
161     return out, ver, d["observe_identical"], d["scenarios"]
162
163
164 def main():
165     ap = argparse.ArgumentParser()
166     ap.add_argument("--cp-rows", required=True); ap.add_argument("--planetlab", required=True)
167     ap.add_argument("--live", required=True); ap.add_argument("--out", required=True)
168     a = ap.parse_args()
169     live = json.load(open(a.live))
170     hrows, tally, hobs, abl, native, hmean = held_out()
171     crows, cobs, cn = cp_study(a.cp_rows)
172     prows, pver, pobs, pn = planetlab(a.planetlab)
173     verify = (ROOT / "results/VERIFY_LOG.txt").read_text().splitlines()
174     n_pass = sum(1 for l in verify if l.startswith("PASS")); n_fail = sum(1 for l in verify if l.startswith(
"FAIL"))
175     E = lambda arm: hmean(arm, "energy_kwh")
176     L = []
177     w = L.append
178     w("# Omni-Compass: Kubernetes alone, Kubernetes + Omni-Compass, and Omni-Compass direct")
179     w("")
180     w("Benchmark report, 26 September 2026. Repository: Omni-Compass/The-Omni-Compass-Control-Core-Engine (
private), branch "
181     "claude/kubernetes-clusters-docker-stack-gp26ve. Every number below is produced by code in that repository
and can be "
182     "regenerated; section 21 gives the commands. Each result states whether it was **measured on a live
Kubernetes control plane** "
183     "or **computed in simulation**.")
184     w("")
185     w("## 1. Summary")
186     w("")
187     w("Omni-Compass is a single control engine that senses the whole compute stack and drives its actuators (its
\"muscles\": "
188     "replica counts, node pools, power caps and others) from one six-state dynamical model, with a safety
shield before every "
189     "action and a kill switch that hands control back. It was compared in three architectures:")
190     w("")
191     w("- **A. Kubernetes alone.** Kubernetes' own controllers decide: Horizontal Pod Autoscaler (HPA) for
replicas, Cluster Autoscaler "
192     "for nodes; other managers act on their own proposals.")
193     w("- **B. Kubernetes + Omni-Compass.** Kubernetes' controllers keep running; Omni-Compass governs on top of
them (sets the HPA "
194     "target, gates and sizes the node pool, caps power) as the single authority over their settings.")
195     w("- **C. Omni-Compass direct.** Omni-Compass is the only decision-maker and actuates the muscles directly;
the separate managers "
196     "no longer decide.")
197     w("")
198     w(f"***Main result (pre-registered, 1,000 held-out scenarios, simulation).** Against Kubernetes alone, B used
"
199     f"{pct(E(A), E(B))} energy and C {pct(E(A), E(C))}; time healthy rose from {hmean(A,'time_healthy'):.0%}
to "
200     f"{hmean(B,'time_healthy'):.0%} (B) and {hmean(C,'time_healthy'):.0%} (C); recovery time fell from "
201     f"{hmean(A,'recovery_minutes'):.0f} to {hmean(B,'recovery_minutes'):.0f} and {hmean(C,'recovery_minutes'):.
0f} minutes; "
202     f"contradictory commands, pages and human interventions went to zero in both; safety-rule violations fell
from "
203     f"{hmean(A,'invariant_violations'):.1f} to {hmean(B,'invariant_violations'):.1f} (B) and {hmean(C,
'invariant_violations'):.1f} (C). "
204     f"Of {len([h for h in HELD if h[1]])} gauges, B is significantly better on {tally['B']['better']} and
worse on {tally['B']['worse']}; "
205     f"C is better on {tally['C']['better']} and worse on {tally['C']['worse']}")
206     w("")
207     lv = live["side_by_side"]
208     w(f"***Live Kubernetes result (measured).** Two identical Kubernetes clusters (1 control plane + 6 workers)
ran the same load at the "
209     f"same time, one without Omni-Compass and one with it. With Omni-Compass: worker nodes in service {lv[
'nodes_native']:.1f} to "
210     f"{lv['nodes_omni']:.1f}; utilisation of the workers in service {lv['util_native']:.3f} to {lv[
'util_omni']:.3f}. **Energy:** with the "
211     f"parked workers kept on standby, powered and ready ({lv['standby_note']}), energy was {lv[
'energy_native_wh']:.0f} vs "
212     f"{lv['energy_omni_wh']:.0f} Wh ({pct(lv['energy_native_wh'], lv['energy_omni_wh'])}: parking alone saves
essentially nothing; "
213     f"energy per unit of work {lv['epc_change']}. The -43% first reported for this run holds only if parked
workers are powered "
214     f"off. Waiting pods and HPA shortfall were not significantly different. The kill switch restored the
original HPA target (50) and all 6 workers. "
215     f"***With every live muscle switched on (section 7.2) the application got slower: p95 response time 486 to
802 ms, energy per "
216     f"unit of work +46% (significant).** The causes were found and fixed across runs 2-4 (section 7.4):
response time went from +65% to a tie at p95.")
217     w("")
218     w("***Where Omni-Compass costs something.** In the pre-registered study both B and C keep more node-hours

```

```

powered than "
219 "Kubernetes alone and start and stop machines more often (more wear), and move the power cap more; C also
lets more work wait "
220 "in the queue and flips scale direction more often. In that study the energy saving comes from power
capping and load "
221 "shaping, not from switching machines off. On the live cluster the saving came from switching machines off.
Other limits: "
222 "the small-cluster release band (section 11); power and heat on the live cluster are modelled, not metered.
")
223 w("")
224 w("***Trade-off in one line:** B is the strongest all-round result in simulation (energy, health, recovery,
queue, coordination and safety "
225 "all better; wear and node-hours worse); C is the strongest on peak power, heat and safety (zero invariant
violations) at the "
226 "cost of queue length, wear and flip-flops. These are the gauges to tune next.")
227 w("")
228 w("## 2. What Omni-Compass is")
229 w("")
230 w("### 2.1 The engine")
231 w("The engine is a six-state ordinary differential equation system, state x = (E, U, I_U, S, B, B_dot):
error E, coherence U, "
232 "pressure I_U, stress S and a damped bath B. The shipped core (`omnicompass/core.py`) integrates it with
fourth-order "
233 "Runge-Kutta and holds the control input constant across the four stages:")
234 w("")
235 w("```")
236 for ln in ["(1) dE/dt = -alpha_E E + beta_int + beta_ext + v_eff",
237 "(2) dU/dt = mu U (1 - U^2) - (dE/dt)/E_max - lambda_U U + u, |u| <= 25",
238 "(3) dI_U/dt = (1 - U) - sigma_I E - delta S - lambda_I I_U",
239 "(4) v_eff = cos(omega_B t / 2) c tanh(lambda_0 + lambda_1 (U - 0.5) + lambda_2 S)",
240 "(5) Phi(S) = alpha_s S^2/2 + beta_s S^3/4 - delta S",
241 "(6) dS/dt = -dPhi/dS",
242 "(7) dB/dt = B_dot; dB_dot/dt = gamma_c delta S - (omega_B/Q_B) B_dot - omega_B^2 B",
243 "(8) R_B[n] = finite-difference audit of (7), never fed back",
244 "Controller: u = clip(-f_U(x,t) + 12 (sigma - U), -25, +25)"]:
245 w(ln)
246 w("```")
247 w("Telemetry (load, queue, power, heat, network, drift, staleness, security) is assimilated into the state
each decision; an "
248 "allocation law turns the state into a demand target rho* (the HPA target), a node change, and a power cap.
Release of capacity "
249 "is gated by equation (2): capacity is only released once the control push has converged.")
250 w("")
251 w("### 2.2 The nervous system (muscles)")
252 w("Each muscle has five parts: afferent (pull: sense), the shared engine, efferent (push: act), reflex (the
shield checks every push) "
253 "and kill (hand the muscle back to its own controller). Status in this repository: **wired live on
Kubernetes** (sense and "
254 "push executed through kubectl, each with shield and kill): nodes, HPA target, power cap (in-place CPU
limits, enforced by "
255 "the kernel), security hold, deployment rollouts (pause, resume, undo), batch queue (admit held Jobs);
**sensed**: heat "
256 "(harness heat law on live power, or GPU temperature), network; **hardware connectors** (built and tested
with fake "
257 "hardware, off on CI machines): CPU power states (RAPL read, cpufreq ceiling) and GPU (nvidia-smi power
and temperature "
258 "read, power limit); **open** (registered, no plant yet): memory, storage, cooling, grid, training,
inference, agent "
259 "containment and the rest of the 52-muscle domain map (`docs/DOMAIN_MAP.md`). AI value alignment is
explicitly not an "
260 "Omni-Compass muscle. Code: `omni_controller/controller.py`, `omni_controller/muscles.py`, `omnicompass/
nervous.py`.")
261 w("")
262 w("### 2.3 The shield and the kill switch")
263 w("Before any action the shield (`omnicompass/shield.py`) enforces invariants I1 to I5: no expansion during
a security block, node "
264 "count within bounds, step limits, never below the capacity running and pending work needs, and the site
power limit. The kill "
265 "switch (a file or `OMNI_KILL=1`) restores every HPA target Omni-Compass changed from the recorded
original, returns the node pool "
266 "to its native size and drops to observe mode. Every decision and action is written to an append-only
audit log.")
267 w("")
268 w("### 2.4 Operating modes")
269 w("Observe (compute and log, write nothing), target (write HPA targets), nodepool (also size a node pool).
Laws: power_protect "
270 "(enforce the site power envelope), throughput (power envelope not enforced; capacity sized at full power)
and fleet variants "
271 "tuned on the 15-second fleet harness.")
272 w("")
273 w("## 3. The three architectures, and how each study realises them")
274 w("")
275 w("| Study | A. Kubernetes alone | B. Kubernetes + Omni-Compass | C. Omni-Compass direct | Live or simulated
|")
276 w("|---|---|---|---|---|")
277 w("| Pre-registered held-out stack benchmark (2 x 500 scenarios) | `k8s_ref_70`: documented HPA law (target

```

```

0.7, 10% tolerance, 300 s "
278 "stabilisation) and Cluster Autoscaler (scale-up on backlog, remove after 10 min under 50%), other
managers act on their own "
279 "proposals | `omni_k8s_throughput`: the same Kubernetes loops, Omni governs on top | `omni_direct`: Omni
senses and actuates "
280 "the stack directly | simulation |")
281 w(f"| Control-plane replica (24 scenarios) | `hpa70_ca`: metrics-server, HPA and Cluster Autoscaler replicas
at 15 s | "
282 "`omni_target_gate_hpa70_ca`: Omni writes the HPA target and gates Cluster Autoscaler scale-down |
`omni_throughput_full`: Omni "
283 "writes the HPA target, owns node scale-down and power cap; the Cluster Autoscaler may only add nodes |
simulation |")
284 w(f"| PlanetLab-shaped demand, fleet plant ({pn} scenarios) | `k8s_hpa70_ca` | `omni_target`: Omni writes
the HPA target | "
285 "`omni_fleet`: Omni is the node-pool authority, Cluster Autoscaler off | simulation on recorded traces |")
286 w("| Live kind cluster, side by side | native: HPA only, 6 workers always on, Omni not running | Omni on top:
sets the HPA target "
287 "and is the sole node-pool authority (cordon, drain, uncordon); Kubernetes' scheduler, kubelet and HPA
still execute | "
288 "not yet built live (section 16) | **live** |")
289 w("")
290 w("## 4. Method and why the comparison is fair")
291 w("")
292 w("- **Frozen before testing.** The allocation law, engine parameters, shield limits, modes, baselines and
benchmark code were "
293 "selected on development seeds (1000, 2000) and frozen, with SHA-256 hashes and program fingerprints in "
294 "`results/PREREGISTRATION.json`, before any run on the held-out seeds 346410161 and 360555127. `verify.py`
re-checks every hash.")
295 w("- **Same scenarios for every arm.** Each scenario is run under every architecture; comparisons are paired
scenario by scenario.")
296 w(f"- **Observe-identity check.** Omni-Compass in observe mode must produce a trajectory bit-identical to
the arm it observes, "
297 f"or the run is invalid. Held-out: {hobs[0][0]} and {hobs[1][0]} of 500 identical to the native stack, {
hobs[0][1]} and "
298 f"{hobs[1][1]} of 500 identical to Kubernetes; control-plane replica: {cobs} of {cn}; PlanetLab: {pobs} of
{pn}; live: 0 writes "
299 "while observing.")
300 w("- **Statistics.** Differences are paired (Omni minus Kubernetes on the same scenario) with 95% bootstrap
confidence intervals. "
301 "In the held-out study a difference is called better or worse only if the interval excludes zero in the
same direction on both "
302 "independent seeds; 'better in N' counts scenarios. Live results use 2-minute blocks and a bootstrap over
blocks.")
303 w("- **Ablations** show the engine, not an accident of tuning, produces the result (section 5.3).")
304 w("- **Independent implementations.** A C++ engine, governor, shield and HPA law are checked against the
Python ones (section 8).")
305 w("")
306 w("## 5. Results: pre-registered held-out benchmark (1,000 scenarios, simulation)")
307 w("")
308 w("Mean over two held-out seeds x 500 scenarios; verdicts require both seeds to agree.")
309 w("")
310 L.extend(hrows)
311 w("")
312 w("### 5.1 What the numbers say")
313 w(f"- **Energy:** B {pct(E(A), E(B))}, C {pct(E(A), E(C))} versus Kubernetes alone, although both keep
slightly more "
314 "node-hours powered. The saving comes from power capping and load shaping: time over the power limit
halves (B) or nearly "
315 "vanishes (C), and C cuts peak power by a quarter. Releasing idle machines, which the live cluster showed,
is not what drives "
316 "this study; idle node-hours are a gauge to tune.")
317 w("- **Reliability:** time healthy and recovery improve sharply in both B and C; the share of incidents
recovered rises.")
318 w("- **Coordination:** separate managers issue contradictory commands (A); a single authority issues none (B,
C). Pages and human "
319 "interventions go to zero because Omni-Compass acts on the conditions that would have paged someone.")
320 w("- **Heat:** C cuts time over the heat limit the most, because it governs power caps and load together.")
321 w("- **Costs:** node start/stop events rise (B +66%, C +102%), machine round trips rise, the power cap moves
more, and C's "
322 "queue and scale reversals are higher than Kubernetes alone. The 24-scenario control-plane study (section
6) shows the "
323 "opposite for wear (fewer machine stops), so wear depends on the plant and law and is a primary tuning
target. These are "
324 "reported, not tuned away.")
325 w("")
326 w("### 5.2 Architecture A has a hidden cost: the fragmented stack without Kubernetes")
327 w(f"For reference, the stack with each manager acting alone and no Kubernetes loops (`native`) used {native[
'energy_kwh']:.1f} kWh, "
328 f"was healthy {native['time_healthy']:.0%} of the time, took {native['recovery_minutes']:.0f} minutes to
recover and produced "
329 f"{native['contradictions']:.1f} contradictory commands and {native['invariant_violations']:.1f} invariant
violations per "
330 "scenario. Kubernetes already improves on that; Omni-Compass improves on Kubernetes.")
331 w("")
332 w("### 5.3 Ablations: is it the engine?")
333 w("| Comparison | Metric | Mean difference | 95% CI (seed 1) | 95% CI (seed 2) |")

```

```

334     w("---|---|---|---|---|---|")
335     for (x, k), es in sorted(abl.items()):
336         if k in ('energy_kwh', 'time_healthy', 'recovery_minutes', 'invariant_violations', 'mean_queue',
"violation_heat"):
337             w(f"| {x.replace('omni_direct_vs_', 'Omni direct vs ')} | {k} | {np.mean([e['mean_delta'] for e in
es]).+.3f} | "
338                 f"[{es[0]['ci95'][0]:+.3f}, {es[0]['ci95'][1]:+.3f}] | "
339                 f"[{es[-1]['ci95'][0]:+.3f}, {es[-1]['ci95'][1]:+.3f}] | ")
340         w("")
341     w("'no_dynamics' runs the same allocation law with equations (1) to (7) frozen; 'no_shield' removes the
shield. The difference "
342     "is full engine minus ablation: negative energy means the evolving dynamics save energy the frozen engine
does not; the "
343     "shield ablation shows what the shield prevents.")
344     w("")
345     w("## 6. Results: Kubernetes control-plane replica (24 scenarios, simulation)")
346     w("")
347     w("Replicas of the documented metrics-server, HPA (15 s) and Cluster Autoscaler (10 s) loops, with Karpenter-
lite as a second "
348     "Kubernetes reference; Omni-Compass decides every 300 s.")
349     w("")
350     L.extend(crows)
351     w("")
352     w("## 7. Results: live Kubernetes (measured)")
353     w("")
354     w("### 7.1 Side by side, two identical clusters at the same time")
355     w("kind clusters (real Kubernetes API server, scheduler, kubelet, HPA and metrics-server), 1 control plane +
6 workers each; the "
356     "official php-apache HPA workload (target 50); the same stepped load (1, 2, 3, 1, 2, 1 load generators
over 20 minutes after a "
357     "2-minute warm-up); capture every 15 s. Native: Omni-Compass not started. Omni: nodepool mode, HPA target,
node pool, power "
358     "sensing. Run 36209933218, 2026-09-26.")
359     w("")
360     w("| Gauge | Kubernetes alone | Kubernetes + Omni-Compass | Change |")
361     w("---|---|---|---|---|---|")
362     for lab, n, o, ch in live["gauges"]:
363         w(f"| {lab} | {n} | {o} | {ch} |")
364     w("")
365     w("| Metric (per unit of work, 2-minute blocks) | Kubernetes alone | Kubernetes + Omni-Compass | Change |
95% CI | Verdict |")
366     w("---|---|---|---|---|---|---|---|---|---|")
367     for r in live["significance"]:
368         w("| " + " | ".join(r) + " |")
369     w("")
370     w(f"Omni-Compass decisions: nodes per minute {live['nodes_decided']}; HPA target {live['hpa_target']}. Node-
pool resizes: 3 "
371     "(three workers cordoned and drained, their pods rescheduled by Kubernetes). Kill switch: HPA target
restored to 50, 6 of 6 "
372     "workers back in service.")
373     w("")
374     aml = json.loads((ROOT / "results/live/LIVE_ALLMUSCLE_1.json").read_text())
375     w("### 7.2 Live, side by side, every live muscle (run " + aml["run"] + ")")
376     w("Same two-cluster setup; the Omni arm drives " + aml["muscles"] + ". Real response times: HTTP requests to
the app timed every 5 s on both clusters. **This run made the application slower.**")
377     w("")
378     w("| Gauge | Kubernetes alone | Kubernetes + Omni-Compass | Change |")
379     w("---|---|---|---|---|---|")
380     for r in aml["gauges"]:
381         w("| " + " | ".join(r) + " |")
382     w("")
383     w("| Metric (2-minute blocks) | Kubernetes alone | + Omni-Compass | Change | 95% CI | Verdict |")
384     w("---|---|---|---|---|---|---|---|---|---|")
385     for r in aml["significance"]:
386         w("| " + " | ".join(r) + " |")
387     w("")
388     w("Actions: " + " | ".join(f"{k.replace('_', ' ')} {v}" for k, v in aml["actions"].items()) + ". Kill switch:
" + aml["kill_switch"] + ".")
389     w("")
390     w("***Diagnosis:** " + aml["diagnosis"])
391     w("")
392     w("### 7.3 Other live runs")
393     w("| Run | Setup | Result |")
394     w("---|---|---|---|")
395     for r in live["other_runs"]:
396         w("| " + " | ".join(r) + " |")
397     w("")
398     w("### 7.4 Live runs after the first all-muscle run")
399     w("")
400     w("| Run | Change tested | Validity | Result |")
401     w("---|---|---|---|---|")
402     for f in ("LIVE_ALLMUSCLE_2_LATENCY.json", "LIVE_ALLMUSCLE_3_SLO.json", "LIVE_ALLMUSCLE_4_FAILSAFE.json"):
403         pth = ROOT / "results/live" / f
404         if not pth.exists():
405             continue
406         r = json.loads(pth.read_text())

```

```

407         g = {row[0]: row for row in r["gauges"]}
408         def gv(name):
409             row = g.get(name)
410             return f"{row[1]} to {row[2]} ({row[3]})" if row else "n/a"
411         res = (f"p95 {gv('Response time (ms)', 95th percentile)}; p99 {gv('Response time (ms)', 99th
percentile)}");
412         f"energy {gv('Energy (Wh), standby counted')}")
413         w(f"| {r['run']} | {r['change']} | {r.get('validity', 'valid').split(':')[0]} | {res} |")
414         w("")
415         w("Run 3 is recorded as invalid: a refused Kubernetes call stopped the controller after 3 of 20 decisions
and the cluster held ")
416         "Omni's last settings with no engine running. The fix is a fail-safe: a failed decision is skipped; three
in a row restore "
417         "native settings and stop the controller. In run 4 the fail-safe did exactly that at minute 16, and the
audit showed why: the "
418         "least-privilege role allowed writing a pod's CPU limit but not reading it first, which kubectl does. In
runs 3 and 4 no power "
419         "cap was ever applied. Over run 4's 16 governed minutes node-hours per unit of work fell 24% and
utilisation rose 36% (both "
420         "significant); median response time -10%, p95 a tie, p99 +6% (not significant). The permission is fixed (
with a receipt) and a "
421         "run in which the fail-safe fires is now rejected; the re-run is in progress.")
422         w("")
423         w("## 8. Engineering verification")
424         w("")
425         w(f"python verify.py` runs {n_pass} checks ({n_fail} failures in the recorded log `results/VERIFY_LOG.txt`),
including:")
426         w("")
427         for key in ("SHA-256", "core parity", "C++ core vs 500 fixtures", "C++ governor vs Python governor (
throughput)",
428         "C++ shield vs Python shield", "negative control", "independent C++ HPA", "soak test (
throughput), 100,000,000",
429         "observe", "kill"):
430             for l in verify:
431                 if l.startswith("PASS") and key in l:
432                     w(f"- {l[4:].strip()}"); break
433         w("")
434         w("Also tested: the live controller against a fake cluster (observe writes nothing; targets bounded; kill
restores from the ")
435         "HPA annotation after a restart; node pool bounded and dry-run safe), the full-engine options (parked and
control-plane nodes ")
436         "excluded; the kill switch restores the node pool exactly once) and pilot scoring (detects a real gain,
reports no gain on ")
437         "identical clusters, detects a service regression).")
438         w("")
439         w("## 9. Physics and models")
440         w("")
441         w("- **Stack plant power:** each node draws idle 0.38 kW plus 1.12 kW x utilisation; a power cap throttles
delivered capacity.")
442         w("- **Heat:** thermal state follows a first-order lag toward 0.34 + 0.62 x power stress (time constant
about 7 steps); heat above ")
443         "0.82 throttles capacity; 'over the heat limit' means thermal above 1.03.")
444         w("- **Live kind cluster:** kind nodes have no power meter, so power is a declared model: 100 W idle + 150 W
x CPU utilisation per ")
445         "worker in service, plus a standby power for each parked (cordoned and drained) worker. Standby defaults
to the idle power ")
446         "(the worker stays powered and ready); lower values apply only to a declared sleep state, zero only to
machines really powered ")
447         "off. The first live reports counted parked workers as zero; section 7.1 gives both. The same constants
drive the governor's ")
448         "power sense and the energy score, so they cannot disagree. On real hardware this is replaced by metered
power (RAPL, PDU or BMC).")
449         w("- **Savings model** (`results/SAVINGS.csv`): a 1,000-node web cluster at 0.4 kW per node, PUE 1.4, $0.12/
kWh and 0.4 kg CO2/kWh; ")
450         "reduction versus HPA 0.7 + Karpenter-lite of 14% to 20% (fleet plant) gives roughly 710 to 960 MWh, $85,
000 to $115,000 and ")
451         "280 to 380 t CO2 per year.")
452         w("- **Engine overhead:** about 2.9 microseconds per decision in C++, memory flat over 100 million decisions.
")
453         w("")
454         w("## 10. What is proven, what is simulated, what is not claimed")
455         w("")
456         w("| Claim | Status | Evidence |")
457         w("|---|---|---|")
458         for r in [{"Engine equations are finite and deterministic; C++ equals Python", "Proven", "verify.py, 500
fixtures, max error 3.6e-15"},
459         ("Observe mode changes nothing", "Proven (simulation and live)", "bit-identical trajectories; 0
writes live"),
460         ("Kill switch restores native control", "Proven (simulation and live)", "HPA target 50 and all
workers restored live"),
461         ("Omni-Compass acts on a real Kubernetes control plane (HPA target, node pool)", "Proven live",
"section 7"),
462         ("Fewer nodes in service than Kubernetes with a fixed node pool, same load served", "Measured
live", "section 7.1"),
463         ("Lower energy on the live cluster", "Not shown while parked nodes stay on standby; -22% to -43%
only if parked nodes sleep or power off", "section 7.1"),

```

```

464         ("Better energy, health, recovery, coordination than Kubernetes (HPA + CA)", "Pre-registered
simulation", "section 5"),
465         ("Better than Karpenter-lite on energy", "Simulation", "sections 6, and PlanetLab fleet plant"),
466         ("Better than upstream Karpenter or Cluster Autoscaler binaries, live", "Not yet tested", "section
12"),
467         ("Metered energy savings on physical servers", "Not yet tested", "power is modelled"),
468         ("GPU power-limit muscle saves 15-19% energy with <1% slower responses", "Simulation calibrated to
metered H100 data (MLPerf)", "section 12"),
469         ("CPU frequency muscle", "Simulation (uncalibrated): about -3.5% energy, -67% heat", "section 12"),
470         ("Tuned laws remove wear and node-hour negatives (C-throughput)", "Pre-registered amendment, new
held-out data", "section 13"),
471         ("Cooling, grid, memory, storage and the other open muscles", "Not claimed", "no plant or
connector yet"),
472         ("Makes AI models aligned or trustworthy", "Not claimed", "value alignment is outside Omni-
Compass"))]:
473     w("| " + " | ".join(r) + " |")
474     w("")
475     w("## 11. Limits and threats to validity")
476     w("")
477     for s in ["Simulated studies use documented-behaviour replicas of Kubernetes controllers, not the upstream
binaries; the "
478         "Kubernetes reference omits Karpenter consolidation, VPA, scheduling constraints and disruption
budgets.",
479         "The live cluster is kind: nodes are containers on one CI machine; the two live arms ran on two
machines at the same "
480         "time, so machine-to-machine variation is part of the noise; each live arm is 20 minutes, one
repetition.",
481         "Live power and heat are modelled; parked kind workers are drained containers. If parked machines
must stay on "
482         "standby, parking reduces nodes in service but not energy; live energy savings then have to come
from power caps, CPU "
483         "power states and heat control, which are not yet wired live.",
484         "The live native arm had no node autoscaler, so its node pool was always full; the fair live
opponent is Karpenter or "
485         "Cluster Autoscaler (section 16).",
486         "The live significance for energy per core-hour with standby power is computed from minute
averages (10 two-minute "
487         "blocks), not from the 15-second capture.",
488         "The fleet law releases a node only when the pool has more than three nodes of slack; a three-
worker pool cannot scale "
489         "down (observed live, reproduced offline). Small clusters need a pool-size-aware release band.",
490         "Architecture C was measured in simulation only; the live C (Kubernetes' controllers parked, Omni-
Compass as the only "
491         "brain) is not built yet.",
492         "Service quality differences in the live runs (pending pods, HPA shortfall) are not statistically
significant at "
493         "this run length."):
494         w(f"- {s}")
495         w("")
496         # ----- device plant (GPU power limit, CPU frequency) -----
497         hw = json.loads((ROOT / "results/hardware/SUMMARY.json").read_text()); cal = json.loads((ROOT / "results/
hardware/CALIBRATION_MLPERF.json").read_text())
498         w("## 12. GPU and CPU muscles: device plant calibrated to metered hardware (simulation)")
499         w("")
500         w("The GPU power-limit and CPU frequency muscles cannot be actuated on CI machines (no GPU; the hypervisor
hides RAPL and "
501         "cpufreq). Their connectors are built (section 2.2) and this plant shows what they do. **GPU calibration:
** the "
502         "performance-versus-power-limit exponent is fitted to MLPerf Inference v4.0 results for an NVIDIA DGX-H100
(8 x H100-SXM, "
503         "700 W TDP), MaxQ (power-limited with `nvidia-smi -pl`, the same command the Omni-Compass GPU connector
sends) versus "
504         "MaxP, system power metered by a Yokogawa WT333E; Apache 2.0. Fitted exponent " + f"{cal['gamma_min']:.3f}
to {cal['gamma_max']:.3f} (median {cal['gamma_median']:.3f}); "
505         "all three are run. **CPU:** a standard first-order model (dynamic power ~ frequency cubed) against a
schedutil-style "
506         "governor; not calibrated to metered data yet. **S** is a fixed manual 70% cap, what an operator could do
by hand. "
507         "24 scenarios, 8 load families; * = paired 95% interval excludes zero.")
508         w("")
509         w("| Vessel | Gauge | A. Native | S. Fixed 70% cap | B. + Omni-Compass | C. Omni-Compass direct |")
510         w("| ---|---|---|---|---|---|---|")
511         labels = [("energy_kwh", "Energy (kWh)", ("kwh_per_work", "Energy per unit of work"), ("p95_latency_x",
"95th-pct response time (x baseline)"),
512             ("slo_breach_min", "Minutes over service target"), ("heat_over_min", "Minutes over heat limit"), ("
peak_kw", "Peak power (kW)",
513             ("power_over_min", "Minutes near full power"))]
514         for vn, v in hw["vessels"].items():
515             m = v["means"]
516             for k, lab in labels:
517                 a0 = m["A"][k]
518                 cell = lambda arm: f"{m[arm][k]:.3f} ({pct(a0, m[arm][k])}{'*' if v['paired_vs_A'][arm][k]
'significant' else ''})"
519                 w(f"| {vn} | {lab} | {a0:.3f} | {cell('S')} | {cell('B')} | {cell('C')} |")
520         w("")
521         w("***Reading:** on GPUs Omni-Compass saves 15-19% energy across the measured range with response time 0.6-0.

```

```

8% slower; a "
522     "fixed cap saves about the same energy but slows responses 30-63% and misses the service target. On CPUs
the native "
523     "governor already tracks demand, so frequency control alone saves about 3.5%; its main effect is heat (-
67%). CPU "
524     "frequency control is below the 10% bar and is a physical ceiling of that muscle, not a tuning gap.")
525     w("")
526     # ----- amendment -----
527     am = json.loads((ROOT / "tuning/HELDOUT_AMENDMENT.json").read_text())
528     w("## 13. Amendment: tuned laws, frozen, then tested on new held-out data (simulation)")
529     w("")
530     w("Two law variants were tuned on the development seeds only (1000, 2000; `tuning/SEARCH*.json`), frozen
with SHA-256 hashes "
531     "in `tuning/PREREGISTRATION_AMENDMENT_2026-09-26.json` and pushed (commit b79d1dd, 03:17 UTC) before a
single run on new "
532     "held-out seeds 731001 and 731002 (2 x 500 scenarios). The frozen engine files are unchanged; the original
pre-registered "
533     "result (section 5) stays the primary result. B-wear: node release held longer. C-throughput: Omni-Compass
direct under "
534     "the throughput law with a wider release band and engine-gated cap. + better, ! worse (both seeds agree),
blank not significant.")
535     w("")
536     names = ["B_frozen", "B_wear", "C_frozen", "C_throughput"]
537     w("| Gauge | A. Kubernetes alone | " + " | ".join(names) + " |")
538     w("|---|---|:" + "----:" * len(names))
539     A0 = am["means"]["A_k8s"]
540     for k in A0:
541         cells = []
542         for n in names:
543             vv = am["means"][n][k]; vd = am["paired"][n][k]["verdict"]
544             cells.append(f'{vv:.3f} ({pct(A0[k], vv)})' + '!' if vd == 'better' else '!' if vd == 'worse' else
''')
545         w(f"| {k} | {A0[k]:.3f} | " + " | ".join(cells) + " |")
546     w("")
547     w("Worse gauges: " + " | ".join(f'{n} {sum(v["verdict"] == "worse" for v in am["paired"][n].values())}' for n
in names) + ".")
548     w("")
549     # ----- response time pass -----
550     w("## 14. Response time: can one mode beat Kubernetes on every gauge? (simulation, development seeds)")
551     w("")
552     w("The fleet-mode constants that drove the live runs were selected (before this work) by a rule that scored
energy, finished "
553     "work and backlog, but never waiting time. A response-time gauge was added to the fleet plant (fleet/
sim_slo.py: M/M/c queueing "
554     "delay per workload plus backlog drain; the frozen trace is reproduced exactly). Measured with it, fleet
mode cut energy 31% on "
555     "web services while p95 response time went from 132 ms to about 20 s. That is where the live slowdown came
from: its HPA target "
556     "floor (rho_min = 76.9%) packs pods too hot for latency-sensitive services.")
557     w("")
558     w("A speed-first law (omnicompass/speed.py; the engine equations unchanged) separates the two jobs the
frozen law gave one number: "
559     "pods run at a latency-safe target while machines are sized to what the pods request. 400 configurations
were searched against "
560     "HPA 70% + Cluster Autoscaler with the rule that no gauge may be worse in any development scenario. None
passed, for three measured "
561     "reasons: (1) several gauges are already at their physical floor (100% work done, zero violations, batch
p95 = pure service "
562     "time), so a tie is the best any controller can do; (2) on the GPU vessel demand exceeds the hardware, so
every arm is saturated; "
563     "(3) energy, response time and machine churn trade two-of-three, because the autoscaler's idle slack is
both where the energy "
564     "is and what absorbs the 90-second boot delay. Examples, mean change vs Kubernetes: fleet mode energy -31%
with far worse response "
565     "time; speed #176 energy -1%, p99 -41%, churn 3.6x; speed #159 p99 -48%, churn -27%, energy +16%. Source: "
566     "tuning/SPEED_FINDINGS_2026-09-26.md.")
567     w("")
568     # ----- named products -----
569     vc = json.loads((ROOT / "tuning/VENDOR_COMPARE_DEV.json").read_text())
570     w("## 15. Named products: OpenShift, Google GKE, Azure AKS, IBM Turbonomic (simulation, development seeds)")
571     w("")
572     w("Each product is emulated from its documented behaviour on the same fleet plant (not the vendors'
binaries): OpenShift's "
573     "documented ClusterAutoscaler example (threshold 0.4, unneeded 5 min, delay after add 10 min); GKE
optimize-utilization "
574     "(MostAllocated packing, more aggressive scale-down; declared as threshold 0.65, 2 min, because Google
publishes no numbers); "
575     "AKS node auto-provisioning (Karpenter, WhenEmptyOrUnderutilized, consolidateAfter 0 s); Turbonomic (
container requests resized "
576     "every 10 min to p99 per-pod usage, its default aggressiveness; nodes suspended toward 0.7 packing). Means
over web services, "
577     "4 development seeds:")
578     w("")
579     arms = ["k8s_hpa70_ca", "openshift", "gke_optimize", "aks_nap", "turbonomic", "omni_fleet", "speed159"]
580     names = {"k8s_hpa70_ca": "Kubernetes (GKE balanced)", "openshift": "OpenShift", "gke_optimize": "GKE
optimize",

```

```

581         "aks_nap": "AKS NAP", "turbonomic": "Turbonomic", "omni_fleet": "Omni fleet mode", "speed159":
"Omni speed #159"}
582     w("| Gauge (web) | " + " | ".join(names[a] for a in arms) + " |")
583     w("|---|" + "----:" * len(arms))
584     for m, lab in [("energy_kwh", "energy (kWh)", ("p95_ms", "p95 (ms)"), ("p99_ms", "p99 (ms)"), ("start_stop",
"machine starts+stops"),
585                  ("node_reversals", "scale reversals"))]:
586         w(f"| {lab} | " + " | ".join(f"{vc['web'][a][m]:.4g}" for a in arms) + " |")
587     w("")
588     w("Every product sits on the same energy / response-time / churn triangle; none wins all three. Carpenter-
style consolidation ")
589     w("AKS NAP) saves about 20% energy with a worse p99 tail and about twice the machine churn; Turbonomic's
p99 resizing interacts ")
590     w("with the HPA and inflates the tail; OpenShift's documented settings sit close to upstream. Omni's speed
mode leads on response ")
591     w("time and churn at an energy cost; its fleet mode leads on energy at a large response-time cost.")
592     w("")
593     # ----- problem-map muscles -----
594     w("## 16. The problem-map muscles: all nine built and tested on held-out data (simulation)")
595     w("")
596     w("Each open row of the industry problem map (section 18) is now a muscle in omnilib/, with the strongest
native tool as opponent, ")
597     w("Omni-Compass, and Omni-Compass with its engine equations not evolved (to show what the equations
themselves add). Constants ")
598     w("were chosen on development seeds 1-8, every file was frozen by SHA-256 (results/muscles/PREREGISTRATION.
json), then 30 held-out ")
599     w("seeds were run once. Verdicts from a paired bootstrap 95% interval.")
600     w("")
601     w("| Muscle (map row) | Strongest native opponent | Omni better | Omni worse | Tie / not significant |")
602     w("|---|---|---|---|---|")
603     best = {"rightsize": "native_hpa_vpa", "coldstart": "native_keda", "gpupack": "native_binpack",
"powersmooth": "native_floor_safe",
604     "health": "native_detect", "cooling": "native_reset", "inference": "native_keda", "containment":
"native_static",
605     "vmenergy": "native_ratio"}
606     label = {"rightsize": "Right-sizing, HPA+VPA conflict (1, 2)", "coldstart": "Cold start (5)", "gpupack":
"GPU packing (6)",
607     "powersmooth": "Training power swings (7)", "health": "GPU failures and stragglers (8)", "cooling":
"Cooling (10)",
608     "vmenergy": "Energy attribution in VMs (11)", "inference": "LLM inference, KV cache (12)",
"containment": "Runaway AI agents (13)"}
609     eng_gap = []
610     for m, nat in best.items():
611         r = json.loads((ROOT / f"results/muscles/{m.upper()}_HELDOUT.json").read_text())
612         o = r["vs"][nat]["omni"]; ne = r["vs"][nat]["omni_no_engine"]
613         def ch(g):
614             a0, b0 = o[g]["native"], o[g]["omni"]
615             return f"{g.replace('_', ' ')} {{(b0 - a0) / abs(a0) * 100:+.0f}%}" if abs(a0) > 1e-12 else f"{g.
replace('_', ' ')} {{a0:.3g}} to {{b0:.3g}}"
616         b = [ch(g) for g in o if o[g]["verdict"] == "better"]
617         ww = [ch(g) for g in o if o[g]["verdict"] == "worse"]
618         t = [g.replace('_', ' ') for g in o if o[g]["verdict"] in ("tie", "not significant")]
619         w(f"| {label[m]} | {nat.replace('native_', ' ')} | {'; '.join(b) or '-'} | {'; '.join(ww) or '-'} | {'; '.
join(t) or '-'} |")
620         eng_gap.append(sum(abs(o[g]['better_by'] - ne[g]['better_by']) for g in o) / len(o))
621     w("")
622     w("Each entry is the plain change of Omni-Compass against the opponent (for example p99 -99% means the slow
tail is 99% ")
623     w("shorter; goodput +10% means 10% more useful training time). Worse entries are real costs, not rounding:
right-sizing ")
624     w("trades a slightly slower typical response (p95 +15%, both far under the 500 ms target) and a few more
memory kills for a 99% ")
625     w("shorter tail; cold start keeps fewer idle instances than KEDA's 5-minute cooldown, so more requests meet
a cold start, but ")
626     w("users wait far less because KEDA polls every 30 s; GPU packing powers idle GPUs off sooner and jobs wait
a few seconds longer; ")
627     w("containment throttles honest agents in their legitimate bursts (about 7% of their work) and stops about
two honest agents a ")
628     w("day, against a 96% cut in runaway spend and containment in minutes instead of hours. Against the other
native arms (the ")
629     w("default tools most teams run) the wins are larger; every comparison is in results/muscles/*_HELDOUT.json.
")
630     w("")
631     w(f"***What the engine itself contributes.** Across the nine muscles the full engine and the engine-not-
evolved arm differ by ")
632     w(f"{100*sum(eng_gap)/len(eng_gap):.1f} percentage points on average per gauge. The mapping from engine
state to action carries most ")
633     w("of each result; the evolved dynamics mainly smooth. In power smoothing an arm where the bath equation (7)
alone sets the site's ")
634     w("draw, with no ramp rule, cut the steepest ramp 63% but did not hold the grid's limit: the human sets the
boundary, the engine ")
635     w("operates inside it.")
636     w("")
637     # ----- every negative -----
638     w("## 17. Every negative, its cause and its status")
639     w("")

```

```

640     w("| Negative | Where | Status | Cause | What would fix it |")
641     w("----|----|----|----|----|")
642     for r in [
643         ("Node start/stop cycles (wear)", "B, pre-registered", "reduced (+65% to +34%), not removed", "node
release thresholds are fixed numbers", "hybrid: Kubernetes serves the queue, Omni releases nodes with the C-
throughput law (removed wear there: -50%)"),
644         ("Node start/stop, round trips, reversals", "C, pre-registered", "fixed in C-throughput (new held-out)",
"power-protect law released nodes too eagerly", "adopted in C-throughput"),
645         ("Node-hours and idle node-hours", "B and C, pre-registered", "fixed in C-throughput (-10%, -23%); not
in B", "power capping trades lower watts for more servers on", "hybrid as above; node-aware cap"),
646         ("Queue / backlog / availability -0.3%", "C (both variants)", "UNRESOLVED", "not the replica law and not
the sizing constants (both tested); likely the delayed observation or direct-mode proposals", "trace one
scenario step by step; candidate: feed the queue into the engine without the extra delay"),
647         ("Power-cap movement", "B and C", "inherent", "moving the cap is how capping saves energy; freeing it
removed most of the saving (tested: -28% to -2%/-7%)", "none needed: electronic setting, no physical wear;
reported"),
648         ("CPU frequency saving 3.5%", "device plant", "physical ceiling", "the native Linux governor already
follows demand", "value is in heat (-67%) and in combining with power caps"),
649         ("GPU node on/off saving 3%", "fleet plant", "superseded", "training nodes cannot be switched off", "the
GPU power-limit muscle (15-19%)"),
650         ("Live response time +65% (p95) and energy per unit of work +46% with every muscle", "live kind, run
36213152881", "reduced to p95 +25% (run 2), tie in run 4",
651         "no response-time afferent; cap held at the law floor; HPA target floor 76.9%", "latency afferent, cap
reflex, SLO reflex; speed-first law (section 14)"),
652         ("Fleet mode p95 20 s vs 132 ms on web services", "fleet plant with response-time gauge", "cause found;
speed-first law built",
653         "fleet constants were selected without a response-time gauge; HPA target floor 76.9%", "choose the mode
per service: speed-first where latency matters"),
654         ("No mode better than Kubernetes on every gauge in every scenario", "fleet plant, 400 configurations",
"not achievable as posed",
655         "gauges at physical floors; hardware-bound GPU vessel; energy / response time / churn trade two-of-
three", "anticipation (pre-adding machines before the daily rise) is the one untested mechanism that could
break the trade"),
656         ("Controller stopped after one refused call; power cap never applied under least privilege", "live runs
3 and 4", "fixed",
657         "one exception ended the loop; the role lacked get on pods/resize", "fail-safe restore after 3 failures;
permission and receipt added; runs with a fail-safe rejected"),
658         ("Right-sizing p95 +15%, more memory kills than VPA", "omnilab rightsize, held-out", "open, small",
659         "Omni sizes CPU closer to demand; VPA's 8-hour p90 memory keeps more slack", "larger memory margin (
costs memory-hours)"),
660         ("More cold starts than KEDA", "omnilab coldstart, held-out", "trade-off",
661         "shorter keep-alive than KEDA's 5-minute cooldown", "longer keep-alive where cold starts matter more
than instance-hours"),
662         ("GPU jobs wait +16% (seconds) vs Volcano binpack", "omnilab gpupack, held-out", "trade-off", "idle GPUs
powered off sooner", "longer power-off delay"),
663         ("Honest agents throttled (-7% work) and ~2 false stops a day vs static caps", "omnilab containment,
held-out", "trade-off",
664         "throttling on the engine's integrated need catches legitimate bursts", "per-agent declared burst
budgets; human approval before stop"),
665         ("Engine dynamics add little beyond the mapping", "all nine muscles", "reported", "the mapping from
state to action carries the effect", "wire the engine's control effort (equation 2) directly as the actuator
command and test it"),
666         ("Live energy with parked servers on standby ~0%", "live kind", "open", "parked servers still draw
standby power", "live power cap and CPU/GPU muscles; sleep states where hardware allows"),
667         ("Small clusters (<= 3 workers) never release", "live kind", "open", "fleet law release band of 3 nodes",
"pool-size-aware release band (law change)"),
668         ("Pages in the 24-scenario replica", "control-plane replica", "open", "longer queue triggers the page
rule", "same as queue"),
669     ]:
670         w("| " + " | ".join(r) + " |")
671         w("")
672         w("## 18. The industry problem map: what Omni-Compass is aimed at")
673         w("")
674         w("One engine; the vessel (the plant it sits on) is the only thing that changes. Industry figures are
approximate, from the ")
675         "public sources named, and are context, not results of this report. Status: **live** = measured on a real
Kubernetes control "
676         "plane; **sim** = demonstrated in this repository's simulations; **open** = mapped, connector not built.")
677         w("")
678         w("| Problem | Scale in the industry (approximate, source) | Best software today | Omni-Compass vessel and
muscles | Gauge that shows it | Status |")
679         w("----|----|----|----|----|")
680         for r in [
681             ("Data-centre electricity growth", "about 415 TWh in 2024, about 1.5% of world electricity, projected
near 945 TWh by 2030 (IEA, Energy and AI, 2025)",
682             "Karpenter, Cluster Autoscaler, CAST AI, Spot Ocean; Kepler for metering", "compute vessel: nodes, HPA,
power cap", "energy, node-hours, idle node-hours",
683             "sim; live only where parked nodes can sleep or power off"),
684             ("Idle and over-provisioned capacity", "Kubernetes clusters commonly run near 10-15% average CPU
utilisation (CAST AI and Datadog industry reports); roughly a quarter to a third of cloud spend reported as
waste (Flexera State of the Cloud)",
685             "VPA, Goldilocks, StormForge, Kubecost/OpenCost", "compute vessel: nodes, HPA; requests and memory (
omnilab/rightsize.py)", "utilisation, node-hours per core-hour, core- and GiB-hours", "live + sim"),
686             ("Controllers fighting each other", "documented conflicts, e.g. HPA and VPA on the same CPU metric (
Kubernetes documentation advises against it)",
687             "none: each tool decides alone", "single authority over replicas and requests", "contradictory commands,

```

```

scale reversals, OOM kills", "sim, held-out (section 16)"),
688 ("Outages and slow recovery", "most significant outages cost over $100,000 (Uptime Institute annual
outage analysis)",
689 "Argo Rollouts, Flagger, SRE runbooks, AIOps (Dynatrace, Datadog)", "compute vessel + deployments (
partial)", "time healthy, recovery time, SLA breaches", "sim"),
690 ("On-call load and alert fatigue", "widely reported burnout in SRE surveys", "PagerDuty, alert tuning",
"all muscles: act before the page", "pages, human interventions", "sim"),
691 ("Heat and cooling limits", "cooling is a large share of facility energy; average PUE about 1.5 (Uptime
Institute survey)",
692 "DCIM (Schneider EcoStruxure), DeepMind cooling AI (reported about 40% less cooling energy)", "heat (
sensed), cooling plant (omnilab/cooling.py)", "PUE, cooling energy, inlet violations", "sim, held-out (section
16)"),
693 ("Site power and grid-connection limits", "multi-year waits for new grid connections are widely
reported", "Meta Dynamo power capping, Intel RAPL",
694 "power cap (wired), batteries and demand response (open)", "peak power, time over power limit", "sim"),
695 ("GPU energy and power limits", "GPU fleets widely reported well below full utilisation; H100 TDP 700 W",
"NVIDIA DCGM and MIG, Run:ai, Kueue; manual MaxQ power limits",
696 "GPU power-limit muscle (hardware connector)", "energy per unit of work, response time, heat", "sim,
calibrated to MLPerf metered H100 data: -15% to -19% energy, +0.6-0.8% response time"),
697 ("Batch deadlines and fair sharing", "", "Kueue, Volcano, Slurm", "batch muscle (live: admit held jobs);
GPU packing (omnilab/gpupack.py)", "queue wait, fragmentation", "live wired + sim, held-out"),
698 ("Hardware wear", "power cycling and churn shorten component life", "none as a governed objective",
"nodes: start/stop cycles and reversals",
699 "machines started and stopped, round trips", "mixed: better in the 24-scenario study, worse in the held-
out study; tuning target"),
700 ("Carbon reporting and reduction", "regulatory disclosure is expanding", "Google carbon-aware computing,
Kepler", "carbon-aware placement (open)", "kWh and CO2 per unit of work", "sim (modelled)"),
701 ("Runaway AI agents and spend", "~$10,000 overnight examples (Dark Reading)", "per-tool quotas and
permissions", "agent containment muscle (omnilab/containment.py)", "rogue spend, time to contain, false stops",
"sim, held-out (section 16)"),
702 ]:
703     w("| " + " | ".join(r) + " |")
704     w("")
705     w("## 19. What comes next")
706     w("")
707     for s in ["Live architecture C: park HPA, VPA, Cluster Autoscaler and Karpenter; Omni-Compass sets replicas,
resources, "
708         "placement, priorities and quotas directly; Kubernetes keeps execution and reflexes (restarts,
rescheduling); the kill "
709         "switch wakes the parked controllers.",
710         "Live opponent at full strength: Karpenter (kwok provider) and Cluster Autoscaler in architecture
A.",
711         "24 live scenarios (traffic, failures, power and heat limits, batch and AI, growth, mixed) with
repetitions.",
712         "More muscles two-way: CPU power states, memory, batch queues, network, security, then GPU and
cooling on hardware.",
713         "Metered power on physical machines."]:
714         w(f"- {s}")
715         w("")
716         w("## 20. Questions and answers")
717         w("")
718         qa = [
719             ("Does Omni-Compass replace Kubernetes?", "No. Kubernetes keeps running containers, placing pods,
restarting failures and "
720             "networking. Omni-Compass replaces the separate decision loops (how many replicas, how many nodes, what
power) with one "
721             "authority. In architecture C Kubernetes becomes one muscle."),
722             ("What happens if Omni-Compass crashes or is switched off?", "The kill switch restores every setting it
changed and returns "
723             "control to Kubernetes' own controllers; this was exercised live and in simulation. A crashed
controller writes nothing further."),
724             ("Can it make things worse?", "Every action passes the shield first; it never goes below the capacity
running and pending work "
725             "needs, and never changes more than the step limit. In the held-out benchmark it had fewer safety
violations than Kubernetes. "
726             "Its real costs are listed in sections 1 and 5.1."),
727             ("How fast does it decide, and what does it cost to run?", "One decision per 60 s on live Kubernetes (
300 s in the replica), "
728             "with a 15 s fast path that adds nodes for pending pods. The engine takes about 2.9 microseconds per
decision."),
729             ("Why does it save energy?", "In the pre-registered stack study mainly by power capping and load shaping
(fewer minutes over "
730             "the power limit, lower peak), while keeping slightly more machines on. On the live cluster it took
machines out of service "
731             "(6 to 3 workers) and packed replicas more densely; that saves energy only if the parked machines sleep
or power off. With "
732             "parked machines on standby the live saving was about zero, so live savings must come from power caps
and CPU power states."),
733             ("Does it slow applications down?", "It can, in the energy-first fleet mode: its HPA target floor packs
pods too hot for "
734             "latency-sensitive services (section 14). Live, the first all-muscle run was slower (p95 +65%); after
the fixes run 4 was a tie at "
735             "p95 and 10% faster at the median. For latency-sensitive services the speed-first mode is the right
setting."),
736             ("Can Omni-Compass control AI agents?", "It controls what an agent can touch, spend and do, and how fast;
not what the model ")

```

```

737     "thinks. In the containment muscle (section 16) it cut runaway spend 96% and contained runaways in
minutes instead of hours, "
738     "with a least-privilege identity that blocks forbidden actions outright; the cost is some throttling of
honest agents' bursts."),
739     ("Does the engine hold everything in its basin by itself?", "Not against an outside limit it is not told.
The bath equation alone "
740     "smoothed training power ramps 63% but did not keep them under the grid's limit; with the human-set
limit as the boundary it "
741     "held it with zero violations. The human sets the boundaries; the engine operates inside them."),
742     ("Is this tuned to the test?", "Parameters were selected on development seeds and frozen with hashes
before the held-out seeds "
743     "were run; verify.py fails if any frozen file changes."),
744     ("How many scenarios and how certain?", "1,000 pre-registered held-out scenarios, 24 control-plane
scenarios, PlanetLab traces "
745     "and live runs; 95% bootstrap intervals, two independent seeds must agree."),
746     ("What is measured versus modelled?", "Live: node counts, replicas, pods, CPU, the controller's actions
and the kill switch. "
747     "Modelled: power and heat everywhere, and everything in the simulated studies."),
748     ("Is the mathematics sound?", "The core is a closed six-state system integrated with RK4; the C++ and
Python implementations "
749     "agree to 3.6e-15 on 500 reference trajectories; 100 million decisions ran without a non-finite value.
"),
750     ("Who owns it and how can it be used?", "The Omni-Compass LLC. Free for evaluation, research and non-
commercial use; commercial "
751     "use requires a paid licence; protected by copyright and by patents and patent applications (see
LICENSE and NOTICE)."),
752     ("What is not claimed?", "Superiority over upstream Karpenter or Cluster Autoscaler live, metered
savings on physical hardware, "
753     "GPU or facility control on real hardware, the vendor products' own binaries (they are emulated from
documentation), and anything about AI value alignment."),
754     ("How do I check it myself?", "Run the commands in section 21; the live runs are GitHub Actions
workflows in the repository."),
755 ]
756 for q, ans in qa:
757     w(f"***{q}** {ans}")
758     w("")
759 w("## 21. Reproduce")
760 w("")
761 w("```")
762 for c in ["pip install -r requirements.txt",
763          "python verify.py",
764          "python benchmarks/stack_benchmark.py ...",
765          "python -m k8s_controlplane.benchmark --scenarios 24 --seed 424242",
766          "python -m fleet.planetlab --dir fleet/traces/planetlab --scenarios 8 --out /tmp/pl",
767          "GitHub Actions: benchmark (live side by side), live-kind-full, live-kind",
768          "python tools/full_report.py ... && python pilot/bench_pdf.py docs/BENCHMARK_REPORT.md docs/
BENCHMARK_REPORT.pdf"]:
769     w(c)
770     w("```")
771     w("")
772 w("## 22. Glossary")
773 w("")
774 for t, d in [("HPA", "Horizontal Pod Autoscaler: Kubernetes controller that sets replica counts from CPU
utilisation versus a target."),
775             ("Cluster Autoscaler, Karpenter", "Kubernetes add-ons that add and remove nodes."),
776             ("Node, worker", "a machine (here a container in kind) that runs pods."),
777             ("Cordon, drain", "mark a node unschedulable, then move its pods elsewhere."),
778             ("kind", "Kubernetes in Docker: a real Kubernetes control plane whose nodes are containers."),
779             ("Observe mode", "Omni-Compass computes and logs but writes nothing."),
780             ("Invariant", "a safety rule the shield enforces before any action."),
781             ("Paired bootstrap CI", "resampling the per-scenario differences to get a 95% interval for the
mean difference."),
782             ("Pre-registration", "freezing code and parameters, with hashes, before running the test data.
")]:
783     w(f"- **{t}**: {d}")
784     w("")
785     Path(a.out).write_text("\n".join(L))
786     print(f"wrote {a.out} ({len(L)} lines)")
787
788
789 if __name__ == "__main__":
790     main()
791

```

tools/gpu_baseline.py

71 lines SHA-256 a6811496c803c27ed276d5be795d48244b65f66925722f95470d6b7ada40d4de

```
1  """The speed lock's baseline: response time without Omni, as a function of arrival rate.
2
3  python3 tools/gpu_baseline.py baseline.json native-run-1/latency.csv [native-run-2/latency.csv ...]
4
5  Each file is a run without Omni (elapsed_seconds,latency_ms,ok, as tools/gpu_workload.py writes). Windows of
6  --window-s seconds, one every --every-s, give (arrival rate, mean, p95, p99); windows are grouped by rate in bins
7  of
8  --bin requests per second and each bin keeps the median of its windows. omni_controller/gpu_governor.py
9  --baseline-file reads the result with the same window, so the lock compares like with like. Use runs that are not
10 the
11 ones the comparison is made on (other seeds), so the lock is not tuned on the arms it is judged by.
12 """
13 from __future__ import annotations
14 import argparse, json, statistics, sys
15 from pathlib import Path
16
17 def windows(rows, window_s, every_s):
18     """rows: [(t, ms, ok)] sorted by t. Yields (rate, mean, p95, p99) per window with at least 20 requests."""
19     if not rows:
20         return
21     t0, t1 = rows[0][0], rows[-1][0]
22     t, i0 = t0 + window_s, 0
23     while t <= t1:
24         while rows[i0][0] < t - window_s:
25             i0 += 1
26         ms = sorted(m for (tt, m, ok) in rows[i0:] if tt < t and ok)
27         if len(ms) >= 20:
28             yield len(ms) / window_s, sum(ms) / len(ms), ms[int(0.95 * len(ms))], ms[min(len(ms) - 1, int(0.99 *
29 len(ms)))]
30         t += every_s
31
32 def build(runs, window_s=60.0, every_s=5.0, bin_w=1.0):
33     """runs: list of row lists. Returns the baseline dict."""
34     groups = {}
35     for rows in runs:
36         for rate, mean, p95, p99 in windows(sorted(rows), window_s, every_s):
37             groups.setdefault(round(rate / bin_w), []).append((rate, mean, p95, p99))
38     bins = []
39     for k in sorted(groups):
40         w = groups[k]
41         bins.append({"rate": statistics.median(x[0] for x in w), "mean": statistics.median(x[1] for x in w),
42 "p95": statistics.median(x[2] for x in w), "p99": statistics.median(x[3] for x in w),
43 "windows": len(w)})
44     return {"window_s": window_s, "every_s": every_s, "bin": bin_w, "runs": len(runs), "bins": bins}
45
46 def read(path):
47     out = []
48     for line in Path(path).read_text().splitlines()[1:]:
49         if line.strip():
50             t, ms, ok = line.split(",")[3:]
51             out.append((float(t), float(ms), ok.strip() == "1"))
52     return out
53
54
55 def main(argv=None):
56     ap = argparse.ArgumentParser(description=__doc__, formatter_class=argparse.RawDescriptionHelpFormatter)
57     ap.add_argument("out")
58     ap.add_argument("latency", nargs="+")
59     ap.add_argument("--window-s", type=float, default=60.0)
60     ap.add_argument("--every-s", type=float, default=5.0)
61     ap.add_argument("--bin", type=float, default=1.0)
62     a = ap.parse_args(argv)
63     b = build([read(p) for p in a.latency], a.window_s, a.every_s, a.bin)
64     Path(a.out).write_text(json.dumps(b, indent=1))
65     print(f"{len(b['bins'])} rate bins from {b['runs']} run(s): " + ", ".join(f"{x['rate']:.1f}/s p95 {x['p95']:.
66 0f} ms" for x in b["bins"]))
67     return 0
68
69 if __name__ == "__main__":
70     sys.exit(main())
71
```

tools/gpu_physics_sim.py

177 lines SHA-256 918040d9868e07fb130975d44d55dc162f66776a8126e109e3a0f8d77b27f8a5

```
1  """The GPU test without a GPU: Omni-Compass's own GPU governor (omni_controller/gpu_governor.py, its command-
2  line defaults)
3  against a modelled NVIDIA card, in virtual time. A preview, not evidence: every energy number here comes from
4  the model
5  below, not from a meter. The hardware answer is scripts/gpu_paired.sh on a real card.
6
7  5      python3 tools/gpu_physics_sim.py                                # 10 paired repetitions, both card models, table to
8  stdout
9  6      python3 tools/gpu_physics_sim.py --reps 5 --out results/gpu/sim --gov "---min-share 0 --util-gate 0" # other
10 settings
11
12 7  The card (an assumption, stated):
13 8  idle draw      40 W while no request is being served
14 9  busy draw      P(s) = 80 + 220 s^gamma W at speed s (1 = full clock); 300 W at full clock, the start limit
15 10 power limit L  below 300 W the card slows to the speed where P(s) = L, as real cards do by lowering clocks
16 11 heat          first order toward 30 C + 0.12 C/W x draw, time constant 20 s
17 12 gamma is how much the top of the clock range costs: 3 when voltage falls with clock (good for capping), 1.5 when
18 the
19 card already sits near its voltage floor (hard for capping). Both are run; a real card lies somewhere between.
20
21 16 The workload is tools/gpu_workload.py's: Poisson arrivals at 30, 60, 80, 30, 60, 30% of full-speed capacity, seed
22 20260928 + repetition, one server first come first served, 50 ms per request at full speed, 600 s plus 30 s
23 drain; the
24 response-time target is ten service times (500 ms), as scripts/gpu_paired.sh sets it. The governor decides at
25 its own
26 interval (default 2 s) from the modelled card's nvidia-smi readings and the response-time file, exactly as on
27 hardware.
28 """
29 from __future__ import annotations
30
31 23 import argparse, json, math, os, random, shlex, statistics, sys, tempfile
32 24 from pathlib import Path
33
34 26 ROOT = Path(__file__).resolve().parents[1]; sys.path.insert(0, str(ROOT))
35 27 from omni_controller.gpu_governor import GpuGovernor, parser as gov_parser
36
37 29 START, MIN = 300.0, 100.0
38 30 P_IDLE, P_STATIC, P_DEMAND = 40.0, 80.0, 300.0
39 31 SERVICE_MS, SLO_MS = 50.0, 500.0
40 32 PHASES = [0.3, 0.6, 0.8, 0.3, 0.6, 0.3]
41 33 DURATION, DRAIN = 600, 30
42 34 T95 = {1: 12.706, 2: 4.303, 3: 3.182, 4: 2.776, 5: 2.571, 6: 2.447, 7: 2.365, 8: 2.306, 9: 2.262, 10: 2.228,
43 35 14: 2.145, 19: 2.093, 29: 2.045}
44
45 37 SMI = '"/usr/bin/env python3
46 38 import json, os, sys
47 39 p = os.environ["PHYS_STATE"]; st = json.load(open(p)); a = sys.argv[1:]
48 40 if "-pl" in a:
49 41     st["limit"] = float(a[a.index("-pl") + 1]); json.dump(st, open(p, "w")); print("ok"); sys.exit(0)
50 42 f = [x for x in a if x.startswith("--query-gpu=")][0].split("=")[1].split(",")
51 43 if any("reasons" in x for x in f): sys.exit(2)
52 44 v = {"index": 0, "power.draw": st["draw"], "temperature.gpu": st["temp"], "utilization.gpu": round(100 * st[
53 "util"]),
54 45 "power.limit": st["limit"], "power.min_limit": %s, "clocks.sm": round(1800 * st["speed"]),
55 46 "enforced.power.limit": st["limit"], "power.management": "Enabled", "persistence_mode": "Enabled",
56 47 "power.default_limit": 300, "power.max_limit": 300}
57 48 print(", ".join(str(v[x]) for x in f))
58 49 ''' % MIN
59
60 52 def speed_at(limit, gamma):
61 53     return 1.0 if limit >= P_DEMAND else max(0.05, (limit - P_STATIC) / (P_DEMAND - P_STATIC)) ** (1.0 / gamma)
62
63 56 def busy_power(s, gamma):
64 57     return P_STATIC + (P_DEMAND - P_STATIC) * s ** gamma
65
66 60 def arrivals(seed):
67 61     rng = random.Random(seed); out = []; per = DURATION / len(PHASES)
68 62     for i, share in enumerate(PHASES):
69 63         rate = share / (SERVICE_MS / 1000.0); t = i * per
70 64         while True:
71 65             t += rng.expovariate(rate)
72 66             if t >= (i + 1) * per:
73 67                 break
74 68             out.append(t)
75 69     return out
76
77 72 def run(arm, seed, gamma, work, gov_argv):
78 73     """One arm: native (the start limit throughout) or omni (the governor in cap mode). Returns its gauges."""
79 74     tag = f"{arm}-{seed}-{gamma}"
```

```

75 state, lat = work / f"state-{tag}.json", work / f"latency-{tag}.csv"
76 state.write_text(json.dumps({"limit": START, "draw": P_IDLE, "temp": 35.0, "util": 0.0, "speed": 1.0}))
77 lat.write_text("elapsed_seconds,latency_ms,ok\n")
78 os.environ["PHYS_STATE"] = str(state)
79 gov = None
80 if arm == "omni":
81     a = gov_parser().parse_args(["--mode", "cap", "--smi", str(work / "smi.py"), "--audit", str(work /
f"audit-{tag}.jsonl"), "--kill-file", str(work / "no-kill"),
82                                     "--latency-file", str(lat), "--slo-ms", str(SLO_MS)] + gov_argv)
83     gov = GpuGovernor(a); every = max(1, int(round(a.interval)))
84     arr = arrivals(seed); ai = 0; queue = []; energy = 0.0; temp = 35.0; lats = []; limits = []; peak = 0.0
85     for sec in range(DURATION + DRAIN):
86         st = json.loads(state.read_text()); limit = st["limit"]; limits.append(limit)
87         s = speed_at(limit, gamma); pb = busy_power(s, gamma)
88         t, end, busy, rows = float(sec), sec + 1.0, 0.0, []
89         while t < end:
90             while ai < len(arr) and arr[ai] <= t:
91                 queue.append([arr[ai], SERVICE_MS / 1000.0]); ai += 1
92             nxt = arr[ai] if ai < len(arr) else math.inf
93             if not queue:
94                 t = min(end, nxt); continue
95             job = queue[0]
96             step = min(job[1] / s, end - t, nxt - t)
97             job[1] -= step * s; busy += step; t += step
98             if job[1] <= 1e-12:
99                 queue.pop(0); ms = (t - job[0]) * 1000.0; lats.append(ms); rows.append(f"{t:.3f},{ms:.2f},1\n")
100             draw = busy * pb + (1 - busy) * P_IDLE
101             energy += draw; temp += (30 + 0.12 * draw - temp) / 20.0; peak = max(peak, temp)
102             with open(lat, "a") as f:
103                 f.writelines(rows)
104             st.update(draw=round(draw, 2), temp=round(temp, 1), util=busy, speed=s); state.write_text(json.dumps(st))
105             if gov and sec < DURATION and sec % every == every - 1:
106                 gov.step()
107         if gov:
108             gov.restore()
109         lats.sort()
110         return {"requests served per kJ": len(lats) / (energy / 1000.0), "GPU energy (J)": energy,
111                 "GPU mean power (W)": energy / (DURATION + DRAIN), "requests served": len(lats),
112                 "requests not served": len(arr) - len(lats), "response time, mean (ms)": statistics.mean(lats),
113                 "response time, 95th percentile (ms)": lats[int(0.95 * len(lats))],
114                 "response time, 99th percentile (ms)": lats[int(0.99 * len(lats))], "peak temperature (C)": peak,
115                 "power limit, mean (W)": statistics.mean(limits[:DURATION]), "power-limit writes": gov.writes if gov
else 0}
116
117
118 def paired(nat, om, k):
119     d = [b[k] - a[k] for a, b in zip(nat, om)]; n = len(d); m = statistics.mean(d)
120     h = T95.get(n - 1, 2.0) * statistics.stdev(d) / math.sqrt(n) if n > 1 else float("nan")
121     return m, m - h, m + h
122
123
124 def table(gamma, nat, om):
125     out = [f"### Card model gamma {gamma:g} ({'voltage falls with clock' if gamma >= 2.5 else 'near its voltage
floor'})",
126           "", f"{len(nat)} paired repetitions. A change is proven when its 95% interval excludes zero.", "",
127           "| Gauge | Native | Omni | Change | 95% interval of the difference |", "|---|---|---|---|---|"]
128     for k in nat[0]:
129         n = statistics.mean(x[k] for x in nat); o = statistics.mean(x[k] for x in om); m, lo, hi = paired(nat,
om, k)
130         pct = f"{100 * (o - n) / n:+.1f}%" if n else "n/a"
131         out.append(f"| {k} | {n:.4g} | {o:.4g} | {pct} | {lo:+.4g} to {hi:+.4g} |")
132         kj = "requests served per kJ"; p95 = "response time, 95th percentile (ms)"
133         m, lo, hi = paired(nat, om, kj)
134         base95 = statistics.mean(x[p95] for x in nat); _, _, hi95 = paired(nat, om, p95)
135         guard = hi95 <= 0.10 * base95
136         word = "better, proven" if lo > 0 else "worse, proven" if hi < 0 else "not proven"
137         if lo > 0 and not guard:
138             word = "better on energy, fails the service guardrail"
139         out += [ "", f"Guardrail (preregistered): 95th-percentile response time not above +10%: "
140                 f"{'held' if guard else 'FAILED'} (upper bound {100 * hi95 / base95:+.1f}%).",
141                 f"***Model verdict: {word}***", "" ]
142     return "\n".join(out)
143
144
145 def main(argv=None):
146     ap = argparse.ArgumentParser(description=__doc__, formatter_class=argparse.RawDescriptionHelpFormatter)
147     ap.add_argument("--reps", type=int, default=10)
148     ap.add_argument("--gammas", default="3,1.5")
149     ap.add_argument("--gov", default="", help="extra governor arguments, e.g. \"--min-share 0 --util-gate 0\"")
150     ap.add_argument("--out", default="", help="folder for GPU_PHYSICS_SIM.md and .json")
151     a = ap.parse_args(argv)
152     work = Path(tempfile.mkdtemp(prefix="gpu-physics-sim-"))
153     (work / "smi.py").write_text(SMI); (work / "smi.py").chmod(0o755)
154     gov_argv = shlex.split(a.gov)
155     settings = vars(gov_parser().parse_args(["--smi", "x"] + gov_argv))
156     report = [f"# GPU physics simulation: native against Omni-Compass, modelled card", "",

```

```

157         "Every number below comes from the card model in tools/gpu_physics_sim.py, not from a meter. It
shows what",
158         "the governor does to a card that behaves as modelled; the hardware answer is scripts/gpu_paired.
sh.", "",
159         "Governor settings: " + ", ".join(f"{k} {settings[k]}" for k in ("headroom", "min_share",
"util_gate", "util_band", "interval")),
160         "" ]
161     raw = {}
162     for gamma in (float(g) for g in a.gammas.split(", ")):
163         nat = [run("native", 20260928 + r, gamma, work, gov_argv) for r in range(a.reps)]
164         om = [run("omni", 20260928 + r, gamma, work, gov_argv) for r in range(a.reps)]
165         report.append(table(gamma, nat, om)); raw[str(gamma)] = {"native": nat, "omni": om}
166     text = "\n".join(report)
167     print(text)
168     if a.out:
169         Path(a.out).mkdir(parents=True, exist_ok=True)
170         (Path(a.out) / "GPU_PHYSICS_SIM.md").write_text(text + "\n")
171         (Path(a.out) / "GPU_PHYSICS_SIM.json").write_text(json.dumps({"settings": settings, "reps": a.reps,
"raw": raw}, indent=1, default=str))
172     return 0
173
174
175 if __name__ == "__main__":
176     sys.exit(main())
177

```

tools/gpu_pipeline_sim.py

156 lines SHA-256 67dc5a6bd17d83e24702dab72ae840d12bca7dbcf58127103b6d55f11d7370b8

```
1  """Speed won on the CPU, spent on the GPU: a request served by a CPU stage and then a GPU stage, with Omni-
2  Compass's
3  speed lock (omni_controller/gpu_governor.py --baseline-file) on the GPU. A preview on a model, not evidence.
4
5  python3 tools/gpu_pipeline_sim.py                                # 10 paired repetitions, both card models
6  python3 tools/gpu_pipeline_sim.py --cpu-ms 10 --out results/gpu/sim/pipeline
7
8  Each request needs --cpu-ms of CPU work on one core (preparing the input) and then the GPU (tools/
9  gpu_physics_sim.py's
10 card and workload: 50 ms at full clock, Poisson arrivals at 30/60/80/30/60/30% of GPU capacity). The serving
11 pod's
12 CPU limit is enforced as Linux does it (CFS bandwidth): a quota of limit x 100 ms of CPU time per 100 ms period;
13 the
14 pod's work runs at a full core until the period's quota is spent, then waits for the next period. The operator's
15 --native-limit (0.5 core: 50 ms per 100 ms) throttles bursts; Omni's conveyance (omni_controller/muscles.py
16 convey:
17 the machine's idle CPU to the serving pods) raises the limit to --conveyed-limit (1 core: never throttled for one
18 thread). Both stages are first come first served queues.
19
20 Arms, paired on the same arrivals:
21 native operator's CPU limit, GPU at its start limit (no Omni)
22 speed CPU conveyed, GPU at its start limit: the speed the CPU muscle wins, handed to the customer
23 lock CPU conveyed, GPU under the speed lock: the speed won is spent on GPU watts, keeping every gauge (mean,
24 p95, p99) at least --speed-gain faster than native
25 The lock's baseline comes from native runs on other seeds (tools/gpu_baseline.py), never from the arms judged.
26 CPU work is the same core-seconds in every arm (conveyance moves no work, it removes waiting); only GPU energy
27 differs.
28 """
29 from __future__ import annotations
30
31 import argparse, json, math, os, shlex, statistics, sys, tempfile
32 from collections import deque
33 from pathlib import Path
34
35 ROOT = Path(__file__).resolve().parents[1]; sys.path.insert(0, str(ROOT))
36 from omni_controller.gpu_governor import GpuGovernor, parser as gov_parser
37 from tools.gpu_physics_sim import SMI, START, P_IDLE, SERVICE_MS, SLO_MS, DURATION, DRAIN, speed_at, busy_power,
38 arrivals, paired
39 from tools.gpu_baseline import build as build_baseline
40
41 DT = 0.001
42
43 def run(seed, gamma, work, cpu_ms, cpu_limit, lock_argv=None, tag=""):
44     """One arm. lock_argv None: GPU held at its start limit. Returns gauges and the latency rows."""
45     state, lat = work / f"state-{tag}.json", work / f"latency-{tag}.csv"
46     state.write_text(json.dumps({"limit": START, "draw": P_IDLE, "temp": 35.0, "util": 0.0, "speed": 1.0}))
47     lat.write_text("elapsed_seconds,latency_ms,ok\n")
48     os.environ["PHYS_STATE"] = str(state)
49     gov, every = None, 1
50     if lock_argv is not None:
51         a = gov_parser().parse_args(["--mode", "cap", "--smi", str(work / "smi.py"), "--audit", str(work /
52 f"audit-{tag}.jsonl"),
53                                     "--kill-file", str(work / "no-kill"), "--latency-file", str(lat),
54                                     "--slo-ms", str(SLO_MS)] + lock_argv)
55     gov = GpuGovernor(a); every = max(1, int(round(a.interval))); vt = {"t": 0.0}; gov.clock = lambda: vt[
56 "t"]
57     arr = arrivals(seed); ai = 0
58     cpu_q, gpu_q = deque(), deque()
59     quota, used = cpu_limit * 0.1, 0.0 # CFS: seconds of CPU per 100 ms period
60     per = int(round(0.1 / DT))
61     energy = 0.0; temp = 35.0; lats = []; rows_all = []; limits = []
62     steps = int(round(1.0 / DT))
63     for sec in range(DURATION + DRAIN):
64         st = json.loads(state.read_text()); limit = st["limit"]; limits.append(limit)
65         s = speed_at(limit, gamma); pb = busy_power(s, gamma)
66         busy = 0.0; rows = []
67         for k in range(steps):
68             t = sec + k * DT
69             while ai < len(arr) and arr[ai] <= t:
70                 cpu_q.append([arr[ai], cpu_ms / 1000.0]); ai += 1
71             if k % per == 0:
72                 used = 0.0 # a new CFS period
73             if cpu_q and used < quota - 1e-12:
74                 j = cpu_q[0]; j[1] -= DT; used += DT
75                 if j[1] <= 1e-12:
76                     cpu_q.popleft(); gpu_q.append([j[0], SERVICE_MS / 1000.0])
77             if gpu_q:
78                 j = gpu_q[0]; j[1] -= DT * s; busy += DT
79                 if j[1] <= 1e-12:
80                     gpu_q.popleft(); ms = (t + DT - j[0]) * 1000.0; lats.append(ms)
81                     rows.append(f"{t + DT:.3f},{ms:.2f},1\n"); rows_all.append((t + DT, ms, True))
82         draw = busy * pb + (1 - busy) * P_IDLE
```

```

75         energy += draw; temp += (30 + 0.12 * draw - temp) / 20.0
76         with open(lat, "a") as f:
77             f.writelines(rows)
78         st.update(draw=round(draw, 2), temp=round(temp, 1), util=busy, speed=s); state.write_text(json.dumps(st))
79         if gov and sec < DURATION and sec % every == every - 1:
80             vt["t"] = sec + 1.0; gov.step()
81         if gov:
82             gov.restore()
83             lats.sort()
84             g = {"requests served per kJ (GPU)": len(lats) / (energy / 1000.0), "GPU energy (J)": energy,
85                 "requests served": len(lats), "requests not served": len(arr) - len(lats),
86                 "response time, mean (ms)": statistics.mean(lats), "response time, 95th percentile (ms)": lats[int(0.95
* len(lats))],
87                 "response time, 99th percentile (ms)": lats[int(0.99 * len(lats))],
88                 "GPU power limit, mean (W)": statistics.mean(limits[:DURATION]), "power-limit writes": gov.writes if
gov else 0}
89             return g, rows_all
90
91
92 def compare(title, nat, arm, speed_gain, lock=False):
93     out = [f"#### {title}", "", "| Gauge | Native | This arm | Change | 95% interval of the difference |", "|---
|---:|---:|---:|---:|"]
94     for k in nat[0]:
95         n = statistics.mean(x[k] for x in nat); o = statistics.mean(x[k] for x in arm); m, lo, hi = paired(nat,
arm, k)
96         pct = f"{100 * (o - n) / n:+.1f}%" if n else "n/a"
97         out.append(f"| {k} | {n:.4g} | {o:.4g} | {pct} | {lo:+.4g} to {hi:+.4g} |")
98         if lock:
99             held = []
100             for k in ("response time, mean (ms)", "response time, 95th percentile (ms)", "response time, 99th
percentile (ms)":):
101                 n = statistics.mean(x[k] for x in nat); _, _, hi = paired(nat, arm, k)
102                 held.append((k.split(", ")[1].split(" ")[0], 100 * hi / n))
103                 ok = all(h <= -100 * speed_gain for _, h in held)
104                 out += [f"Speed lock (every gauge at least {100 * speed_gain:g}% faster than native, upper end of
its 95% interval): ",
105                         + ", ".join(f"{name} {h:+.1f}%" for name, h in held) + f" -> **{'HELD' if ok else 'NOT
HELD'}**."]
106             return "\n".join(out + [""])
107
108
109 def main(argv=None):
110     ap = argparse.ArgumentParser(description=__doc__, formatter_class=argparse.RawDescriptionHelpFormatter)
111     ap.add_argument("--reps", type=int, default=10)
112     ap.add_argument("--gammas", default="3,1.5")
113     ap.add_argument("--cpu-ms", type=float, default=20.0, help="CPU work per request on one core (ms)")
114     ap.add_argument("--native-limit", type=float, default=0.5, help="operator's CPU limit of the serving pod (
cores)")
115     ap.add_argument("--conveyed-limit", type=float, default=1.0, help="the pod's CPU limit under conveyance (
cores)")
116     ap.add_argument("--calib-runs", type=int, default=3)
117     ap.add_argument("--gov", default="", help="extra governor arguments for the lock arm")
118     ap.add_argument("--out", default="")
119     a = ap.parse_args(argv)
120     work = Path(tempfile.mkdtemp(prefix="gpu-pipeline-sim-"))
121     (work / "smi.py").write_text(SMI); (work / "smi.py").chmod(0o755)
122     # the baseline: native runs on seeds the comparison never uses (latency at the start limit does not depend
on gamma)
123     calib = [run(900001 + i, 3.0, work, a.cpu_ms, a.native_limit, tag=f"calib-{i}") for i in range(a.
calib_runs)]
124     base = build_baseline(calib); bf = work / "baseline.json"; bf.write_text(json.dumps(base))
125     lock_argv = ["--baseline-file", str(bf)] + shlex.split(a.gov)
126     settings = vars(gov_parser().parse_args(["--smi", "x"] + lock_argv))
127     report = [f"# Speed won on the CPU, spent on the GPU (model)", "",
128              f"CPU work {a.cpu_ms:g} ms per request; pod CPU limit {a.native_limit:g} core native, {a.
conveyed_limit:g} "
129              f"conveyed. GPU as tools/gpu_physics_sim.py. Baseline from {a.calib_runs} native runs on other
seeds.",
130              "Lock settings: " + ", ".join(f"{k} {settings[k]}" for k in ("speed_gain", "lock_margin",
"lock_step",
131                                     "lock_window_s", "lock_floor",
132                                     "interval")),
133              "Every number comes from the model in this file and tools/gpu_physics_sim.py, not from a meter.",
134              ""]
135     raw = {}
136     for gamma in (float(x) for x in a.gammas.split(", ")):
137         nat, spd, lck = [], [], []
138         for r in range(a.reps):
139             seed = 20260928 + r
140             nat.append(run(seed, gamma, work, a.cpu_ms, a.native_limit, tag=f"n-{gamma}-{r}")[0])
141             spd.append(run(seed, gamma, work, a.cpu_ms, a.conveyed_limit, tag=f"s-{gamma}-{r}")[0])
142             lck.append(run(seed, gamma, work, a.cpu_ms, a.conveyed_limit, lock_argv, tag=f"l-{gamma}-{r}")[0])
143             report.append(f"### Card model gamma {gamma:g}, {a.reps} paired repetitions\n")
144             report.append(compare("CPU muscle only (speed handed to the customer)", nat, spd, settings[
"speed_gain"])))
145             report.append(compare("CPU muscle + GPU speed lock (speed spent on watts)", nat, lck, settings[

```

```

    "speed_gain"], lock=True))
144     raw[str(gamma)] = {"native": nat, "speed": spd, "lock": lck}
145     text = "\n".join(report); print(text)
146     if a.out:
147         Path(a.out).mkdir(parents=True, exist_ok=True)
148         (Path(a.out) / "GPU_PIPELINE_SIM.md").write_text(text + "\n")
149         (Path(a.out) / "GPU_PIPELINE_SIM.json").write_text(json.dumps({"settings": settings, "args": vars(a),
"baseline": base,
150                                     "raw": raw}, indent=1, default=str))
151     return 0
152
153
154 if __name__ == "__main__":
155     sys.exit(main())
156

```

tools/gpu_reps.py

541 lines SHA-256 d00b185618919164b72147155ecf9d8d9dd37d8101a69c32cfc226176b56340b

```
1  """The GPU bench table (scripts/gpu_paired.sh): native / watch (Omni runs, writes nothing) / omni (Omni writes
   power limits).
2
3  Every energy number comes from the device: joules = the time integral of nvidia-smi power.draw over the measured
   window
4  (trapezoid rule on its own sample timestamps), summed over the measured GPUs; the CPU package joules come from
   RAPL
5  energy counters when the machine has them. Nothing here is a model.
6
7  Per repetition, per arm: energy, joules per served request, mean power, peak and mean temperature, served and not
8  served requests, response time (mean, 95th, 99th percentile), power-limit writes. Across repetitions: for watch
   and omni
9  against native, the paired mean difference with a t-based 95% interval. A difference is proven when that interval
10 excludes zero; otherwise the table says not proven.
11
12 Three contrasts, each paired by repetition: observation = watch - native (does Omni's presence alone move the
   plant),
13 authority = omni - watch (what granting Omni the write adds), total = omni - native (the preregistered
   comparison).
14 Three receipts, kept apart: A the governor's decisions (audit.jsonl), B its actuator records (requested, return
   code,
15 read-back, enforced limit, delay), C the bench's own nvidia-smi samples and the workload's requests. From A and
   B the
16 table also reports actuator fidelity, control effort, and how well the engine's six-state projection matches
   what the
17 telemetry shows next (the intertwining residual and directional accuracy), defined in docs/GPU_PREREGISTRATION.
   md.
18 The result label is chosen by rule (label()), never by hand. A meter that was not fitted prints UNAVAILABLE.
19
20 The run is INVALID (exit 2) when: a native or watch arm saw a power limit (or enforced limit) other than the one
   read at
21 start, the watch arm executed a write, an arm ended with a limit that differs from the start (the kill switch
   did not
22 restore), a governor exited nonzero (a refused write, a refused start), a write was refused, Omni's frozen files
23 changed, a confirmation ran on uncommitted code, or smoke and confirmation repetitions were mixed.
24 Usage: python tools/gpu_reps.py RUN_DIR (RUN_DIR holds rep-<n>/<arm>/) -> RUN_DIR/GPU_REPS.md, RUN_DIR/
   GPU_REPS.json"""
25 from __future__ import annotations
26
27 import csv, datetime as dt, json, math, statistics, sys
28 from pathlib import Path
29
30 U_AUTHORITY = 25.0          # omniconpass/core.py: the U-channel command's bound (saturation)
31 DEADBAND = 0.01            # directional accuracy: a movement under this (state units) is not a direction
32 W_INT = {"E": 1.0, "U": 1.0, "I_U": 1.0, "S": 1.0, "B": 1.0} # the residual's weights, declared before any
   trial
33 MARGIN_SERVED, MARGIN_P95, MARGIN_LOST = 0.01, 0.10, 0.01      # guardrails (docs/GPU_PREREGISTRATION.md)
34 SMI_DEFAULT = "timestamp,index,power.draw,temperature.gpu,utilization.gpu,power.limit,clocks.sm"
35
36 T95 = {1: 12.706, 2: 4.303, 3: 3.182, 4: 2.776, 5: 2.571, 6: 2.447, 7: 2.365, 8: 2.306, 9: 2.262, 10: 2.228, 11:
   2.201,
37        12: 2.179, 13: 2.160, 14: 2.145, 15: 2.131, 16: 2.120, 17: 2.110, 18: 2.101, 19: 2.093, 20: 2.086}
38 LOWER = {"energy, whole machine at the wall (J)", "energy, GPU (J)", "energy per served request (J)", "power,
   GPU mean (W)", "temperature, peak (C)",
39          "temperature, mean (C)", "requests not served", "response time, mean (ms)", "response time, 95th
   percentile (ms)",
40          "response time, 99th percentile (ms)", "energy, CPU package (J)", "energy, rest of the machine (J)",
41          "energy, DRAM (J)", "energy, platform psys (J)", "energy, GPU device counter (J)"}
42 PRIMARY = "work per energy (served requests per kJ)"
43 WALL = "work per wall energy (served requests per kJ, whole machine)"
44 KEYS = [PRIMARY, WALL, "energy, whole machine at the wall (J)", "energy, GPU (J)", "energy per served request (
   J)", "power, GPU mean (W)", "requests served", "requests not served",
45          "response time, mean (ms)", "response time, 95th percentile (ms)", "response time, 99th percentile (ms)",
46          "temperature, peak (C)", "temperature, mean (C)", "energy, CPU package (J)", "energy, DRAM (J)", "energy,
   platform psys (J)",
47          "energy, rest of the machine (J)", "energy, GPU device counter (J)", "power-capped share of samples"]
48 PHYSICAL = {"energy, whole machine at the wall (J)", "energy, CPU package (J)", "energy, rest of the machine (
   J)", WALL,
49             "energy, DRAM (J)", "energy, platform psys (J)", "energy, GPU device counter (J)"}
50
51
52 def smi_time(s):
53     s = s.strip()
54     for f in ("%Y/%m/%d %H:%M:%S.%f", "%Y/%m/%d %H:%M:%S"):
55         try:
56             return dt.datetime.strptime(s, f).replace(tzinfo=dt.timezone.utc).timestamp()
57         except ValueError:
58             pass
59     raise ValueError(s)
60
61
62 def pct(xs, q):
63     xs = sorted(xs)
```

```

64     return xs[min(len(xs) - 1, int(q * len(xs)))] if xs else float("nan")
65
66
67 def arm(d, gpus):
68     t0 = float((d / "window_start.txt").read_text().split()[0]); t1 = float((d / "window_end.txt").read_text().
69     split()[0])
70     fields = ((d / "smi_fields.txt").read_text().strip() if (d / "smi_fields.txt").exists() else SMI_DEFAULT).
71     split(",")
72     ix = {f: i for i, f in enumerate(fields)}
73     ie = ix.get("enforced.power.limit"); ir = ix.get("clocks_event_reasons.active", ix.get(
74     "clocks_throttle_reasons.active"))
75     series, enforced, capped, nre = {}, set(), 0, 0
76     for r in csv.reader(open(d / "smi.csv")):
77         if len(r) < len(fields) or not r[1].strip().isdigit():
78             continue
79         g = int(r[1])
80         if g not in gpus:
81             continue
82         t = smi_time(r[0])
83         if t0 <= t <= t1:
84             series.setdefault(g, []).append((t, float(r[2]), float(r[3]), float(r[5])))
85             if ie is not None:
86                 enforced.add(float(r[ie]))
87             if ir is not None:
88                 try:
89                     capped += bool(int(r[ir].strip(), 16) & 0x4); nre += 1 # bit 2: software power cap
90                 except ValueError:
91                     pass
92     joules, temps, limits, peak = 0.0, [], set(), float("nan")
93     for g, s in series.items():
94         s.sort()
95         joules += sum((b[0] - a[0]) * (a[1] + b[1]) / 2.0 for a, b in zip(s, s[1:]))
96         temps += [x[2] for x in s]; limits |= {x[3] for x in s}
97         peak = max(temps) if temps else float("nan")
98         req = list(csv.DictReader(open(d / "requests.csv")))
99         ms = [float(r["latency_ms"]) for r in req if r["ok"] == "1"]
100         served, lost = len(ms), sum(1 for r in req if r["ok"] != "1")
101         g = {PRIMARY: served / (joules / 1000.0) if joules > 0 else float("nan"), "energy, GPU (J)": joules, "energy
102         per served request (J)": joules / served if served else float("nan"),
103         "power, GPU mean (W)": joules / max(1e-9, t1 - t0), "requests served": float(served), "requests not
104         served": float(lost),
105         "response time, mean (ms)": sum(ms) / served if served else float("nan"),
106         "response time, 95th percentile (ms)": pct(ms, 0.95), "response time, 99th percentile (ms)": pct(ms, 0.
107         99),
108         "temperature, peak (C)": peak, "temperature, mean (C)": sum(temps) / len(temps) if temps else float(
109         "nan"),
110         "energy, CPU package (J)": float("nan"), "energy, whole machine at the wall (J)": float("nan"), WALL:
111         float("nan"),
112         "energy, rest of the machine (J)": float("nan"), "energy, DRAM (J)": float("nan"), "energy, platform
113         psys (J)": float("nan"),
114         "power-capped share of samples": capped / nre if nre else float("nan")}
115     if (d / "wall.csv").exists():
116         w = [(float(r["epoch_s"]), float(r["watts"])) for r in csv.DictReader(open(d / "wall.csv")) if r[
117         "watts"]]
118         w = [x for x in w if t0 <= x[0] <= t1]
119         # a gap in the plug's readings longer than 5 s leaves the arm without a wall number, never an
120         interpolated one
121         if len(w) > 1 and w[0][0] - t0 < 5 and t1 - w[-1][0] < 5 and all(b[0] - a[0] < 5 for a, b in zip(w, w[1:
122         ])):
123             wj = sum((b[0] - a[0]) * (a[1] + b[1]) / 2.0 for a, b in zip(w, w[1:]))
124             g["energy, whole machine at the wall (J)"] = wj
125             g[WALL] = served / (wj / 1000.0) if wj > 0 else float("nan")
126             g.update(rapl_energy(d))
127             g["energy, GPU device counter (J)"] = float("nan")
128             dv0, dv1 = _kv(d / "device_start.txt"), _kv(d / "device_end.txt")
129             try:
130                 g["energy, GPU device counter (J)"] = (float(dv1["energy_mj"]) - float(dv0["energy_mj"])) / 1000.0
131             except (KeyError, ValueError):
132                 pass
133             # the rest of the machine only when the wall, the GPU and the CPU package measured this arm: wall - GPU -
134             package
135             # (- DRAM where measured), never modelled
136             if not math.isnan(g["energy, whole machine at the wall (J)"]) and not math.isnan(g["energy, CPU package (
137             J)"]):
138                 dram = g["energy, DRAM (J)"]
139                 g["energy, rest of the machine (J)"] = (g["energy, whole machine at the wall (J)"] - joules - g["energy,
140                 CPU package (J)"]
141                 - (0.0 if math.isnan(dram) else dram))
142     writes = would = 0
143     if (d / "audit.jsonl").exists():
144         for line in open(d / "audit.jsonl"):
145             rec = json.loads(line)
146             writes += "write" in rec; would += "would_write" in rec
147     restored = (d / "limit_end.txt").read_text().split() == (d / "limit_start.txt").read_text().split() \
148     if (d / "limit_end.txt").exists() else False
149     gx = (d / "governor_exit.txt").read_text().strip() if (d / "governor_exit.txt").exists() else None

```

```

135     return g, {"limits_seen": sorted(limits), "enforced_seen": sorted(enforced), "writes": writes, "would_write":
would,
136         "restored": restored, "governor_exit": gx, "samples": sum(map(len, series.values())), "health":
health(d),
137         **receipts(d)}
138
139
140 def _kv(f):
141     return dict(x.split(" ", 1) for x in f.read_text().splitlines() if " " in x) if f.exists() else {}
142
143
144 def health(d):
145     """Card health over the arm (bench only): change in uncorrected and corrected ECC errors, pages pending
retirement."""
146     a, b = _kv(d / "device_start.txt"), _kv(d / "device_end.txt"); out = {}
147     for k in ("ecc.errors.uncorrected.volatile.total", "ecc.errors.corrected.volatile.total", "retired_pages.
pending"):
148         try:
149             out[k] = float(b[k]) - float(a[k]) if k != "retired_pages.pending" else b[k]
150         except (KeyError, ValueError):
151             out[k] = None
152     return out
153
154
155 def rapl_delta(e0, e1, rng):
156     """A wrapping microjoule counter's increase: modular when the range is known; unknown (None) if it went
backwards
157     without one."""
158     if e1 >= e0:
159         return e1 - e0
160     return e1 - e0 + rng if rng > 0 else None
161
162
163 def rapl_energy(d):
164     """CPU package and DRAM joules (summed over sockets, separately) and platform (psys) joules from rapl_start.
tsv /
165     rapl_end.tsv; NaN (printed UNAVAILABLE) where the counters were missing or unreadable. package + dram is the
CPU
166     side; psys is a wider, vendor-defined scope and is never added to them."""
167     out = {"energy, CPU package (J)": float("nan"), "energy, DRAM (J)": float("nan"), "energy, platform psys (
J)": float("nan")}
168     f0, f1 = d / "rapl_start.tsv", d / "rapl_end.tsv"
169     if f0.exists() and f1.exists():
170         rd = lambda f: {x.split()[0]: x.split() for x in f.read_text().splitlines() if len(x.split()) == 4}
171         a, b = rd(f0), rd(f1)
172         sums = {}
173         for dom, (_, name, e0, rng) in a.items():
174             if dom not in b:
175                 continue
176             kind = "package" if name.startswith("package") else name
177             dl = rapl_delta(int(e0), int(b[dom][2]), int(rng))
178             sums.setdefault(kind, []).append(dl)
179             for kind, key in (("package", "energy, CPU package (J)"), ("dram", "energy, DRAM (J)"), ("psys", "energy,
platform psys (J)")):
180                 if sums.get(kind) and None not in sums[kind]:
181                     out[key] = sum(sums[kind]) / 1e6
182             elif (d / "rapl_start.txt").exists() and (d / "rapl_end.txt").exists(): # runs recorded before the
named format
183                 a = [int(x) for x in (d / "rapl_start.txt").read_text().split()]; b = [int(x) for x in (d / "rapl_end.
txt").read_text().split()]
184                 rng = [int(x) for x in (d / "rapl_range.txt").read_text().split()] if (d / "rapl_range.txt").exists()
else [0] * len(a)
185                 dl = [rapl_delta(x, y, r) for x, y, r in zip(a, b, rng)]
186                 if a and len(a) == len(b) and None not in dl:
187                     out["energy, CPU package (J)"] = sum(dl) / 1e6
188     return out
189
190
191 def credit(root, runs_dirs):
192     """Per write, credited against the same moments of the same seeded request stream in the native arm (and the
watch
193     arm): GPU joules and requests finished in the interval from this write to the next, omni minus native. The
194     intervals tile the arm, so the credits add up to the whole difference they cover. Also who decided each
write."""
195     rows = []
196     for rep, arms in runs_dirs.items():
197         if "omni" not in arms or "native" not in arms:
198             continue
199         do = arms["omni"]
200         if not (do / "audit.jsonl").exists():
201             continue
202         recs = [json.loads(x) for x in open(do / "audit.jsonl") if x.strip()]
203         t_start = float((do / "window_start.txt").read_text().split()[0]); t_end = float((do / "window_end.txt").
read_text().split()[0])
204         last_dec, writes = {}, []
205         for r in recs:

```

```

206         if "decision" in r:
207             for gk, v in r["decision"].items():
208                 if "decided_by" in v:
209                     last_dec[gk] = v["decided_by"]
210             if "actuator" in r and r["actuator"].get("rc") == 0:
211                 a = r["actuator"]
212                 writes.append((a["t_requested"] - t_start, a["requested_w"], last_dec.get(str(a["gpu"]), "?")))
213             if not writes:
214                 continue
215             L = t_end - t_start
216             edges = [w[0] for w in writes] + [L]
217             series = {a: _series(arms[a], t0=float((arms[a] / "window_start.txt").read_text().split()[0])) for a in (
"omni", "native", "watch")} if a in arms}
218             done = {a: _done(arms[a]) for a in series}
219             for (t, w, who), t_next in zip(writes, edges[1:]):
220                 if t_next <= t:
221                     continue
222                 row = {"rep": rep, "t_s": round(t, 1), "limit_w": w, "decided_by": who, "interval_s": round(t_next -
t, 1)}
223                 for a in ("native", "watch"):
224                     if a in series:
225                         row[f"dJ_vs_{a}"] = _energy(series["omni"], t, t_next) - _energy(series[a], t, t_next)
226                         row[f"dServed_vs_{a}"] = _count(done["omni"], t, t_next) - _count(done[a], t, t_next)
227                 rows.append(row)
228             by = {}
229             for r in rows:
230                 b = by.setdefault(r["decided_by"], {"writes": 0, "dJ_vs_native": 0.0, "dServed_vs_native": 0, "seconds":
0.0})
231                 b["writes"] += 1; b["dJ_vs_native"] += r.get("dJ_vs_native", 0.0); b["dServed_vs_native"] += r.get(
"dServed_vs_native", 0)
232                 b["seconds"] += r["interval_s"]
233             return rows, by
234
235
236 def _series(d, t0):
237     fields = ((d / "smi_fields.txt").read_text().strip() if (d / "smi_fields.txt").exists() else SMI_DEFAULT).
split(",")
238     s = []
239     for r in csv.reader(open(d / "smi.csv")):
240         if len(r) >= len(fields) and r[1].strip().isdigit():
241             try:
242                 s.append((smi_time(r[0]) - t0, float(r[2])))
243             except ValueError:
244                 pass
245     return sorted(s)
246
247
248 def _energy(s, a, b):
249     pts = [x for x in s if a <= x[0] <= b]
250     return sum((q[0] - p[0]) * (p[1] + q[1]) / 2.0 for p, q in zip(pts, pts[1:]))
251
252
253 def _done(d):
254     """Finish times of served requests on the arm's own window clock (the workload's t0_epoch against
window_start)."""
255     off = 0.0
256     if (d / "summary.json").exists():
257         t0e = json.loads((d / "summary.json").read_text()).get("t0_epoch")
258         if t0e:
259             off = t0e - float((d / "window_start.txt").read_text().split()[0])
260     return [float(r["done_s"]) + off for r in csv.DictReader(open(d / "requests.csv")) if r["ok"] == "1" and r[
"done_s"]]
261
262
263 def _count(ts, a, b):
264     return sum(1 for t in ts if a <= t < b)
265
266
267 def receipts(d):
268     """Receipts A and B from the governor's audit: actuator fidelity, control effort, representation fidelity."""
269     if not (d / "audit.jsonl").exists():
270         return {}
271     recs = [json.loads(x) for x in open(d / "audit.jsonl") if x.strip()]
272     snap = next((r["snapshot"] for r in recs if "snapshot" in r), {})
273     start = {g: v.get("limit_w") for g, v in snap.items()}
274     # B: actuator
275     act = [r["actuator"] for r in recs if "actuator" in r]
276     late = {(x["gpu"], x["requested_w"]): x for x in (r["actuator_realized"] for r in recs if
"actuator_realized" in r)}
277     r_act, delay, realized, over = [], [], {}, 0
278     for a in act:
279         if a.get("rc") != 0:
280             continue
281         rb = a.get("readback_w"); lt = late.get((a["gpu"], a["requested_w"]))
282         if not a.get("realized") and lt:
283             rb = lt["readback_w"]; delay.append(lt["delay_s"])

```

```

284         elif a.get("realized"):
285             delay.append(a["delay_s"])
286         if rb is not None:
287             r_act.append(rb - a["requested_w"]); realized.setdefault(str(a["gpu"]), []).append(rb)
288         over += bool(a.get("override"))
289         tv, rev = 0.0, 0
290         for g, seq in realized.items():
291             seq = ([start[g]] if start.get(g) is not None else []) + seq
292             steps = [b - a for a, b in zip(seq, seq[1:])] if abs(b - a) >= 1.0]
293             tv += sum(abs(x) for x in steps); rev += sum(1 for a, b in zip(steps, steps[1:]) if (a > 0) != (b > 0))
294             rest = next((r for r in reversed(recs) if "restored" in r), None)
295             # A: governor decisions and control effort
296             dec = [(r["time"], g, v) for r in recs if "decision" in r for g, v in r["decision"].items() if "telemetry"
in v]
297             u = [abs(v.get("u_push", 0.0)) for _, _, v in dec if "u_push" in v]
298             times = sorted({t for t, _, _ in dec}); dts = [b - a for a, b in zip(times, times[1:])]
299             dt0 = statistics.median(dts) if dts else 0.0
300             J_u = sum(abs(v.get("u_push", 0.0)) * dt0 for _, _, v in dec)
301             bounds = [v.get("shield_bound") for _, _, v in dec]
302             # representation: the engine's projection against what the telemetry shows at the next decision
303             by = {}
304             for t, g, v in dec:
305                 by.setdefault(g, []).append(v)
306             res2, n_res, dirs, bins = 0.0, 0, {k: [0, 0] for k in W_INT}, {"0.01-0.03": [0, 0], "0.03-0.1": [0, 0], ">=0.
1": [0, 0]}
307             for g, vs in by.items():
308                 for a, b in zip(vs, vs[1:]):
309                     if "state_measured" not in a or "state_measured" not in b:
310                         continue
311                     res2 += sum(W_INT[k] * (b["state_measured"][k] - a["state_projected_next"][k]) ** 2 for k in W_INT);
n_res += 1
312                 for k in W_INT:
313                     pred = a["state_projected_next"][k] - a["state_observed"][k]
314                     real = b["state_measured"][k] - a["state_measured"][k]
315                     if abs(pred) >= DEADBAND and abs(real) >= DEADBAND:
316                         hit = (pred > 0) == (real > 0); dirs[k][0] += hit; dirs[k][1] += 1
317                         key = "0.01-0.03" if abs(pred) < 0.03 else "0.03-0.1" if abs(pred) < 0.1 else ">=0.1"
318                         bins[key][0] += hit; bins[key][1] += 1
319                     hits, valid = sum(x[0] for x in dirs.values()), sum(x[1] for x in dirs.values())
320                     return {"actuator": {"writes_ok": sum(1 for a in act if a.get("rc") == 0), "writes_refused": sum(1 for a in
act if a.get("rc") != 0)},
321                             "r_act_mean_abs_w": statistics.mean(abs(x) for x in r_act) if r_act else None,
322                             "r_act_max_abs_w": max((abs(x) for x in r_act), default=None),
323                             "delay_median_s": statistics.median(delay) if delay else None, "delay_max_s": max(delay,
default=None),
324                             "enforced_under_requested": over, "total_variation_w": tv, "reversals": rev,
325                             "restored_ok": rest.get("ok") if rest else None},
326                             "governor": {"decisions": len(dec), "J_u": J_u, "u_mean_abs": statistics.mean(u) if u else None,
327                             "u_max_abs": max(u, default=None), "saturations": sum(1 for x in u if x >= U_AUTHORITY -
1e-9),
328                             "shield_interventions": sum(1 for b in bounds if b not in ("engine", "start_ceiling")),
329                             "bounds": {b: bounds.count(b) for b in sorted(set(filter(None, bounds)))},
330                             "holds": sum(1 for r in recs if "decision" in r for v in r["decision"].values() if
"hold" in v)},
331                             "representation": {"pairs": n_res, "R_int": math.sqrt(res2 / n_res) if n_res else None,
332                             "directional_accuracy": hits / valid if valid else None, "directional_valid":
valid,
333                             "by_state": {k: (x[0] / x[1] if x[1] else None, x[1]) for k, x in dirs.items()},
334                             "by_predicted_size": {k: (x[0] / x[1] if x[1] else None, x[1]) for k, x in bins.
items()}}}}
335
336
337 def label(prim, g_served, g_p95, g_lost, valid, observation=None):
338     """The result label, by rule (docs/GPU_PREREGISTRATION.md), from the primary verdict, the three guardrails,
and the
339     observation contrast: if Omni watching (writing nothing) already differs from native on the primary outcome,
proven
340     either way, the effect cannot be credited to Omni's authority and no omni result is published."""
341     if not valid:
342         return "INVALID"
343     if observation in ("better, proven", "worse, proven"):
344         return "NOT ATTRIBUTABLE: WATCH DIFFERS FROM NATIVE"
345     ok = g_served and g_p95 and g_lost
346     if prim == "better, proven":
347         return "SUPERIOR WITHIN GUARDRAILS" if ok else "ENERGY IMPROVEMENT WITH SERVICE TRADEOFF"
348     if prim == "worse, proven":
349         return "WORSE"
350     return "NONINFERIOR / INCONCLUSIVE" if ok else "NOT ESTABLISHED"
351
352
353 def verdict(k, mean, half, n):
354     if n < 2 or math.isnan(half):
355         return "not proven (too few repetitions)"
356     if abs(mean) < 1e-12 and half < 1e-12:
357         return "equal"
358     better = (mean < 0) == (k in LOWER)

```

```

359     proven = mean - half > 0 or mean + half < 0
360     return ("better" if better else "worse") + (" , proven" if proven else " , not proven")
361
362
363 def main(root):
364     root = Path(root)
365     receipt = json.loads((root / "receipt.json").read_text()) if (root / "receipt.json").exists() else {}
366     gpus = [int(x) for x in str(receipt.get("gpus", "0")).split(",")]
367     if not receipt and sorted(root.glob("rep-*/receipt.json")):
368         receipt = json.loads(sorted(root.glob("rep-*/receipt.json"))[0].read_text())
369     gpus = [int(x) for x in str(receipt.get("gpus", "0")).split(",")]
370     start0 = (root / "snapshot.txt").read_text().split() if (root / "snapshot.txt").exists() else []
371     start = start0
372     runs, checks, problems, dirs = {}, {}, [], {}
373     for d in sorted(root.glob("rep-*/")):
374         if not (d / "requests.csv").exists():
375             continue
376         rep, a = d.parent.name.split("-", 1)[1], d.name
377         start = (d.parent / "snapshot.txt").read_text().split() if (d.parent / "snapshot.txt").exists() else
start0
378         g, c = arm(d, gpus); dirs.setdefault(rep, {})[a] = d
379         runs.setdefault(a, {})[rep] = g; checks.setdefault(a, {})[rep] = c
380         lim = [f"{x:.2f}" for x in c["limits_seen"]]
381         if a in ("native", "watch") and start and any(all(abs(x - float(s)) >= 1.0 for s in start) for x in c[
"limits_seen"]):
382             problems.append(f"{a} rep {rep}: power limit {lim} differs from the start limit {start}")
383             if a == "watch" and c["writes"]:
384                 problems.append(f"watch rep {rep}: {c['writes']} power-limit writes executed (the control arm must
write nothing)")
385             if a == "native" and c["writes"] + c["would_write"]:
386                 problems.append(f"native rep {rep}: Omni records present")
387             if not c["restored"]:
388                 problems.append(f"{a} rep {rep}: the limit at the end differs from the start (kill switch did not
restore)")
389             rc = json.loads((d.parent / "receipt.json").read_text()) if (d.parent / "receipt.json").exists() else
receipt
390             enf0 = str(rc.get("power_limit_enforced_w", "unsupported"))
391             if a in ("native", "watch") and c["enforced_seen"] and enf0 not in ("", "unsupported") \
and any(abs(x - float(enf0)) >= 1.0 for x in c["enforced_seen"]):
392                 problems.append(f"{a} rep {rep}: enforced power limit {c['enforced_seen']} differs from the snapshot
{enf0}")
393             if a in ("watch", "omni") and c["governor_exit"] not in (None, "0"):
394                 problems.append(f"{a} rep {rep}: the governor exited {c['governor_exit']} (a refused write or a
refused start)")
395             if c.get("actuator", {}).get("writes_refused"):
396                 problems.append(f"{a} rep {rep}: {c['actuator']['writes_refused']} power-limit writes refused by the
device")
397             fz = [json.loads((root / f).read_text()) for f in ("FREEZE.json", "FREEZE_END.json") if (root / f).exists()]
398             fz += [json.loads(f.read_text()) for f in sorted(root.glob("rep-*/FREEZE*.json"))]
399             if fz and any(x["files"] != fz[0]["files"] for x in fz):
400                 problems.append("Omni changed during the run: the frozen file hashes differ between the start, the end,
or the machines")
401             if fz and fz[0].get("phase") == "confirm" and fz[0].get("dirty"):
402                 problems.append("confirmation run on uncommitted Omni code")
403             if len({x.get("phase") for x in fz}) > 1:
404                 problems.append("smoke and confirmation repetitions mixed: smoke data never enter the confirmation")
405             out = {"freeze": fz[0] if fz else None, "receipt": receipt, "start_limit_w": start, "checks": checks,
"problems": problems, "means": {}, "paired": {}}
406             names = {"native": "Native", "watch": "Omni watches only", "omni": "Omni governs"}
407             cols = [a for a in ("native", "watch", "omni") if a in runs]
408             L = ["# GPU bench: native vs Omni-Compass, metered by the device", ""]
409             if receipt:
410                 L += [f"GPU {receipt.get('gpu_name', '?')}, driver {receipt.get('driver', '?')}, persistence mode "
411 f"{receipt.get('persistence', '?')}, start power limit {' '.join(start)} W, workload {receipt.get(
'workload', '?')}.",
412 f"{receipt.get('reps', '?')} repetitions, order rotated, {receipt.get('duration_s', '?')} s per
arm plus "
413 f"{receipt.get('drain_s', '?')} s drain, {receipt.get('cooldown_s', '?')} s idle before each arm.",
414 ""]
415             if fz:
416                 L += [f"Phase: **{fz[0].get('phase')}**. Omni frozen at commit {fz[0].get('commit', '?')[:12]}"
417 f"{f' (uncommitted changes present)' if fz[0].get('dirty') else ''}; file hashes in FREEZE.json,
rechecked at the end.", ""]
418             if problems:
419                 L += ["## INVALID RUN", ""] + [f"- {p}" for p in problems] + [""]
420             for a in cols:
421                 out["means"][a] = {k: _mean([g[k] for g in runs[a].values()]) for k in KEYS}
422             L += ["## Every column, mean over repetitions", " ", "| Gauge | " + " | ".join(names[a] for a in cols) + " |",
423 "|---|" + "---|" * len(cols)]
424             L += [f"| {k} | " + " | ".join("UNAVAILABLE" if k in PHYSICAL and math.isnan(out["means"][a][k]) else _f(out[
"means"][a][k]) for a in cols) + " |"
425 for k in KEYS if k in PHYSICAL or not all(math.isnan(out["means"][a][k]) for a in cols)]
426             L.append("")
427             # three contrasts: total (the preregistered one), observation, authority
428             for a, b, key, what in (("omni", "native", "omni", "Total: Omni governs against native"),
429 ("watch", "native", "watch", "Observation: Omni watches only against native"),

```

```

430         ("omni", "watch", "omni_vs_watch", "Authority: Omni governs against Omni watching")):
431     if a not in runs or b not in runs:
432         continue
433     reps = sorted(set(runs[a]) & set(runs[b]), key=int)
434     out["paired"][key] = {}
435     L += [f"## {what}, {len(reps)} paired repetitions", "",
436         f"| Gauge | {names[b]} | {names[a]} | Change | 95% interval of the difference | Verdict |", "|---|
--|---|---|---|---|"]
437     for k in KEYS:
438         dd = [runs[a][r][k] - runs[b][r][k] for r in reps if not math.isnan(runs[a][r][k]) and not math.
isnan(runs[b][r][k])]
439         if not dd:
440             if k in PHYSICAL:
441                 L.append(f"| {k} | UNAVAILABLE | UNAVAILABLE | | | no meter |")
442             continue
443             n = len(dd); m = sum(dd) / n
444             sd = math.sqrt(sum((x - m) ** 2 for x in dd) / (n - 1)) if n > 1 else float("nan")
445             half = T95.get(n - 1, 1.96) * sd / math.sqrt(n) if n > 1 else float("nan")
446             nb = _mean([runs[b][r][k] for r in reps]); ob = nb + m
447             v = verdict(k, m, half, n)
448             ch = f"({ob - nb) / abs(nb) * 100:+.1f}%" if abs(nb) > 1e-12 else f"{m:+.3g}"
449             out["paired"][key][k] = {"native" if b == "native" else b: nb, a: ob, "base": nb, "diff": m, "ci95":
[m - half, m + half], "verdict": v}
450             L.append(f"| {'*' + k + ' (primary)*' if k == PRIMARY else k} | {_f(nb)} | {_f(ob)} | {ch} | {m -
half:+.4g} to {m + half:+.4g} | {v} |")
451             L.append("")
452             wr = {a: sum(c["writes"] for c in checks.get(a, {}).values()) for a in cols}
453             # the preregistered verdict (docs/GPU_PREREGISTRATION.md): primary outcome with the two service guardrails
454             po = out["paired"].get("omni", {})
455             if PRIMARY in po:
456                 prim = po[PRIMARY]["verdict"]
457                 sv, p95 = po.get("requests served"), po.get("response time, 95th percentile (ms)")
458                 lost = po.get("requests not served")
459                 g_served = sv is not None and sv["ci95"][0] >= -MARGIN_SERVED * abs(sv["native"])
460                 g_p95 = p95 is not None and p95["ci95"][1] <= MARGIN_P95 * abs(p95["native"])
461                 g_lost = lost is None or sv is None or lost["ci95"][1] <= MARGIN_LOST * abs(sv["native"])
462                 ww = out["paired"].get("watch", {}).get(PRIMARY, {}).get("verdict")
463                 head = label(prim, g_served, g_p95, g_lost, not problems, ww)
464                 av = out["paired"].get("omni_vs_watch", {}).get(PRIMARY, {}).get("verdict")
465                 out["headline"] = {"primary": prim, "guardrail_served": g_served, "guardrail_p95": g_p95,
"guardrail_not_served": g_lost,
466                     "verdict": head, "watch_primary": ww, "authority_primary": av,
467                     "wall": po.get(WALL, {}).get("verdict"), "valid": not problems}
468                 c = po[PRIMARY]
469                 L += [f"## Verdict on the preregistered question", "",
470                     f"Work per energy under Omni against native: {_f(c['native'])} -> {_f(c['omni'])} served requests
per kJ, "
471                     f"difference {c['diff']:+.4g} (95% interval {c['ci95'][0]:+.4g} to {c['ci95'][1]:+.4g}).",
472                     f"Guardrails: requests served {'held' if g_served else 'FAILED'} (not below -1%), "
473                     f"95th-percentile response time {'held' if g_p95 else 'FAILED'} (not above +10%), "
474                     f"requests not served {'held' if g_lost else 'FAILED'} (not above +1% of native served).",
475                     f"Observation (watch against native) on the same outcome: {ww or 'n/a'}. "
476                     f"Authority (omni against watch): {av or 'n/a'}.",
477                     *([f"Whole machine at the wall (smart plug Omni never reads): {po[WALL]['verdict']}, "
478                     f"{_f(po[WALL]['native'])} -> {_f(po[WALL]['omni'])} served requests per kJ." if WALL in po
else []),
479                     f"***Result, by rule: {head}.*", "")]
480             L += receipt_lines(checks)
481             crows, cby = credit(root, dirs)
482             out["credit"] = {"writes": crows, "by_decider": cby}
483             if cby:
484                 L += [f"## Credit per write (against native at the same moments of the same request stream)", "",
485                     "Each write owns the interval until the next write. Joules and requests finished in it, Omni
governing minus native; "
486                     "the intervals tile the arm, so the credits add up. The decider is the rule that set the written
limit "
487                     "(engine, a floor, the busy gate, a reflex, heat, the speed lock).", "",
488                     "| Decided by | Writes | Seconds owned | GPU joules vs native | Requests finished vs native |", "|-
--|---|---|---|---|"]
489                 L += [f"| {k} | {v['writes']} | {v['seconds']:.0f} | {v['dJ_vs_native']:+.4g} | {v['dserved_vs_native']
+d} |" for k, v in sorted(cby.items())]
490                 L.append("")
491                 L += [f"## The control", "", f"- Power-limit writes executed: " + " ".join(f"{names[a]} {wr[a]}" for a in
cols) + " ",
492                     "- Every arm ended at the start limit." if not any("restore" in p for p in problems) else "- An arm
did NOT end at the start limit.",
493                     "- Energy is the device's own power.draw integrated over time; no number here is modelled.", ""]
494                 (root / "GPU_REPS.json").write_text(json.dumps(out, indent=1)); (root / "GPU_REPS.md").write_text("\n".join(
L))
495                 print("\n".join(L))
496                 return 2 if problems else 0
497
498
499 def receipt_lines(checks):
500     """Receipts A and B, per arm and repetition, as the table's own sections."""
501     L = []

```

```

502     rows = [(a, r, c) for a in ("omni", "watch") for r, c in sorted(checks.get(a, {}).items(), key=lambda x: int(
503         x[0])) if c.get("governor")]
504     if not rows:
505         return L
506     fmt = lambda x, n=3: "n/a" if x is None else f"{x:.{n}g}" if isinstance(x, float) else str(x)
507     L += ["## Actuator fidelity (receipt B: requested against read back from the device)", "",
508         "| Arm | Rep | Writes | Refused | mean abs(read back - requested) W | max W | delay median s | delay
509         max s | enforced under requested | total variation W | reversals | restored |",
510         "|---|---|---|---|---|---|---|---|---|---|---|---|",
511         L += [f"| {a} | {r} | {c['actuator']['writes_ok']} | {c['actuator']['writes_refused']} | {fmt(c['actuator']
512         'r_act_mean_abs_w')} | "
513         f"{fmt(c['actuator']['r_act_max_abs_w'])} | {fmt(c['actuator']['delay_median_s'])} | {fmt(c[
514         'actuator']['delay_max_s'])} | "
515         f"{c['actuator']['enforced_under_requested']} | {fmt(c['actuator']['total_variation_w'])} | {c[
516         'actuator']['reversals']} | "
517         f"{c['actuator']['restored_ok']} |" for a, r, c in rows]
518     L += ["", "## Control effort (receipt A)", "",
519         "| Arm | Rep | Decisions | J_u = sum abs(u) dt | mean abs(u) | max abs(u) | saturations | shield
520         interventions | holds |",
521         "|---|---|---|---|---|---|---|---|",
522     L += [f"| {a} | {r} | {c['governor']['decisions']} | {fmt(c['governor']['J_u'])} | {fmt(c['governor']
523         'u_mean_abs')} | "
524         f"{fmt(c['governor']['u_max_abs'])} | {c['governor']['saturations']} | {c['governor']
525         'shield_interventions']} | "
526         f"{c['governor']['holds']} |" for a, r, c in rows]
527     L += ["", "## Representation fidelity (the engine's projection against the telemetry at the next decision)",
528         "",
529         f"R_int: weighted root-mean-square of h(z next) - F_h(h(z), u) over E, U, I_U, S, B (weights 1).
530         Directional "
531         f"accuracy: the share of movements the projection called in the right direction, movements under {
532         DEADBAND} excluded.", "",
533         "| Arm | Rep | Pairs | R_int | directional accuracy | valid directions | by predicted size |", "|---|---
534         -|---|---|---|---|---|",
535     L += [f"| {a} | {r} | {c['representation']['pairs']} | {fmt(c['representation']['R_int'])} | "
536         f"{fmt(c['representation']['directional_accuracy'])} | {c['representation']['directional_valid']} | "
537         + ", ".join(f"{k}: {fmt(v[0])} of {v[1]}" for k, v in c['representation']['by_predicted_size'].items(
538         )) + " |"
539         for a, r, c in rows]
540     return L + [""]
541
542 def _mean(xs):
543     xs = [x for x in xs if not math.isnan(x)]
544     return sum(xs) / len(xs) if xs else float("nan")
545
546 def _f(x):
547     return "n/a" if math.isnan(x) else f"{x:.4g}"
548
549 if __name__ == "__main__":
550     sys.exit(main(sys.argv[1]))
551

```

tools/gpu_workload.py

145 lines SHA-256 flf1d19a347db982e6c3381ffb4796ddb3694df049eaccbe253ce859d4cad7bf

```
1  """The pinned GPU workload for scripts/gpu_paired.sh: a fixed stream of inference-like requests served by one
   GPU.
2
3  Each request is --iters fp16 matrix products of size --n x --n on the GPU (synchronised, so its time is the
   GPU's time).
4  Requests arrive open-loop on a schedule fixed by --seed: Poisson arrivals whose rate steps through --phases,
   each phase a
5  share of the GPU's own capacity measured by `calibrate` at full power before any arm runs. Every arm therefore
   receives
6  exactly the same requests at exactly the same moments; what differs is only how the GPU is governed.
7
8  calibrate times --calib requests back to back at the current (start) power limit and writes calib.json:
9  iters (chosen so one request takes about --target-ms), service_ms, and every setting below
10 run serves the schedule for --duration seconds, then --drain seconds for requests still queued, and
   writes
11     latency.csv (elapsed_seconds, latency_ms, ok) live, one row per finished request
12     requests.csv (arrival_s, start_s, done_s, latency_ms, ok)
13     summary.json (requests, served, not served, t0_epoch: the wall clock at time 0, the settings)
14 Any other workload (vLLM, TensorRT-LLM, an MLPerf inference harness) plugs into the bench through WORKLOAD_CMD
15 (scripts/gpu_paired.sh) by writing the same three files; tools/gpu_reps.py checks them.
16 A request still queued when the drain ends is not served (ok = 0).
17 --sim replaces the GPU by a sleep of service_ms (for the harness's own tests on machines without a GPU).
18 """
19 from __future__ import annotations
20
21 import argparse, json, queue, random, sys, threading, time
22 from pathlib import Path
23
24
25 def gpu_kernel(n, device):
26     import torch
27     a = torch.randn(n, n, device=device, dtype=torch.float16)
28     b = torch.randn(n, n, device=device, dtype=torch.float16)
29
30     def run(iters):
31         for _ in range(iters):
32             torch.matmul(a, b)
33             torch.cuda.synchronize(device)
34         run(3)
35     return run
36
37
38 def sim_kernel(service_ms):
39     def run(iters):
40         time.sleep(service_ms / 1000.0 * iters / 10.0)
41     return run
42
43
44 def calibrate(a):
45     run = sim_kernel(a.sim_ms) if a.sim else gpu_kernel(a.n, a.device)
46     iters = 10
47     t = time.perf_counter(); run(iters); one = (time.perf_counter() - t) * 1000.0
48     iters = max(1, round(iters * a.target_ms / max(one, 1e-3)))
49     ts = []
50     for _ in range(a.calib):
51         t = time.perf_counter(); run(iters); ts.append((time.perf_counter() - t) * 1000.0)
52     ts.sort()
53     c = {"iters": iters, "service_ms": ts[len(ts) // 2], "n": a.n, "seed": a.seed, "phases": a.phases,
54         "target_ms": a.target_ms, "sim": bool(a.sim)}
55     Path(a.out).mkdir(parents=True, exist_ok=True)
56     (Path(a.out) / "calib.json").write_text(json.dumps(c, indent=1))
57     print(json.dumps(c))
58
59
60 def arrivals(c, duration):
61     """The fixed schedule: Poisson arrivals, rate = phase share / service time, from the seed alone."""
62     rng = random.Random(c["seed"])
63     shares = [float(x) for x in str(c["phases"]).split(",")]
64     per = duration / len(shares)
65     out, t = [], 0.0
66     for i, f in enumerate(shares):
67         rate = f / (c["service_ms"] / 1000.0)
68         t = max(t, i * per)
69         while True:
70             t += rng.expovariate(rate)
71             if t >= (i + 1) * per:
72                 t = (i + 1) * per
73                 break
74         out.append(t)
75     return out
76
77
78 def serve(a):
```

```

79     c = json.loads(Path(a.calib_file).read_text())
80     run = sim_kernel(c["service_ms"] * 10.0 / c["iters"]) if c["sim"] else gpu_kernel(c["n"], a.device)
81     sched = arrivals(c, a.duration)
82     out = Path(a.out); out.mkdir(parents=True, exist_ok=True)
83     lat = open(out / "latency.csv", "w"); lat.write("elapsed_seconds,latency_ms,ok\n"); lat.flush()
84     q, rows, t0 = queue.Queue(), [], time.perf_counter()
85     t0_epoch = time.time() # the wall clock at t0, so the bench can line requests up with its
meters
86     end = a.duration + a.drain
87
88     def worker():
89         while True:
90             arr = q.get()
91             if arr is None:
92                 return
93             now = time.perf_counter() - t0
94             if now >= end:
95                 rows.append((arr, None, None, None, 0)); continue
96             run(c["iters"])
97             done = time.perf_counter() - t0
98             ms = (done - arr) * 1000.0
99             rows.append((arr, now, done, ms, 1))
100            lat.write(f"{done:.3f},{ms:.3f},1\n"); lat.flush()
101
102    th = threading.Thread(target=worker, daemon=True); th.start()
103    for arr in sched:
104        wait = arr - (time.perf_counter() - t0)
105        if wait > 0:
106            time.sleep(wait)
107        q.put(arr)
108    q.put(None)
109    th.join(timeout=max(0.0, end - (time.perf_counter() - t0)) + 5.0)
110    left = end - (time.perf_counter() - t0)
111    if left > 0:
112        time.sleep(left) # every arm is measured over the same window length
113    served = sum(1 for r in rows if r[4] == 1)
114    with open(out / "requests.csv", "w") as f:
115        f.write("arrival_s,start_s,done_s,latency_ms,ok\n")
116        for r in rows:
117            f.write(",".join(" " if v is None else f"{v:.4f}" if isinstance(v, float) else str(v)) for v in r) +
"\n")
118    s = {"requests": len(sched), "served": served, "not_served": len(sched) - served, "duration_s": a.duration,
"t0_epoch": t0_epoch,
119        "drain_s": a.drain, **{k: c[k] for k in ("iters", "service_ms", "n", "seed", "phases", "sim")}}
120    (out / "summary.json").write_text(json.dumps(s, indent=1))
121    print(json.dumps(s))
122
123
124    def main(argv=None):
125        ap = argparse.ArgumentParser(description=__doc__, formatter_class=argparse.RawDescriptionHelpFormatter)
126        ap.add_argument("cmd", choices=["calibrate", "run"])
127        ap.add_argument("--out", default="gpu_run")
128        ap.add_argument("--calib-file", default="")
129        ap.add_argument("--device", default="cuda:0")
130        ap.add_argument("--n", type=int, default=4096)
131        ap.add_argument("--target-ms", type=float, default=50.0)
132        ap.add_argument("--calib", type=int, default=40)
133        ap.add_argument("--seed", type=int, default=20260928)
134        ap.add_argument("--phases", default="0.3,0.6,0.8,0.3,0.6,0.3", help="shares of full-power capacity, one per
phase")
135        ap.add_argument("--duration", type=float, default=600.0)
136        ap.add_argument("--drain", type=float, default=30.0)
137        ap.add_argument("--sim", action="store_true")
138        ap.add_argument("--sim-ms", type=float, default=20.0)
139        a = ap.parse_args(argv)
140        calibrate(a) if a.cmd == "calibrate" else serve(a)
141
142
143    if __name__ == "__main__":
144        main()
145

```

tools/live_reps.py

94 lines SHA-256 3c3e587aa10f7eef42271fa29e082223a300540134eafcd04b925acade4adc5

```
1  """Repeated live runs: aggregate native / watch (Omni runs, writes nothing) / omni (B) / strict (C) arms over
   repetitions of scripts/kind_bench.sh.
2  Per repetition the gauges come from pilot/bench_report.py (capture.csv, latency.csv); across repetitions this
   reports
3  means and, for each Omni arm against native, the paired mean difference with a t-based 95% interval (n
   repetitions).
4  Usage: python tools/live_reps.py DIR (DIR holds bench-<arm>-<rep>/) -> DIR/LIVE_REPS.md, DIR/LIVE_REPS.json"""
5  import csv, json, math, sys
6  from pathlib import Path
7  import numpy as np
8  ROOT = Path(__file__).resolve().parents[1]; sys.path.insert(0, str(ROOT))
9  from pilot.bench_report import gauges, latency, pod_starts, LOWER_BETTER
10
11  T95 = {1: 12.706, 2: 4.303, 3: 3.182, 4: 2.776, 5: 2.571, 6: 2.447, 7: 2.365, 8: 2.306, 9: 2.262}
12  KEYS = ["worker nodes in service, mean", "node-hours", "energy, parked workers still on at idle power (Wh)",
   "energy (Wh)", "response time (ms), mean", "response time (ms), 95th percentile",
13  "response time (ms), 99th percentile", "failed requests (%)", "pending pods, pod-minutes", "utilisation (
   used / allocatable)",
14  "CPU used (cores), mean", "Omni's own CPU (cores), mean", "CPU used with Omni's own (cores), mean",
15  "energy per core-hour (Wh)", "HPA replicas, mean",
16  "pods started", "pod start wait, total (s)", "pod start wait, mean (s)"]
17
18
19  LABEL = {"energy, parked workers still on at idle power (Wh)": "energy, parked workers still on at idle power (Wh,
   declared model)",
20  "energy (Wh)": "energy, parked workers at 25 W standby (Wh, declared model; kind never does this)",
21  "energy per core-hour (Wh)": "energy per core-hour (Wh, the 25 W standby model)"}
22  NEUTRAL = {"CPU used (cores), mean", "utilisation (used / allocatable)", "Omni's own CPU (cores), mean",
23  "CPU used with Omni's own (cores), mean"} # more is not better or worse by itself
24  NOTE = ["**Energy on kind is a declared model, not a meter.** Every worker stays powered and Ready in every arm;
   the first",
25  "energy row counts a parked worker at its full idle power, which is what kind does. The second counts it
   at the",
26  "declared standby power, which needs a node autoscaler that really removes the machine; this run has none.
   ", ""]
27
28
29  def arm_gauges(d):
30  rows = list(csv.DictReader(open(d / "capture.csv")))
31  g = gauges(rows); g.update(latency(str(d / "latency.csv"))); g.update(pod_starts(d))
32  # the controller's own cost (its process and every command it ran; it runs beside the cluster, not in it),
   from its
33  # audit: counted so a CPU saving in the cluster is never reported without what Omni itself spent. Native: 0.
34  own = 0.0
35  if "-native-" not in d.name and (d / "audit.jsonl").exists():
36  for line in open(d / "audit.jsonl"):
37  if "overhead" in line:
38  own = float(json.loads(line)["overhead"].get("cores_mean", 0.0))
39  if own == 0.0:
40  own = float("nan") # an Omni arm without its cost record: unknown, never zero
41  elif "-native-" not in d.name:
42  own = float("nan")
43  g["Omni's own CPU (cores), mean"] = own
44  g["CPU used with Omni's own (cores), mean"] = g.get("CPU used (cores), mean", float("nan")) + own
45  return g
46
47
48  def main(root):
49  root = Path(root); runs = {}
50  for d in sorted(root.glob("bench-*-*")):
51  _, arm, rep = d.name.split("-", 2)
52  if (d / "capture.csv").exists() and not (d / "INVALID").exists():
53  runs.setdefault(arm, {})[rep] = arm_gauges(d)
54  out = {"repetitions": {a: sorted(r) for a, r in runs.items()}, "means": {}, "paired": {}}
55  for a, r in runs.items():
56  out["means"][a] = {k: float(np.nanmean([g.get(k, np.nan) for g in r.values()])) for k in KEYS}
57  L = ["# Repeated live runs on kind (native vs Omni watching only vs Omni on top vs Omni alone)", ""]
58  cols = [a for a in ("native", "watch", "omni", "strict") if a in runs]
59  names = {"native": "Native", "watch": "Omni watches only", "omni": "Omni on top", "strict": "Omni alone"}
60  L += ["## All columns, mean over repetitions", "", "| Gauge | " + " | ".join(names[a] for a in cols) + " |",
61  "|---|" + "----:" * len(cols)]
62  L += [f"| {LABEL.get(k, k)} | " + " | ".join(f"{out['means'][a][k]:.4g}" for a in cols) + " |" for k in KEYS]
63  L += [""] + NOTE
64  for a in ("watch", "omni", "strict"):
65  if a not in runs or "native" not in runs:
66  continue
67  reps = sorted(set(runs[a]) & set(runs["native"]))
68  out["paired"][a] = {}
69  title = {"watch": "W: Omni-Compass watches only (dry run: the cost of being there)", "omni": "B: Omni-
   Compass on top",
70  "strict": "C: Omni-Compass decides (strict)"}[a]
71  L += [f"## {title} vs native, {len(reps)} paired repetitions", "",
72  "| Gauge | Native | Omni | Change | 95% interval of the difference | Significant |", "|---|---|---|"]
```

```

73 |---:|---:|---|"]
74     for k in KEYS:
75         d = np.array([runs[a][r].get(k, np.nan) - runs["native"][r].get(k, np.nan) for r in reps], float)
76         d = d[~np.isnan(d)]
77         if len(d) == 0:
78             continue
79         nb = float(np.nanmean([runs["native"][r].get(k, np.nan) for r in reps])); ob = nb + float(d.mean())
80         half = T95.get(len(d) - 1, 1.96) * (d.std(ddof=1) / math.sqrt(len(d))) if len(d) > 1 else float("nan")
81         sig = len(d) > 1 and (d.mean() - half > 0 or d.mean() + half < 0)
82         better = (d.mean() < 0) == (k in LOWER_BETTER)
83         ch = (ob - nb) / abs(nb) * 100 if abs(nb) > 1e-12 else None
84         out["paired"][a][k] = {"native": nb, "omni": ob, "diff": float(d.mean()), "ci95": [float(d.mean() -
85 half), float(d.mean() + half)], "significant": bool(sig)}
86         ch_s = f"{ch:+.1f}%" if ch is not None else f"{ob - nb:+.3g} (native is 0)"
87         L.append(f"| {LABEL.get(k, k)} | {nb:.4g} | {ob:.4g} | {ch_s} | {d.mean() - half:+.4g} to {d.mean() +
88 half:+.4g} | "
89             f"({'yes, more' if d.mean() > 0 else 'yes, less'}) if k in NEUTRAL else ('yes, better' if
90 better else 'yes, worse')) if sig else 'no' |")
91         L.append("")
92     (root / "LIVE_REPS.json").write_text(json.dumps(out, indent=1)); (root / "LIVE_REPS.md").write_text("\n".join(
93 L))
94     print("\n".join(L))
95
96 if __name__ == "__main__":
97     main(sys.argv[1])
98

```

tools/mechanism.py

113 lines SHA-256 dd0683e64dee77278b3ea0e2b9e62aa33b30b96e1f80f140f09463253bb93874

```
1  """Mechanism of action: what the Omni-Compass equations actually do inside the governor, measured.
2
3  1. linearisation    Jacobian of equations (1)-(7) at the operating point the governor lives at; eigenvalues
   give each
4
5                      mode's time constant, in engine time and in governor decisions (the adapter advances engine
   time
6
7                      0.1 per decision: omnicompass/adapter.py, Governor.step)
8
9  2. who moves the   per decision the state changes twice: assimilation (the observation blended in with weight
   0.339)
10
11                     state          and evolution (one macro step of the equations). Their sizes are measured on a real
   observation
12
13                     stream (fleet plant, web vessel, fleet mode)
14
15  3. the controller  whether equation (2)'s control input u is applied during governance
16
17  4. compute         microseconds per governor decision
18
19  Writes results/MECHANISM_OF_ACTION.json.
20  """
21
22  from __future__ import annotations
23
24  import json, math, sys, time
25  from dataclasses import replace
26  from pathlib import Path
27
28  import numpy as np
29
30  ROOT = Path(__file__).resolve().parents[1]
31  sys.path.insert(0, str(ROOT))
32  from omnicompass import adapter as A
33  from omnicompass.core import State, derivatives, macro_step
34
35  NAMES = ["E", "U", "I_U", "S", "B", "B_dot"]
36
37  def jacobian(x, p, t=0.0, h=1e-6):
38      f0 = np.array(derivatives(x, p, t).as_tuple())
39      J = np.zeros((6, 6))
40      for j in range(6):
41          v = list(x.as_tuple()); v[j] += h
42          J[:, j] = (np.array(derivatives(State(*v), p, t).as_tuple()) - f0) / h
43      return J
44
45  def record_stream(vessel="web", seed=101):
46      """Observation stream a fleet-mode governor sees on the fleet plant (captured from a real run)."""
47      from fleet import sim_slo
48      from fleet.harness import make_scenario
49      seen = []
50      orig = A.Governor.step
51
52      def spy(self, obs, conflicts):
53          seen.append(dict(obs))
54          return orig(self, obs, conflicts)
55      A.Governor.step = spy
56
57      try:
58          sim_slo.run(make_scenario(vessel, seed), "omni_fleet")
59      finally:
60          A.Governor.step = orig
61      return seen
62
63  def split(stream):
64      p = replace(A.STACK_PARAMS)
65      x = State(**vars(A.INITIAL_STATE)); t = 0.0
66      d_obs, d_math = [], []
67      for obs in stream:
68          o = A.observe_vector(obs, 0)
69          xa = A.assimilate(x, o, p)
70          xn, _ = macro_step(xa, p, t, target=None)
71          d_obs.append(np.abs(np.array(xa.as_tuple()) - np.array(x.as_tuple())))
72          d_math.append(np.abs(np.array(xn.as_tuple()) - np.array(xa.as_tuple())))
73          x, t = xn, t + 0.1
74      return np.mean(d_obs, 0), np.mean(d_math, 0), x, p
75
76  def main():
77      stream = record_stream()
78      obs_move, math_move, x_op, p_op = split(stream)
79      J = jacobian(x_op, p_op)
80      ev = np.linalg.eigvals(J)
81      modes = []
82      for e in sorted(ev, key=lambda z: z.real):
83          tau = (-1.0 / e.real) if e.real < 0 else float("inf")
84          modes.append({"eigenvalue": [float(e.real), float(e.imag)], "time_constant_engine": tau,
85                      "time_constant_decisions": tau / 0.1 if math.isfinite(tau) else None,
```

```

80         "period_decisions": (2 * math.pi / abs(e.imag) / 0.1) if abs(e.imag) > 1e-9 else None})
81     g = A.Governor(law=A.mode_law("fleet")); g.set_mode(A.AUTOPILLOT)
82     t0 = time.perf_counter(); n = 0
83     for obs in stream * 3:
84         g.step(obs, 0); n += 1
85     us = (time.perf_counter() - t0) / n * 1e6
86     share = {k: {"assimilation": float(a), "equations": float(m), "equations_share": float(m / max(a + m, 1e-
12))}
87             for k, a, m in zip(NAMES, obs_move, math_move)}
88     out = {"operating_point": dict(zip(NAMES, map(float, x_op.as_tuple()))),
89           "forcing_at_operating_point": {"beta_int": p_op.beta_int, "beta_ext": p_op.beta_ext},
90           "jacobian": J.tolist(), "modes": modes,
91           "per_decision_movement": share,
92           "engine_time_per_decision": 0.1, "assimilation_weight": A.ASSIMILATION,
93           "controller_applied_in_governance": False,
94           "controller_note": "Governor.step evolves with target=None: equation (2)'s input u is zero; the
control law is "
95               "evaluated only to produce 'push', which gates capacity release",
96               "microseconds_per_decision_python": us, "decisions_measured": n}
97     (ROOT / "results/MECHANISM_OF_ACTION.json").write_text(json.dumps(out, indent=1))
98     print("operating point:", {k: round(v, 3) for k, v in out["operating_point"].items()})
99     print("modes (engine time constant -> decisions):")
100     for m in modes:
101         print(f"    lambda = {m['eigenvalue'][0]:+.3f}{m['eigenvalue'][1]:+.3f}i    tau = {m[
'time_constant_engine']:.3g} "
102               f"-> {m['time_constant_decisions'] if m['time_constant_decisions'] is None else round(m[
'time_constant_decisions'], 1)} decisions"
103               + (f", period {m['period_decisions']:.0f} decisions" if m["period_decisions"] else ""))
104     print("mean movement per decision (assimilation vs equations):")
105     for k, v in share.items():
106         print(f"    {k:6} observation {v['assimilation']:.4f} equations {v['equations']:.4f} -> equations do {
100*v['equations_share']:.0f}%")
107     print(f"controller u applied during governance: {out['controller_applied_in_governance']}")
108     print(f"compute: {us:.1f} microseconds per decision (Python)")
109
110
111 if __name__ == "__main__":
112     main()
113

```

tools/mechanism_identity.py

165 lines SHA-256 1fa5edcbf447c488abc92936529d28158e8bba2e95a9f3c5a407d67234dea389

```
1  """The mechanism's identity: M = (F, Theta, C, h, G, M_act, dt, A), one fingerprint per engine configuration.
2
3  python3 tools/mechanism_identity.py          # write results/MECHANISM_IDENTITY.json
4  python3 tools/mechanism_identity.py --check  # compare the code with the recorded identity (verify.py runs
5  this)
6  python3 tools/mechanism_identity.py --id     # the canonical mechanism id of the code as it stands (the GPU
7  bench records it)
8
9  Each component is fingerprinted (SHA-256) from the source text of the functions that compute it, and Theta from
10 the
11 values of every constant and parameter. The mechanism id is the SHA-256 of the eight component fingerprints in
12 order,
13 so it changes exactly when a component changes, and changes in unrelated code (reports, docs) leave it alone.
14
15 F      state law          dx/dt = F(x, d, u; Theta): derivatives (+ v_eff, s_flow), per configuration
16 Theta parameters        Params defaults, PARAMETER_RANGES, KP, U_AUTHORITY, steps, windows,
17 ASSIMILATION
18 C      internal controller u = clip(-f_U + KP (sigma - U), +-U_AUTHORITY), per configuration
19 h      observation map     observe_vector, assimilate (omnicompass/adaptor.py)
20 G      authority law       Governor.step, AllocationLaw, mode_law (omnicompass/adaptor.py)
21 M_act  authority -> actuator GPU power limit (omni_controller/gpu_governor.py: shield_limit, busy_gate,
22 lock_decide, GpuGovernor.set_limit, GpuGovernor.step); Kubernetes (
23 omni_controller/muscles.py)
24 dt      execution contract rk4_step, macro_step (u computed before each RK4 micro step, held through
25 its four
26 stages), MACRO_DT, MICRO_STEPS; the GPU decision interval
27 A      admissibility / shield omnicompass/shield.py (violations, enforce, ShieldLimits), shield_limit
28
29 Also records the frozen 500-fixture comparison of the two configurations (benchmarks/core_evidence.py population,
30 seed 223387268): executable both, not interchangeable. No new population is drawn.
31 """
32 from __future__ import annotations
33
34 import argparse, hashlib, inspect, json, math, sys
35 from dataclasses import asdict
36 from pathlib import Path
37
38 ROOT = Path(__file__).resolve().parents[1]
39 sys.path.insert(0, str(ROOT))
40 import numpy as np
41 from omnicompass import core as C, configurations as K, adapter as A, shield as SH
42 from omni_controller import gpu_governor as GG, muscles as MU
43
44 OUT = ROOT / "results" / "MECHANISM_IDENTITY.json"
45 SEED, RUNS = 223387268, 500
46
47 def h_src(*objs):
48     return hashlib.sha256("\n".join(inspect.getsource(o) for o in objs).encode()).hexdigest()
49
50 def h_val(v):
51     return hashlib.sha256(json.dumps(v, sort_keys=True).encode()).hexdigest()
52
53 def theta():
54     return {"Params_defaults": asdict(C.Params()), "PARAMETER_RANGES": C.PARAMETER_RANGES, "KP": C.KP,
55            "U_AUTHORITY": C.U_AUTHORITY, "MACRO_STEPS": C.MACRO_STEPS, "MICRO_STEPS": C.MICRO_STEPS,
56            "MACRO_DT": C.MACRO_DT, "BASIN_TOL": C.BASIN_TOL, "CONVEY_WINDOW": C.CONVEY_WINDOW,
57            "CERT_WINDOW": C.CERT_WINDOW, "ASSIMILATION": A.ASSIMILATION,
58            "gpu_governor_defaults": {k: v for k, v in vars(GG.parser().parse_args(["--smi", "x"])).items()
59                                     if k not in ("smi", "audit", "kill_file", "latency_file", "baseline_file",
60 "gpus")}}
61
62 def components(name):
63     F = h_src(C.derivatives, C.v_eff, C.s_flow) if name == K.SYMMETRIC else h_src(K.printed_derivatives, K.
64 printed_v_eff)
65     Cc = h_src(C.control_command) if name == K.SYMMETRIC else h_src(K.control_command_for, K.derivatives_for)
66     return {"F": F, "Theta": h_val(theta()), "C": Cc,
67            "h": h_src(A.observe_vector, A.assimilate, A.clamp),
68            "G": h_src(A.Governor, A.AllocationLaw, A.mode_law),
69            "M_act": hashlib.sha256((h_src(GG.shield_limit, GG.busy_gate, GG.lock_decide, GG.GpuGovernor)
70 + hashlib.sha256((ROOT / "omni_controller" / "muscles.py").read_bytes()).
71 hexdigest()).encode()).hexdigest(),
72            "dt": h_src(C.rk4_step, C.axy, C.macro_step),
73            "A": h_src(SH.ShieldLimits, SH.violations, SH.enforce, GG.shield_limit)}
74
75 def mech_id(comp):
76     return hashlib.sha256("".join(comp[k] for k in ("F", "Theta", "C", "h", "G", "M_act", "dt", "A")).encode()).
77 hexdigest()
```

```

73
74 def trajectory(name, x0, p):
75     """The canonical 20-macro-step run for a configuration, with the same predicates as core.simulate. For
76     symmetric_verified it is checked equal to core.simulate on every fixture."""
77     deriv = K.derivatives_for(name)
78     tgt = C.nearest_basin_sign(x0.U); h = C.MICRO_DT
79     b0 = C.basin_sign(x0.U); streak = 1 if b0 else 0; ssign = b0; convey = cert = 0
80     J = peak = 0.0; sat = 0; x = x0; finite = True
81     for m in range(1, C.MACRO_STEPS + 1):
82         t = (m - 1) * C.MACRO_DT # as core.macro_step: each macro interval starts on its own
83     clock
84     for _ in range(C.MICRO_STEPS):
85         drift = deriv(x, p, t, 0.0).U
86         raw = -drift + C.KP * (tgt - x.U); u = max(-C.U_AUTHORITY, min(C.U_AUTHORITY, raw))
87         J += abs(u) * h; peak = max(peak, abs(u)); sat += abs(raw) > C.U_AUTHORITY + C.EPS
88         k1 = deriv(x, p, t, u); k2 = deriv(C._axpy(x, k1, 0.5 * h), p, t + 0.5 * h, u)
89         k3 = deriv(C._axpy(x, k2, 0.5 * h), p, t + 0.5 * h, u); k4 = deriv(C._axpy(x, k3, h), p, t + h, u)
90         x = C.State(*(a + (h / 6.0) * (b1 + 2 * b2 + 2 * b3 + b4) for a, b1, b2, b3, b4 in
91             zip(x.as_tuple(), k1.as_tuple(), k2.as_tuple(), k3.as_tuple(), k4.as_tuple())))
92         t += h
93         finite = finite and all(math.isfinite(v) for v in x.as_tuple())
94         b = C.basin_sign(x.U)
95         streak, ssign = ((streak + 1, b) if b == ssign else (1, b)) if b else (0, 0)
96         convey = convey or int(streak >= C.CONVEY_WINDOW); cert = cert or int(streak >= C.CERT_WINDOW)
97     return {"finite": finite, "convey": convey, "cert": cert, "err": abs(x.U - tgt), "J_u": J, "peak": peak,
98     "sat": sat, "final": x}
99
100 def comparison():
101     rng = np.random.default_rng(SEED)
102     pop = [(C.sample_state(rng), C.sample_params(rng)) for _ in range(RUNS)]
103     out = {}
104     for name in (K.SYMMETRIC, K.PRINTED):
105         rs = [trajectory(name, x0, p) for x0, p in pop]
106         if name == K.SYMMETRIC:
107             for (x0, p), r in zip(pop, rs):
108                 ref = C.simulate(x0, p)
109                 assert ref.final.as_tuple() == r["final"].as_tuple() and ref.convey == r["convey"] and ref.cert
110                 == r["cert"], \
111                     "the comparison loop differs from core.simulate"
112         out[name] = {"finite": sum(r["finite"] for r in rs), "convey5": sum(r["convey"] for r in rs),
113             "cert10": sum(r["cert"] for r in rs), "runs": RUNS,
114             "mean_final_target_error": float(np.mean([r["err"] for r in rs])),
115             "mean_integrated_abs_u": float(np.mean([r["J_u"] for r in rs])),
116             "max_actuator_magnitude": float(max(r["peak"] for r in rs)),
117             "saturated_microsteps": int(sum(r["sat"] for r in rs))}
118     return out
119
120 def build():
121     cfigs = {}
122     for name in (K.SYMMETRIC, K.PRINTED):
123         comp = components(name)
124         cfigs[name] = {"mechanism_id": mech_id(comp), "components": comp, "role": K.describe(name),
125             "canonical": name == K.DEFAULT}
126     return {"definition": "M = (F, Theta, C, h, G, M_act, dt, A); state x = (E, U, I_U, S, B, B_dot); dx/dt = F(
127     x, d, u; Theta)",
128         "canonical": K.DEFAULT, "configurations": cfigs, "theta": theta(),
129         "fixture_comparison": {"population": f"benchmarks/core_evidence.py, seed {SEED}, {RUNS} fixtures (
130     frozen)",
131             "results": comparison(),
132             "reading": "both configurations are executable; they are not experimentally
133     interchangeable "
134             "(different final error and control effort). Evidence names the
135     mechanism id that produced it.}}"}
136
137 def check():
138     if not OUT.exists():
139         return ["no results/MECHANISM_IDENTITY.json: run python3 tools/mechanism_identity.py"]
140     rec = json.loads(OUT.read_text()); bad = []
141     for name, c in rec["configurations"].items():
142         now = components(name)
143         for k, v in now.items():
144             if c["components"][k] != v:
145                 bad.append(f"{name}: component {k} changed since the identity was recorded")
146         if mech_id(now) != c["mechanism_id"]:
147             bad.append(f"{name}: mechanism id changed")
148     return bad
149
150 def main(argv=None):
151     ap = argparse.ArgumentParser(description=__doc__, formatter_class=argparse.RawDescriptionHelpFormatter)
152     ap.add_argument("--check", action="store_true")
153     ap.add_argument("--id", action="store_true", help="print the canonical mechanism id of the code as it
154     stands")

```

```

151     a = ap.parse_args(argv)
152     if a.id:
153         print(mech_id(components(K.DEFAULT))); return 0
154     if a.check:
155         bad = check(); print("\n".join(bad) if bad else "mechanism identity matches the code"); return 1 if bad
156     else 0
157     rec = build(); OUT.write_text(json.dumps(rec, indent=1) + "\n")
158     for n, c in rec["configurations"].items():
159         print(f"{n}: {c['mechanism_id']}")
160     print(json.dumps(rec["fixture_comparison"]["results"], indent=1))
161     return 0
162
163 if __name__ == "__main__":
164     sys.exit(main())
165

```

tools/protocol_bench.py

199 lines SHA-256 54660e345eab6316fdb2d0e49f8f9b563c10806ebddf0f74b426bbb63ecd3c75

```
1  """OmniCompass Runtime Benchmark Protocol (benchmark_protocol.docx), stage 1, on the fleet plant (fleet/sim_slo.
2  py).
3  Same scenarios, same faults, every architecture:
4  A  each platform alone      k8s_hpa70_ca, openshift, gke_optimize, aks_nap, turbonomic, cast_ai, spot_ocean
5  B  Omni-Compass on top of it omniB:<platform> with the frozen per-platform settings (tuning/
6  B_SETTINGS_PER_PLATFORM.json)
7  C  Omni-Compass alone      C-hpa: the held-out C league pick (HPA kept); C-strict: the closure law, HPA
8  replaced                    (settings: tuning/CLOSURE_SEARCH2_DEV.json, else CLOSURE_SEARCH_DEV.json, else
9  defaults)
10 Stress levels 1-5 (protocol section 6). Level L injects L fault events, times drawn uniformly in [120, 1200)
11 ticks, each
12 one of (drawn with the seed, identical for every arm):
13 node_crash      L x 10% of the cluster's running machines die at once (at least one)
14 load_spike      every workload in the cluster: demand x (1 + 0.5 L) for 20-80 ticks      (queue overload,
15 scaling)
16 service_crash   L x 20% of the cluster's services lose half their pods (crash loop, restart churn) and queue
17 two ticks of demand as backlog
18 cpu_pressure    a noisy neighbour: extra demand of L x 8% of the cluster's allocatable cores for 40-160 ticks,
19 spread over the workloads
20 Admissible basin, per tick (declared before any run): worst response time <= 1000 ms (10 x service time),
21 queue ratio < 0.28, pending pods <= 5% of requested, site power <= 1.02 x limit, thermal < 0.96.
22 Gauges (protocol section 7):
23 conveyed        every fault recovered within 20 min (80 ticks) AND all clusters in the basin for the last 10
24 min
25 recovery time    from a fault to the first tick that starts 2 min (8 ticks) in the basin; unrecovered = to
26 the end
27 failure persist. minutes outside the basin
28 pages           out-of-basin episodes lasting 5 min or longer (the Prometheus `for: 5m` rule): each one
29 calls a
30 human (the protocol's human intervention count; nobody intervenes in the simulation)
31 max error / CPU / queue / latency, timeouts (demand share with response > 2 s), work completed, energy
32 Usage: python tools/protocol_bench.py [runs_per_level=100] -> results/protocol/{raw_results.csv,
33 summary_results.csv,
34 PROTOCOL_SUMMARY.json}
35 """
36 import csv, json, math, sys
37 from multiprocessing import Pool
38 from pathlib import Path
39 import numpy as np
40 ROOT = Path(__file__).resolve().parents[1]; sys.path.insert(0, str(ROOT))
41 from fleet import sim_slo
42 from fleet.harness import make_scenario, STEPS, TICK
43 from fleet.sim_slo import BLaw, DirectLaw
44 from omnicompass.speed import SpeedLaw
45 from omnicompass.closure import ClosureLaw
46 from tuning.league import candidates
47
48 PLATFORMS = ["k8s_hpa70_ca", "openshift", "gke_optimize", "aks_nap", "turbonomic", "cast_ai", "spot_ocean"]
49 VESSELS = ["web", "multi", "batch", "gpu"]
50 LEVELS = [1, 2, 3, 4, 5]
51 SEED0 = 900001
52 R_LIM, Q_LIM, PEND_LIM, P_LIM, TH_LIM = 1000.0, 0.28, 0.05, 1.02, 0.96
53 REC_OK, REC_LIMIT, FINAL, PAGE = 8, 80, 40, 20
54 KINDS = ["node_crash", "load_spike", "service_crash", "cpu_pressure"]
55 B_SET = json.load(open(ROOT / "tuning/B_SETTINGS_PER_PLATFORM.json"))
56 C_PICK = json.load(open(ROOT / "tuning/LEAGUE_HELDOUT.json"))["omni_per_vessel"]
57 CANDS = candidates()
58 _cs = [ROOT / "tuning/CLOSURE_FINAL_DEV.json", ROOT / "tuning/CLOSURE_SEARCH2_DEV.json", ROOT / "tuning/
59 CLOSURE_SEARCH_DEV.json"]
60 CLOSURE = next((json.load(open(p)) for p in _cs if p.exists()), {})
61 _g = ROOT / "tuning/GLOBAL_LEAGUE_PREREGISTRATION.json"
62 if _g.exists():
63     # one global setting for every workload (the confirmatory
64     setting)
65     _s = json.load(open(_g))["setting"]
66     CLOSURE = {v: {"closure": dict(_s["closure"], site=(v == "multi")), "direct": _s["direct"], "speed": _s[
67     "speed"]}} for v in VESSELS}
68 SL = dict(rho0=0.7, rho_min=0.7, kI=0.0, kE=0.0)
69
70 def plan(vessel, seed, level):
71     rng = np.random.default_rng(seed * 10 + level)
72     n_cl = 4 if vessel == "multi" else 1
73     ev = []
74     for _ in range(level):
75         ev.append(dict(t=int(rng.integers(120, 1200)), ci=int(rng.integers(0, n_cl)), kind=KINDS[int(rng.
76         integers(0, 4))],
77         dur=int(rng.integers(20, 81)), dur2=int(rng.integers(40, 161)), pick=float(rng.uniform(
78         )))
79     return sorted(ev, key=lambda e: e["t"])
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
```

```

69
70 def injector(events, level):
71     from fleet.sim import ALLOC
72     by_t = {}
73     for e in events:
74         by_t.setdefault(e["t"], []).append(e)
75
76     def on_tick(t, scn):
77         for e in by_t.get(t, []):
78             c = scn.clusters[e["ci"]]; p = c.pool
79             if e["kind"] == "node_crash":
80                 p.nodes = max(0, p.nodes - max(1, int(math.floor(0.10 * level * p.nodes))))
81             elif e["kind"] == "load_spike":
82                 for w in c.workloads:
83                     w.demand[t:t + e["dur"]] *= (1.0 + 0.5 * level)
84             elif e["kind"] == "service_crash":
85                 svc = [w for w in c.workloads if w.hpa] or c.workloads
86                 k = max(1, int(round(0.2 * level * len(svc))))
87                 start = int(e["pick"] * len(svc))
88                 for j in range(k):
89                     w = svc[(start + j) % len(svc)]
90                     if w.hpa:
91                         w.replicas = max(1, w.replicas // 2)
92                         w.backlog += 2.0 * float(w.demand[t])
93             elif e["kind"] == "cpu_pressure":
94                 extra = 0.08 * level * p.nodes * p.cores * ALLOC / max(1, len(c.workloads))
95                 for w in c.workloads:
96                     w.demand[t:t + e["dur2"]] += extra
97         return on_tick
98
99
100 def arms(vessel):
101     a = [{"A", pl, pl, {}} for pl in PLATFORMS]
102     a += [{"B", pl, "omniB:" + pl, dict(b_low=BLaw(**B_SET[vessel][pl]))} for pl in PLATFORMS]
103     pick = C_PICK[vessel]; cfg = CANDS[pick]
104     if cfg is None:
105         a.append(("C", "C-hpa", "omni_fleet", {}))
106     else:
107         a.append(("C", "C-hpa", "omni_speed", dict(speed_low=SpeedLaw(**cfg["law"]), omni_every=cfg["every"],
108                                                     lat_gain=cfg["lat_gain"], slo_mult=cfg["slo_mult"])))
109     cl = CLOSURE.get(vessel, {})
110     a.append(("C", "C-strict", "omni_closure", dict(speed_low=SpeedLaw(**cl.get("speed", SL)), omni_every=1,
111                                                     direct_low=DirectLaw(**cl.get("direct", {})),
112                                                     closure_low=ClosureLaw(**cl.get("closure", {}))))))
113     return a
114
115
116 def gauge(r, events):
117     s = r["series"]
118     ok = np.array([(x[0] <= R_LIM and x[1] < Q_LIM and x[2] <= PEND_LIM and x[3] <= P_LIM and x[4] < TH_LIM) for
119 x in s])
120     # run of in-basin ticks starting at each tick
121     run = np.zeros(len(ok) + 1, int)
122     for i in range(len(ok) - 1, -1, -1):
123         run[i] = run[i + 1] + 1 if ok[i] else 0
124     recs, unrec = [], 0
125     for e in events:
126         t0 = e["t"]
127         hit = next((t for t in range(t0, len(ok)) if run[t] >= REC_OK or (run[t] > 0 and t + run[t] == len(ok))),
128 None)
129         if hit is None:
130             recs.append(len(ok) - t0); unrec += 1
131         else:
132             recs.append(hit - t0)
133     pages, i = 0, 0
134     while i < len(ok):
135         if not ok[i]:
136             j = i
137             while j < len(ok) and not ok[j]:
138                 j += 1
139             pages += (j - i) >= PAGE; i = j
140         else:
141             i += 1
142     conveyed = all(x <= REC_LIMIT for x in recs) and bool(ok[-FINAL:].all())
143     mins = TICK / 60.0
144     return {"conveyed": int(conveyed), "recovery_min": float(np.mean(recs)) * mins if recs else 0.0,
145             "unrecovered": unrec, "first_conveyance_min": (float(np.argmax(run >= REC_OK)) * mins if (run >=
146 REC_OK).any() else STEPS * mins),
147             "persistence_min": float((-ok).sum()) * mins, "pages": pages,
148             "max_latency_ms": float(max(x[0] for x in s)), "max_queue": float(max(x[1] for x in s)),
149             "max_pending": float(max(x[2] for x in s)), "max_power": float(max(x[3] for x in s)),
150             "timeouts": r["timeouts"], "work_completed": r["work_completed"], "energy_kwh": r["energy_kwh"],
151             "p99_ms": r["p99_ms"], "start_stop": r["machines_started"] + r["machines_stopped"]}
152
153 def job(args):

```

```

152     vessel, seed, level = args
153     ev = plan(vessel, seed, level)
154     sc = make_scenario(vessel, seed)
155     rows = []
156     for col, name, arm, kw in arms(vessel):
157         r = sim_slo.run(sc, arm, on_tick=injector(ev, level), **kw)
158         rows.append(dict(vessel=vessel, level=level, seed=seed, column=col, system=name, **gauge(r, ev)))
159     return rows
160
161
162 GAUGES = ["conveyed", "recovery_min", "unrecovered", "persistence_min", "pages", "max_latency_ms", "max_queue",
163           "max_pending", "timeouts", "work_completed", "energy_kwh", "p99_ms", "start_stop"]
164
165
166 def main():
167     n = int(sys.argv[1]) if len(sys.argv) > 1 else 100
168     jobs = [(v, SEED0 + i, L) for v in VESSELS for L in LEVELS for i in range(n)]
169     with Pool(4) as p:
170         rows = [r for rs in p.imap_unordered(job, jobs, chunksize=4) for r in rs]
171         rows.sort(key=lambda r: (r["vessel"], r["level"], r["seed"], r["column"], r["system"]))
172         out = ROOT / "results/protocol"; out.mkdir(parents=True, exist_ok=True)
173         with open(out / "raw_results.csv", "w", newline="") as f:
174             wr = csv.DictWriter(f, fieldnames=list(rows[0])); wr.writeheader(); wr.writerows(rows)
175         summ = []
176         for v in VESSELS:
177             for L in LEVELS + ["all"]:
178                 for sysname in sorted({r["system"] for r in rows}, key=lambda s: (s.startswith("C"), s)):
179                     for col in ("A", "B", "C"):
180                         rs = [r for r in rows if r["vessel"] == v and (L == "all" or r["level"] == L) and r[
181                             "system"] == sysname and r["column"] == col]
182                         if not rs:
183                             continue
184                         d = dict(vessel=v, level=L, column=col, system=sysname, runs=len(rs))
185                         for g in GAUGES:
186                             d[g] = float(np.sum([r[g] for r in rs])) if g in ("conveyed", "pages", "unrecovered")
187                             else float(np.mean([r[g] for r in rs]))
188                         summ.append(d)
189                         with open(out / "summary_results.csv", "w", newline="") as f:
190                             wr = csv.DictWriter(f, fieldnames=list(summ[0])); wr.writeheader(); wr.writerows(summ)
191                             (out / "PROTOCOL_SUMMARY.json").write_text(json.dumps({"runs_per_level": n, "seeds": [SEED0, SEED0 + n - 1],
192                                     "closure_settings": CLOSURE, "summary": summ},
193                                     indent=1))
194                             for d in summ:
195                                 if d["level"] == "all":
196                                     print(f"{d['vessel']:6} {d['column']:13} {d['system']:13} conveyed {int(d['conveyed']):4}/{d['runs']}
197                                     pages {int(d['pages']):5} "
198                                     f"recovery {d['recovery_min']:6.1f} min persist {d['persistence_min']:6.1f} min timeouts {
199                                     100*d['timeouts']:.2f}% energy {d['energy_kwh']:.2f}")
200
201 if __name__ == "__main__":
202     main()
203

```

tools/release_manifest.py

124 lines SHA-256 a3e9874f10e62e0b9e47252d0fcfb671ebd6969e466fd4e4a863713795cd8fff

```
1  """The release's one identity: RELEASE_MANIFEST.json.
2
3  python3 tools/release_manifest.py          # write the manifest, run verify.py --quick, record its receipt
4  python3 tools/release_manifest.py --check  # compare the files with the manifest (verify.py runs this)
5
6  The manifest names the commit it was written at and fingerprints (SHA-256) everything that identifies the
7  release: the
8  canonical engine and the reference engine, the C++ twins (through the seal), the GPU protocol, the live evidence
9  (execution commit, run id and raw-file digests of sets 20 and 21), the canonical engine declaration, and the
10 receipt
11 of the verification run made when it was written. One object, so nobody has to reconstruct provenance from
12 history.
13 """
14 from __future__ import annotations
15
16 import argparse, datetime, hashlib, json, subprocess, sys
17 from pathlib import Path
18
19 ROOT = Path(__file__).resolve().parents[1]
20 MANIFEST = ROOT / "RELEASE_MANIFEST.json"
21 RECEIPT = ROOT / "results" / "VERIFY_RECEIPT.txt"
22
23 GROUPS = {
24     "canonical_engine": ["omnicompass/core.py", "cpp/src/core.cpp", "docs/CANONICAL_ENGINE.md"],
25     "reference_engine": ["reference/omni_compass_reference_engine.py", "reference/PROVENANCE.json"],
26     "cpp_twins_seal": ["results/SEAL.json"],
27     "gpu_protocol": ["docs/GPU_PREREGISTRATION.md", "scripts/gpu_paired.sh", "tools/gpu_workload.py", "tools/
28 gpu_reps.py",
29                     "omni_controller/gpu_governor.py", "omnicompass/adapters.py", "omni_controller/muscles.py"],
30     "live_set_20": ["results/live/LIVE_REPS_20.md", "results/live/raw/run-36366603505/SHA256SUMS_ALL.txt"],
31     "live_set_21": ["results/live/LIVE_REPS_21.md", "results/live/SET21_ARTIFACTS.json"],
32     "live_set_22": ["results/live/LIVE_REPS_22.md", "results/live/SET22_ARTIFACTS.json",
33                     "results/live/reaggregated/LIVE_REPS_36488547793.md"],
34     "preregistration": ["results/PREREGISTRATION.json", "results/LOCK_AMENDMENTS.json"],
35     "mechanism": ["results/MECHANISM_IDENTITY.json", "tools/mechanism_identity.py", "docs/TRACKING_THEOREM.md",
36                  "tools/tracking_bounds.py", "results/TRACKING_BOUNDS.json", "docs/EVIDENCE_LEDGER.md",
37                  "RECEIPT.md"],
38     "license": ["LICENSE", "NOTICE"],
39 }
40
41 # engine components checked identical to this release's at each execution commit (tools/mechanism_identity.py
42 # components F, C, h, G, dt, recomputed from the source at that commit); the Kubernetes actuator map and shield
43 # are as
44 # at the execution commit
45 SAME = ["F", "C", "h", "G", "dt"]
46 LIVE = {
47     "set_20": {"run_id": 36366603505, "execution_commit": "18220d4", "repetitions": 10, "plant": "kind (GitHub
48 Actions)",
49               "engine_components_identical_to_release": SAME,
50               "raw": "results/live/raw/run-36366603505/ (SHA256SUMS_ALL.txt covers every file)"},
51     "set_21": {"run_id": 36466558583, "execution_commit": "9e64f7b", "repetitions": 10, "plant": "kind (GitHub
52 Actions)",
53               "engine_components_identical_to_release": SAME,
54               "raw": "GitHub artifacts, digests in results/live/SET21_ARTIFACTS.json; recomputed by reaggregate
55 run 36485672281"},
56     "set_22": {"run_id": 36488547793, "execution_commit": "cfdc17c", "repetitions": 10, "plant": "kind (GitHub
57 Actions)",
58               "load": "open-loop, fixed rate (equal work)", "engine_components_identical_to_release": SAME,
59               "raw": "GitHub artifacts, digests in results/live/SET22_ARTIFACTS.json; recomputed by the
60 reaggregate workflow"},
61 }
62
63 def sha_files(paths):
64     h = hashlib.sha256()
65     for p in paths:
66         h.update(p.encode() + b"\0" + (ROOT / p).read_bytes())
67     return h.hexdigest()
68
69 def sha(p):
70     return hashlib.sha256((ROOT / p).read_bytes()).hexdigest()
71
72 def git_head():
73     return subprocess.run(["git", "-C", str(ROOT), "rev-parse", "HEAD"], capture_output=True, text=True,
74                           check=True).stdout.strip()
75
76 def check():
77     """Returns a list of mismatches (empty when every fingerprinted file matches)."""
78     if not MANIFEST.exists():
79         return ["no RELEASE_MANIFEST.json: run python3 tools/release_manifest.py"]
80     m = json.loads(MANIFEST.read_text()); bad = []
```

```

72     for group, files in m["files"].items():
73         for p, h in files.items():
74             if not (ROOT / p).exists():
75                 bad.append(f"{group}: {p} is missing")
76             elif sha(p) != h:
77                 bad.append(f"{group}: {p} changed since the manifest (python3 tools/release_manifest.py)")
78     r = m.get("verification_receipt")
79     if r and (not RECEIPT.exists() or sha(RECEIPT.relative_to(ROOT)) != r["sha256"]):
80         bad.append("verification receipt: results/VERIFY_RECEIPT.txt differs from the manifest")
81     return bad
82
83
84 def main(argv=None):
85     ap = argparse.ArgumentParser(description=__doc__, formatter_class=argparse.RawDescriptionHelpFormatter)
86     ap.add_argument("--check", action="store_true")
87     a = ap.parse_args(argv)
88     if a.check:
89         bad = check()
90         print("\n".join(bad) if bad else "release manifest matches the files")
91         return 1 if bad else 0
92     mi = json.loads((ROOT / "results" / "MECHANISM_IDENTITY.json").read_text())
93     canon = mi["configurations"][mi["canonical"]]
94     tests = sorted(str(p.relative_to(ROOT)) for p in (ROOT / "tests").glob("test_*.py")) + ["verify.py"]
95     head = git_head()
96     m = {"release": "Omni-Compass", "release_id": f"omni-compass-{head[:12]}", "written_at_commit": head,
97          "written_at": datetime.datetime.now(datetime.timezone.utc).isoformat(timespec="seconds"),
98          "canonical_engine": "symmetric_verified (docs/CANONICAL_ENGINE.md)",
99          "mechanism_id": canon["mechanism_id"], "mechanism_configuration": mi["canonical"],
100         "mechanism_components": canon["components"],
101         "alternative_embodiments": {n: c["mechanism_id"] for n, c in mi["configurations"].items() if not c["canonical"]},
102         "fingerprints": {"engine_sha256": sha("omnicompass/core.py"), "cpp_sha256": sha("cpp/src/core.cpp"),
103                          "controller_sha256": sha("omni_controller/controller.py"),
104                          "observation_map_sha256": canon["components"]["h"], "authority_map_sha256": canon["components"]["G"],
105                          "actuator_map_sha256": canon["components"]["M_act"], "shield_sha256": sha("omnicompass/shield.py"),
106                          "gpu_protocol_sha256": sha_files(GROUPS["gpu_protocol"]),
107                          "preregistration_sha256": sha_files(GROUPS["preregistration"] + ["docs/GPU_PREREGISTRATION.md"])},
108         "test_suite_sha256": sha_files(tests)},
109     "files": {g: {p: sha(p) for p in fs} for g, fs in GROUPS.items()},
110     "live_evidence": LIVE,
111     "physical_meter_results": "none yet: the GPU bench (scripts/gpu_paired.sh) has not been run on a card"}
112     MANIFEST.write_text(json.dumps(m, indent=1) + "\n")
113     r = subprocess.run([sys.executable, str(ROOT / "verify.py"), "--quick"], capture_output=True, text=True,
114                        cwd=ROOT)
115     RECEIPT.write_text(f"verify.py --quick at commit {m['written_at_commit']}, {m['written_at']}\n\n" + r.stdout)
116     ok = r.returncode == 0 and "VERIFICATION: PASS" in r.stdout
117     m["verification_receipt"] = {"file": "results/VERIFY_RECEIPT.txt", "sha256": sha("results/VERIFY_RECEIPT.txt"), "passed": ok}
118     MANIFEST.write_text(json.dumps(m, indent=1) + "\n")
119     print(f"RELEASE_MANIFEST.json written at {m['written_at_commit'][:12]}; verification {'PASS' if ok else 'FAILED'}")
120     return 0 if ok else 1
121
122 if __name__ == "__main__":
123     sys.exit(main())
124

```

tools/seal.py

106 lines SHA-256 2eb3e46b22443212a3bad8e5dba875bcl2b4e6a15293ef4caa3c8fa4385786c

```
1  """The seal: every law that exists in Python and in C++, locked together.
2
3  python3 tools/seal.py          # build the C++, run every twin's parity test, and only if all pass write
  SEAL.json
4  python3 tools/seal.py --check  # compare the files with SEAL.json (verify.py runs this)
5
6  SEAL.json (results/SEAL.json) records, for each twin, the SHA-256 of its Python file and of its C++ files, the
  parity
7  test that proved them equal, and when. A change to either side of a twin breaks the seal until its parity test
  passes
8  again and the seal is rewritten, so the Python and the C++ can never drift apart unnoticed. Mechanisms that
  exist in
9  Python only are listed too, so nobody takes them for twinned.
10 """
11 from __future__ import annotations
12
13 import argparse, datetime, hashlib, json, shutil, subprocess, sys, tempfile
14 from pathlib import Path
15
16 ROOT = Path(__file__).resolve().parents[1]; sys.path.insert(0, str(ROOT))
17 SEAL = ROOT / "results" / "SEAL.json"
18
19 TWINS = {
20     "core engine (six-state equations)": {"python": ["omnicompass/core.py"],
21     "cpp": ["cpp/src/core.cpp", "cpp/include/omnicompass/core.hpp"], "proof": "cpp/scripts/compare_results.
  py on fixtures/current_500_*.csv"},
22     "governor (allocation laws)": {"python": ["omnicompass/adaptor.py"],
23     "cpp": ["cpp/src/governor.cpp", "cpp/include/omnicompass/governor.hpp"], "proof": "tests/
  test_cpp_governor_parity.py"},
24     "safety shield": {"python": ["omnicompass/shield.py"],
25     "cpp": ["cpp/src/shield.cpp", "cpp/include/omnicompass/shield.hpp"], "proof": "tests/
  test_cpp_shield_parity.py"},
26     "HPA replica law": {"python": ["fleet/harness.py"],
27     "cpp": ["cpp/src/hpa.cpp", "cpp/include/omnicompass/hpa.hpp"], "proof": "tests/test_cpp_hpa_parity.py"},
28     "closure law (machines)": {"python": ["omnicompass/closure.py"],
29     "cpp": ["cpp/src/closure.cpp", "cpp/include/omnicompass/closure.hpp"], "proof": "tests/
  test_cpp_closure_parity.py"},
30     "conveyance law (the conserved budget)": {"python": ["omnicompass/conveyance.py"],
31     "cpp": ["cpp/src/conveyance.cpp", "cpp/include/omnicompass/conveyance.hpp"], "proof": "tests/
  test_cpp_conveyance_parity.py"},
32     "nervous system (authority per organ, living band)": {"python": ["omnicompass/nervous_system.py"],
33     "cpp": ["cpp/src/nervous_system.cpp", "cpp/include/omnicompass/nervous_system.hpp"], "proof": "tests/
  test_cpp_twins_parity.py"},
34     "compass and ledger (composite storage)": {"python": ["omnicompass/compass.py", "omnicompass/storage.py"],
35     "cpp": ["cpp/src/compass.cpp", "cpp/include/omnicompass/compass.hpp"], "proof": "tests/
  test_cpp_twins_parity.py"},
36     "GPU governor rules (shield limit, busy gate, speed lock, window, baseline)": {"python": ["omni_controller/
  gpu_governor.py"],
37     "cpp": ["cpp/src/gpu_rules.cpp", "cpp/include/omnicompass/gpu_rules.hpp"], "proof": "tests/
  test_cpp_twins_parity.py"},
38 }
39 PYTHON_ONLY = {
40     "omni_controller/gpu_governor.py (device I/O)": "reading nvidia-smi and writing power limits; its decision
  rules are twinned above",
41     "omni_controller/muscles.py": "Kubernetes and hardware muscles (convey, rightsize, coldstart, batch, contain,
  cooling, CPU/GPU connectors)",
42     "omni_controller/controller.py": "Kubernetes controller (HPA target, pod reflex, machines, strict replicas)",
43     "hardware/node_exchange.py": "CPU + GPU on one power budget (simulation harness; the law itself is twinned
  above)",
44 }
45
46
47 def sha(p):
48     return hashlib.sha256((ROOT / p).read_bytes()).hexdigest()
49
50
51 def prove():
52     """Build the C++ and run every twin's parity test; raises on any failure."""
53     import verify
54     tmp = Path(tempfile.mkdtemp()); b = verify.build_cpp(tmp)
55     subprocess.run([str(b / "oc_run_fixture"), str(ROOT / "fixtures/current_500_inputs.csv"), str(tmp / "cpp.
  csv")], check=True)
56     subprocess.run([sys.executable, str(ROOT / "cpp/scripts/compare_results.py"), str(ROOT / "fixtures/
  current_500_expected.csv"),
57     str(tmp / "cpp.csv"), str(tmp / "cpp.json")], check=True, capture_output=True)
58     assert json.loads((tmp / "cpp.json").read_text())["pass"], "core engine: C++ differs from the 500 fixtures"
59     from tests import (test_cpp_governor_parity, test_cpp_shield_parity, test_cpp_hpa_parity,
  test_cpp_closure_parity,
60     test_cpp_conveyance_parity)
61     for m in ("power_protect", "throughput", "fleet", "fleet_balanced", "fleet_wear"):
62         test_cpp_governor_parity.main(str(b / "oc_governor"), 20, m)
63     test_cpp_shield_parity.main(str(b / "oc_shield"), 8)
64     test_cpp_hpa_parity.main(str(b / "oc_hpa"), (901,))
```

```

65     test_cpp_closure_parity.main(str(b / "oc_closure"), 1)
66     assert test_cpp_conveyance_parity.main(str(b / "oc_conveyance"), 200) < 1e-9
67     from tests import test_cpp_twins_parity
68     test_cpp_twins_parity.main(str(b / "oc_twins"), 2000)
69
70
71 def check():
72     """Returns a list of broken twins (empty when every sealed file is unchanged)."""
73     if not SEAL.exists():
74         return ["no seal: run python3 tools/seal.py"]
75     seal = json.loads(SEAL.read_text()); broken = []
76     for name, t in TWINS.items():
77         rec = seal["twins"].get(name)
78         if rec is None:
79             broken.append(f"{name}: not in the seal"); continue
80         for p in t["python"] + t["cpp"]:
81             if rec["sha256"].get(p) != sha(p):
82                 broken.append(f"{name}: {p} changed since the seal ({t['proof']}) must pass, then python3 tools/
seal.py)")
83     return broken
84
85
86 def main(argv=None):
87     ap = argparse.ArgumentParser(description=__doc__, formatter_class=argparse.RawDescriptionHelpFormatter)
88     ap.add_argument("--check", action="store_true")
89     a = ap.parse_args(argv)
90     if a.check:
91         broken = check()
92         print("\n".join(broken) if broken else f"seal intact: {len(TWINS)} Python/C++ twins")
93         return 1 if broken else 0
94     prove()
95     out = {"sealed_at": datetime.datetime.now(datetime.timezone.utc).isoformat(timespec="seconds"),
96           "rule": "each twin's Python and C++ files were proven equal by its parity test before these hashes
were written",
97           "twins": {name: {"proof": t["proof"], "sha256": {p: sha(p) for p in t["python"] + t["cpp"]}} for name,
t in TWINS.items()},
98           "python_only": PYTHON_ONLY}
99     SEAL.write_text(json.dumps(out, indent=1) + "\n")
100     print(f"sealed {len(TWINS)} Python/C++ twins after every parity test passed; {len(PYTHON_ONLY)} mechanisms
Python only")
101     return 0
102
103
104 if __name__ == "__main__":
105     sys.exit(main())
106

```

tools/source_book.py

135 lines SHA-256 303b336123f3d795796fd63539294c75767847d69353f730a7c1010bb41dc097

```
1  """The source code as a PDF, for reading, archiving and copyright deposit.
2
3  python3 tools/source_book.py cpp      # the C++ engine: docs/source/OMNI_COMPASS_CPP_SOURCE.pdf
4  python3 tools/source_book.py python   # the Python engine, controllers and harnesses: docs/source/
   OMNI_COMPASS_PYTHON_SOURCE.pdf
5
6  Each PDF has a title page (title, owner, version and date), a contents list with page numbers, and every file in
   full
7  with its path, line count and SHA-256 at its head, numbered lines, and the path and page number on every page.
   The
8  version is the git commit the files were read from, so the PDF names exactly the code it holds.
9  """
10 from __future__ import annotations
11
12 import argparse, datetime, hashlib, subprocess, sys
13 from pathlib import Path
14
15 from reportlab.lib.pagesizes import letter
16 from reportlab.lib.styles import ParagraphStyle
17 from reportlab.lib.units import inch
18 from reportlab.pdfbase import pdfmetrics
19 from reportlab.pdfbase.ttfonts import TTFont
20 from reportlab.platypus import BaseDocTemplate, Frame, PageBreak, PageTemplate, Paragraph, Preformatted, Spacer,
   Table
21
22 ROOT = Path(__file__).resolve().parents[1]
23 OWNER = "The Omni-Compass LLC"
24 MONO = "/usr/share/fonts/truetype/dejavu/DejaVuSansMono.ttf"
25 SANS = "/usr/share/fonts/truetype/dejavu/DejaVuSans.ttf"
26 SANS_B = "/usr/share/fonts/truetype/dejavu/DejaVuSans-Bold.ttf"
27
28 BOOKS = {
29     "cpp": {"title": "Omni-Compass C++20 Engine: Source Code", "out": "docs/source/OMNI_COMPASS_CPP_SOURCE.pdf",
30            "about": "The C++20 engine of Omni-Compass: the six-state core, the governor and its allocation laws,
   the safety "
31                  "shield, the HPA replica law, the closure law, the conveyance law, the nervous system, the
   compass and its "
32                  "ledger, and the GPU governor's rules, with the command-line tools that run them. Each law
   is the twin of "
33                  "its Python original and proven equal to it by a parity test (results/SEAL.json).",
34            "globs": ["cpp/CMakeLists.txt", "cpp/include/omnicompass/*.hpp", "cpp/src/*.cpp", "cpp/tools/*.cpp",
   "cpp/scripts/*.py"]},
35     "python": {"title": "Omni-Compass Python Engine, Controllers and Harnesses: Source Code",
36               "out": "docs/source/OMNI_COMPASS_PYTHON_SOURCE.pdf",
37               "about": "The Python engine of Omni-Compass (omnicompass/), its Kubernetes and GPU controllers
   and muscles "
38                     "(omni_controller/), the hardware and power-budget harnesses (hardware/), the tools (
   tools/), the "
39                     "benchmarks (benchmarks/) and the verification entry point (verify.py).",
40               "globs": ["omnicompass/*.py", "omni_controller/*.py", "hardware/*.py", "tools/*.py", "benchmarks/
   *.py", "verify.py"]},
41 }
42
43
44 def files(book):
45     out = []
46     for g in BOOKS[book]["globs"]:
47         out += sorted(p for p in ROOT.glob(g) if p.is_file() and p.name != "__init__.py" or (p.name == "__init__
   py" and p.stat().st_size > 0))
48     seen, uniq = set(), []
49     for p in out:
50         if p not in seen:
51             seen.add(p); uniq.append(p)
52     return uniq
53
54
55 def commit():
56     try:
57         return subprocess.run(["git", "-C", str(ROOT), "rev-parse", "HEAD"], capture_output=True, text=True,
   check=True).stdout.strip()
58     except (subprocess.SubprocessError, OSError):
59         return "unknown"
60
61
62 def build(book):
63     pdfmetrics.registerFont(TTFont("Mono", MONO)); pdfmetrics.registerFont(TTFont("Sans", SANS)); pdfmetrics.
   registerFont(TTFont("SansB", SANS_B))
64     mono_face = TTFont("MonoCheck", MONO).face
65     cfg = BOOKS[book]; fs = files(book); ver = commit(); today = datetime.date.today().isoformat()
66     code = ParagraphStyle("code", fontName="Mono", fontSize=6.6, leading=8.0)
67     head = ParagraphStyle("head", fontName="SansB", fontSize=10, leading=13, spaceAfter=2)
68     meta = ParagraphStyle("meta", fontName="Mono", fontSize=6.6, leading=8.5, spaceAfter=6)
69     body = ParagraphStyle("body", fontName="Sans", fontSize=10, leading=14, spaceAfter=8)
```

```

70 title = ParagraphStyle("title", fontName="SansB", fontSize=20, leading=26, spaceAfter=16)
71 missing = set()
72 texts = {}
73 for p in fs:
74     t = p.read_text(encoding="utf-8").replace("\t", "    ")
75     for ch in set(t):
76         if ord(ch) > 126 and ord(ch) not in mono_face.charToGlyph:
77             missing.add(ch)
78     texts[p] = t
79 if missing:
80     raise SystemExit(f"font lacks glyphs for {sorted(missing)}; the PDF would show boxes")
81
82 total = sum(t.count("\n") + 1 for t in texts.values())
83 state = {"file": ""}
84
85 def footer(canvas, doc):
86     canvas.saveState(); canvas.setFont("Sans", 7)
87     canvas.drawString(0.6 * inch, 0.4 * inch, cfg["title"] + (f" | {state['file']}" if state["file"] else
""))
88     canvas.drawRightString(letter[0] - 0.6 * inch, 0.4 * inch, f"page {doc.page}")
89     canvas.drawString(0.6 * inch, 0.27 * inch, f"Copyright (c) 2026 {OWNER}. All rights reserved. Version {
ver[:12]}")
90     canvas.restoreState()
91
92 out = ROOT / cfg["out"]; out.parent.mkdir(parents=True, exist_ok=True)
93 pages = {}
94
95 class Doc(BaseDocTemplate):
96     def afterFlowable(self, f):
97         if getattr(f, "_file", None):
98             state["file"] = f._file; pages.setdefault(f._file, self.page)
99
100 def story(page_map):
101     s = [Spacer(1, 1.2 * inch), Paragraph(cfg["title"], title),
102         Paragraph(f"Copyright © 2026 {OWNER}. All rights reserved.", body),
103         Paragraph(cfg["about"], body),
104         Paragraph(f"Version: git commit {ver}<br/>Date: {today}<br/>Files: {len(fs)}<br/>Lines: {total:,.}",
body),
105         Paragraph("Use of this code is governed by the LICENSE file distributed with it.", body), PageBreak(
),
106         Paragraph("Contents", title)]
107     rows = [[str(p.relative_to(ROOT)), f"{texts[p].count(chr(10)) + 1:,.} lines", str(page_map.get(str(p.
relative_to(ROOT)), ""))] for p in fs]
108     tbl = Table(rows, colWidths=[4.6 * inch, 1.2 * inch, 0.8 * inch])
109     tbl.setStyle([(("FONT", (0, 0), (-1, -1), "Mono", 7.5), ("ALIGN", (1, 0), (-1, -1), "RIGHT"))])
110     s += [tbl, PageBreak()]
111     for p in fs:
112         rel = str(p.relative_to(ROOT)); t = texts[p]
113         h = Paragraph(rel, head); h._file = rel
114         s += [h, Paragraph(f"{t.count(chr(10)) + 1:,.} lines    SHA-256 {hashlib.sha256(p.read_bytes()).
hexdigest()}", meta)]
115         lines = t.split("\n")
116         w = len(str(len(lines)))
117         numbered = "\n".join(f"{i + 1:>{w}} {l}" for i, l in enumerate(lines))
118         s += [Preformatted(numbered, code, maxLineLength=118, newLineChars="    "), PageBreak()]
119     return s
120
121 for _ in range(2):
122     # pass 1 finds each file's page; pass 2 prints them in the contents
123     doc = Doc(str(out), pagesize=letter, leftMargin=0.6 * inch, rightMargin=0.6 * inch, topMargin=0.6 * inch,
bottomMargin=0.6 * inch,
124         title=cfg["title"], author=OWNER, subject=f"source code, version {ver}")
125     doc.addPageTemplates([PageTemplate(frames=[Frame(doc.leftMargin, doc.bottomMargin, doc.width, doc.
height)], onPageEnd=footer)])
126     found = dict(pages); pages.clear(); state["file"] = ""
127     doc.build(story(found))
128     print(f"{out.relative_to(ROOT)}: {len(fs)} files, {total:,.} lines, version {ver[:12]}")
129     return out
130
131 if __name__ == "__main__":
132     ap = argparse.ArgumentParser(description=__doc__, formatter_class=argparse.RawDescriptionHelpFormatter)
133     ap.add_argument("book", choices=sorted(BOOKS))
134     build(ap.parse_args().book)
135

```

tools/strip_reference.py

34 lines SHA-256 1dbab6494b976ec7f934b53fca971691a2f9e494c4bdea4f5609f185a17c

```
1  """Remove comments and docstrings from a Python file; assert the program is unchanged."""
2  import ast, io, sys, tokenize
3  sys.path.insert(0, __import__("os").path.dirname(__file__))
4  from code_fingerprint import fingerprint
5
6
7  def strip(src):
8      tree = ast.parse(src)
9      lines = src.splitlines(keepends=True)
10     edits = []
11     for node in ast.walk(tree):
12         if isinstance(node, (ast.Module, ast.FunctionDef, ast.AsyncFunctionDef, ast.ClassDef)):
13             b = node.body
14             if b and isinstance(b[0], ast.Expr) and isinstance(getattr(b[0], "value", None), ast.Constant) \
15                 and isinstance(b[0].value.value, str):
16                 d = b[0]
17                 edits.append((d.lineno, d.end_lineno, d.col_offset, len(b) == 1))
18     for start, end, col, only in sorted(edits, reverse=True):
19         repl = [" " * col + "pass\n"] if only else []
20         lines[start - 1:end] = repl
21     src = "".join(lines)
22     toks = [t for t in tokenize.generate_tokens(io.StringIO(src).readline) if t.type != tokenize.COMMENT]
23     return tokenize.untokenize(toks)
24
25
26 if __name__ == "__main__":
27     src_path, dst_path = sys.argv[1], sys.argv[2]
28     open(dst_path, "w", encoding="utf-8").write(strip(open(src_path, encoding="utf-8").read()))
29     a, b = fingerprint(src_path), fingerprint(dst_path)
30     print("program fingerprint original:", a)
31     print("program fingerprint stripped:", b)
32     assert a == b, "stripped file is not the same program"
33     print("IDENTICAL PROGRAM")
34
```

tools/three_way.py

93 lines SHA-256 b7b937e09a4039fca94b6a74fba69cdb2068d5ccc5d8b62cbb2cf28ec04950ed

```
1 """Three columns: Kubernetes alone (HPA 70% + Cluster Autoscaler) | Omni-Compass on top of it (B) | Omni-Compass
   alone (C),
2 on fresh scenarios never used before (seeds 713001-713030 per workload), every gauge, paired bootstrap vs
   Kubernetes.
3 Settings are all frozen before this run: B = tuning/B_SETTINGS_TONE.json (Kubernetes column), C = the global
   setting
4 of tuning/GLOBAL_LEAGUE_PREREGISTRATION.json with the nervous-system coordination on (tuning/
   COORD_PREREGISTRATION.json).
5 THEORETICAL SIMULATION (fleet/sim_slo.py). Usage: python tools/three_way.py -> results/THREE_WAY.md, results/
   THREE_WAY.json"""
6 import json, sys
7 from multiprocessing import Pool
8 from pathlib import Path
9 import numpy as np
10 ROOT = Path(__file__).resolve().parents[1]; sys.path.insert(0, str(ROOT))
11 from fleet import sim_slo
12 from fleet.harness import make_scenario
13 from fleet.sim_slo import DirectLaw, BLaw
14 from omnicompass.speed import SpeedLaw
15 from omnicompass.closure import ClosureLaw
16 from tuning.speed_search import gauges
17 from tuning.league import VESSELS, LOWER, HIGHER
18
19 K8S = "k8s_hpa70_ca"
20 C2 = "--c2" in sys.argv
21 SEEDS = list(range(714001, 714031)) if C2 else list(range(713001, 713031))
22 TONE = json.load(open(ROOT / "tuning/B_SETTINGS_TONE.json"))
23 SET = json.load(open(ROOT / ("tuning/C2_PREREGISTRATION.json" if "--c2" in sys.argv else "tuning/
   GLOBAL_LEAGUE_PREREGISTRATION.json")))[("setting")]
24 NAMES = {"energy_kwh": "Energy (kWh)", "node_hours": "Machine-hours", "p95_ms": "Response time p95 (ms)",
25          "p99_ms": "Response time p99 (ms)", "mean_ms": "Response time mean (ms)", "violation_backlog": "Time
   over backlog limit",
26          "violation_power": "Time over power limit", "violation_heat": "Time over heat limit", "start_stop":
   "Machine starts+stops",
27          "node_reversals": "Scale reversals", "pod_changes": "Pod changes", "work_completed": "Work completed",
28          "time_healthy": "Time healthy", "contradictions": "Controllers fighting (per day)"}
29
30
31 def job(a):
32     v, s = a
33     sc = make_scenario(v, s)
34     r0 = sim_slo.run(sc, K8S)
35     rb = sim_slo.run(sc, "omniB:" + K8S, b_law=BLaw(**TONE[v][K8S]))
36     cl = dict(SET["closure"], site=(v == "multi"), coord=True)
37     rc = sim_slo.run(sc, "omni_closure", speed_law=SpeedLaw(**SET["speed"]), omni_every=1,
38                    direct_law=DirectLaw(**SET["direct"]), closure_law=ClosureLaw(**cl))
39     out = []
40     for r in (r0, rb, rc):
41         g = gauges(r); g["contradictions"] = r.get("contradictions", 0); out.append(g)
42     return (v, s), out
43
44
45 def main():
46     with Pool(4) as p:
47         res = dict(p.map(job, [(v, s) for v in VESSELS for s in SEEDS], chunksize=2))
48         rng = np.random.default_rng(20260927)
49         keys = [k for k in NAMES if k in res[(VESSELS[0], SEEDS[0])][0]]
50         out = {"seeds": [SEEDS[0], SEEDS[-1]], "workloads": {}}
51         L = ["# Three columns: Kubernetes alone | Omni-Compass on top | Omni-Compass alone", "",
52             f"Fresh scenarios {SEEDS[0]}-{SEEDS[-1]} (30 per workload, never used before). THEORETICAL SIMULATION (
   fleet/sim_slo.py).",
53             "Settings frozen before the run. Marks: **better** / *worse* = paired 95% interval vs Kubernetes
   excludes 0 "
54             "and the difference is over 0.5%; otherwise equal.", ""]
55         for v in VESSELS:
56             L += [f"## {v}", "", "| Gauge | Kubernetes alone | Omni on top | Omni alone |", "|---|---|---|---|"]
57             out["workloads"][v] = {}
58             for k in keys:
59                 cols = [np.array([res[(v, s)][i][k] for s in SEEDS], float) for i in range(3)]
60                 cells = [f"{cols[i].mean():.4g}"
61                         rec = {"k8s": float(cols[0].mean())}
62                 for i, name in ((1, "on_top"), (2, "alone")):
63                     d = cols[i] - cols[0]
64                     bs = d[rng.integers(0, len(d), (5000, len(d)))].mean(1); lo, hi = np.percentile(bs, [2.5, 97.5])
65                     low_better = k not in HIGHER and k not in ("work_completed", "time_healthy")
66                     pct = 100 * d.mean() / max(abs(cols[0].mean()), 1e-9)
67                     sig = (lo > 0 or hi < 0) and abs(pct) > 0.5
68                     better = (d.mean() < 0) == low_better
69                     mark = ("**better**" if better else "**worse**") if sig else "equal"
70                     cells.append(f"{cols[i].mean():.4g} ({pct:+.0f}%, {mark})")
71                     rec[name] = {"mean": float(cols[i].mean()), "pct": float(pct), "verdict": mark.strip("**")}
72             out["workloads"][v][k] = rec
73             L.append(f"| {NAMES[k]} | " + " | ".join(cells) + " |")
```

```

74         L.append("")
75         tally = {c: {"better": 0, "equal": 0, "worse": 0} for c in ("on_top", "alone")}
76         for v in out["workloads"].values():
77             for rec in v.values():
78                 for c in tally:
79                     tally[c][rec[c]["verdict"]] += 1
80         out["tally"] = tally
81         L += ["## Tally across all workloads and gauges", "", "| | better | equal | worse |", "|---|---:|---:|---:|",
82             f"| Omni on top | {tally['on_top']['better']} | {tally['on_top']['equal']} | {tally['on_top']['worse']} |",
83             f"| Omni alone | {tally['alone']['better']} | {tally['alone']['equal']} | {tally['alone']['worse']} |"]
84         tag = "_C2" if C2 else ""
85         if C2:
86             L[0] += " (Omni alone v2, frozen in tuning/C2_PREREGISTRATION.json)"
87         (ROOT / f"results/THREE_WAY{tag}.md").write_text("\n".join(L) + "\n"); (ROOT / f"results/THREE_WAY{tag}.
88         json").write_text(json.dumps(out, indent=1))
89         print("\n".join(L))
90
91 if __name__ == "__main__":
92     main()
93

```

tools/tower_off_on.py

19 lines SHA-256 ecfb5b9d574bf99798db33f69c3325ab7f9d46fb4682cd35b01eced5

```
1 #!/usr/bin/env python3
2 """Paired same-tower experiment: native stack / Omni observe / Omni authority ON."""
3 import argparse, json, sys
4 from pathlib import Path
5 ROOT=Path(__file__).resolve().parents[1]; sys.path.insert(0,str(ROOT))
6 from benchmarks.stack_benchmark import run
7 from omnicompass.adapter import AllocationLaw
8 KEYS=['availability','time_healthy','recovery_minutes','recovered','energy_kwh','sla_violation_physical',
9       'sla_violation_total','violation_backlog','violation_power','violation_heat','contradictions','peak_power_kw',
10      'node_hours','idle_node_hours','machine_round_trips','scale_reversals']
11 def main():
12     p=argparse.ArgumentParser(); p.add_argument('--runs',type=int,default=100); p.add_argument('--seed',type=int,
13     default=360555127); p.add_argument('--out',default='results/tower_off_on'); a=p.parse_args()
14     out=Path(a.out); out.mkdir(parents=True,exist_ok=True)
15     s=run(a.runs,a.seed,out,AllocationLaw(),arms=['k8s_ref_70','omni_k8s_observe','omni_k8s_throughput'])
16     b=s['means']['k8s_ref_70']; o=s['means']['omni_k8s_observe']; c=s['means']['omni_k8s_throughput']
17     ident={k:abs(float(b[k])-float(o[k])) <= 1e-12 for k in KEYS}
18     receipt={'runs':a.runs,'seed':a.seed,'same_tower_observe_identity':all(ident.values()),'identity_by_metric':
19     ident,'off':{'k:b[k] for k in KEYS},'on':{'k:c[k] for k in KEYS}}
20     (out/'TOWER_OFF_ON_RECEIPT.json').write_text(json.dumps(receipt,indent=2))
21     print(json.dumps(receipt,indent=2))
22 if __name__=='__main__': main()
23
```

tools/tracking_bounds.py

104 lines SHA-256 c5fc54965f14a4757f1c6c47d53f2a1011f8b1660d66e21b7b6c2e81aac6e9e1

```
1  """Computational checks behind docs/TRACKING_THEOREM.md, on the frozen 500-fixture population of
2  benchmarks/core_evidence.py (seed 223387268). No new population is drawn. Evidence class of every number here: V
3  (finite computational verification over these fixtures) – not a theorem, not a physical result.
4
5  python3 tools/tracking_bounds.py      # writes results/TRACKING_BOUNDS.json
6
7  Checks, canonical engine (symmetric_verified), both with the target locked to the nearest basin and aimed wrong:
8  drift bound      max |f_U| along every micro step against the analytic bound F_bar of theorem 2 (must be <
9  U_AUTHORITY)
10 saturation      micro steps with |raw| > U_AUTHORITY (occupancy of the saturated region)
11 sampled data     e_(k+1) = rho e_k + d_k with rho = 1 - h KP (theorem 3): the largest |d_k| (epsilon_h), the
12                  ultimate bound epsilon_h / (1 - rho), and whether |e_(k+1)| <= rho |e_k| + epsilon_h held at
13                  every step
14 invariance       every micro step: x_k in A implies x_(k+1) in A, A the box of theorem 4 for that fixture
15 """
16 from __future__ import annotations
17 import json, math, sys
18 from pathlib import Path
19 ROOT = Path(__file__).resolve().parents[1]
20 sys.path.insert(0, str(ROOT))
21 import numpy as np
22 from omnicompass import core as C
23
24 SEED, RUNS = 223387268, 500
25 # theorem 2 over the whole declared box (PARAMETER_RANGES, E0 in [0, 1]): mu <= 6, |E| <= 7/4.2, |dE/dt| <= 4.2
26 # |E| + 2 + 5,
27 # E_max >= 1, lambda_U <= 0.5, |U| <= 1.1
28 F_GLOBAL = 6 * 2 / (3 * math.sqrt(3)) + (4.2 * 7 / 4.2 + 7) / 1.0 + 0.5 * 1.1
29 OUT = ROOT / "results" / "TRACKING_BOUNDS.json"
30
31 def f_bar(p, E0):
32     """Theorem 2's drift bound on the box |U| <= 1.1, E in the E box of this fixture:
33     |f_U| <= mu * 2/(3 sqrt 3) + max|dE/dt| / E_max + 1.1 lambda_U."""
34     lo, hi = e_box(p, E0)
35     dE = max(abs(-p.alpha_E * lo + p.beta_int + p.beta_ext + p.c), abs(-p.alpha_E * hi + p.beta_int + p.beta_ext
36 - p.c),
37             abs(-p.alpha_E * lo + p.beta_int + p.beta_ext - p.c), abs(-p.alpha_E * hi + p.beta_int + p.beta_ext
38 + p.c))
39     return p.mu * 2.0 / (3.0 * math.sqrt(3.0)) + dE / p.E_max + 1.1 * p.lambda_U
40
41 def e_box(p, E0):
42     """The E interval that equation (1) cannot leave: [(beta - c)/alpha_E, (beta + c)/alpha_E], widened to hold
43     E0."""
44     b = p.beta_int + p.beta_ext
45     lo, hi = (b - p.c) / p.alpha_E, (b + p.c) / p.alpha_E
46     return (lo, hi) if E0 is None else (min(lo, E0), max(hi, E0))
47
48 def run(x0, p, tgt):
49     h = C.MICRO_DT; rho = 1.0 - h * C.KP
50     e_lo, e_hi = e_box(p, x0.E)
51     s_minus, s_plus = C.s_roots(p)
52     s_lo, s_hi = min(x0.S, s_plus), max(x0.S, s_plus)
53     e0 = x0.U - tgt
54     out = {"fU_max": 0.0, "sat": 0, "d_max": 0.0, "steps": 0, "inv_fail": 0, "sign_flip": 0, "e_grew": 0}
55     x = x0
56     for m in range(C.MACRO_STEPS):
57         t = m * C.MACRO_DT
58         for _ in range(C.MICRO_STEPS):
59             fU = C.derivatives(x, p, t, 0.0).U
60             u, raw = C.control_command(x, p, t, tgt)
61             out["fU_max"] = max(out["fU_max"], abs(fU)); out["sat"] += abs(raw) > C.U_AUTHORITY + C.EPS
62             xn = C.rk4_step(x, p, t, h, u)
63             e, en = x.U - tgt, xn.U - tgt
64             out["d_max"] = max(out["d_max"], abs(en - rho * e))
65             out["sign_flip"] += (e != 0 and en != 0 and (e > 0) != (en > 0))
66             out["e_grew"] += abs(en) > abs(e) + 1e-15
67             inside = lambda z: (e_lo - 1e-12 <= z.E <= e_hi + 1e-12 and abs(z.U - tgt) <= abs(e0) + 1e-12
68 and s_lo - 1e-12 <= z.S <= s_hi + 1e-12)
69             out["inv_fail"] += inside(x) and not inside(xn)
70             out["steps"] += 1
71             x = xn; t += h
72     return out
73
74 def main():
75     rng = np.random.default_rng(SEED)
76     pop = [(C.sample_state(rng), C.sample_params(rng)) for _ in range(RUNS)]
77     h = C.MICRO_DT; rho = 1.0 - h * C.KP
```

```

78     res = {"population": f"benchmarks/core_evidence.py, seed {SEED}, {RUNS} fixtures", "evidence_class": "V",
79            "rho": rho, "h": h, "KP": C.KP, "U_AUTHORITY": C.U_AUTHORITY}
80     for case, flip in (("nearest_target", 1), ("wrong_target", -1)):
81         rs, fb = [], []
82         for x0, p in pop:
83             tgt = flip * C.nearest_basin_sign(x0.U)
84             rs.append(run(x0, p, tgt)); fb.append(f_bar(p, x0.E))
85         eps = max(r["d_max"] for r in rs)
86         res[case] = {
87             "drift_max_observed": max(r["fU_max"] for r in rs),
88             "drift_bound_F_bar_max": max(fb), "F_bar_global": F_GLOBAL,
89             "drift_within_bound_all": all(r["fU_max"] <= b + 1e-9 for r, b in zip(rs, fb)),
90             "F_bar_below_U_AUTHORITY_all": all(b < C.U_AUTHORITY for b in fb),
91             "saturated_microsteps": sum(r["sat"] for r in rs), "microsteps": sum(r["steps"] for r in rs),
92             "epsilon_h": eps, "ultimate_bound": eps / (1.0 - rho), "basin_tol": C.BASIN_TOL,
93             "ultimate_bound_inside_basin": eps / (1.0 - rho) < C.BASIN_TOL,
94             "error_sign_changes": sum(r["sign_flip"] for r in rs), "error_growth_steps": sum(r["e_grew"] for r
in rs),
95             "invariance_failures": sum(r["inv_fail"] for r in rs)}
96     OUT.write_text(json.dumps(res, indent=1) + "\n")
97     if __name__ == "__main__":
98         print(json.dumps(res, indent=1))
99     return res
100
101
102 if __name__ == "__main__":
103     main()
104

```

tools/wall_meter.py

66 lines SHA-256 27114160a1ffa8adfb7ca2fa88c92a02779127cd88082eb876e2ed94af457fe

```
1  """Whole-machine watts from the wall, read from a smart plug that Omni-Compass never reads or writes.
2
3  The plug sits between the wall socket and the machine under test. Its reading covers everything the machine draws:
4  GPU, CPU, memory, fans, power-supply losses. It is the measurement a buyer understands at once, and it is
   independent
5  of Omni by construction: the governor is never given the plug's address.
6
7  SPEC
8  shelly1:<ip>      Shelly Gen1 plug (Plug S, Plug US):  http://<ip>/meter/0      -> "power"
9  shelly2:<ip>      Shelly Gen2/Gen3 plug (Plus Plug):  http://<ip>/rpc/Switch.GetStatus?id=0 -> "apower"
10  tasmota:<ip>      any plug flashed with Tasmota:      http://<ip>/cm?cmd=Status%208 -> StatusSNS.ENERGY.Power
11  cmd:<command>     any shell command that prints watts (a USB meter, a PDU script, IPMI, a lab instrument)
12
13  Usage
14  python tools/wall_meter.py SPEC --once                print one reading (check the plug before a run)
15  python tools/wall_meter.py SPEC OUT.csv [--interval 1] write epoch_s,watts until stopped
16  """
17
18  from __future__ import annotations
19
20
21  import argparse, json, shlex, signal, subprocess, sys, time, urllib.request
22
23  def read(spec, timeout=3.0):
24      kind, _, where = spec.partition(":")
25      if kind == "cmd":
26          out = subprocess.run(where, shell=True, capture_output=True, text=True, timeout=timeout + 5).stdout
27          return float(out.split()[0])
28      url = {"shelly1": f"http://{where}/meter/0", "shelly2": f"http://{where}/rpc/Switch.GetStatus?id=0",
29            "tasmota": f"http://{where}/cm?cmd=Status%208"}[kind]
30      # a plug on the home network: never through a proxy
31      opener = urllib.request.build_opener(urllib.request.ProxyHandler({}))
32      d = json.loads(opener.open(url, timeout=timeout).read())
33      if kind == "shelly1":
34          return float(d["power"])
35      if kind == "shelly2":
36          return float(d["apower"])
37      return float(d["StatusSNS"]["ENERGY"]["Power"])
38
39  def main(argv=None):
40      ap = argparse.ArgumentParser(description=__doc__, formatter_class=argparse.RawDescriptionHelpFormatter)
41      ap.add_argument("spec")
42      ap.add_argument("out", nargs="?")
43      ap.add_argument("--once", action="store_true")
44      ap.add_argument("--interval", type=float, default=1.0)
45      a = ap.parse_args(argv)
46      if a.once or not a.out:
47          print(f"{read(a.spec):.2f}")
48          return 0
49      stop = {"now": False}
50      signal.signal(signal.SIGTERM, lambda *_: stop.__setitem__("now", True))
51      with open(a.out, "w") as f:
52          f.write("epoch_s,watts\n"); f.flush()
53          while not stop["now"]:
54              t = time.time()
55              try:
56                  f.write(f"{t:.3f},{read(a.spec):.3f}\n")
57              except Exception as e: # noqa: BLE001 a missed reading is a gap, recorded as such, never a guess
58                  f.write(f"{t:.3f},\n")
59                  print(f"wall meter: {type(e).__name__}: {e}", file=sys.stderr)
60                  f.flush()
61                  time.sleep(max(0.0, a.interval - (time.time() - t)))
62      return 0
63
64  if __name__ == "__main__":
65      sys.exit(main())
66
```

benchmarks/core_evidence.py

63 lines SHA-256 5013dcc81db1c3eb6b05043675ea8e5bc3a670eb07daf7bbfc49a747d4b7ee3c

```
1  """Canonical 500-run population and counterfactual variants.
2
3  Variants on the identical population (seed 223387268):
4  controlled      canonical: target = nearest basin of U0
5  wrong_target    controller aimed at the opposite basin
6  uncontrolled    u = 0 throughout
7  controlled_mu0  canonical controller, double-well coefficient mu set to 0
8  Also: S-channel admissibility S0 >= S_minus(p), bath residual RMS, RK4 step refinement.
9  """
10 import json, sys, math
11 from dataclasses import replace
12 from pathlib import Path
13 import numpy as np
14 ROOT = Path(__file__).resolve().parents[1]; sys.path.insert(0, str(ROOT))
15 from omnicompass import core as C
16
17 def main(out=ROOT / "results" / "core_evidence.json", seed=223387268, runs=500):
18     rng = np.random.default_rng(seed)
19     pop = [(C.sample_state(rng), C.sample_params(rng)) for _ in range(runs)]
20     res = {}
21     for name in ("controlled", "wrong_target", "uncontrolled", "controlled_mu0"):
22         cv = ce = 0; peak = 0.0; sat = 0; confirm = []
23         for x0, p in pop:
24             pp = replace(p, mu=0.0) if name == "controlled_mu0" else p
25             tgt = -C.nearest_basin_sign(x0.U) if name == "wrong_target" else None
26             r = C.simulate(x0, pp, target=tgt, controlled=(name != "uncontrolled"))
27             cv += r.convey; ce += r.cert; peak = max(peak, r.actuator_peak); sat += r.actuator_saturated_periods
28             if r.convey: confirm.append(r.convey_confirm)
29         res[name] = {"convey5": cv, "cert10": ce, "runs": runs, "actuator_peak": peak, "saturated_microsteps":
30 sat,
31 "convey_confirm_step_counts": {str(k): int(v) for k, v in zip(*np.unique(confirm,
32 return_counts=True))}}
33 # closed-form prediction under exact cancellation: U_k = s + (U0 - s) exp(-KP * 0.1 k)
34 hit = 0
35 for x0, p in pop:
36     s = C.nearest_basin_sign(x0.U); r = C.simulate(x0, p)
37     Uk = [s + (x0.U - s) * math.exp(-C.KP * C.MACRO_DT * k) for k in range(C.MACRO_STEPS + 1)]
38     first = next(k for k, u in enumerate(Uk) if abs(u - s) <= C.BASIN_TOL)
39     hit += int((4 if first == 0 else first + 4) == r.convey_confirm)
40 res["first_order_prediction_of_convey_step"] = {"matched": hit, "runs": runs}
41 adm = [x0.S >= C.s_roots(p)[0] for x0, p in pop]
42 res["s_admissibility"] = {"pass": int(sum(adm)), "runs": runs,
43 "min_margin": float(min(x0.S - C.s_roots(p)[0] for x0, p in pop))}
44
45 rms = []
46 for x0, p in pop[:100]:
47     xs = [x0]; x = x0
48     for m in range(200):
49         u, _ = C.control_command(x, p, m * C.MICRO_DT, C.nearest_basin_sign(x0.U))
50         x = C.rk4_step(x, p, m * C.MICRO_DT, C.MICRO_DT, u); xs.append(x)
51         r = [C.bath_residual(xs[i - 1], xs[i], xs[i + 1], p, C.MICRO_DT) for i in range(1, len(xs) - 1)]
52         rms.append(math.sqrt(sum(v * v for v in r) / len(r)))
53 res["bath_residual_rms"] = {"median": float(np.median(rms)), "max": float(max(rms)), "runs": 100}
54 x0, p = pop[0]
55 def final(h):
56     x = x0; n = int(round(2.0 / h))
57     for m in range(n):
58         u, _ = C.control_command(x, p, m * h, C.nearest_basin_sign(x0.U)); x = C.rk4_step(x, p, m * h, h, u)
59     return np.array(x.as_tuple())
60 ref = final(0.000625)
61 res["rk4_refinement_run0"] = {str(h): float(np.max(np.abs(final(h) - ref))) for h in (0.02, 0.01, 0.005, 0.
62 0025)}
63 Path(out).write_text(json.dumps(res, indent=2)); print(json.dumps(res, indent=1))
64
65 if __name__ == "__main__":
66     main()
67
```

benchmarks/fleet_overhead.py

23 lines SHA-256 232accd3a645ba881ba71916e13ddb2c882069cbefdd10dbd425c3b42736ceaf

```
1  """Fleet overhead of infrastructure components under the keep-or-remove rule."""
2  import json, sys
3  from pathlib import Path
4  ROOT = Path(__file__).resolve().parents[1]; sys.path.insert(0, str(ROOT))
5  from omnicompass.muscles import Fleet, fleet_report
6
7  def main(out=ROOT / "results" / "fleet_overhead.json"):
8      res = {}
9      for util in (0.08, 0.30):
10         for mesh in ("sidecar", "ambient"):
11             res[f"util{int(util*100)}_{mesh}"] = fleet_report(Fleet(mean_cpu_utilization=util, mesh=mesh))
12             Path(out).write_text(json.dumps(res, indent=2))
13             r = res[f"util{int(util*100)}_{mesh}"]
14             print(f"fleet: {r['fleet']['nodes']:,} nodes, {r['clusters']} clusters, {r['fleet_vcpu']:,} vCPU, {r['fleet_mwh_per_year']:,} MWh/yr (util {int(util*100)}%)")
15             for c in r["components"]:
16                 print(f"  {c['name'][:46]:46s} {c['role']:13s} {c['status']:10s} vCPU {c['vcpu']:,} ({100*c['share_of_fleet_vcpu']:5.2f}%) energy {100*c['share_of_fleet_energy_low']:5.2f}-{100*c['share_of_fleet_energy_high']:5.2f}%)")
17                 for k, v in res.items():
18                     print(k, {role: f"{100*g['share_of_fleet_vcpu']:.2f}% vCPU, {100*g['share_of_fleet_energy_low']:.2f}-{100*g['share_of_fleet_energy_high']:.2f}% energy" for role, g in v["by_role"].items()},
19                           f"| idle vCPU {100*v['idle_share_of_fleet_vcpu']:.0f}% of fleet, idle draw {100*v['idle_vcpu_share_of_fleet_energy']:.0f}% of energy")
20
21 if __name__ == "__main__":
22     main()
23
```

benchmarks/multiplicity.py

65 lines SHA-256 5ef32b53e6f36f97d09d56783c32cc8b6b704804e569fcf1747b0259d50676e4

```
1  """Primary-endpoint and multiplicity analysis of the held-out comparison.
2
3  Primary comparison: omni_k8s_throughput vs k8s_ref_70. Primary metric: energy_kwh.
4  Secondary metrics are tested with a paired sign-flip permutation test (two-sided, 10,000 permutations per test,
5  seeded) and Holm-Bonferroni adjusted across all secondary metrics and both held-out seeds together.
6  """
7  import csv, json, sys
8  from pathlib import Path
9  import numpy as np
10
11  ROOT = Path(__file__).resolve().parents[1]
12  PRIMARY = "energy_kwh"
13  SECONDARY = [{"time_healthy", False}, ("recovery_minutes", True), ("recovered", False), ("sla_violation_physical",
14      True),
15      ("violation_backlog", True), ("violation_power", True), ("violation_heat", True), ("availability",
16      False),
17      ("invariant_violations_ex_power", True), ("scale_reversals", True), ("machine_round_trips", True),
18      ("thermal_travel", True)]
19  BASE, CAND, PERMS = "k8s_ref_70", "omni_k8s_throughput", 10000
20
21  def diffs(sd, metric):
22      rows = list(csv.DictReader(open(ROOT / "results" / f"heldout_seed_{sd}" / "RUNS.csv")))
23      a = {r["scenario_id"]: float(r[metric]) for r in rows if r["arm"] == BASE}
24      b = {r["scenario_id"]: float(r[metric]) for r in rows if r["arm"] == CAND}
25      return np.array([b[k] - a[k] for k in sorted(a)])
26
27  def perm_p(d, rng):
28      obs = abs(d.mean())
29      if np.all(d == 0):
30          return 1.0
31      signs = rng.choice([-1.0, 1.0], size=(PERMS, len(d)))
32      null = np.abs((signs * d).mean(axis=1))
33      return float((1 + np.sum(null >= obs - 1e-15)) / (PERMS + 1))
34
35  def main(out=ROOT / "results" / "MULTIPLICITY.json"):
36      seeds = json.loads((ROOT / "results" / "PREREGISTRATION.json").read_text())["held_out_seeds"]
37      rng = np.random.default_rng(20260924)
38      res = {"primary": {}, "secondary": [], "method": __doc__.strip()}
39      for sd in seeds:
40          d = diffs(sd, PRIMARY)
41          res["primary"][str(sd)] = {"metric": PRIMARY, "mean_delta": float(d.mean()), "p": perm_p(d, rng),
42              "better": int((d < 0).sum()), "n": len(d)}
43          tests = []
44          for sd in seeds:
45              for m, lower in SECONDARY:
46                  d = diffs(sd, m)
47                  tests.append({"seed": sd, "metric": m, "mean_delta": float(d.mean()), "p": perm_p(d, rng),
48                      "direction": ("better" if (d.mean() < 0) == lower else "worse") if d.mean() != 0 else
49                      "none"})
50          order = sorted(range(len(tests)), key=lambda i: tests[i]["p"])
51          k = len(tests); running = 0.0
52          for rank, i in enumerate(order):
53              running = max(running, min(1.0, (k - rank) * tests[i]["p"]))
54              tests[i]["p_holm"] = running
55              tests[i]["verdict"] = tests[i]["direction"] if running < 0.05 else "not significant"
56          res["secondary"] = tests
57          Path(out).write_text(json.dumps(res, indent=2))
58          for sd, r in res["primary"].items():
59              print(f"primary seed {sd}: energy delta {r['mean_delta']:+.2f} kWh, p = {r['p']:.1e}, better in {r[
60                  'better']}/{r['n']}")
61          for t in tests:
62              print(f"  {t['seed']} {t['metric']:30s} {t['mean_delta']:+9.4f}  p_holm {t['p_holm']:.1e}  {t[
63                  'verdict']}")
64
65  if __name__ == "__main__":
66      main()
```

benchmarks/savings.py

61 lines SHA-256 f6bb43b9d6bb2638a001731cfaac4274f275ea7cf5bf2fe32c7ab70b54dca93a

```
1  """Savings projection: fractional energy reductions measured in the fleet benchmark and the recorded-trace run,
2  applied to declared fleet profiles. The reductions come from simulation; the profiles are declared examples. The
   output
3  is a projection of annual energy, cost and CO2 differences, not a measurement.
4
5  Reductions are taken against HPA 0.7 + Karpenter-lite (the stronger baseline) for the energy-first setting, with
   low /
6  high bounds from the paired 95% interval of the energy difference.
7  """
8  import csv, json, subprocess, sys
9  from pathlib import Path
10 R00T = Path(__file__).resolve().parents[1]
11 PROFILES = [
12     # name, nodes, average kW per node at the wall before PUE, PUE, USD per kWh, kg CO2 per kWh
13     ("cluster_1k_nodes", 1000, 0.40, 1.40, 0.12, 0.40),
14     ("fleet_100k_nodes", 100000, 0.40, 1.30, 0.08, 0.40),
15 ]
16
17
18 GPU_NODE_KW = 6.0 # declared average draw of an 8-GPU node before PUE, applied to the gpu vessels
19
20
21 def reductions():
22     out = []
23     fs = json.loads((R00T / "results" / "fleet" / "heldout" / "SUMMARY.json").read_text())
24     for v, d in fs["vessels"].items():
25         p = {x["metric"]: x for x in d["paired"] ["omni_fleet_vs_k8s_hpa70_karpenter"]} ["energy_kwh"]
26         b = p["base"]
27         out.append((v, "HPA 0.7 + Karpenter-lite", -p["ci95"][1] / b, -p["delta"] / b, -p["ci95"][0] / b))
28     pl = json.loads((R00T / "results" / "fleet" / "planetlab" / "SUMMARY.json").read_text())
29     p = {x["metric"]: x for x in pl["paired"] ["omni_fleet_vs_k8s_hpa70_karpenter"]} ["energy_kwh"]
30     b = pl["means"] ["k8s_hpa70_karpenter"] ["energy_kwh"]
31     out.append(("planetlab_shapes", "HPA 0.7 + Karpenter-lite", -p["ci95"][1] / b, -p["delta"] / b, -p["ci95"][0]
   / b))
32     return out
33
34
35 def python_reference(row):
36     name, v, base, nodes, kw, pue, usd, co2, lo, mid, hi = row
37     mwh = nodes * kw * pue * 8760.0 / 1000.0
38     return [mwh] + [mwh * r for r in (lo, mid, hi)] + [mwh * 1000.0 * usd * r for r in (lo, mid, hi)] + [mwh *
   co2 * r for r in (lo, mid, hi)]
39
40
41 def main(exe=None, out_dir=R00T / "results"):
42     rows = [(pn, v, base, n, (GPU_NODE_KW if v.startswith("gpu") else kw), pue, usd, co2, lo, mid, hi)
   for (pn, v, base, n, kw, pue, usd, co2) in PROFILES for (v, base, lo, mid, hi) in reductions()]
43     inp = Path(out_dir) / "SAVINGS_INPUTS.csv"
44     with open(inp, "w", newline="") as f:
45         w = csv.writer(f); w.writerow(["profile", "vessel", "baseline", "nodes", "node_avg_kw", "pue",
   "usd_per_kwh", "kg_co2_per_kwh", "reduction_low", "reduction_mid", "reduction_high"])
46         for r in rows: w.writerow([r[0], r[1], r[2]] + [repr(float(x)) for x in r[3:]])
47     exe = exe or str(R00T / "cpp" / "build" / "oc_savings")
48     subprocess.run([exe, str(inp), str(Path(out_dir) / "SAVINGS.csv")], check=True)
49     got = list(csv.DictReader(open(Path(out_dir) / "SAVINGS.csv")))
50     keys = ["baseline_mwh_per_year", "saved_mwh_low", "saved_mwh_mid", "saved_mwh_high", "saved_usd_low",
   "saved_usd_mid", "saved_usd_high", "saved_tco2_low", "saved_tco2_mid", "saved_tco2_high"]
51     worst = max(abs(float(g[k]) - ref) / max(1.0, abs(ref)) for g, r in zip(got, rows) for k, ref in zip(keys,
   python_reference(r)))
52     assert worst < 1e-9, worst
53     for g in got:
54         print(f"{g['profile']}:17s {g['vessel']}:16s saved {float(g['saved_mwh_mid']):>12,.0f} MWh/yr  ${float(g[
   'saved_usd_low']):>14,.0f} to ${float(g['saved_usd_high']):>14,.0f}  {float(g['saved_tco2_mid']):>10,.0f} tCO2")
55     return got
56
57
58
59 if __name__ == "__main__":
60     main(sys.argv[1] if len(sys.argv) > 1 else None)
61
```

benchmarks/stack_benchmark.py

424 lines SHA-256 012cce208ce893f542e55cae70b01a3095bfc5f3fada5c732e6b8ae79376edbb

```
1  """Stack governance benchmark.
2
3  Arms:
4      native          each manager executes its own proposal; a paged human intervenes after the operator delay
5      omni_observe     governor in OBSERVE
6      omni_autopilot   governor in AUTOPILOT
7      omni_no_dynamics AUTOPILOT with equations (1)-(7) not evolved
8      omni_kill        AUTOPILOT, killed at event onset + 2 intervals
9      omni_switch_on   OBSERVE until event onset, then AUTOPILOT
10 """
11 from __future__ import annotations
12
13 import argparse, csv, hashlib, json, math, sys, time
14 from dataclasses import asdict
15 from pathlib import Path
16 from typing import Any, Dict, List
17
18 import numpy as np
19
20 ROOT = Path(__file__).resolve().parents[1]
21 sys.path.insert(0, str(ROOT))
22 from omnicompass import stack_sim as S
23 from omnicompass.shield import enforce, violations, ShieldLimits
24 from omnicompass.adapter import mode_law
25 from omnicompass.adapter import (Governor, AllocationLaw, OBSERVE, AUTOPILOT, directive_to_actions,
26                                 observe_vector, assimilate)
27
28 ARMS = ["native", "k8s_ref_50", "k8s_ref_60", "k8s_ref_70", "k8s_ref_80", "omni_k8s_protect",
29         "omni_k8s_throughput",
30         "omni_k8s_observe", "omni_k8s_kill", "omni_k8s_throughput_no_dynamics", "omni_k8s_throughput_no_gate",
31         "omni_observe", "omni_pipeline", "omni_direct", "omni_direct_no_dynamics",
32         "omni_direct_no_shield", "omni_kill", "omni_switch_on"]
33 K8S_REFS = ("k8s_ref_50", "k8s_ref_60", "k8s_ref_70", "k8s_ref_80")
34 LAYERED = {"omni_k8s_protect", "omni_k8s_throughput", "omni_k8s_observe", "omni_k8s_kill",
35            "omni_k8s_throughput_no_dynamics",
36            "omni_k8s_throughput_no_gate"}
37 THROUGHPUT = {"omni_k8s_throughput", "omni_k8s_throughput_no_dynamics", "omni_k8s_throughput_no_gate"}
38 BASELINES = ("native",) + K8S_REFS
39 DIRECT_SENSING = {*K8S_REFS, *LAYERED, "omni_direct", "omni_direct_no_dynamics", "omni_direct_no_shield",
40                  "omni_kill",
41                  "omni_switch_on"}
42 DIRECT_ACTUATION = {*LAYERED, "omni_direct", "omni_direct_no_dynamics", "omni_direct_no_shield", "omni_kill",
43                    "omni_switch_on"}
44 SHIELDED = {*LAYERED, "omni_pipeline", "omni_direct", "omni_direct_no_dynamics", "omni_kill", "omni_switch_on"}
45
46 class K8sReference:
47     """Documented-behaviour reference model of Kubernetes autoscaling (not the upstream controllers).
48
49     HPA (kubernetes.io, Horizontal Pod Autoscaling): desired = ceil(current * metric / target); no change while
50     |metric/target - 1| <= 0.1; scale-down uses the highest recommendation over the 300 s stabilization window.
51     metric = load ratio; the target is workload configuration, so the benchmark runs targets 0.5, 0.6, 0.7, 0.
52
53     8.
54     Cluster Autoscaler (kubernetes/autoscaler FAQ), simplified: scale-up when work is unschedulable (backlog
55     present),
56     sized to the pending work, without scheduling-feasibility simulation; a node is removed after it has been
57     unneeded (utilization < 0.5) for 10 minutes, not within
58     10 minutes of a scale-up.
59     Alerting: page when queue ratio > 0.35 for 15 minutes (Prometheus 'for: 15m' rule).
60     Terraform is not a live capacity controller (desired size ignored after creation).
61     Power/thermal, rollout, rollback and routing proposals are executed as proposed.
62     Not represented: Karpenter consolidation, VPA, pod scheduling constraints, disruption budgets.
63 """
64
65 def __init__(self, hpa_target: float):
66     self.hpa_target = hpa_target
67     self.rec_hist = []
68     self.under = 0
69     self.since_add = 99
70     self.alert = 0
71
72 def actions(self, props, st, obs, cfg):
73     acts = []
74     load = float(obs.get("load_ratio", 0.0)); q = float(obs.get("queue_ratio", 0.0))
75     ratio = load / self.hpa_target
76     cur = int(st["replicas"])
77     rec = cur if abs(ratio - 1.0) <= 0.1 else max(20, min(500, int(np.ceil(cur * ratio))))
78     self.rec_hist = (self.rec_hist + [rec])[-2:]
79     want = rec if rec >= cur else max(self.rec_hist)
80     if want != cur:
81         acts.append({"manager": "HPA", "action": "replicas", "target": want, "direction": 1 if want > cur
82                     else -1})
83     n = int(st["actual_nodes"])
```

```

76         util = min(1.0, load)
77         self.since_add += 1
78         if q > 0.0:
79             need = int(np.ceil(q * n))
80             if need > 0 and n < cfg.maximum_nodes:
81                 acts.append({"manager": "Cluster Autoscaler", "action": "nodes", "target": min(cfg.maximum_nodes,
n + need), "direction": 1})
82                 self.since_add = 0
83                 self.under = 0
84             else:
85                 self.under = self.under + 1 if util < 0.5 else 0
86                 if self.under >= 2 and self.since_add >= 2 and n > cfg.minimum_nodes:
87                     acts.append({"manager": "Cluster Autoscaler", "action": "nodes", "target": n - 1, "direction": -
1})
88                 self.under = 0
89             for p in native_actions(props):
90                 if p["action"] in ("power_cap", "rollback", "route_shift", "rollout"):
91                     acts.append(p)
92                 if any(p["action"] == "alert" for p in props):
93                     acts.append({"manager": "Alertmanager", "action": "alert", "target": 1, "direction": 1})
94                 self.alert = self.alert + 1 if q > 0.35 else 0
95                 if self.alert >= 3:
96                     acts.append({"manager": "Alertmanager", "action": "page", "target": 1, "direction": 1})
97                 self.alert = 0
98             return acts
99         clamp = S._mom_clamp
100
101
102 def native_actions(proposals: List[Dict[str, Any]]) -> List[Dict[str, Any]]:
103     """Native stack actions."""
104     acts = [p for k in ("replicas", "nodes", "terraform_plan", "rollout") for p in proposals if p["action"] == k]
105     if any(p["action"] == "rollback" for p in proposals):
106         acts.append({"manager": "Argo CD / Rollouts", "action": "rollback", "target": 1.0, "direction": -1})
107     caps = [float(p["target"]) for p in proposals if p["action"] == "power_cap"]
108     if caps:
109         acts.append({"manager": "Power/Thermal", "action": "power_cap", "target": min(caps), "direction": -1})
110     if any(p["action"] == "route_shift" for p in proposals):
111         acts.append({"manager": "Istio/Cilium", "action": "route_shift", "target": 1.0, "direction": 1})
112     acts += [p for p in proposals if p["action"] == "alert"] + [p for p in proposals if p["action"] == "page"]
113     return acts
114
115
116 def layer_over_k8s(omni_acts, d, k8s, props, st, obs, cfg):
117     """Omni-Compass governing Kubernetes: Kubernetes autoscalers run as the execution layer; the governor sets
the
118     HPA target (rho*), keeps Cluster Autoscaler scale-up, owns scale-down, power caps, rollback and routing."""
119     k8s.hpa_target = float(d["demand"])
120     k_acts = k8s.actions(props, st, obs, cfg)
121     n = int(st["actual_nodes"])
122     k_up = max([int(a["target"]) for a in k_acts if a["action"] == "nodes" and a["direction"] > 0], default=n)
123     o_nodes = [a for a in omni_acts if a["action"] == "nodes"]
124     o_tgt = int(o_nodes[0]["target"]) if o_nodes else n
125     tgt = max(k_up, o_tgt) if (k_up > n or o_tgt > n) else o_tgt
126     out = [a for a in omni_acts if a["action"] not in ("nodes", "replicas")]
127     if tgt != n:
128         out.append({"manager": "Omni-Compass/Kubernetes", "action": "nodes", "target": tgt, "direction": 1 if
tgt > n else -1})
129     for a in k_acts:
130         if a["action"] == "replicas" and (a["direction"] > 0 or tgt < n):
131             out.append(a)
132     if any(a["direction"] > 0 and a["action"] in ("nodes", "replicas") for a in out):
133         out = [a for a in out if not (a["action"] == "power_cap" and float(a["target"]) < float(st["power_cap"])
- 1e-9)]
134     return out
135
136
137 def contradictions(actions, state, security_denied) -> int:
138     c = 0
139     nd = [a["direction"] for a in actions if a["action"] in ("nodes", "terraform_plan")]
140     if nd and min(nd) < 0 < max(nd):
141         c += 1
142     rd = [a["direction"] for a in actions if a["action"] == "replicas"]
143     if rd and min(rd) < 0 < max(rd):
144         c += 1
145     kinds = {a["action"] for a in actions}
146     if "rollout" in kinds and "rollback" in kinds:
147         c += 1
148     up = any(a["direction"] > 0 and a["action"] in ("nodes", "terraform_plan", "replicas", "rollout") for a in
actions)
149     tighten = any(a["action"] == "power_cap" and float(a["target"]) < float(state["power_cap"]) - 1e-9 for a in
actions)
150     if up and tighten:
151         c += 1
152     if security_denied and any(a["direction"] > 0 and a["action"] in ("nodes", "terraform_plan", "rollout") for
a in actions):
153         c += 1

```

```

154     if sum(1 for a in actions if a["action"] == "alert") > 1:
155         c += 1
156     return c
157
158
159 def simulate(scn, arm: str, cfg, law: AllocationLaw) -> Dict[str, Any]:
160     st = {"actual_nodes": cfg.initial_nodes, "available_nodes": cfg.initial_nodes,
161          "terraform_desired_nodes": cfg.initial_nodes, "replicas": 100, "pending_queue": 0.0, "thermal": 0.32,
162          "smoothed_load": scn.base_load_fraction, "power_cap": 1.0, "network_recovery": 0.0,
163          "terraform_pending_target": None, "terraform_apply_step": None, "last_node_direction": 0,
164          "last_replica_direction": 0, "operator_page_step": None, "operator_ack_step": None}
165     gov = None
166     if arm in THROUGHPUT:
167         law = mode_law("throughput", law)
168     if arm == "omni_k8s_throughput_no_gate":
169         from dataclasses import replace as _replace
170         law = _replace(law, push_release=9.0)
171     shield_lim = ShieldLimits(power_limit=1e9) if arm in THROUGHPUT else ShieldLimits()
172     k8s = K8sReference(int(arm[-2:]) / 100.0) if arm in K8S_REFS else (K8sReference(0.7) if arm in LAYERED else
173     None)
174     shield_hits = inv_viol = inv_viol_ex_power = 0
175     node_travel = cap_travel = thermal_travel = 0.0
176     node_starts = node_stops = 0
177     if arm not in BASELINES:
178         gov = Governor(law=law, evolve=not arm.endswith("_no_dynamics"))
179         gov.set_mode(OBSERVE if arm in ("omni_observe", "omni_switch_on", "omni_k8s_observe") else AUTOPILOT)
180     hist, qv, rec = [], [], []
181     dem = done = energy = nh = unn = peak = 0.0
182     sla_p = sla_c = adm_n = contra = rev = secv = pages = manual = nact = 0
183     v_queue = v_power = v_heat = 0
184     rec_pending, rec_start, streak = False, None, 0
185     authority_steps = 0
186     trace = []
187     for step in range(cfg.steps):
188         if arm in ("omni_kill", "omni_k8s_kill") and step == scn.event_start + 2:
189             gov.kill()
190         if arm == "omni_switch_on" and step == scn.event_start:
191             gov.set_mode(AUTOPILOT)
192         rng = np.random.default_rng(scn.seed + step * 7919 + 101)
193         src = S._mom_source_signals(scn, step, cfg, rng)
194         if st["terraform_apply_step"] is not None and step >= st["terraform_apply_step"]:
195             st["actual_nodes"] = st["terraform_desired_nodes"] = int(st["terraform_pending_target"])
196             st["terraform_pending_target"] = st["terraform_apply_step"] = None
197             failed = int(round(cfg.initial_nodes * src["failed_fraction"]))
198             st["available_nodes"] = max(cfg.minimum_nodes, st["actual_nodes"] - failed)
199             demand = (src["demand_fraction"] + src["batch_fraction"]) * cfg.initial_nodes * cfg.
200             capacity_units_per_node * src["rollout_multiplier"]
201             dem += demand
202             st["pending_queue"] += demand
203             eff_net = clamp(src["network_efficiency"] + st["network_recovery"], 0.45, 1.0)
204             throttle = clamp(1.0 - max(0.0, st["thermal"] - 0.82) * 0.65, 0.58, 1.0)
205             cap = st["available_nodes"] * cfg.capacity_units_per_node * eff_net * throttle * st["power_cap"]
206             cap *= clamp(st["replicas"] / 100.0, 0.55, 1.35)
207             work = min(st["pending_queue"], max(0.0, cap))
208             st["pending_queue"] -= work
209             done += work
210             util = clamp(work / max(1.0, st["available_nodes"] * cfg.capacity_units_per_node), 0.0, 1.25)
211             pkw = st["actual_nodes"] * (cfg.node_idle_kw + cfg.node_dynamic_kw * min(1.0, util)) * cfg.pue * st[
212             "power_cap"]
213             pstress = pkw / max(1e-9, cfg.site_power_limit_kw * src["power_derate"])
214             st["thermal"] = clamp(0.86 * st["thermal"] + 0.14 * (0.34 + 0.62 * min(1.35, pstress)), 0.0, 1.35)
215             qratio = clamp(st["pending_queue"] / max(1.0, cap * 2.0), 0.0, 2.0)
216             lratio = clamp(demand / max(1.0, cap), 0.0, 2.0)
217             st["smoothed_load"] = 0.78 * st["smoothed_load"] + 0.22 * lratio
218             drift = abs(st["terraform_desired_nodes"] - st["actual_nodes"]) / max(1.0, cfg.maximum_nodes - cfg.
219             minimum_nodes)
220             hist.append({"queue_ratio": qratio, "load_ratio": lratio, "thermal": st["thermal"], "power_stress":
221             pstress,
222             "network_stress": clamp(1 - eff_net, 0, 1), "drift_ratio": drift,
223             "batch_fraction": src["batch_fraction"], "security_block": src["security_block"],
224             "rollout_multiplier": src["rollout_multiplier"], "stale": 0.0})
225         if arm in DIRECT_SENSING:
226             obs = S._mom_delayed_observation(hist, 0)
227         else:
228             delay = scn.telemetry_delay_steps + scn.kafka_lag_steps
229             obs = S._mom_delayed_observation(hist, delay)
230             if delay > 0:
231                 obs["stale"] = 1.0
232             if src.get("telemetry_drop", 0.0) > 0.5:
233                 obs = S._mom_delayed_observation(hist, delay + 2)
234                 obs["stale"] = 1.0
235             props = S._mom_proposals(st, obs, scn, step, cfg, src)
236             nconf, _ = S._mom_conflict_count(props)
237             denied = any(p["action"] == "deny_change" for p in props)
238             if gov is not None:
239                 gov.current_cap = st["power_cap"]

```

```

235     gov.nodes = st["actual_nodes"]
236     d = gov.step(obs, nconf)
237     if gov.has_authority:
238         acts = directive_to_actions(d, st, props, cfg, direct_actuation=arm in DIRECT_ACTUATION)
239         if arm in LAYERED:
240             acts = layer_over_k8s(acts, d, k8s, props, st, obs, cfg)
241         if arm in SHIELDED:
242             acts, h = enforce(acts, st, obs, cfg, shield_lim)
243             shield_hits += h
244             authority_steps += 1
245         elif arm in LAYERED:
246             k8s.hpa_target = 0.70
247             acts = k8s.actions(props, st, obs, cfg)
248         else:
249             acts = native_actions(props)
250     elif k8s is not None:
251         acts = k8s.actions(props, st, obs, cfg)
252     else:
253         acts = native_actions(props)
254     vv = violations(acts, st, hist[-1], cfg)
255     inv_viol += len(vv)
256     inv_viol_ex_power += sum(1 for v in vv if v != "I4")
257     n_before, cap_before = st["actual_nodes"], st["power_cap"]
258     c_step = contradictions(acts, st, denied)
259     contra += c_step
260     nact += len(acts)
261     nd = rd = pg = 0
262     st["network_recovery"] *= 0.55
263     for a in acts:
264         k = a["action"]
265         if k == "replicas":
266             t = int(a["target"]); rd = 1 if t > st["replicas"] else -1 if t < st["replicas"] else 0
267             st["replicas"] = max(20, min(500, t))
268         elif k == "nodes":
269             t = max(cfg.minimum_nodes, min(cfg.maximum_nodes, int(a["target"])))
270             nd = 1 if t > st["actual_nodes"] else -1 if t < st["actual_nodes"] else 0
271             st["actual_nodes"] += int(np.sign(t - st["actual_nodes"])) * min(2, abs(t - st["actual_nodes"]))
272         elif k == "terraform_plan":
273             st["terraform_pending_target"] = max(cfg.minimum_nodes, min(cfg.maximum_nodes, int(a["target"])))
274             st["terraform_apply_step"] = step + cfg.terraform_apply_delay_steps
275         elif k == "power_cap":
276             st["power_cap"] = clamp(float(a["target"]), 0.65, 1.0)
277         elif k == "route_shift":
278             st["network_recovery"] = max(st["network_recovery"], 0.16)
279         elif k == "rollback":
280             st["replicas"] = max(60, int(round(st["replicas"] * 0.88)))
281         elif k == "page":
282             pg += 1
283     pages += pg
284     node_travel += abs(st["actual_nodes"] - n_before)
285     node_starts += max(0, st["actual_nodes"] - n_before)
286     node_stops += max(0, n_before - st["actual_nodes"])
287     cap_travel += abs(st["power_cap"] - cap_before)
288     if len(hist) > 1:
289         thermal_travel += abs(hist[-1]["thermal"] - hist[-2]["thermal"])
290     if pg and st["operator_page_step"] is None:
291         st["operator_page_step"] = step
292         st["operator_ack_step"] = step + scn.operator_delay_steps
293     if st["operator_ack_step"] is not None and step >= st["operator_ack_step"]:
294         manual += 1
295         st["operator_page_step"] = st["operator_ack_step"] = None
296         st["power_cap"] = max(st["power_cap"], 0.90)
297         st["actual_nodes"] = min(cfg.maximum_nodes, st["actual_nodes"] + 1)
298     if nd and st["last_node_direction"] and nd != st["last_node_direction"]:
299         rev += 1
300     if rd and st["last_replica_direction"] and rd != st["last_replica_direction"]:
301         rev += 1
302     if nd: st["last_node_direction"] = nd
303     if rd: st["last_replica_direction"] = rd
304     if src["security_block"] > 0.5 and any(a["action"] in ("nodes", "terraform_plan", "rollout") and a[
"direction"] > 0 for a in acts):
305         secv += 1
306         req = max(cfg.minimum_nodes, min(cfg.maximum_nodes, int(math.ceil((demand + st["pending_queue"] * 0.45) /
cfg.capacity_units_per_node))))
307         unn += max(0, st["actual_nodes"] - req) * cfg.dt_hours
308         nh += st["actual_nodes"] * cfg.dt_hours
309         energy += pkw * cfg.dt_hours
310         peak = max(peak, pkw)
311         qv.append(st["pending_queue"])
312         v_queue += int(qratio > 0.35); v_power += int(pstress > 1.05); v_heat += int(st["thermal"] > 1.03)
313         phys = qratio > 0.35 or pstress > 1.05 or st["thermal"] > 1.03
314         sla_p += int(phys)
315         sla_c += int(phys or c_step > 0)
316         adm = qratio < 0.28 and pstress <= 1.02 and st["thermal"] < 0.96 and c_step == 0
317         adm_n += int(adm)
318         streak = streak + 1 if adm else 0

```

```

319         if S._mom_event_active(scen, step):
320             rec_pending, rec_start = True, scen.event_start + scen.event_duration
321         elif rec_pending and rec_start is not None and step >= rec_start and streak >= cfg.certification_streak:
322             rec.append((step - rec_start + 1) * cfg.interval_minutes); rec_pending = False
323         trace.append((round(qratio, 12), round(pkpw, 12), st["actual_nodes"], st["replicas"], round(st[
"power_cap"], 12)))
324     n = cfg.steps
325     return {"scenario_id": scen.scenario_id, "family": scen.family, "arm": arm,
326            "availability": done / max(1e-9, dem), "sla_violation_physical": sla_p / n,
327            "sla_violation_total": sla_c / n, "violation_backlog": v_queue / n, "violation_power": v_power / n,
328            "violation_heat": v_heat / n, "time_healthy": adm_n / n,
329            "mean_queue": float(np.mean(qv)), "p95_queue": float(np.percentile(qv, 95)),
330            "recovery_minutes": float(np.mean(rec)) if rec else n * cfg.interval_minutes, "recovered": int(bool(
rec)),
331            "energy_kwh": energy, "peak_power_kw": peak, "node_hours": nh, "idle_node_hours": unn,
332            "contradictions": contra, "scale_reversals": rev, "security_violations": secv, "pages": pages,
333            "human_interventions": manual, "actions": nact, "omni_authority_steps": authority_steps,
334            "node_start_stop": node_travel, "machines_started": node_starts, "machines_stopped": node_stops,
335            "machine_round_trips": (node_starts + node_stops - abs(st["actual_nodes"] - cfg.initial_nodes)) / 2.
0, "power_cap_travel": cap_travel, "thermal_travel": thermal_travel,
336            "invariant_violations": inv_viol, "invariant_violations_ex_power": inv_viol_ex_power,
337            "shield_interventions": shield_hits,
338            "trace_hash": hashlib.sha256(repr(trace).encode()).hexdigest()[16]}
339
340
341 LOWER = ["machine_round_trips", "scale_reversals", "node_start_stop", "power_cap_travel", "thermal_travel",
"invvariant_violations", "invariant_violations_ex_power", "sla_violation_physical", "sla_violation_total",
"violation_backlog", "violation_power", "violation_heat", "mean_queue", "p95_queue", "recovery_minutes",
"energy_kwh",
342         "peak_power_kw", "node_hours", "idle_node_hours", "contradictions", "scale_reversals",
"security_violations",
343         "pages", "human_interventions"]
344 HIGHER = ["availability", "time_healthy", "recovered"]
345 NEUTRAL = ["machines_started", "machines_stopped"]
346
347
348 def paired(rows, base, cand, metric, rng):
349     A = {r["scenario_id"]: r[metric] for r in rows if r["arm"] == base}
350     B = {r["scenario_id"]: r[metric] for r in rows if r["arm"] == cand}
351     ids = sorted(A)
352     d = np.array([B[i] - A[i] for i in ids], float)
353     boots = d[rng.integers(0, len(d), (4000, len(d)))].mean(axis=1)
354     lo, hi = np.percentile(boots, [2.5, 97.5])
355     good = d < 0 if metric in LOWER else d > 0
356     bad = d > 0 if metric in LOWER else d < 0
357     return {"metric": metric, "base_mean": float(np.mean([A[i] for i in ids])), "cand_mean": float(np.mean([B[i]
for i in ids])),
358            "mean_delta": float(d.mean()), "ci95": [float(lo), float(hi)], "better": int(good.sum()),
359            "worse": int(bad.sum()), "tied": int(len(d) - good.sum() - bad.sum()), "n": len(d)}
360
361
362 def run(scenarios: int, seed: int, out: Path, law: AllocationLaw, arms=ARMS) -> Dict[str, Any]:
363     cfg = S.ManagerBenchmarkConfig(profile="full", scenarios=scenarios, steps=72)
364     scns = S._mom_generate_scenarios(cfg, seed)[scenarios]
365     rows = [simulate(s, a, cfg, law) for s in scns for a in arms]
366     out.mkdir(parents=True, exist_ok=True)
367     with open(out / "RUNS.csv", "w", newline="") as f:
368         w = csv.DictWriter(f, fieldnames=list(rows[0])); w.writeheader(); w.writerows(rows)
369     rng = np.random.default_rng(12345)
370     means = {a: {m: float(np.mean([r[m] for r in rows if r["arm"] == a])) for m in HIGHER + LOWER + NEUTRAL} for
a in arms}
371     comps = {}
372     for base in [b for b in BASELINES if b in arms]:
373         for cand in [a for a in arms if a not in BASELINES]:
374             comps[f"{cand}_vs_{base}"] = [paired(rows, base, cand, m, rng) for m in HIGHER + LOWER]
375     for x, y in ((("omni_k8s_protect", "omni_direct"), ("omni_k8s_throughput", "omni_k8s_protect"),
376                  ("omni_k8s_throughput", "omni_k8s_throughput_no_dynamics"), ("omni_k8s_throughput",
"omni_k8s_throughput_no_gate"))):
377         if x in arms and y in arms:
378             comps[f"{x}_vs_{y}"] = [paired(rows, y, x, m, rng) for m in HIGHER + LOWER]
379     for abl in ("omni_direct_no_dynamics", "omni_direct_no_shield", "omni_pipeline"):
380         if abl in arms and "omni_direct" in arms:
381             comps[f"omni_direct_vs_{abl}"] = [paired(rows, abl, "omni_direct", m, rng) for m in HIGHER + LOWER]
382     tails = {a: {"time_healthy_p05": float(np.percentile([r["time_healthy"] for r in rows if r["arm"] == a], 5)),
383                  "recovery_minutes_p95": float(np.percentile([r["recovery_minutes"] for r in rows if r["arm"] ==
a], 95)),
384                  "sla_violation_physical_p95": float(np.percentile([r["sla_violation_physical"] for r in rows if
r["arm"] == a], 95)),
385                  "energy_kwh_p95": float(np.percentile([r["energy_kwh"] for r in rows if r["arm"] == a], 95)),
386                  "time_healthy_min": float(min([r["time_healthy"] for r in rows if r["arm"] == a])),
387                  "recovered_fraction": float(np.mean([r["recovered"] for r in rows if r["arm"] == a]))} for a in
arms}
388     obs_identical = None
389     k8s_obs_identical = None
390     if "omni_k8s_observe" in arms and "k8s_ref_70" in arms:
391         tk = {r["scenario_id"]: r["trace_hash"] for r in rows if r["arm"] == "k8s_ref_70"}

```

```

392         to2 = {r["scenario_id"]: r["trace_hash"] for r in rows if r["arm"] == "omni_k8s_observe"}
393         k8s_obs_identical = sum(tk[i] == to2[i] for i in tk)
394     if "omni_observe" in arms:
395         tn = {r["scenario_id"]: r["trace_hash"] for r in rows if r["arm"] == "native"}
396         to = {r["scenario_id"]: r["trace_hash"] for r in rows if r["arm"] == "omni_observe"}
397         obs_identical = sum(tn[i] == to[i] for i in tn)
398         fam = {}
399         for f in sorted({r["family"] for r in rows}):
400             fam[f] = {a: {m: float(np.mean([r[m] for r in rows if r["arm"] == a and r["family"] == f]))
401                        for m in ("time_healthy", "sla_violation_physical", "recovery_minutes", "contradictions",
402                                "energy_kwh")}}
403             for a in ("native", "k8s_ref_70", "omni_k8s_protect", "omni_k8s_throughput") if a in arms:
404                 summary = {"scenarios": len(scns), "steps": cfg.steps, "seed_base": seed, "arms": arms,
405                           "allocation_law": asdict(law), "means": means, "paired": comps, "by_family": fam,
406                           "observe_mode_identical_to_native": obs_identical,
407                           "layered_observe_identical_to_kubernetes": k8s_obs_identical, "tails": tails}
408             (out / "SUMMARY.json").write_text(json.dumps(summary, indent=2))
409         return summary
410
411 def main():
412     ap = argparse.ArgumentParser()
413     ap.add_argument("--scenarios", type=int, default=500)
414     ap.add_argument("--seed", type=int, required=True)
415     ap.add_argument("--out", required=True)
416     a = ap.parse_args()
417     s = run(a.scenarios, a.seed, Path(a.out), AllocationLaw())
418     print(json.dumps({k: {m: round(v, 4) for m, v in s["means"][k].items()} for k in s["arms"]}, indent=1))
419     print("observe identical to native:", s["observe_mode_identical_to_native"], "/", s["scenarios"])
420
421
422 if __name__ == "__main__":
423     main()
424

```

verify.py

315 lines SHA-256 8f819c2021b12df0d2a04eb73d6503ec6d69c62e87db460f2cc57559fcb42d2

```
1  #!/usr/bin/env python3
2  """One-command verification of the Omni-Compass package.
3
4  python verify.py          full verification (about 15-30 minutes, single core)
5  python verify.py --full-replay  additionally re-runs all held-out scenarios for every arm and requires the
6                                regenerated RUNS.csv and SUMMARY.json to equal the shipped files
7  python verify.py --quick    skips the 20,000-case engine comparison, runs a 1,000,000-decision soak
8                                instead of 100,000,000, and replays fewer scenarios
9  """
10 import argparse, csv, hashlib, json, shutil, subprocess, sys, tempfile
11 from pathlib import Path
12 ROOT = Path(__file__).resolve().parent
13 sys.path.insert(0, str(ROOT))
14 ok = True
15
16 def check(name, cond, detail=""):
17     global ok
18     print(("PASS " if cond else "FAIL ") + name + (f" ({detail})" if detail else ""))
19     ok = ok and bool(cond)
20
21 def sha(p): return hashlib.sha256(Path(p).read_bytes()).hexdigest()
22
23 def build_cpp(tmp):
24     src = ROOT / "cpp"; b = tmp / "build"
25     if shutil.which("cmake"):
26         subprocess.run(["cmake", "-S", str(src), "-B", str(b), "-DCMAKE_BUILD_TYPE=Release"], check=True,
27             capture_output=True)
28         subprocess.run(["cmake", "--build", str(b), "-j2"], check=True, capture_output=True)
29     else:
30         b.mkdir(parents=True)
31         cxx = shutil.which("g++") or shutil.which("clang++")
32         srcs = [str(p) for p in (src / "src").glob("*.cpp")]
33         for tool, name in ((("run_fixture.cpp", "oc_run_fixture"), ("run_governor.cpp", "oc_governor"), ("soak.
34             cpp", "oc_soak"),
35             ("smoke.cpp", "oc_smoke"), ("run_shield.cpp", "oc_shield"), ("run_hpa.cpp", "oc_hpa"),
36             ("savings.cpp", "oc_savings"),
37             ("run_closure.cpp", "oc_closure"), ("run_conveyance.cpp", "oc_conveyance"), (
38             "run_twins.cpp", "oc_twins")):
39             subprocess.run([cxx, "-std=c++20", "-O2", "-I", str(src / "include"), *srcs, str(src / "tools" /
40                 tool), "-o", str(b / name)], check=True)
41         return b
42
43 def main():
44     ap = argparse.ArgumentParser(); ap.add_argument("--quick", action="store_true"); ap.add_argument("--full-
45     replay", action="store_true")
46     a = ap.parse_args()
47     sys.path.insert(0, str(ROOT / "tools"))
48     from code_fingerprint import fingerprint
49     prov = json.loads((ROOT / "reference" / "PROVENANCE.json").read_text())
50     ref = ROOT / prov["reference_engine"]
51     import platform
52     pyver = ".".join(platform.python_version_tuple()[2:])
53     check("reference engine SHA-256", sha(ref) == prov["reference_engine_sha256"])
54     reg = json.loads((ROOT / "results" / "PREREGISTRATION.json").read_text())
55     amend = {a["file"]: a for a in json.loads((ROOT / "results" / "LOCK_AMENDMENTS.json").read_text())[
56         "amendments"]}
57
58     def locked(f, h): # unchanged, or changed exactly as a recorded amendment (bug fix) says
59         cur = sha(ROOT / f)
60         if cur == h:
61             return True, ""
62         a = amend.get(f)
63         ok = a is not None and h in a["from_sha256"] and cur == a["to_sha256"]
64         return ok, (f"amended in {a['commit']}: {a['reason'][:60]}..." if ok else "")
65
66     for f, h in reg["sha256"].items():
67         ok, note = locked(f, h)
68         check(f"pre-registered file unchanged or amended on record (SHA-256): {f}", ok, note)
69
70     if pyver == reg.get("fingerprint_python", ""):
71         check("reference engine program fingerprint", fingerprint(ref) == prov["program_fingerprint"])
72         for f, h in reg["code_fingerprint"].items():
73             check(f"program fingerprint: {f}", fingerprint(ROOT / f) == h)
74     else:
75         print(f"SKIP program fingerprints (recorded with Python {reg.get('fingerprint_python')}, running {
76             pyver}; "
77             "ast.dump output is interpreter-version dependent; SHA-256 checks above are authoritative)")
78     from tests import test_stack_sim_provenance, test_native_equivalence, test_cpp_governor_parity
79     test_stack_sim_provenance.main(); check("stack simulator provenance", True)
80     if not a.quick:
81         from tests import test_core_parity
82         test_core_parity.main(); check("core parity vs reference engine and 500 fixtures", True)
83     test_native_equivalence.main(20 if a.quick else 60); check("native arm equals reference simulator", True)
84     tmp = Path(tempfile.mkdtemp()); b = build_cpp(tmp)
85     subprocess.run([str(b / "oc_run_fixture"), str(ROOT / "fixtures" / "current_500_inputs.csv"), str(tmp / "cpp.
```

```

csv"]], check=True)
76 r = subprocess.run([sys.executable, str(ROOT / "cpp" / "scripts" / "compare_results.py"), str(ROOT /
"fixtures" / "current_500_expected.csv"),
77 str(tmp / "cpp.csv"), str(tmp / "cpp.json")], capture_output=True, text=True)
78 rep = json.loads((tmp / "cpp.json").read_text())
79 check("C++ core vs 500 fixtures", rep["pass"], f"max abs {rep['max_abs_error']:.2e}")
80 test_cpp_governor_parity.main(str(b / "oc_governor"), 20 if a.quick else 60); check("C++ governor vs Python
governor (power_protect)", True)
81 test_cpp_governor_parity.main(str(b / "oc_governor"), 20 if a.quick else 60, "throughput"); check("C++
governor vs Python governor (throughput)", True)
82 for m in ("fleet", "fleet_balanced", "fleet_wear"):
83 test_cpp_governor_parity.main(str(b / "oc_governor"), 20 if a.quick else 60, m); check(f"C++ governor
vs Python governor ({m})", True)
84 from tests import test_cpp_shield_parity
85 test_cpp_shield_parity.main(str(b / "oc_shield"), 8 if a.quick else 20); check("C++ shield vs Python shield (
enforce and violations)", True)
86 from tests import test_cpp_shield_adversarial, test_shield_properties
87 test_shield_properties.main(100_000 if a.quick else 1_000_000); check("shield properties: no invariant
violated, idempotent, no invented action, minimal intervention (random + adversarial)", True)
88 test_cpp_shield_adversarial.main(str(b / "oc_shield"), 50_000 if a.quick else 200_000); check("C++ shield vs
Python shield on the adversarial generator", True)
89 from tests import test_cpp_closure_parity
90 test_cpp_closure_parity.main(str(b / "oc_closure"), 1 if a.quick else 3); check("C++ closure law vs Python
closure law, every decision and reserve", True)
91 from tests import test_cpp_conveyance_parity
92 w = test_cpp_conveyance_parity.main(str(b / "oc_conveyance"), 50 if a.quick else 200)
93 check("C++ conveyance law vs Python conveyance law, every allocation (random systems and the CPU+GPU
layout)", w < 1e-9)
94 from tests import test_cpp_twins_parity
95 test_cpp_twins_parity.main(str(b / "oc_twins"), 500 if a.quick else 2000)
96 check("C++ nervous system, compass and ledger, GPU rules vs Python, value for value", True)
97 from tools import seal
98 broken = seal.check()
99 check(f"seal: all {len(seal.TWINS)} Python/C++ twins unchanged since proven equal (results/SEAL.json)", not
broken, "; ".join(broken))
100 from tools import mechanism_identity, tracking_bounds
101 bad = mechanism_identity.check()
102 check("mechanism identity: F, Theta, C, h, G, M_act, dt, A of both configurations match results/
MECHANISM_IDENTITY.json", not bad, "; ".join(bad))
103 tb = tracking_bounds.main() # recomputed; a changed result also shows as a manifest mismatch below
104 nt = tb["nearest_target"]
105 check("tracking theorem checks (V): drift under F_bar < U_AUTHORITY, discrete ultimate bound inside the
basin, no invariance failure",
106 nt["drift_within_bound_all"] and nt["F_bar_below_U_AUTHORITY_all"] and nt[
"ultimate_bound_inside_basin"]
107 and nt["invariance_failures"] == 0 and tb["wrong_target"]["invariance_failures"] == 0)
108 from tools import release_manifest
109 bad = release_manifest.check()
110 check("release manifest: engine, C++ seal, GPU protocol, live evidence and license match RELEASE_MANIFEST.
json", not bad, "; ".join(bad))
111 src = (ROOT / "cpp" / "src" / "shield.cpp").read_text()
112 mut_src = tmp / "shield_mut.cpp"
113 mut_src.write_text(src.replace('return k == "nodes" || k == "terraform_plan" || k == "rollout"; }, 'return
k == "nodes" || k == "terraform_plan"; }'))
114 others = [str(p) for p in (ROOT / "cpp" / "src").glob("*.cpp") if p.name != "shield.cpp"]
115 subprocess.run([shutil.which("g++") or shutil.which("clang++"), "-std=c++20", "-O2", "-I", str(ROOT / "cpp" /
"include"), *others, str(mut_src),
116 str(ROOT / "cpp" / "tools" / "run_shield.cpp"), "-o", str(tmp / "oc_shield_mut")],
check=True)
117 try:
118 test_cpp_shield_parity.main(str(tmp / "oc_shield_mut"), 8)
119 det3 = False
120 except AssertionError:
121 det3 = True
122 check("negative control: parity test detects a C++ shield that no longer blocks rollouts during a security
block", det3)
123 from tests import test_cpp_hpa_parity
124 test_cpp_hpa_parity.main(str(b / "oc_hpa"), (901,) if a.quick else (901, 902)); check("independent C++ HPA
replica law vs fleet harness HPA", True)
125 hsrc = (ROOT / "cpp" / "include" / "omnicompass" / "hpa.hpp").read_text()
126 hm = tmp / "hmut" / "omnicompass"; hm.mkdir(parents=True)
127 for p in (ROOT / "cpp" / "include" / "omnicompass").glob("*.hpp"):
128 (hm / p.name).write_text(p.read_text())
129 (hm / "hpa.hpp").write_text(hsrc.replace("tolerance{0.1}", "tolerance{0.05}"))
130 subprocess.run([shutil.which("g++") or shutil.which("clang++"), "-std=c++20", "-O2", "-I", str(tmp / "hmut"),
str(ROOT / "cpp" / "src" / "hpa.cpp"),
131 str(ROOT / "cpp" / "tools" / "run_hpa.cpp"), "-o", str(tmp / "oc_hpa_mut")], check=True)
132 try:
133 test_cpp_hpa_parity.main(str(tmp / "oc_hpa_mut"), (901,))
134 det4 = False
135 except AssertionError:
136 det4 = True
137 check("negative control: parity test detects a C++ HPA with 5% instead of 10% tolerance", det4)
138 r = subprocess.run([str(b / "oc_smoke")], capture_output=True, text=True)
139 check("C++ governor contract smoke test (bounds, security, determinism, both modes)", r.returncode == 0 and
"SMOKE PASS" in r.stdout)

```

```

140     n_soak = 1000000 if a.quick else 100000000
141     rec_all = json.loads((ROOT / "results" / "SOAK.json").read_text())
142     for mode in ("power_protect", "throughput"):
143         soak = json.loads(subprocess.run([str(b / "oc_soak"), str(n_soak), mode], capture_output=True,
144 text=True).stdout)
145         cps = soak["rss_kb_at_25_50_75_100pct"]
146         check(f"soak test ({mode}), {n_soak:,} decisions, no failures, memory flat after warm-up", soak[
147 "nonfinite_or_out_of_range"] == 0 and max(cps) == min(cps),
148         f"{soak['ns_per_decision']:.0f} ns per decision, RSS checkpoints {cps} kB")
149         if not a.quick:
150             rec = rec_all[mode]
151             same = all(soak[k] == rec[k] for k in ("decisions", "nonfinite_or_out_of_range", "E", "U", "I_U",
152 "S", "B"))
153             check(f"soak test ({mode}) reproduces results/SOAK.json (decisions, failures, state ranges)", same)
154             hdr = (ROOT / "cpp" / "include" / "omnicompass" / "governor.hpp").read_text()
155             mut = tmp / "mut" / "omnicompass"; mut.mkdir(parents=True)
156             for src_h in (ROOT / "cpp" / "include" / "omnicompass").glob("*.hpp"): # every header, so new modules
157                 compile too
158                 (mut / src_h.name).write_text(src_h.read_text())
159                 (mut / "governor.hpp").write_text(hdr.replace("down_dwell{4}", "down_dwell{3}"))
160                 cxx = shutil.which("g++") or shutil.which("clang++")
161                 subprocess.run([cxx, "-std=c++20", "-O2", "-I", str(tmp / "mut"), *[str(p) for p in (ROOT / "cpp" / "src").
162 glob("*.cpp")],
163                             str(ROOT / "cpp" / "tools" / "run_governor.cpp"), "-o", str(tmp / "oc_mut")], check=True)
164             try:
165                 test_cpp_governor_parity.main(str(tmp / "oc_mut"), 20)
166                 detected = False
167             except AssertionError:
168                 detected = True
169             check("negative control: parity test detects a mutated C++ governor (dwell 3 instead of 4)", detected)
170             (mut / "governor.hpp").write_text(hdr.replace("push_release{0.005}", "push_release{9.0}"))
171             subprocess.run([cxx, "-std=c++20", "-O2", "-I", str(tmp / "mut"), *[str(p) for p in (ROOT / "cpp" / "src").
172 glob("*.cpp")],
173                             str(ROOT / "cpp" / "tools" / "run_governor.cpp"), "-o", str(tmp / "oc_mut2")], check=True)
174             try:
175                 test_cpp_governor_parity.main(str(tmp / "oc_mut2"), 20)
176                 detected2 = False
177             except AssertionError:
178                 detected2 = True
179             check("negative control: parity test detects a C++ governor without the equation (2) release gate",
180 detected2)
181             (mut / "governor.hpp").write_text(hdr.replace("push_add{-9.0}", "push_add{9.0}"))
182             subprocess.run([cxx, "-std=c++20", "-O2", "-I", str(tmp / "mut"), *[str(p) for p in (ROOT / "cpp" / "src").
183 glob("*.cpp")],
184                             str(ROOT / "cpp" / "tools" / "run_governor.cpp"), "-o", str(tmp / "oc_mut3")], check=True)
185             try:
186                 test_cpp_governor_parity.main(str(tmp / "oc_mut3"), 20)
187                 detected3 = False
188             except AssertionError:
189                 detected3 = True
190             check("negative control: parity test detects a C++ governor with a wrong equation (2) add gate", detected3)
191             from benchmarks.stack_benchmark import simulate, ARMS
192             from omnicompass import stack_sim as S
193             from omnicompass.adapter import AllocationLaw
194             n = 5 if a.quick else 20
195             for sd in regl["held_out_seeds"]:
196                 rows = list(csv.DictReader(open(ROOT / "results" / f"heldout_seed_{sd}" / "RUNS.csv")))
197                 cfg = S.ManagerBenchmarkConfig(profile="full", scenarios=500, steps=72)
198                 scns = S._mom_generate_scenarios(cfg, sd)[:n]; bad = 0
199                 for s in scns:
200                     for arm in ARMS:
201                         new = simulate(s, arm, cfg, AllocationLaw())
202                         old = next(r for r in rows if int(r["scenario_id"]) == s.scenario_id and r["arm"] == arm)
203                         bad += int(new["trace_hash"] != old["trace_hash"] or abs(new["time_healthy"] - float(old[
204 "time_healthy"]))) > 1e-12)
205             check(f"held-out seed {sd}: {n} scenarios x {len(ARMS)} arms replay identically", bad == 0)
206             summ = json.loads((ROOT / "results" / f"heldout_seed_{sd}" / "SUMMARY.json").read_text())
207             check(f"held-out seed {sd}: observe mode identical to native", summ["observe_mode_identical_to_native"]
208 == summ["scenarios"])
209             for arm in ("omni_k8s_protect", "omni_direct", "omni_direct_no_dynamics"):
210                 check(f"held-out seed {sd}: {arm} invariant violations I1-I5 = 0", summ["means"][arm][
211 "invariant_violations"] == 0.0)
212             check(f"held-out seed {sd}: omni_k8s_throughput invariant violations I1-I3, I5 = 0 (I4 not enforced in
213 throughput mode)",
214                 summ["means"]["omni_k8s_throughput"]["invariant_violations_ex_power"] == 0.0)
215             check(f"held-out seed {sd}: layered observe mode identical to Kubernetes alone",
216                 summ["layered_observe_identical_to_kubernetes"] == summ["scenarios"])
217             freg = json.loads((ROOT / "results" / "fleet" / "PREREGISTRATION.json").read_text())
218             for f, h in freg["sha256"].items():
219                 if f != "omnicompass/adapter.py":
220                     ok, note = locked(f, h)
221                     check(f"fleet pre-registered file unchanged or amended on record (SHA-256): {f}", ok, note)
222             from fleet.sim import run as frun, arms_for
223             from fleet.harness import make_scenario
224             frows = list(csv.DictReader(open(ROOT / "results" / "fleet" / "heldout" / "RUNS.csv")))
225             fsumm = json.loads((ROOT / "results" / "fleet" / "heldout" / "SUMMARY.json").read_text())

```

```

214     for v in ("web", "multi", "batch", "gpu", "gpu_always_on"):
215         seeds = range(freg["held_out_seed_base"], freg["held_out_seed_base"] + (2 if a.quick else 5))
216         bad = 0
217         for s_ in seeds:
218             scn = make_scenario(v, s_)
219             for arm in arms_for(v):
220                 new = frun(scn, arm)
221                 old = next(r for r in frows if r["vessel"] == v and int(r["seed"]) == s_ and r["arm"] == arm)
222                 bad += int(new["trace_hash"] != old["trace_hash"] or abs(float(new["energy_kwh"]) - float(old[
"energy_kwh"]))) > 1e-9)
223     check(f"fleet held-out ({v}): {len(seeds)} scenarios x {len(arms_for(v))} arms replay identically", bad
== 0)
224     check(f"fleet ({v}): observe mode identical to HPA 0.7 + Cluster Autoscaler",
225           fsum["vessels"][v]["observe_identical_to_hpa70_ca"] == fsum["vessels"][v]["scenarios"])
226     from fleet import planetlab as PLm
227     pl_rows = list(csv.DictReader(open(ROOT / "results" / "fleet" / "planetlab" / "RUNS.csv")))
228     tr = PLm.load_dir(ROOT / "fleet" / "traces" / "planetlab"); badp = 0
229     for s_ in range(800000, 800000 + (1 if a.quick else 3)):
230         scn = PLm.make_scenario(tr, s_)
231         for arm in arms_for("web"):
232             new = frun(scn, arm); old = next(r for r in pl_rows if int(r["seed"]) == s_ and r["arm"] == arm)
233             badp += int(new["trace_hash"] != old["trace_hash"])
234     check("recorded-trace (PlanetLab) runs replay identically", badp == 0)
235     tp = json.loads((ROOT / "fleet" / "traces" / "PROVENANCE.json").read_text())
236     check("recorded traces match their registered SHA-256", all(hashlib.sha256((ROOT / "fleet" / "traces" /
"planetlab" / n).read_bytes()).hexdigest() == v["sha256"] for n, v in tp["files"].items()))
237     from tests import test_hpa_three_way
238     test_hpa_three_way.main((901,) if a.quick else (901, 902)); check("three-way HPA agreement (harness, C++,
external implementation with upstream window)", True)
239     import benchmarks.savings as SV
240     sv_tmp = tmp / "savings"; sv_tmp.mkdir()
241     SV.main(str(b / "oc_savings"), sv_tmp)
242     check("C++ savings projector matches Python reference and reproduces results/SAVINGS.csv",
243           (sv_tmp / "SAVINGS.csv").read_text() == (ROOT / "results" / "SAVINGS.csv").read_text())
244     if not a.quick:
245         from tests import test_selfpilot
246         test_selfpilot.main(); check("end-to-end self-pilot: shipped controller, simulated cluster, capture and
scoring (energy and machines no worse, no significant pending-pod increase)", True)
247         from tests import test_pilot_score
248         test_pilot_score.main(); check("pilot scoring: detects a real gain, no false gain on identical clusters,
detects a service regression", True)
249         from tests import test_omni_controller
250         test_omni_controller.main(); check("live controller against a fake cluster: observe writes nothing, target
bounded, kill restores, node pool bounded and dry-run safe", True)
251         from tests import test_muscles
252         test_muscles.main(); check("live muscles: power cap, heat, security, rollout, batch, CPU frequency and GPU
connectors; kill restores; observe writes nothing", True)
253         from tests import test_nervous_system
254         test_nervous_system.main(50_000 if a.quick else 300_000); check("nervous system: observe executes nothing,
kill removes authority, security blocks capacity expansion, stress never grants contraction, envelopes in
hardware range, deterministic", True)
255         from tests import test_conveyance
256         test_conveyance.main(); check("conveyance law: budget conserved, Lyapunov ascent, exact exponential
convergence, bounds, budget never exceeded", True)
257         from tests import test_two_way
258         test_two_way.main(); check("two-way nervous system: blind senses never grant contraction, frozen probe
detected, failures read as pressure, no new order before the last one landed", True)
259         from tests import test_node_release_gate
260         test_node_release_gate.main(); check("machine-organ release gate: own engine view, pods-first coordination,
headroom at the engine's rho", True)
261         from tests import test_hardware_plant
262         test_hardware_plant.main(); check("hardware plant: GPU arms B and C distinct; C = min(engine cap, (want/
rho)^(1/gamma)); no lost work, no extra heat", True)
263         from tests import test_node_exchange
264         test_node_exchange.main(); check("CPU+GPU on one budget: never over the site budget; CPU-measured-only goes
over; more work than today's fixed cap", True)
265         from tests import test_cpufreq_contract
266         test_cpufreq_contract.main(); check("CPU-frequency lever via kernel policy files: schedutil floor, nervous
envelope, never above the operator ceiling, exact restore on kill", True)
267         from tests import test_strict_replicas
268         test_strict_replicas.main(); check("strict C live mode: Omni-Compass sets the replica floor, growth free,
kill restores the range", True)
269         from tests import test_compass
270         test_compass.main(); check("the compass: face, axle, closed circles, inwardness on the boundary, ledger
descent", True)
271         from tests import test_living_band
272         test_living_band.main(); check("living band: every level the nervous system hands out stays inside 5%..95%;
budget kept", True)
273         from tests import test_pod_reflex
274         test_pod_reflex.main(); check("fast pod reflex: the HPA's own rule read from the live queue; calm writes
nothing; hands back; kill restores", True)
275         from tests import test_muscles_levers
276         test_muscles_levers.main(); check("live levers: rightsize, coldstart, batch pace, containment, cooling; kill
restores every one, also from a fresh process", True)
277         from tests import test_pod_starts
278         test_pod_starts.main(); check("exact pod-start timing from the API server's watch stream (creation to Ready,

```

```

window-bounded)", True)
279     from tests import test_gpu_bench
280     test_gpu_bench.main(); check("GPU bench: watch writes nothing, limit never below draw x 1.3, read-back,
blind and SLO reflexes, kill restores; paired run validity", True)
281     from tests import test_convey
282     test_convey.main(); check("energy to where the work is: idle machine CPU conveyed to serving pods in place,
band-bounded, kill restores", True)
283     from tests import test_active_nodes
284     test_active_nodes.main(); check("full-engine controller options: parked and control-plane nodes excluded,
kill switch restores the node pool once", True)
285     from tests import test_schedutil, test_cpufreq_ceiling
286     test_schedutil.main(); check("schedutil model: 1.25 map tips at 80%, OPP snap, uclamp, RT to policy max,
rate limit, iowait boost, Omni ceiling and kill", True)
287     test_cpufreq_ceiling.main(); check("cpufreq ceiling writer: scaling_max_freq on every policy, clamped;
restore puts cpuinfo_max_freq back", True)
288     from tests import test_failsafe
289     test_failsafe.main(); check("no automated fallback: failed decisions are skipped; only the human switch
turns the whole harness off and on", True)
290     from tests import test_fleet_realdata_paths
291     test_fleet_realdata_paths.main(); check("capture replay and PlanetLab vessel on inputs in the real formats",
True)
292     r = subprocess.run(["bash", "-n", str(ROOT / "fleet" / "capture" / "kube_capture.sh")], capture_output=True)
293     check("cluster capture script parses", r.returncode == 0)
294     import benchmarks.multiplicity as MP
295     mp_tmp = tmp / "mult.json"; MP.main(mp_tmp)
296     check("multiplicity analysis reproduces results/MULTIPLICITY.json",
297           json.loads(mp_tmp.read_text()) == json.loads((ROOT / "results" / "MULTIPLICITY.json").read_text()))
298     if a.full_replay:
299         from benchmarks.stack_benchmark import run as run_bench
300         for sd in reg["held_out_seeds"]:
301             out = tmp / f"replay_{sd}"
302             run_bench(reg["held_out_scenarios_per_seed"], sd, out, AllocationLaw())
303             same_runs = (out / "RUNS.csv").read_bytes() == (ROOT / "results" / f"heldout_seed_{sd}" / "RUNS.
csv").read_bytes()
304             new_s = json.loads((out / "SUMMARY.json").read_text()); old_s = json.loads((ROOT / "results" /
f"heldout_seed_{sd}" / "SUMMARY.json").read_text())
305             new_law, old_law = new_s.pop("allocation_law"), old_s.pop("allocation_law")
306             inert = {"push_add": -9.0}
307             law_ok = all(new_law.get(k) == v for k, v in old_law.items()) and all(k in inert and new_law[k] ==
inert[k] for k in set(new_law) - set(old_law))
308             same_summary = new_s == old_s and law_ok
309             check(f"full replay, held-out seed {sd}: all scenarios x all arms regenerate RUNS.csv exactly and
SUMMARY.json exactly (recorded law may add only inert defaults)", same_runs and same_summary)
310             print("\nVERIFICATION:", "PASS" if ok else "FAIL")
311             sys.exit(0 if ok else 1)
312
313 if __name__ == "__main__":
314     main()
315

```